

2025 Workshop on



Remote Memory ~ Local Memory over Reconfigurable Ethernet Fabric^[1]

Vishal Shrivastav

Joint work with Weigao Su

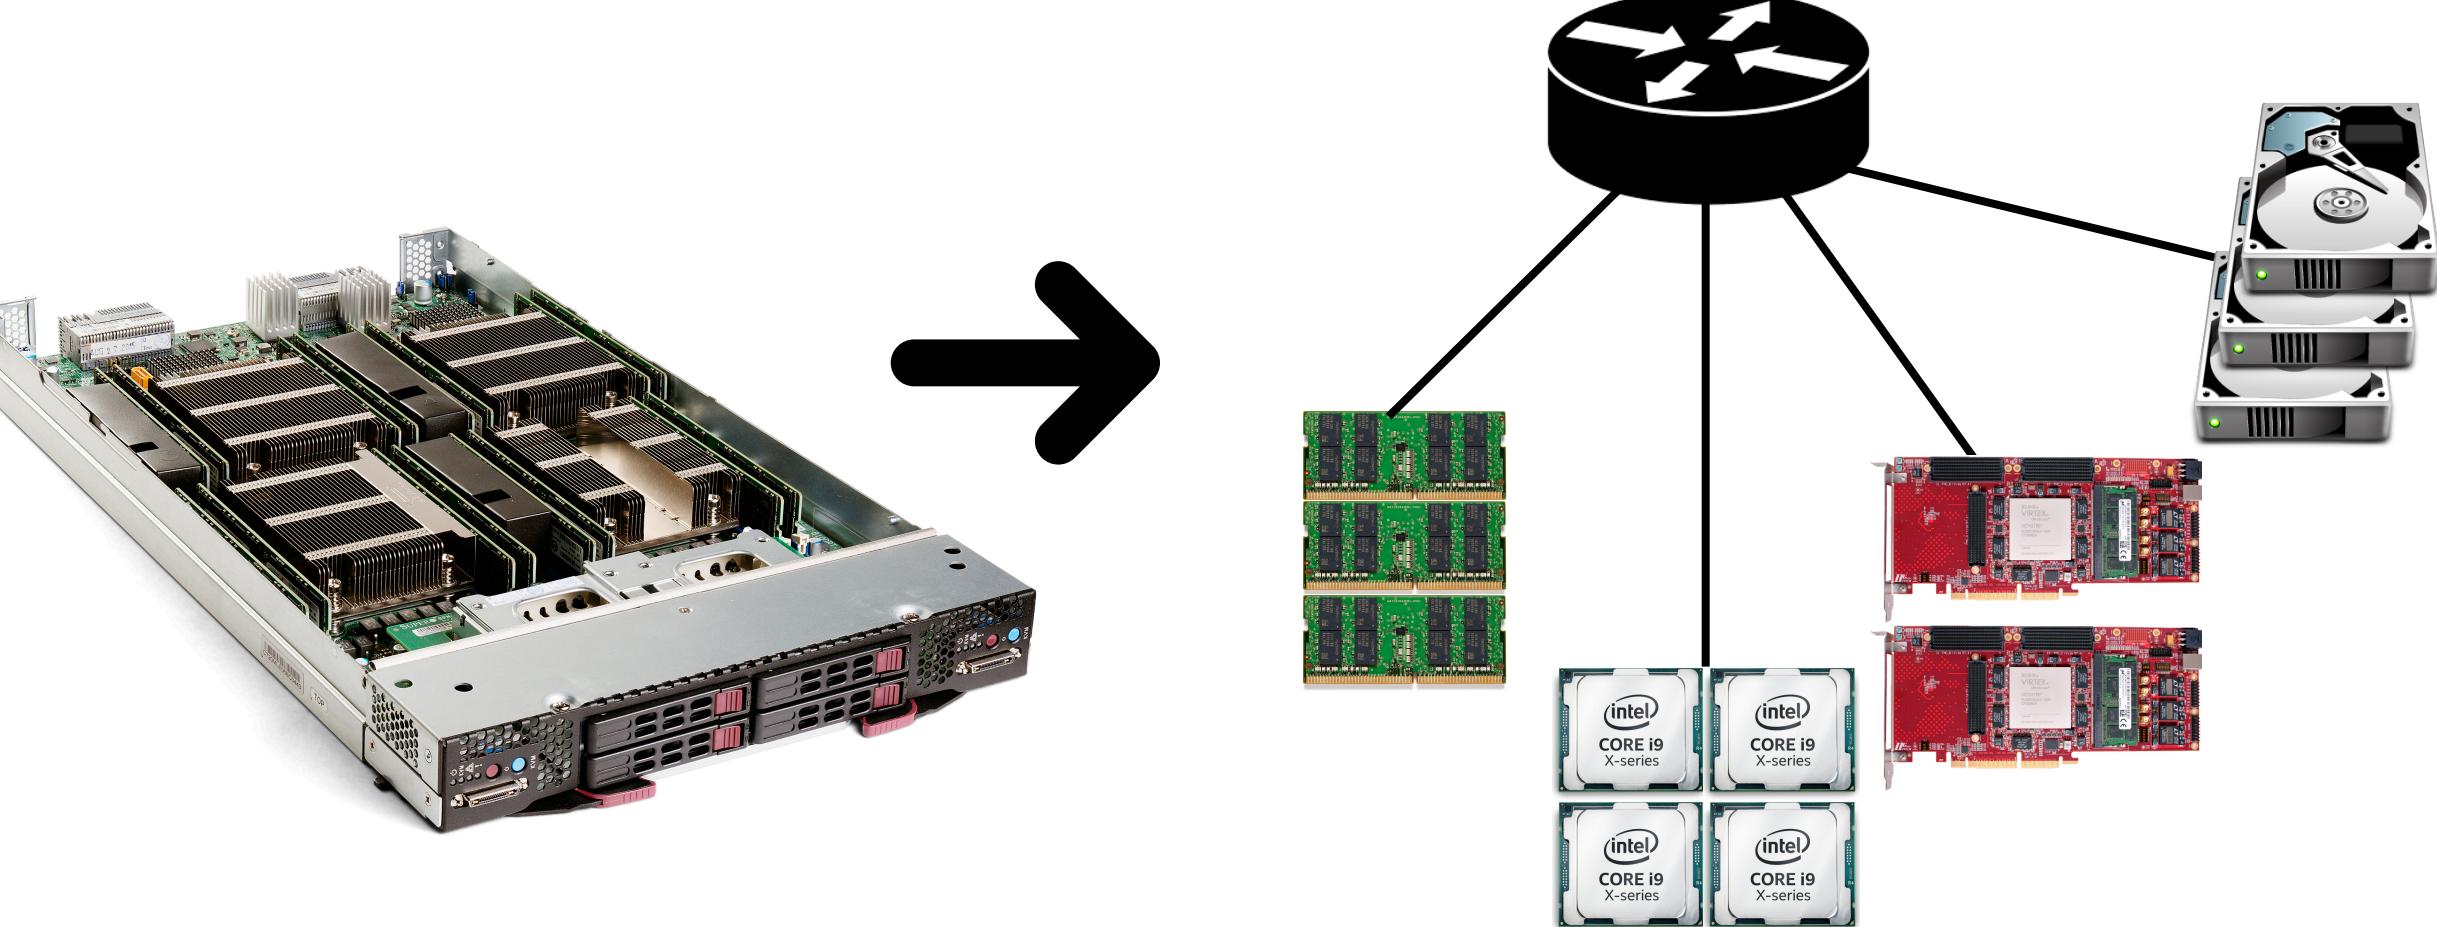


Reference:

- [1] "EDM: An Ultra-Low Latency Ethernet Fabric for Memory Disaggregation". *Weigao Su and Vishal Shrivastav*. ASPLOS 2025 ,

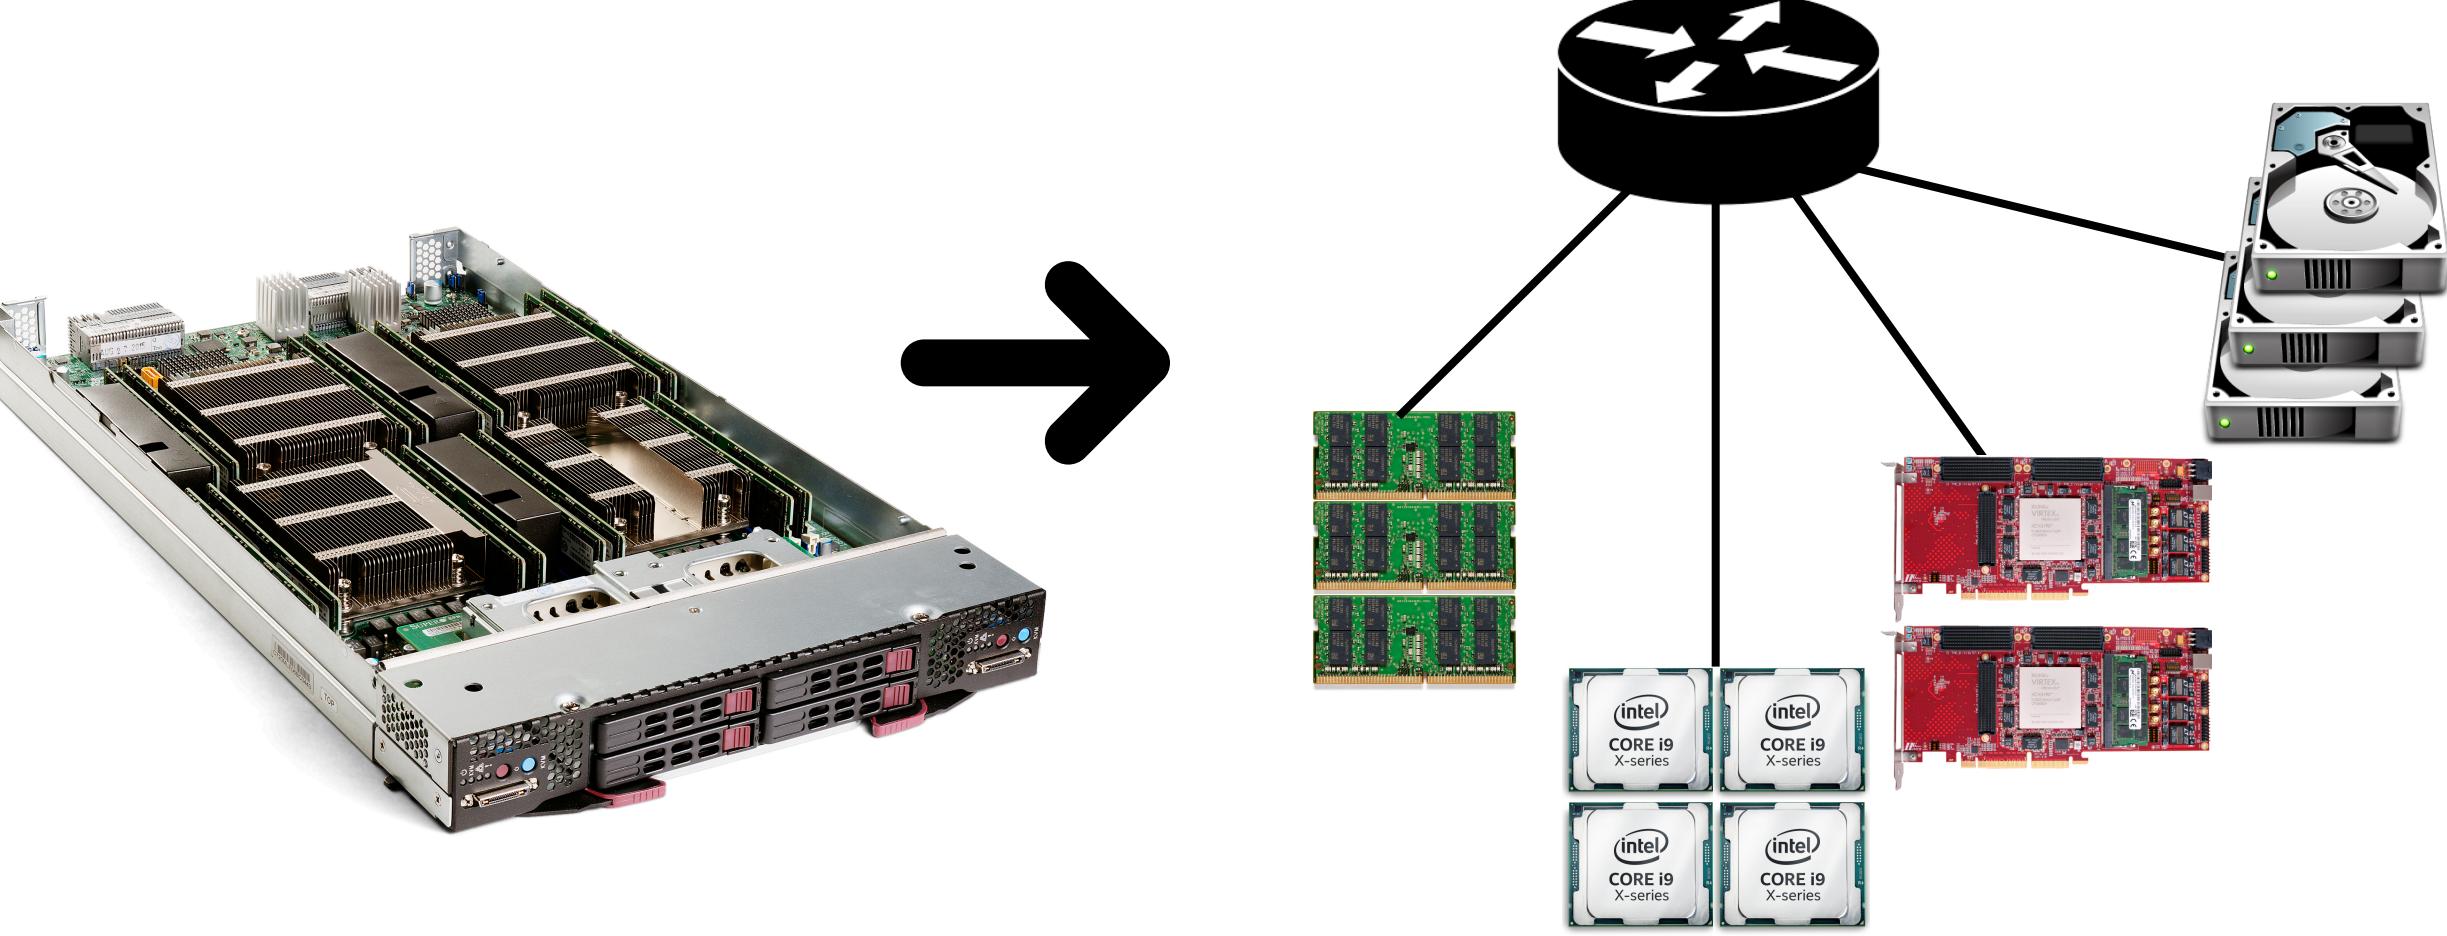
Why memory disaggregation?

- The need for memory is surging
- Constraints of individual servers
- Fine-grained pooling, elastic scaling



Why memory disaggregation?

- The need for memory is surging
- Constraints of individual servers
- Fine-grained pooling, elastic scaling

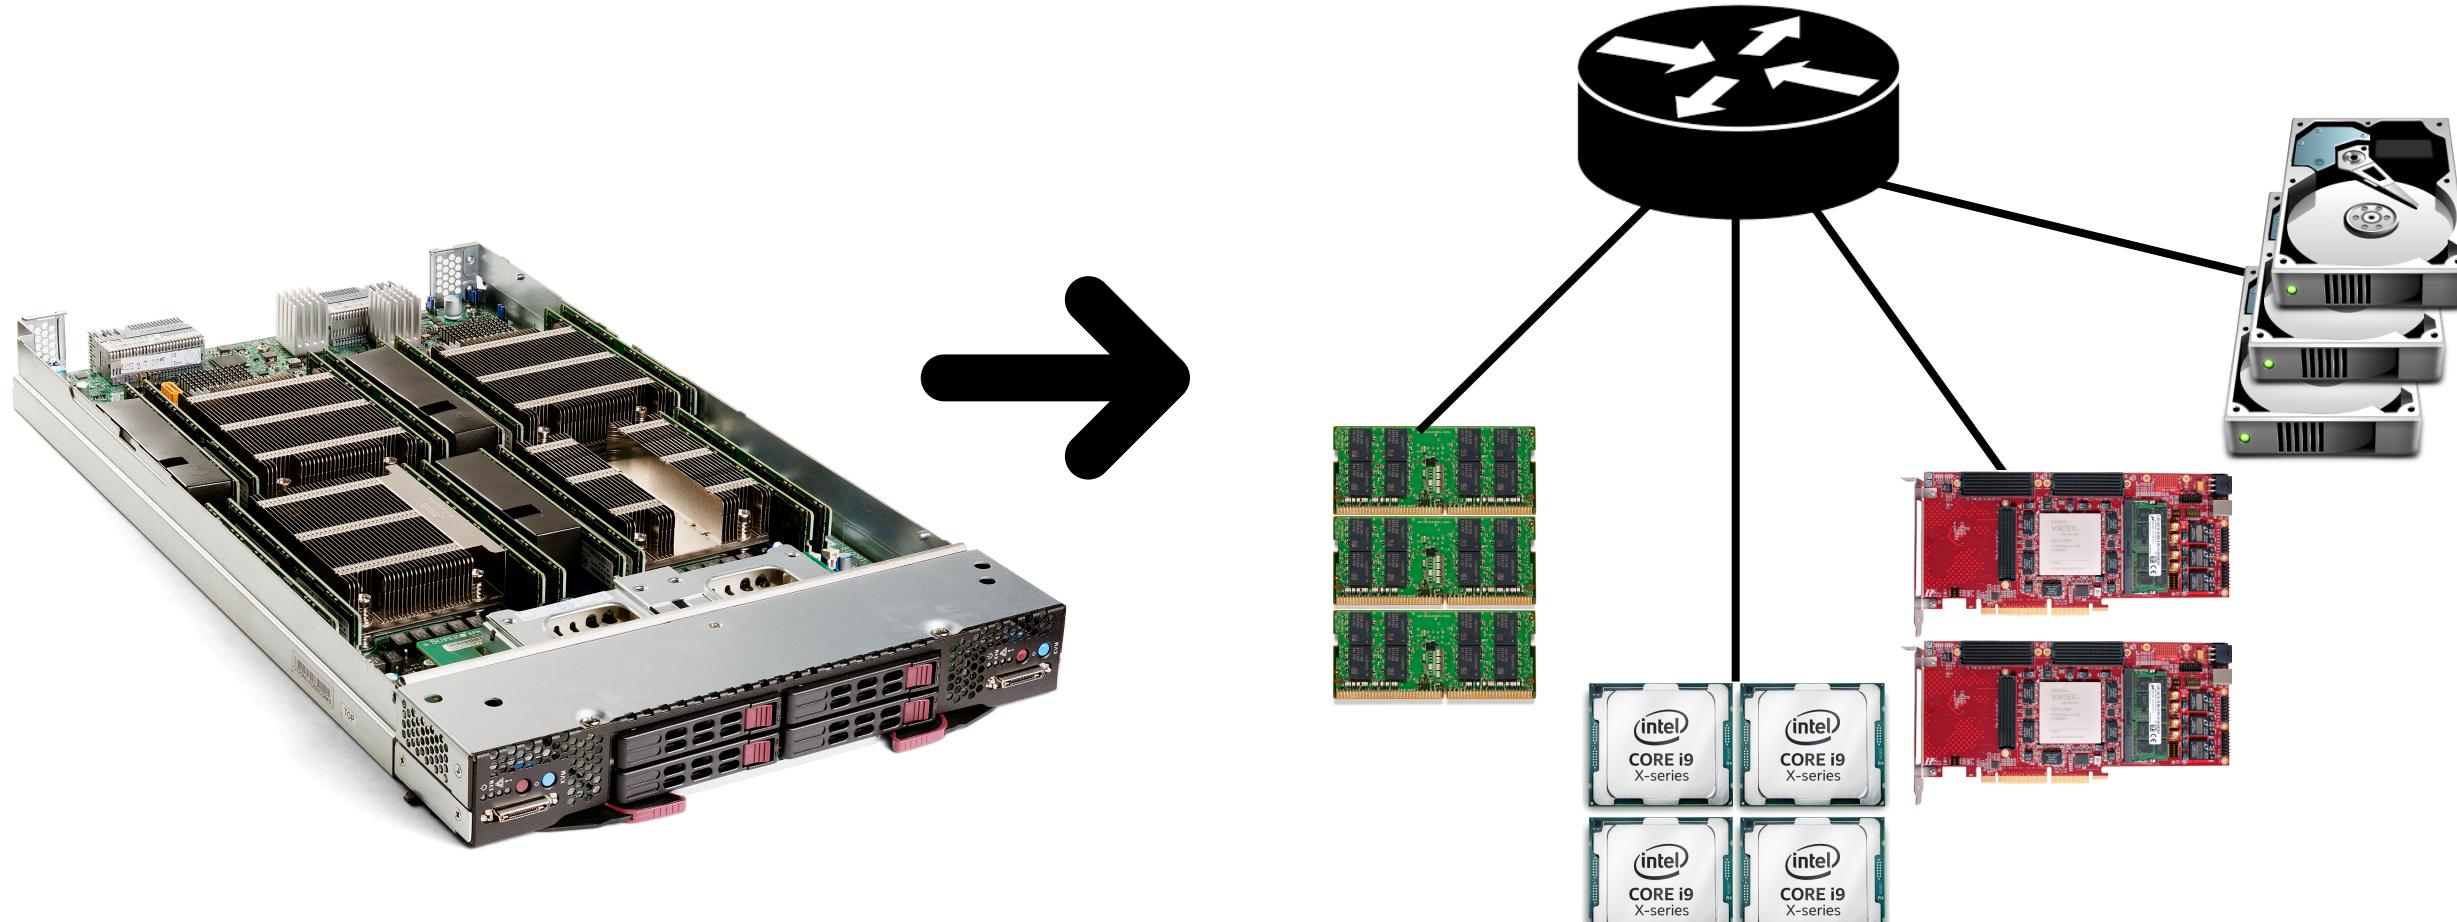


Why is Ethernet promising?

- Dominant datacenter network fabric
 - Low management cost, distance scaling...
- High bandwidth (Terabit Ethernet link)

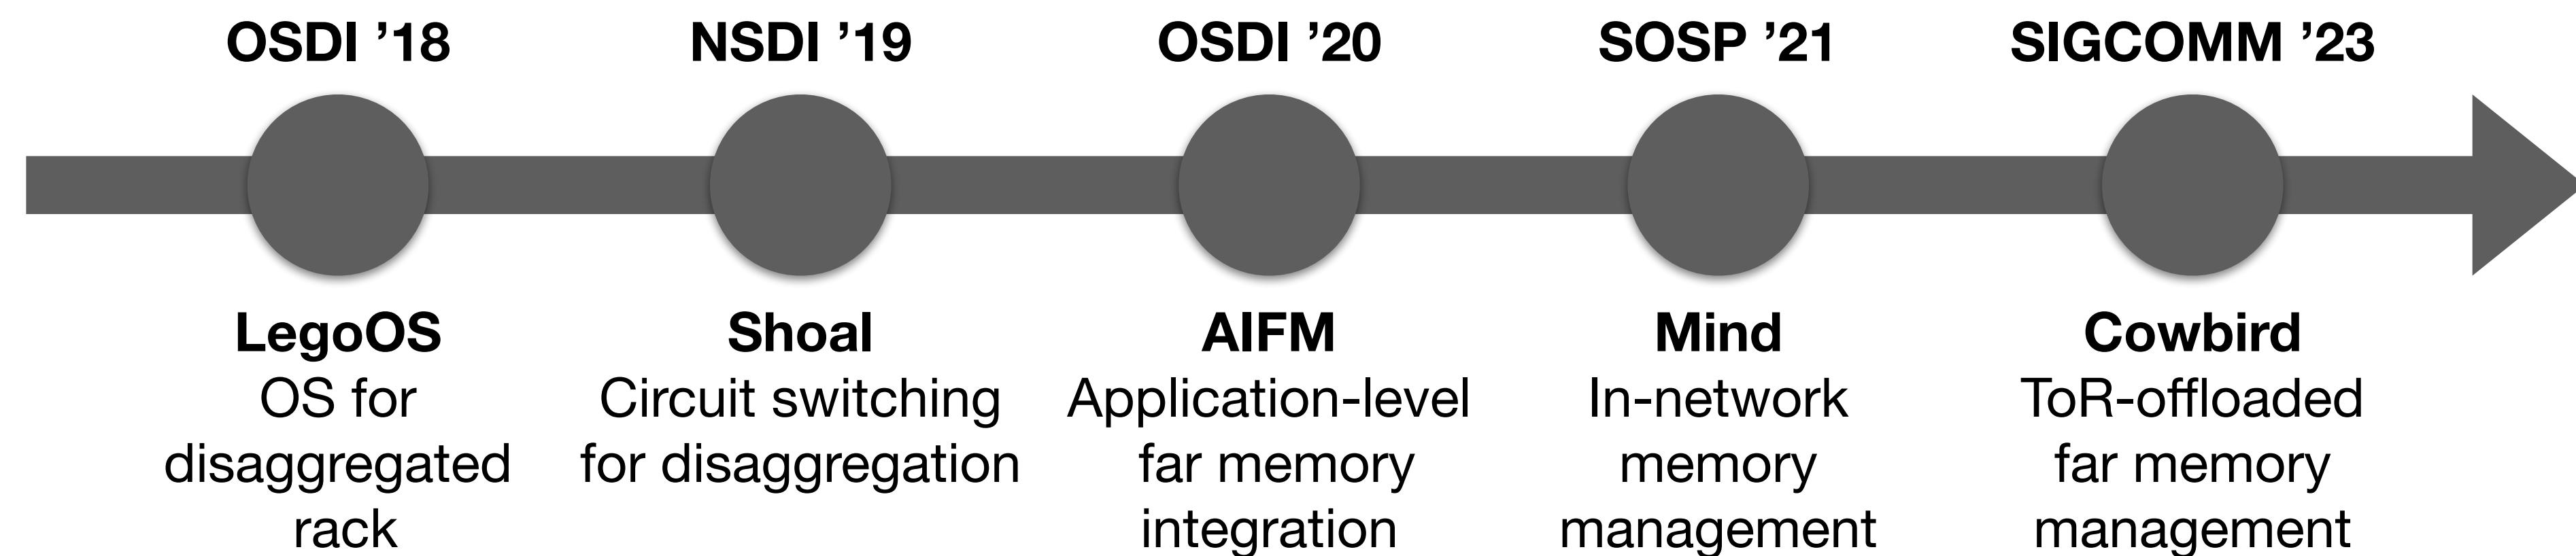
Why memory disaggregation?

- The need for memory is surging
- Constraints of individual servers
- Fine-grained pooling, elastic scaling

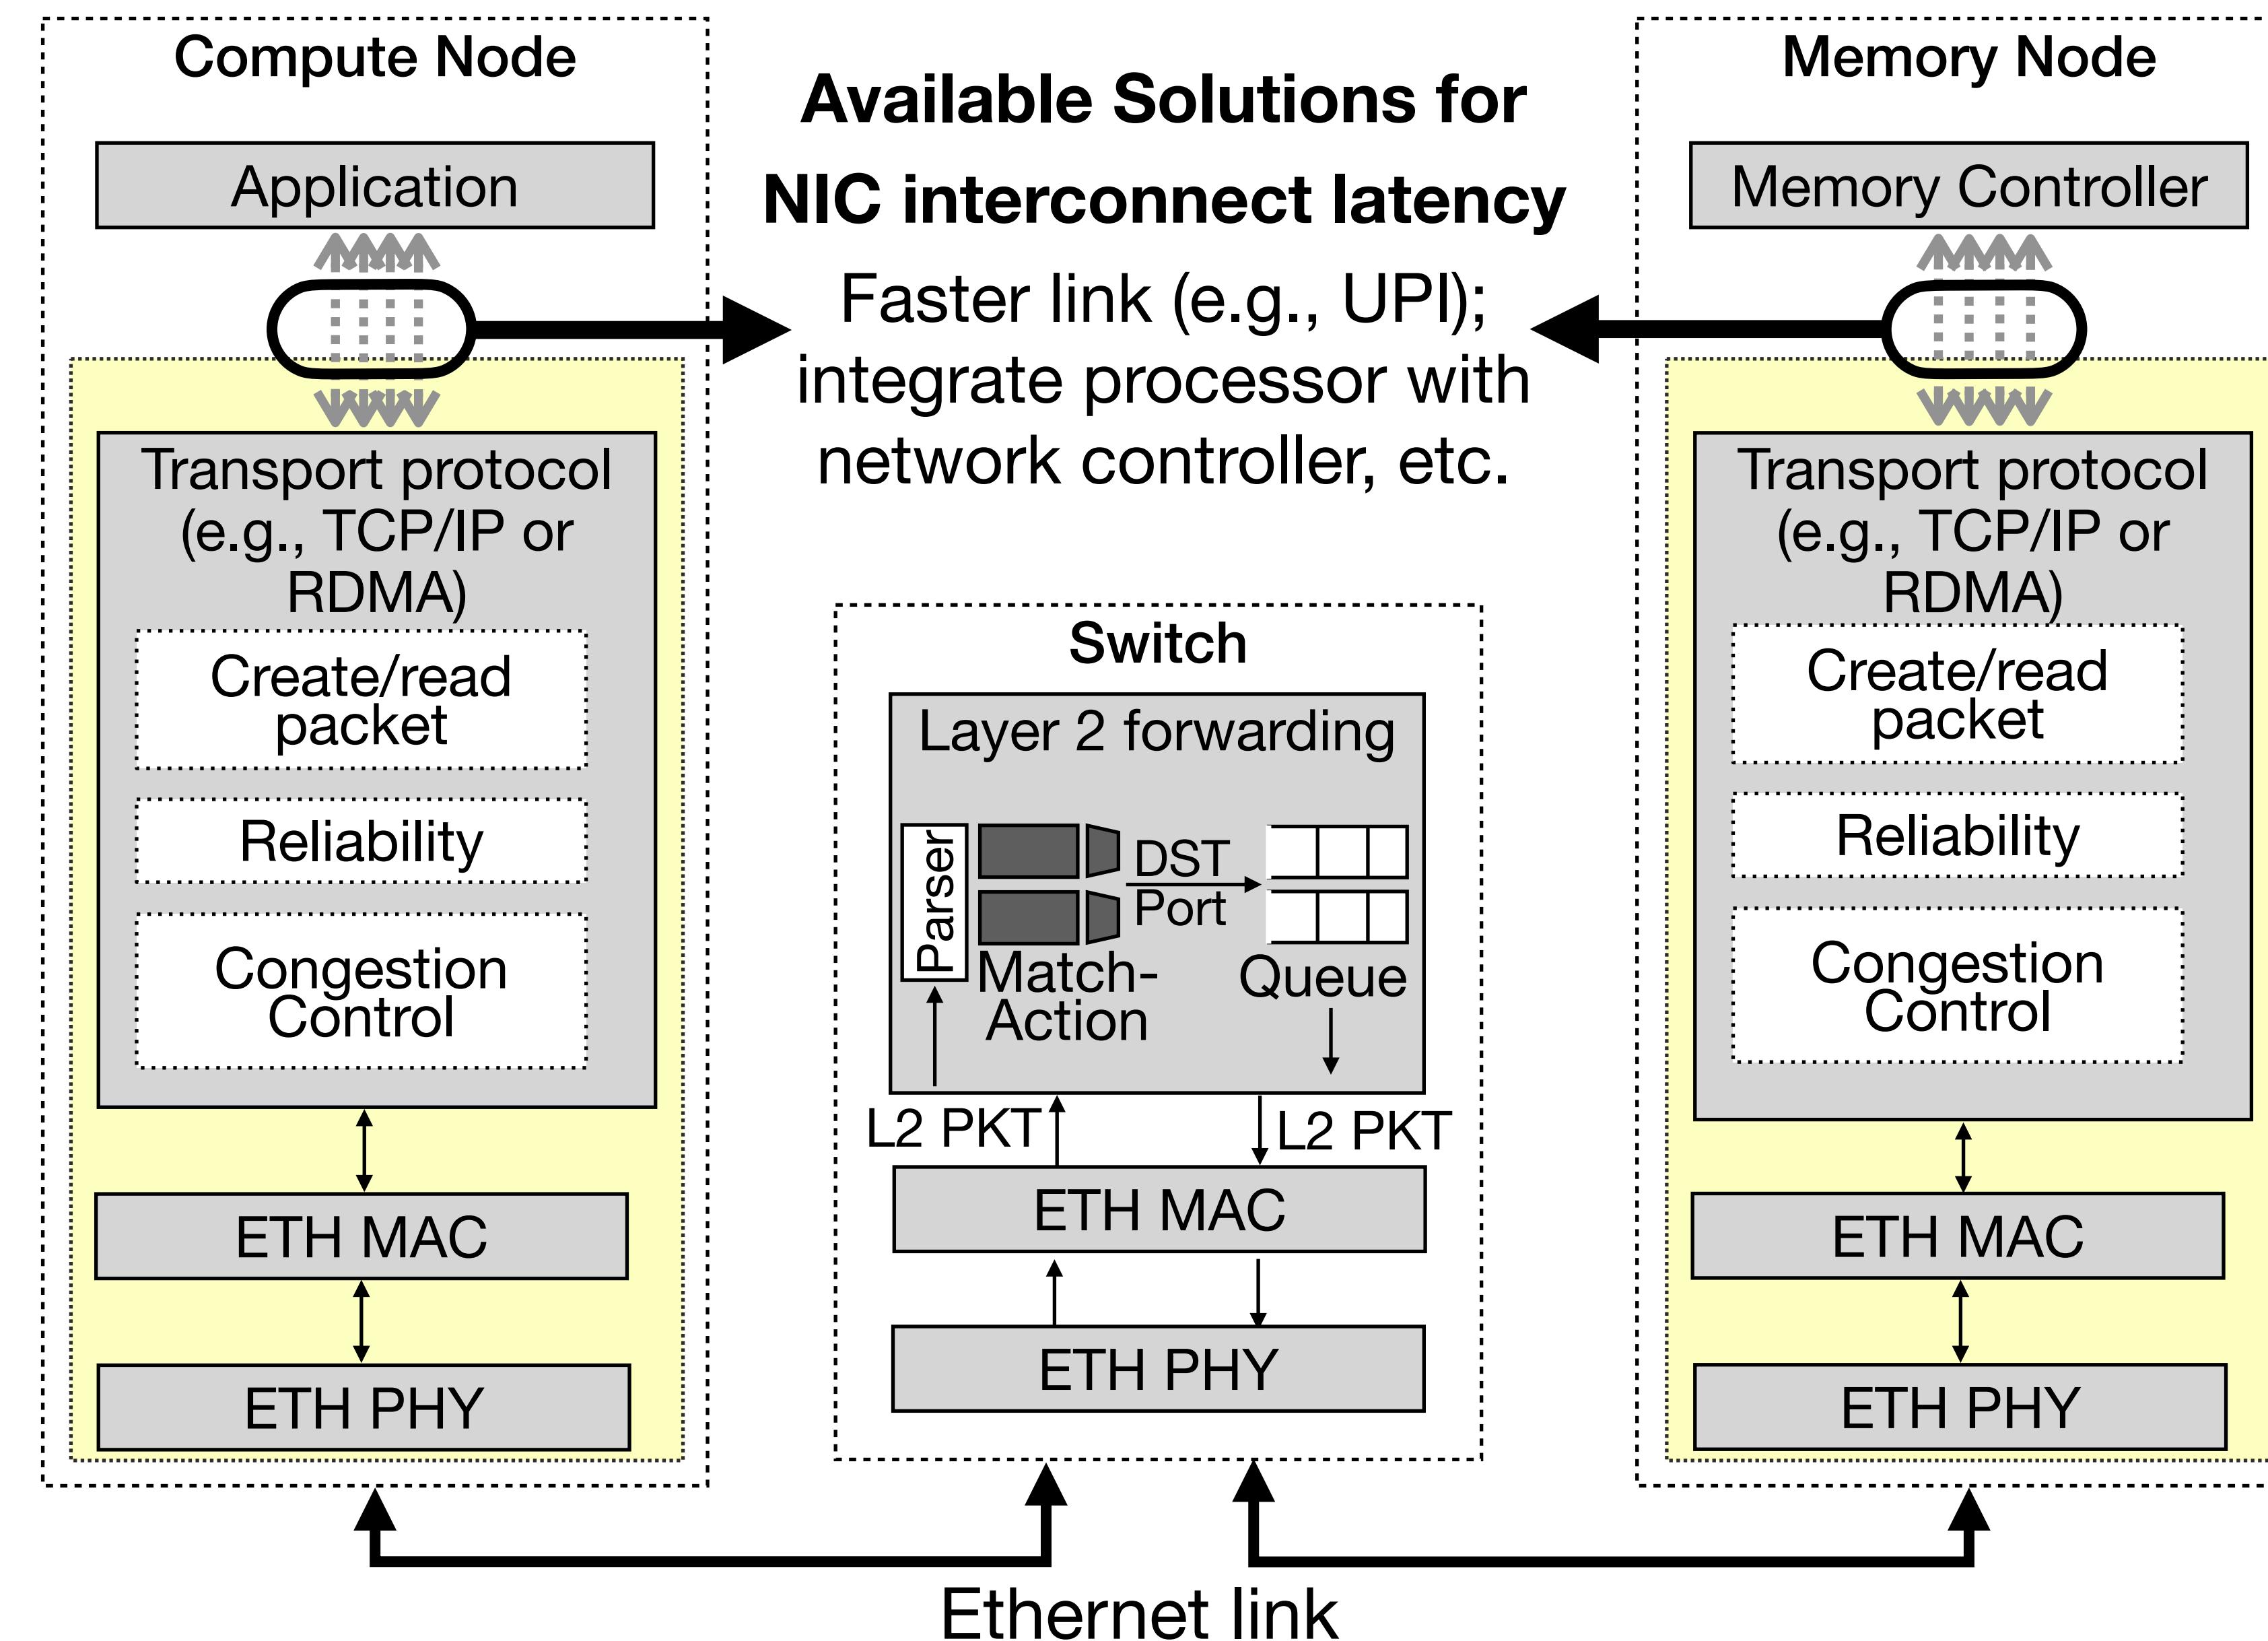


Why is Ethernet promising?

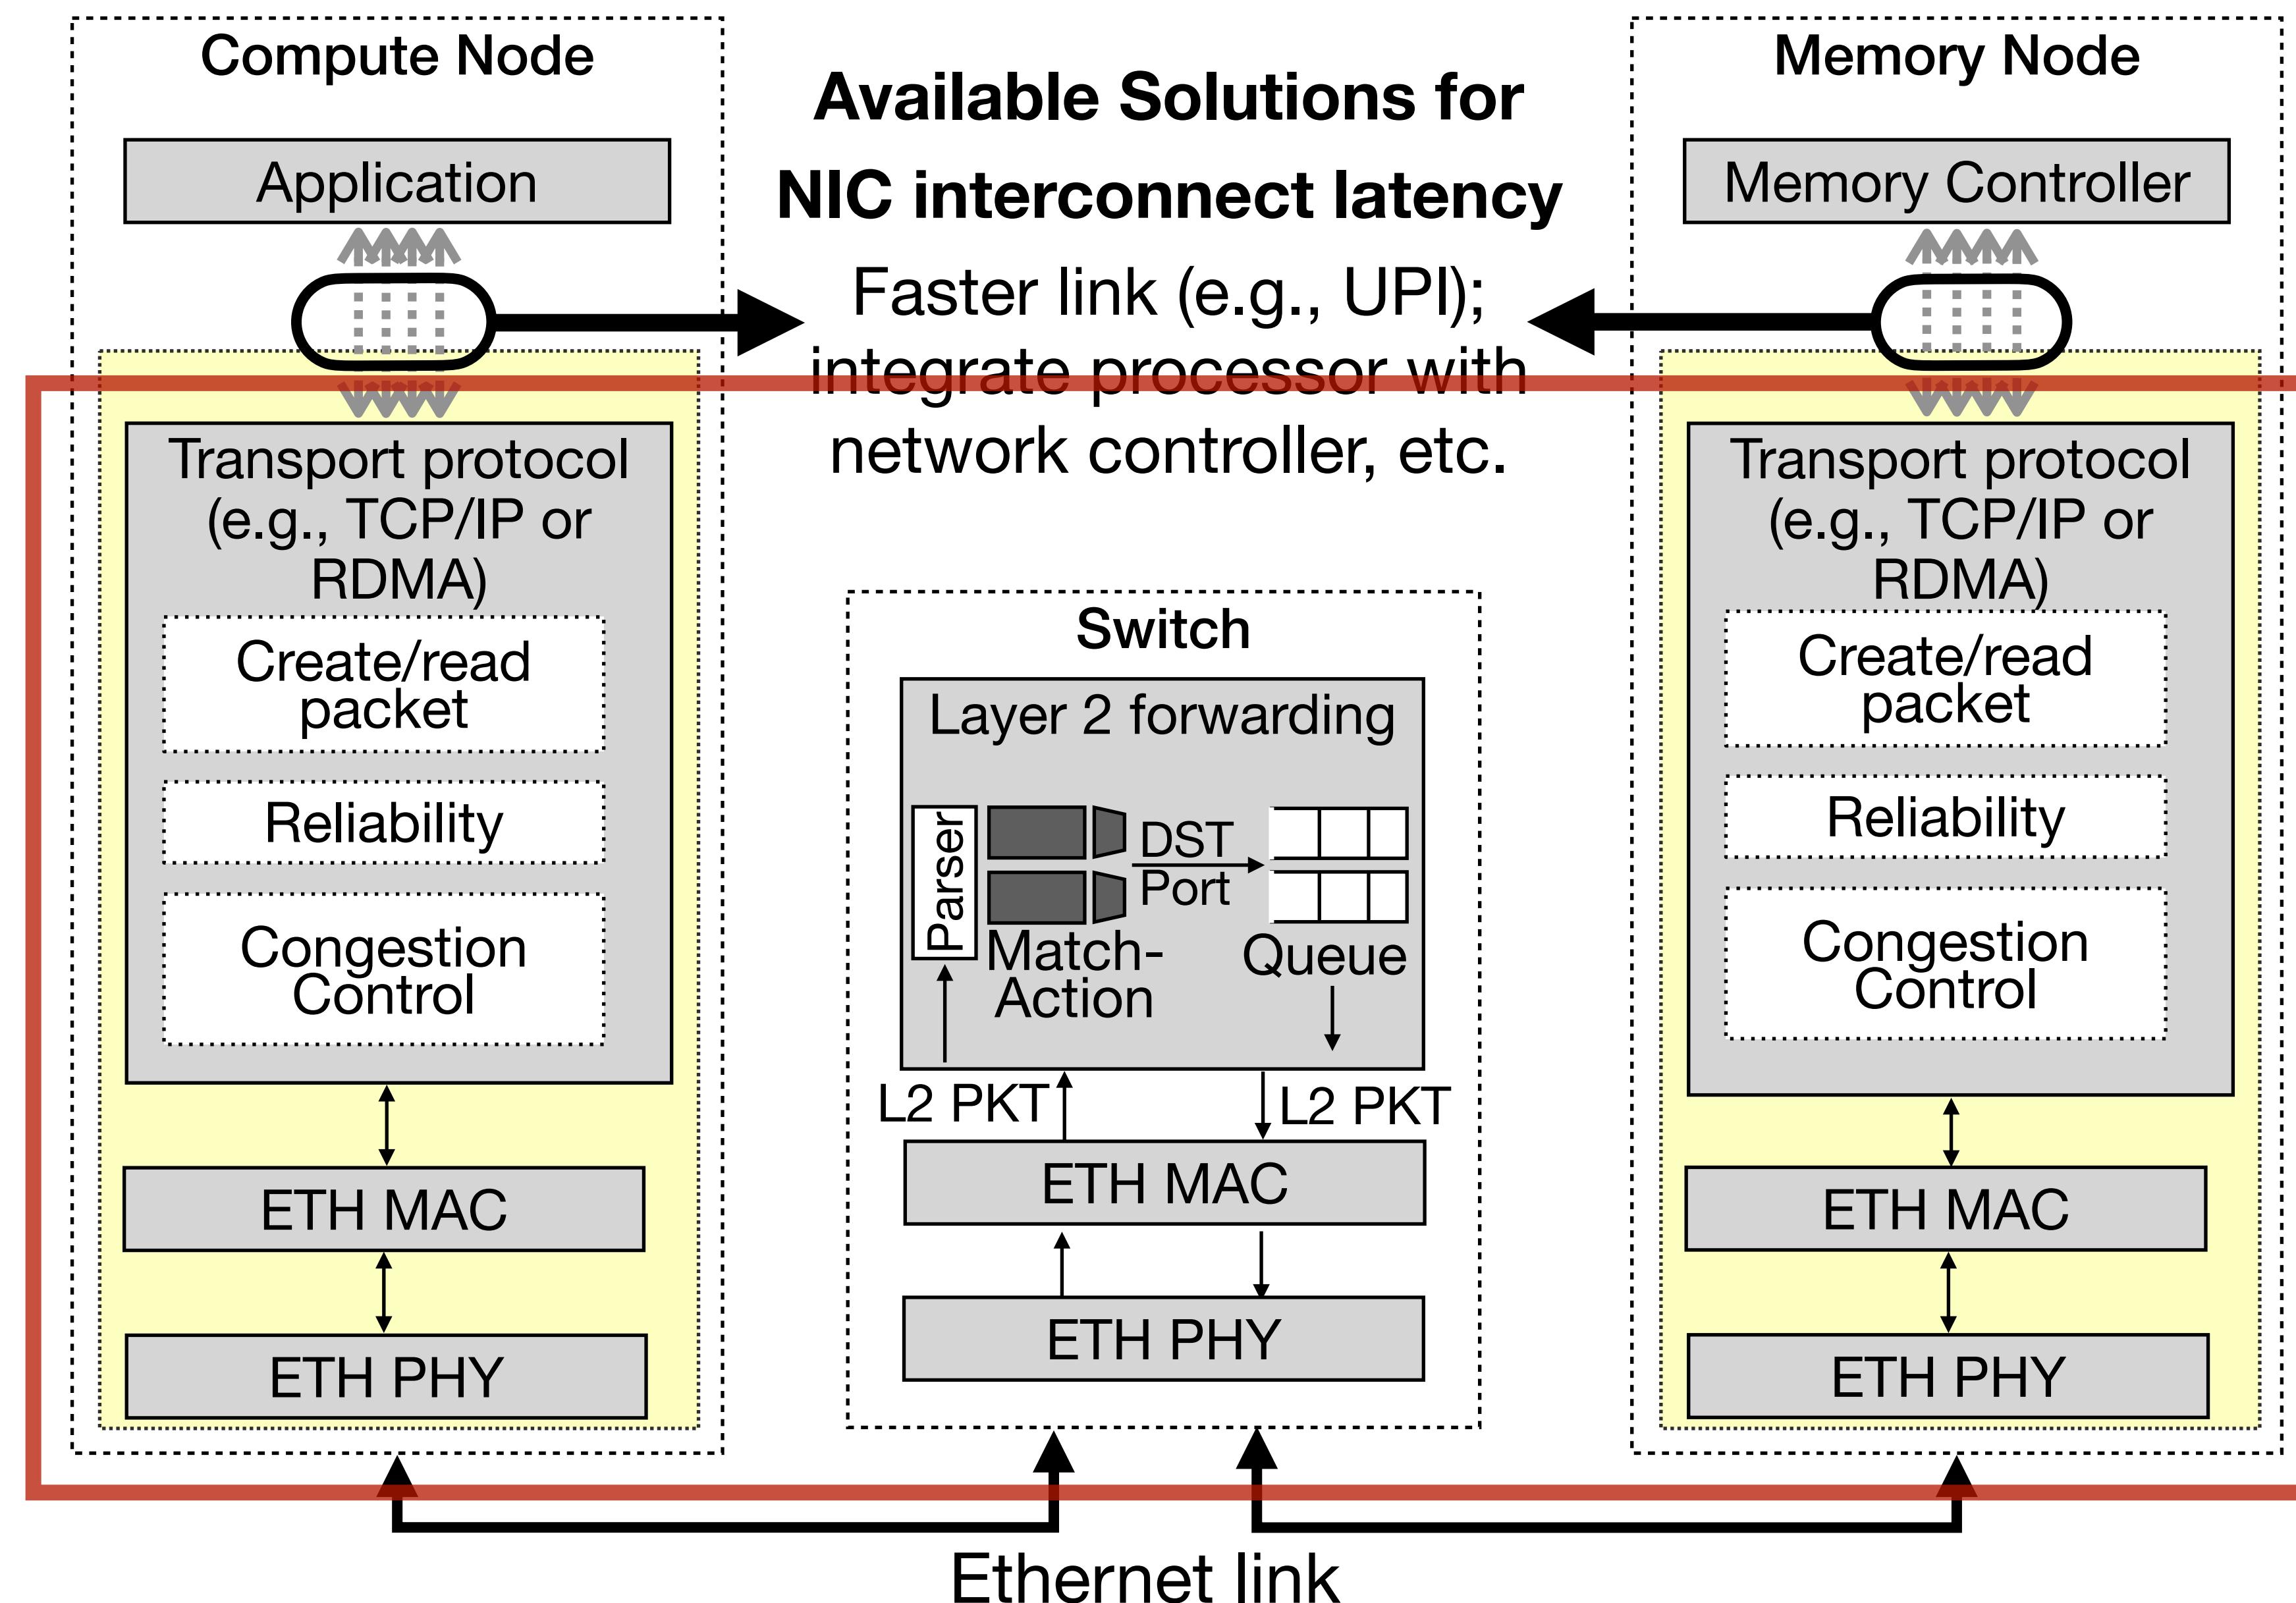
- Dominant datacenter network fabric
 - Low management cost, distance scaling...
- High bandwidth (Terabit Ethernet link)



Memory Disaggregation over Ethernet

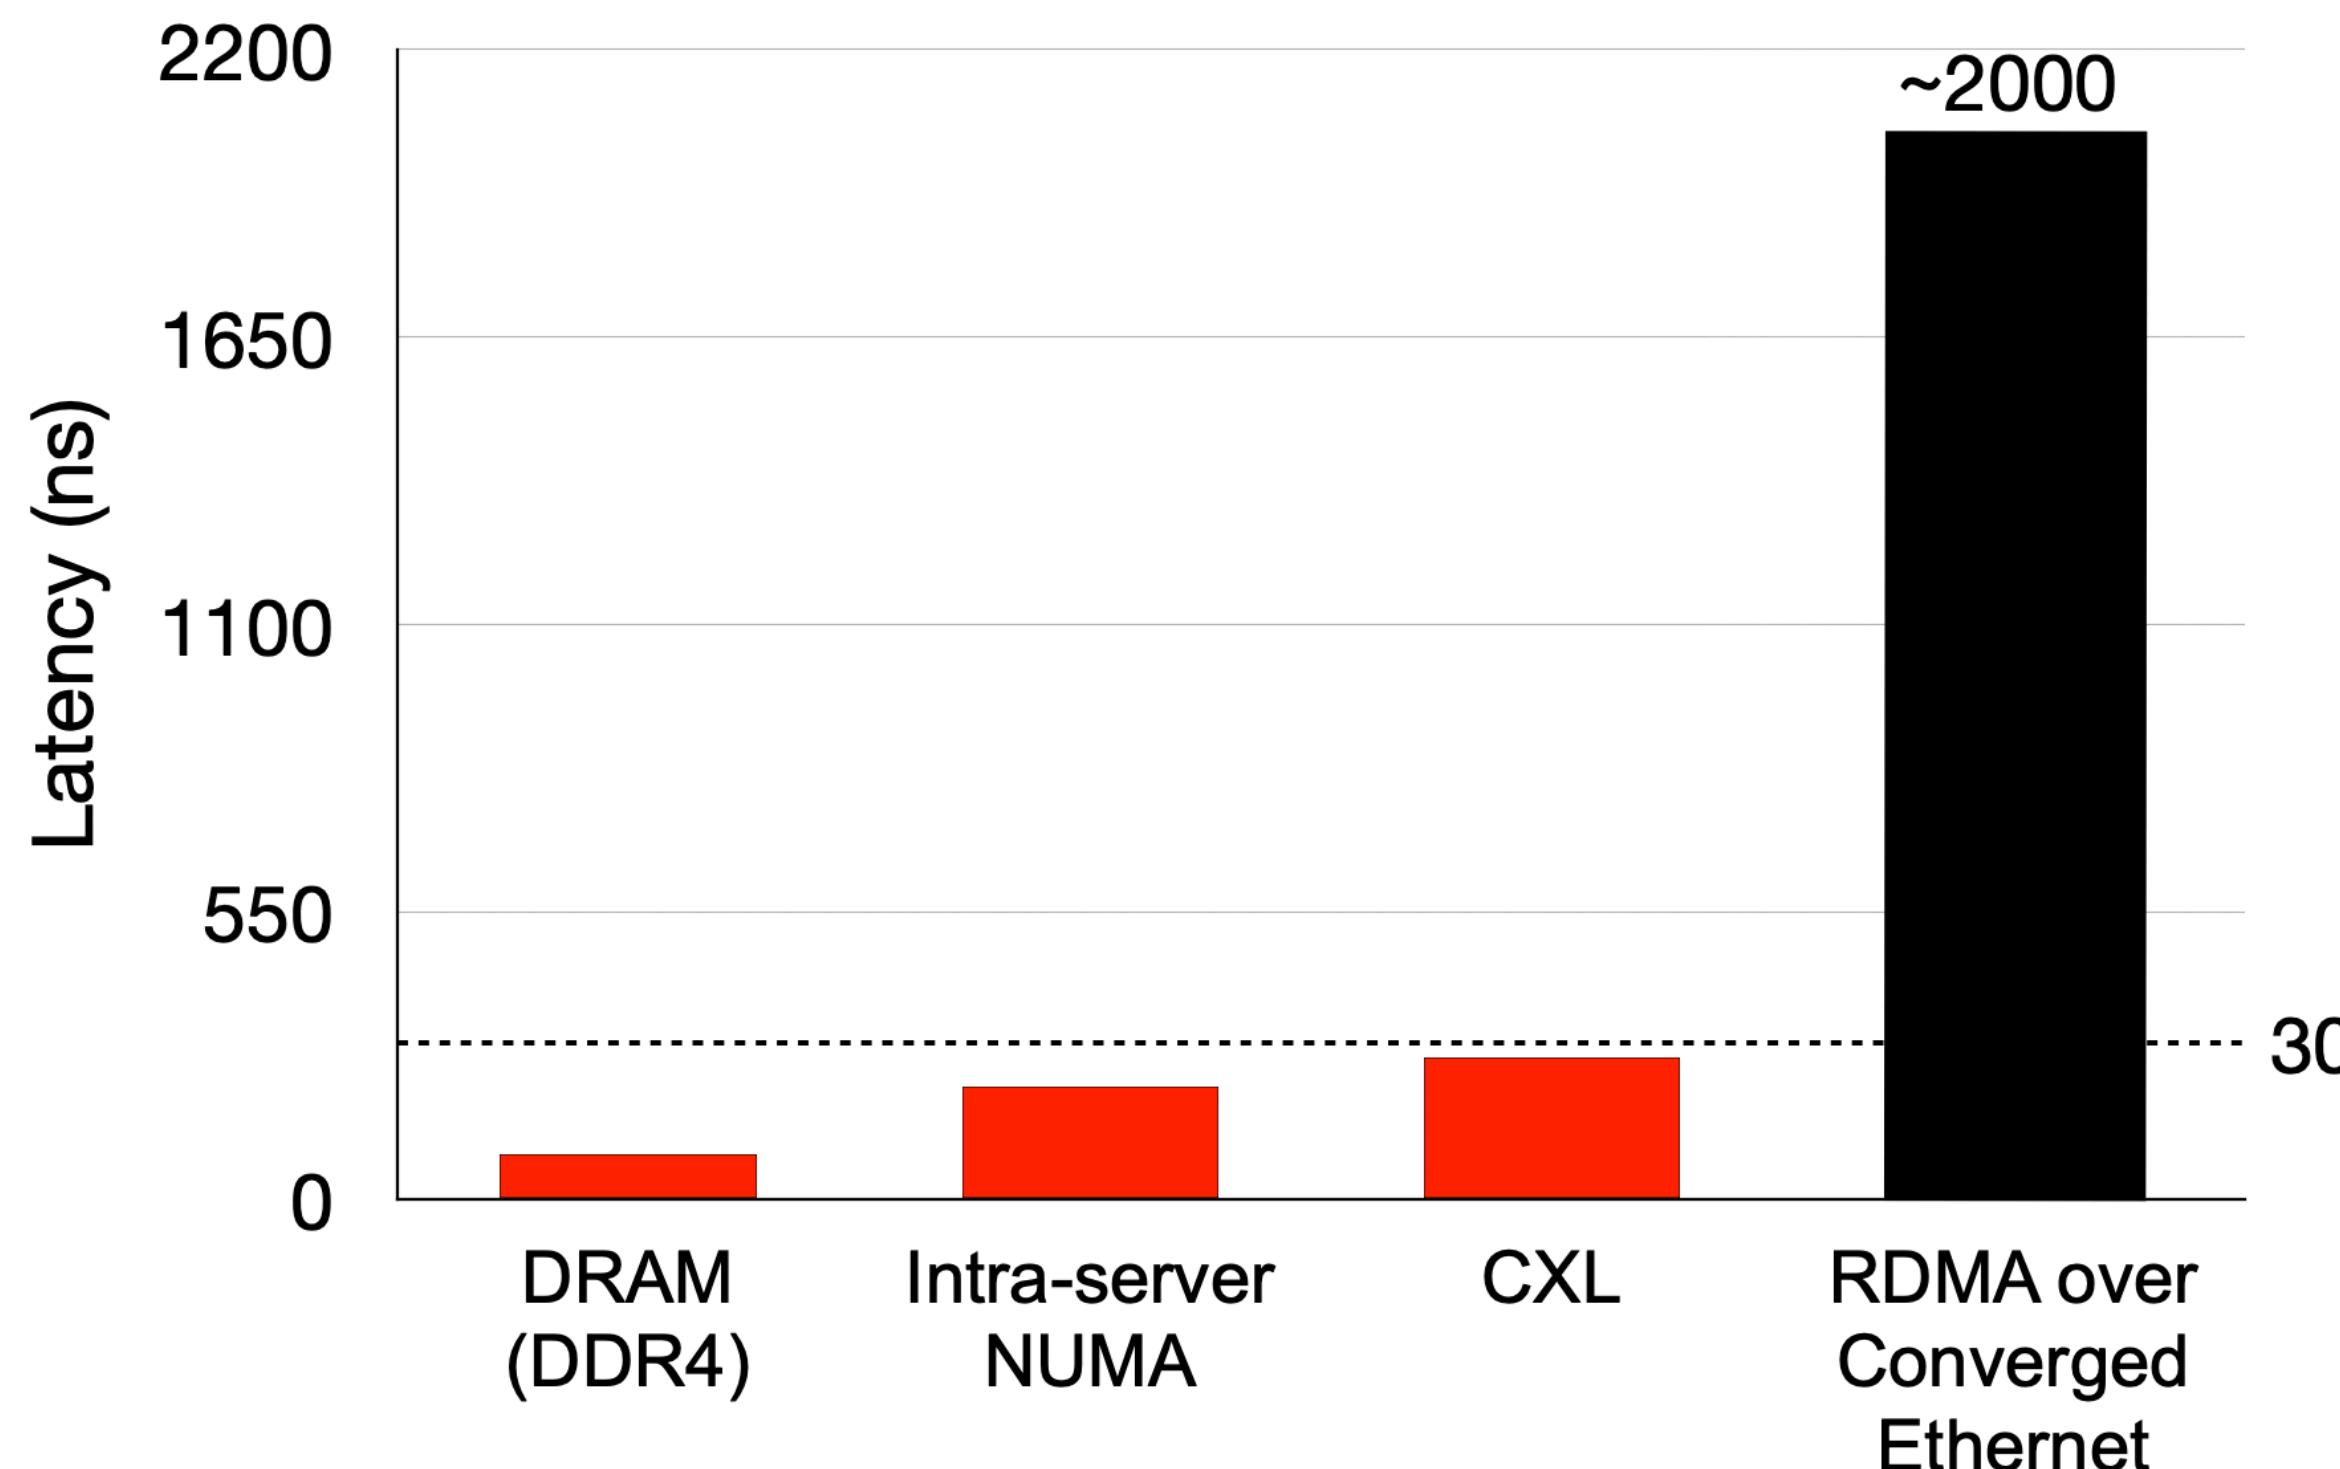


Memory Disaggregation over Ethernet



However, the latency in Ethernet is prohibitive, prompting proposals of separate fabric to carry memory traffic

Custom processor interconnect, PCIe, Infiniband, etc.



Scale-Out NUMA

Stanko Novaković Alexandros Daglis Edouard Bugnion Babak Falsafi Boris Grot†
EcoCloud, EPFL

Pond: CXL-Based Memory Pooling Systems for Cloud Platforms

Huaicheng Li Daniel S. Berger Lisa Hsu
Virginia Tech Microsoft Azure Unaffiliated
Carnegie Mellon University USA University of Washington USA USA

LegoOS: A Disseminated, Distributed OS for Hardware Resource Disaggregation

Yizhou Shan, Yutong Huang, Yilun Chen, Yiyang Zhang
Purdue University

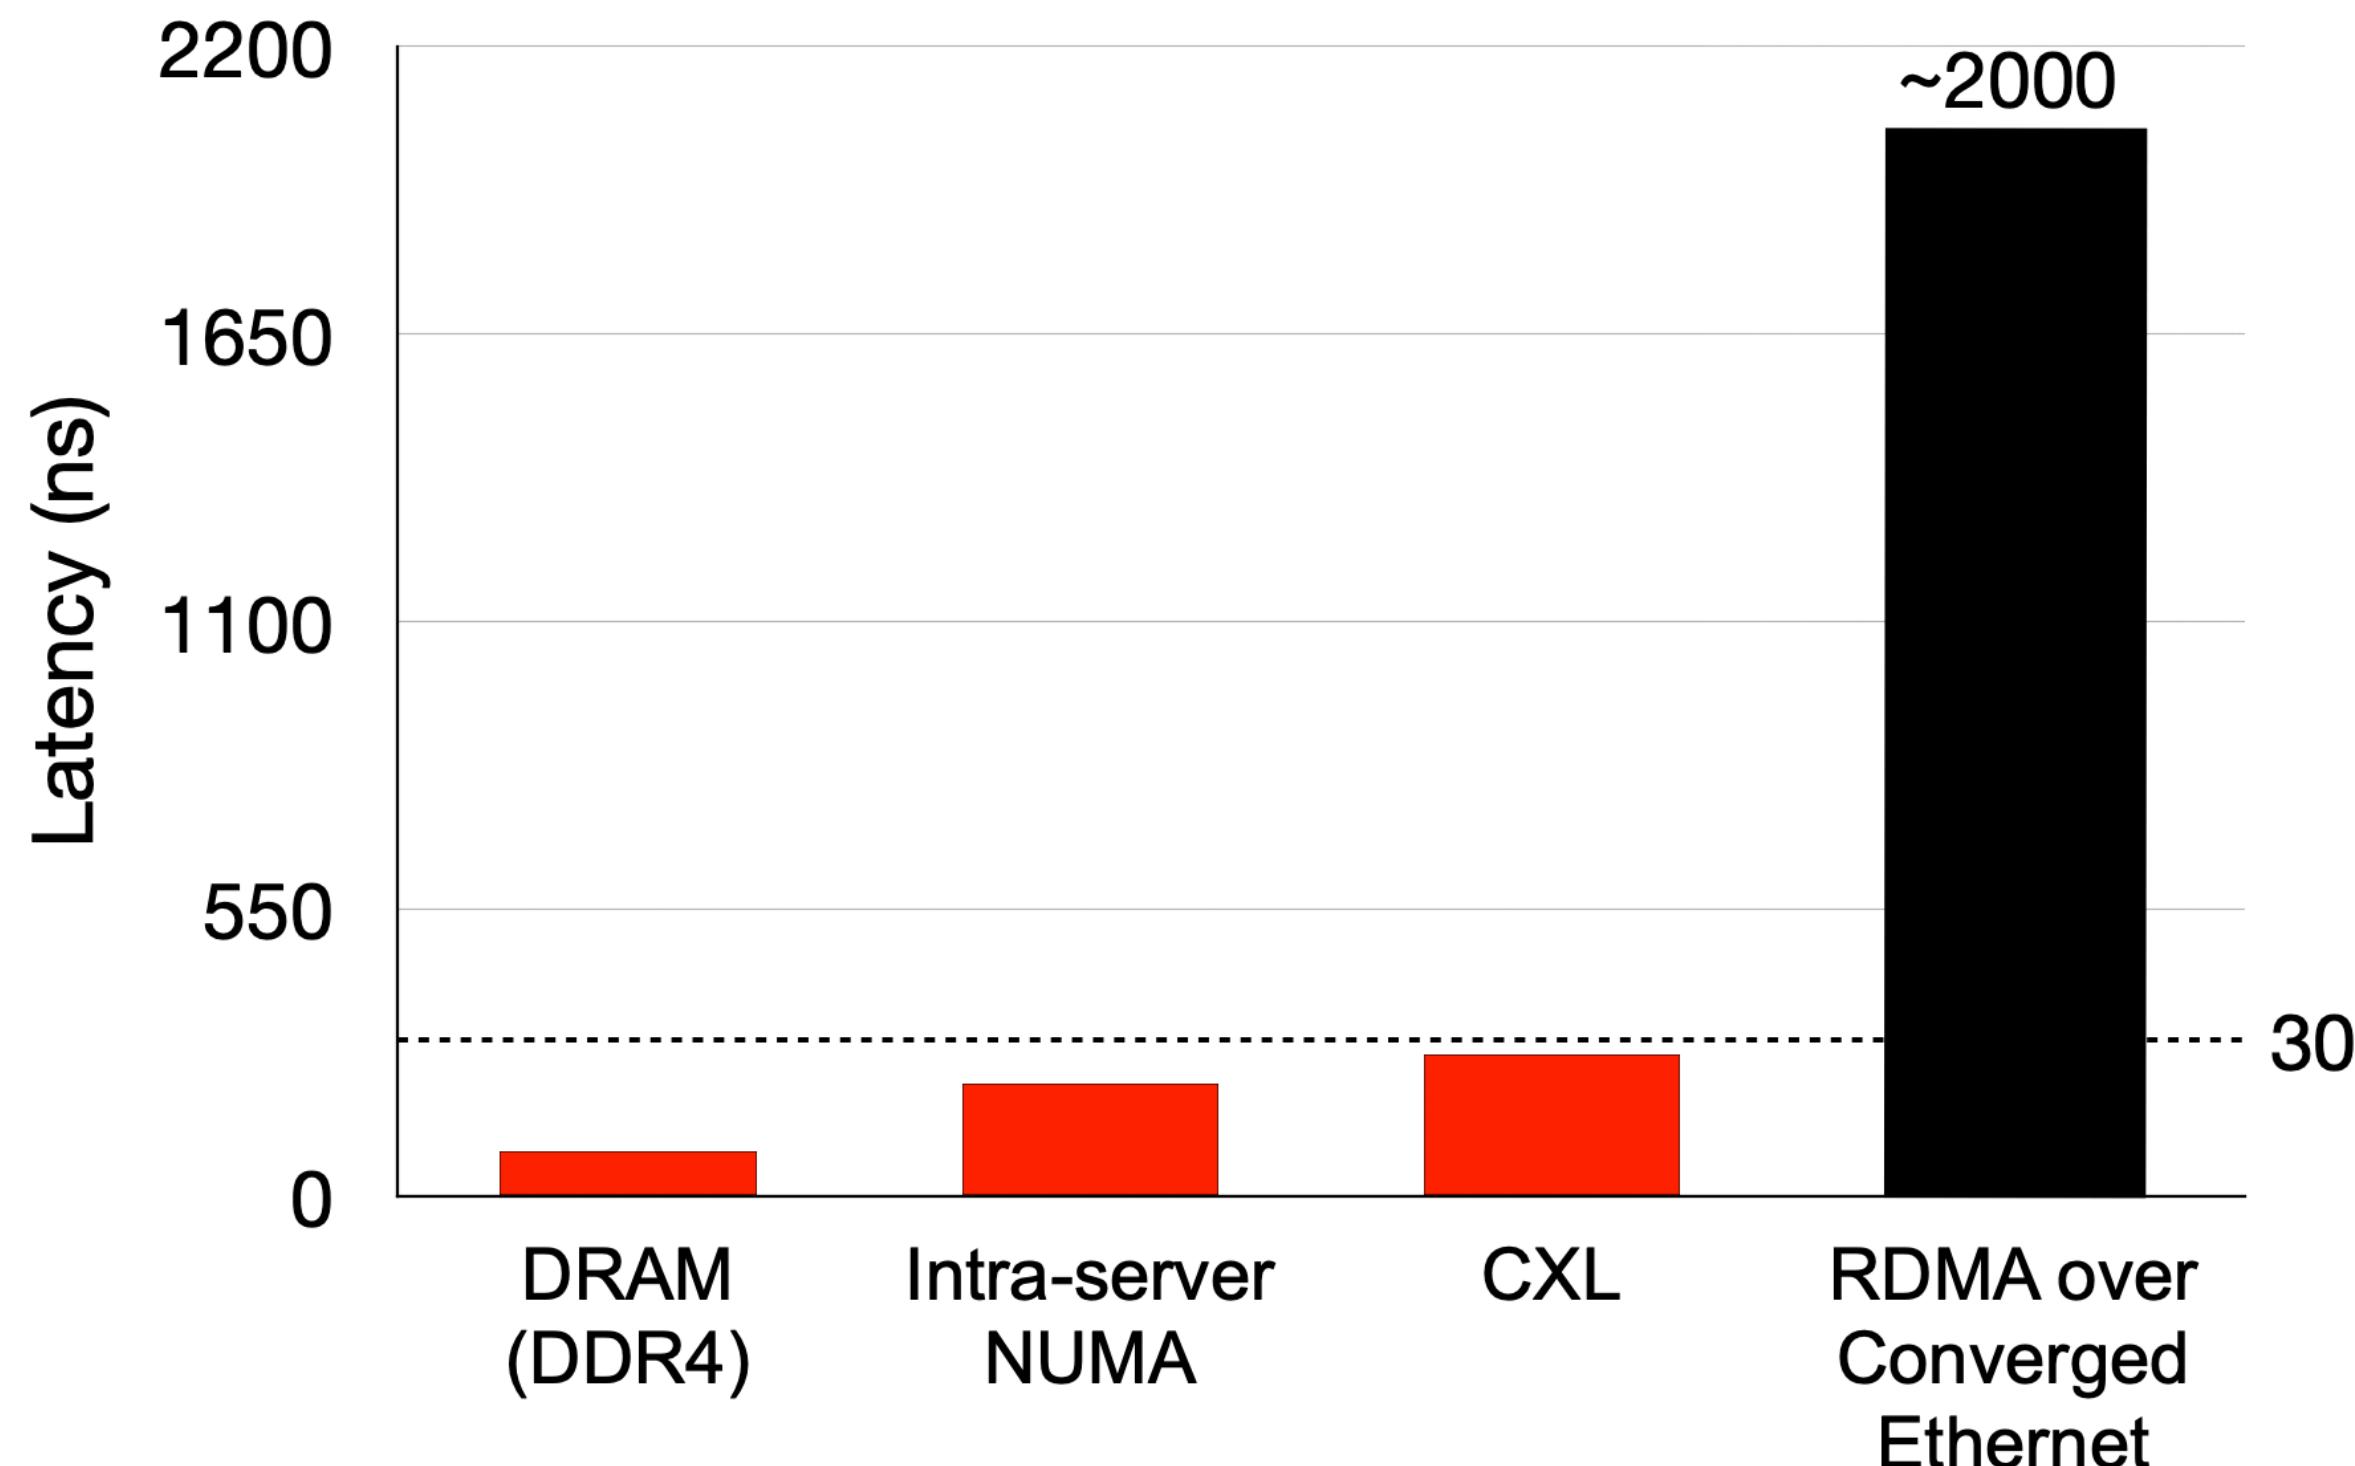
Abstract

The monolithic server model where a server is the unit of deployment, operation, and failure is meeting its limits in the face of several recent hardware and application trends. To improve resource utilization, elasticity, heterogeneity, and reliability, we propose to break up the monolithic server into smaller, more specialized units. This requires a distributed operating system that can fit into monolithic servers and deploying them in datacenters is a painful and cost-ineffective process that often limits the speed of new hardware adoption. We believe that datacenters should break monolithic servers and organize hardware devices like CPU, DRAM, and disks as independent failure isolated units. This is challenging because the hardware devices have different failure characteristics and lifetimes. To address this challenge, we propose a distributed operating system called LegoOS that uses a novel memory pooling mechanism to manage hardware resources. LegoOS provides a unified memory interface to applications, which can access memory from multiple hardware devices. It also provides a distributed file system that can store data across multiple hardware devices. LegoOS is designed to be highly reliable and fault-tolerant, and it can handle hardware failures gracefully. We implemented a prototype of LegoOS and evaluated its performance and reliability. Our results show that LegoOS is able to provide high performance and reliability while maintaining a simple and modular architecture. We believe that LegoOS will be a key component of future datacenter architectures.

Ricardo Bianchini
Microsoft Azure
USA

However, the latency in Ethernet is prohibitive, prompting proposals of separate fabric to carry memory traffic

Custom processor interconnect, PCIe, Infiniband, etc.



Scale-Out NUMA

Stanko Novaković Alexandros Daglis Edouard Bugnion Babak Falsafi Boris Grot†

EcoCloud, EPFL

Pond: CXL-Based Memory Pooling Systems for Cloud Platforms

Huaicheng Li

Virginia Tech

Carnegie Mellon University

USA

Daniel S. Berger

Microsoft Azure

University of Washington

USA

Lisa Hsu

Unaffiliated

USA

LegoOS: A Disseminated, Distributed OS for Hardware Resource Disaggregation

Yizhou Shan, Yutong Huang, Yilun Chen, Yiying Zhang

Purdue University

Abstract

The monolithic server model where a server is the unit of deployment, operation, and failure is meeting its limits in the face of several recent hardware and application trends. To improve resource utilization, elasticity, heterogeneity, and reliability, we propose a new paradigm for server design.

that can fit into monolithic servers and deploying them in datacenters is a painful and cost-ineffective process that often limits the speed of new hardware adoption.

We believe that datacenters should break monolithic servers and organize hardware devices like CPU, DRAM, and disks as *independant failure isolated*

Ricardo Bianchini

Microsoft Azure

USA

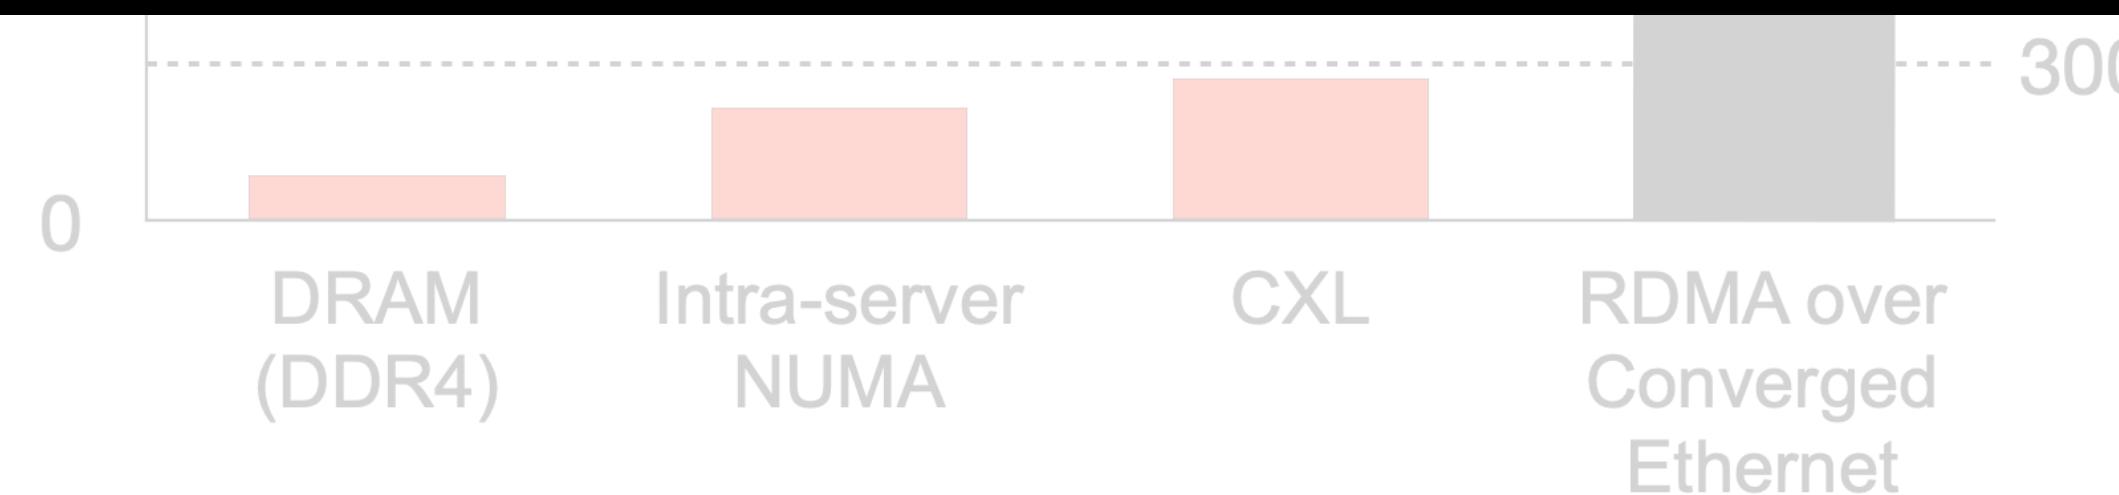
But, separate fabrics for different traffic makes the network **costly** and **harder to manage**

However, the latency in Ethernet is prohibitive, prompting proposals of separate fabric to carry memory traffic

Custom processor interconnect, PCIe, Infiniband, etc.

A **low latency** Ethernet fabric would allow us to have a **single unified** network fabric to carry all kinds of traffic (memory, storage, IP, ...)

... easier to manage, lower cost, statistical bandwidth multiplexing



The monolithic server model where a server is the unit of deployment, operation, and failure is meeting its limits in the face of several recent hardware and application trends. To improve resource utilization, elasticity, heterogeneity, and reliability, we believe that datacenters should break monolithic servers and organize hardware devices like CPU, DRAM, and disks as independent failure isolated components.

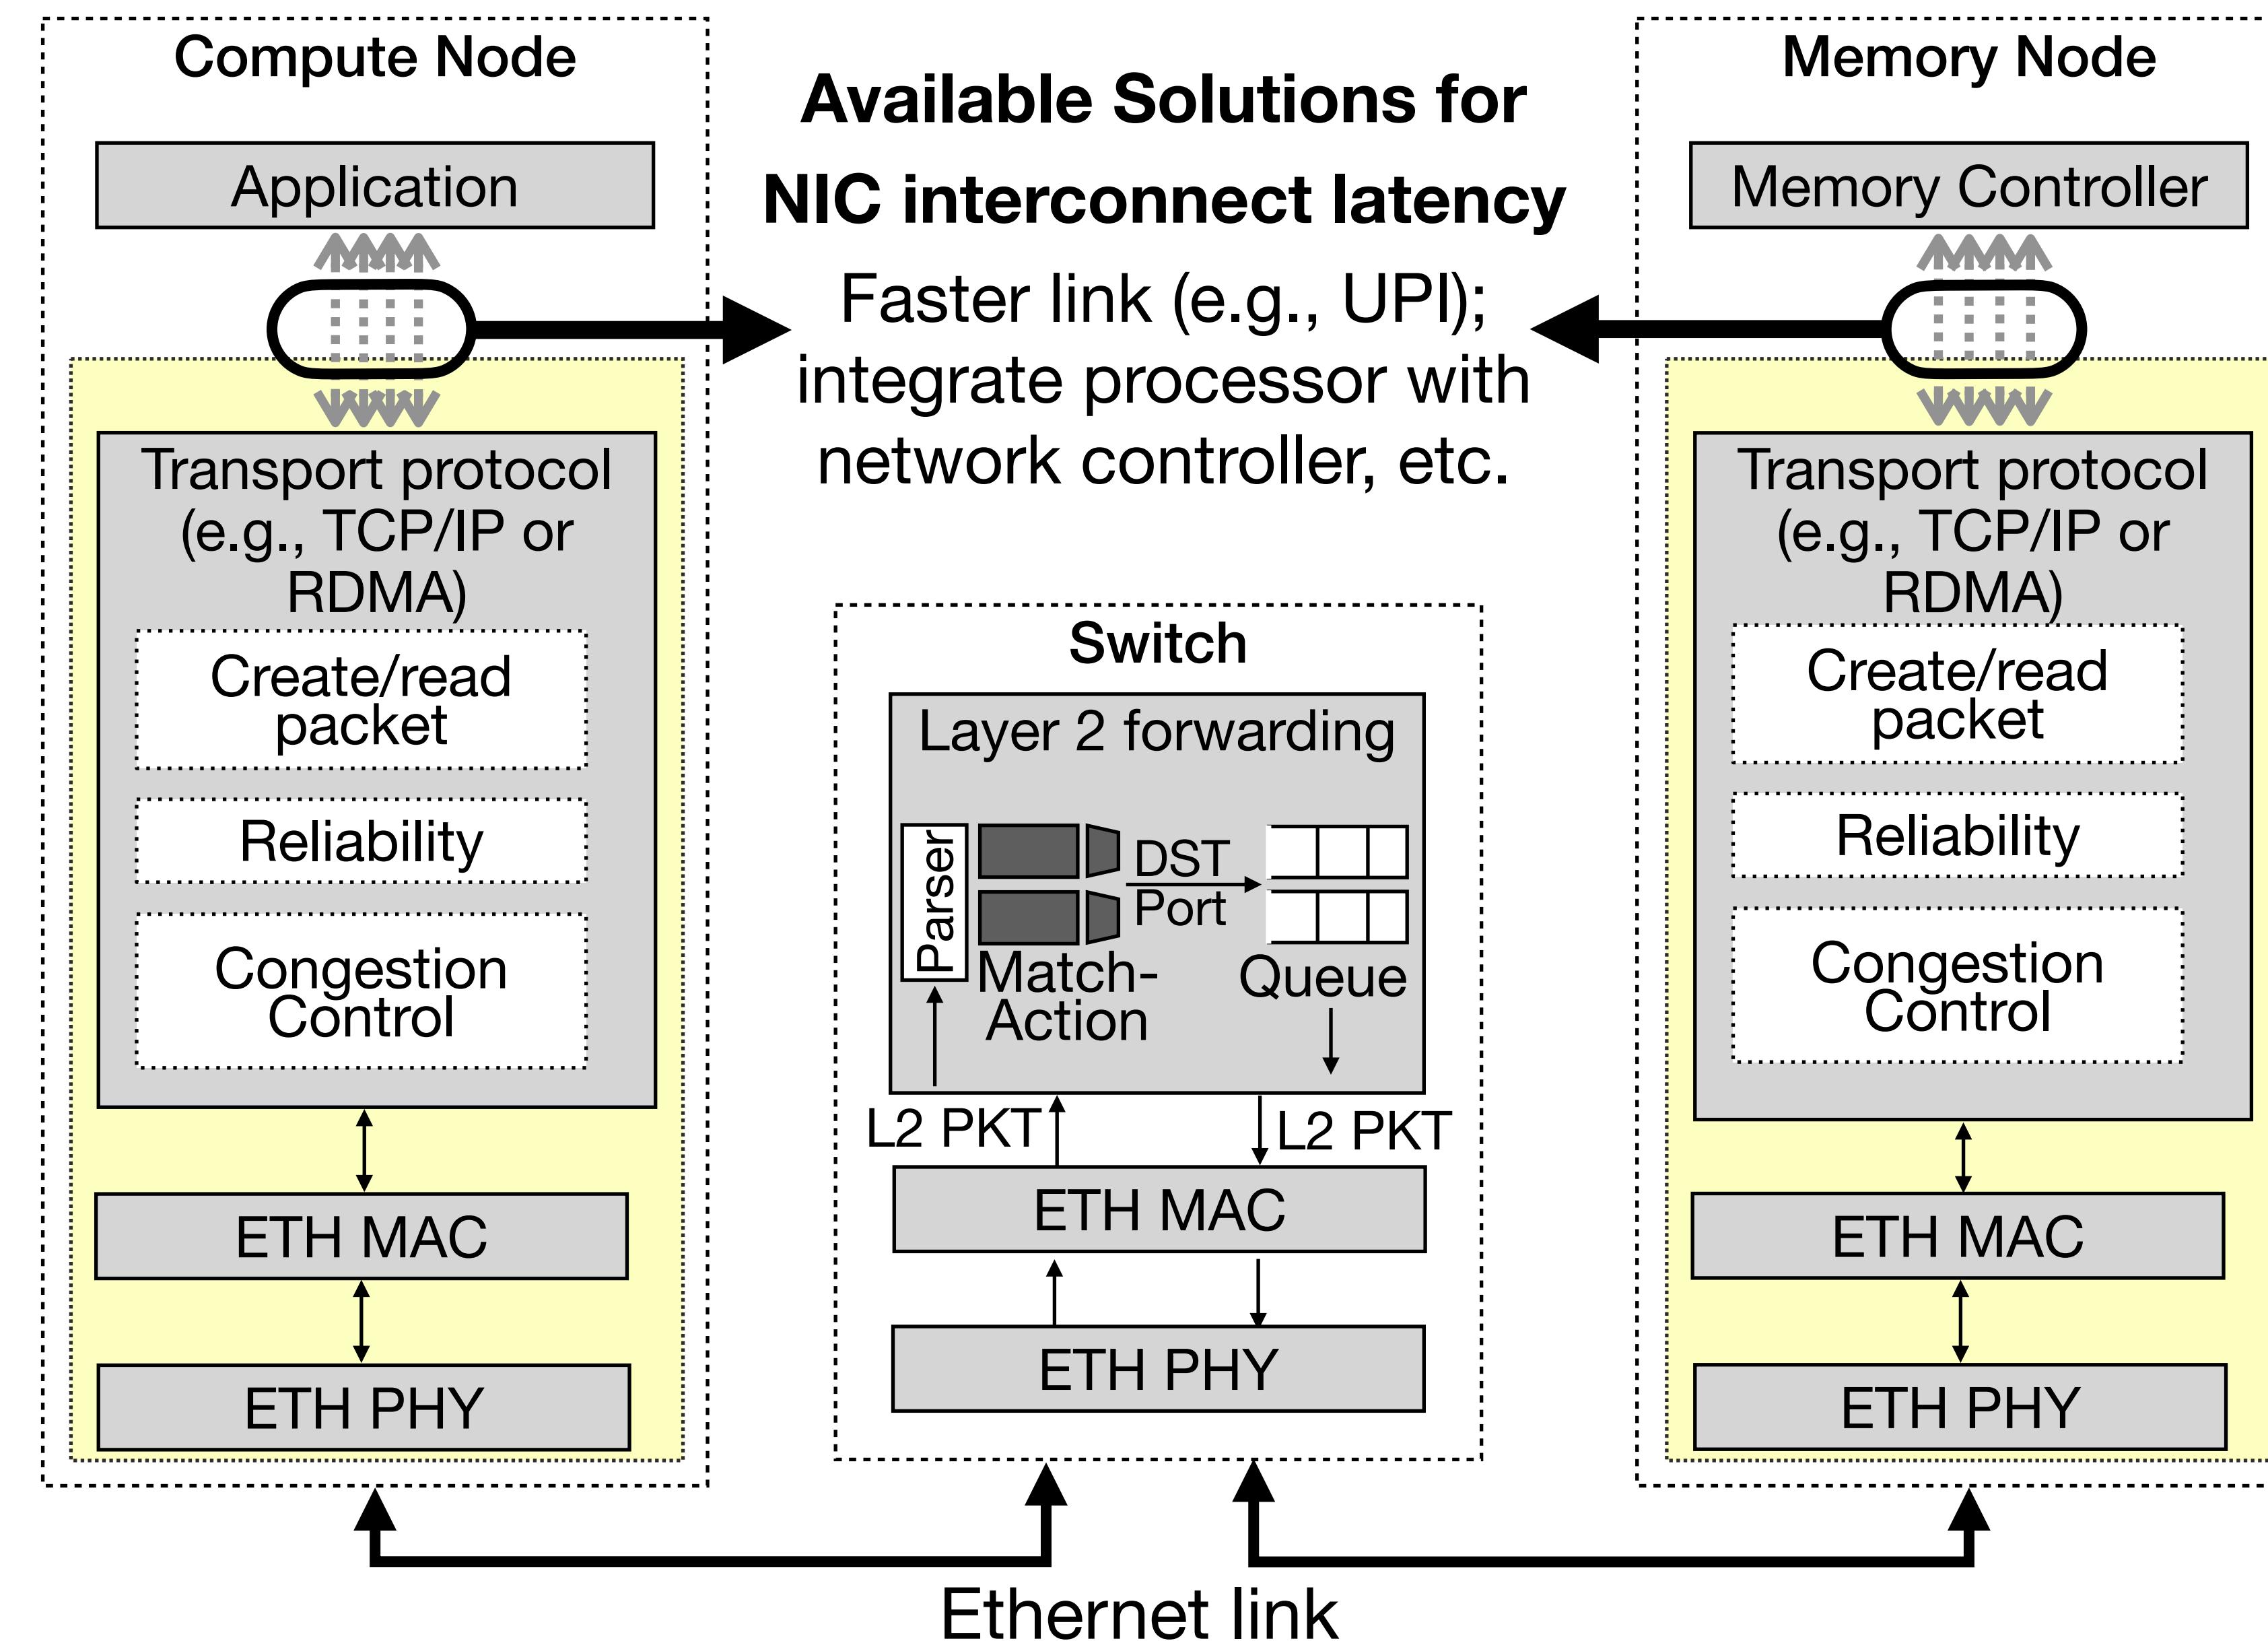
often limits the speed of new hardware adoption.

We believe that datacenters should break monolithic servers and organize hardware devices like CPU, DRAM, and disks as independent failure isolated components.

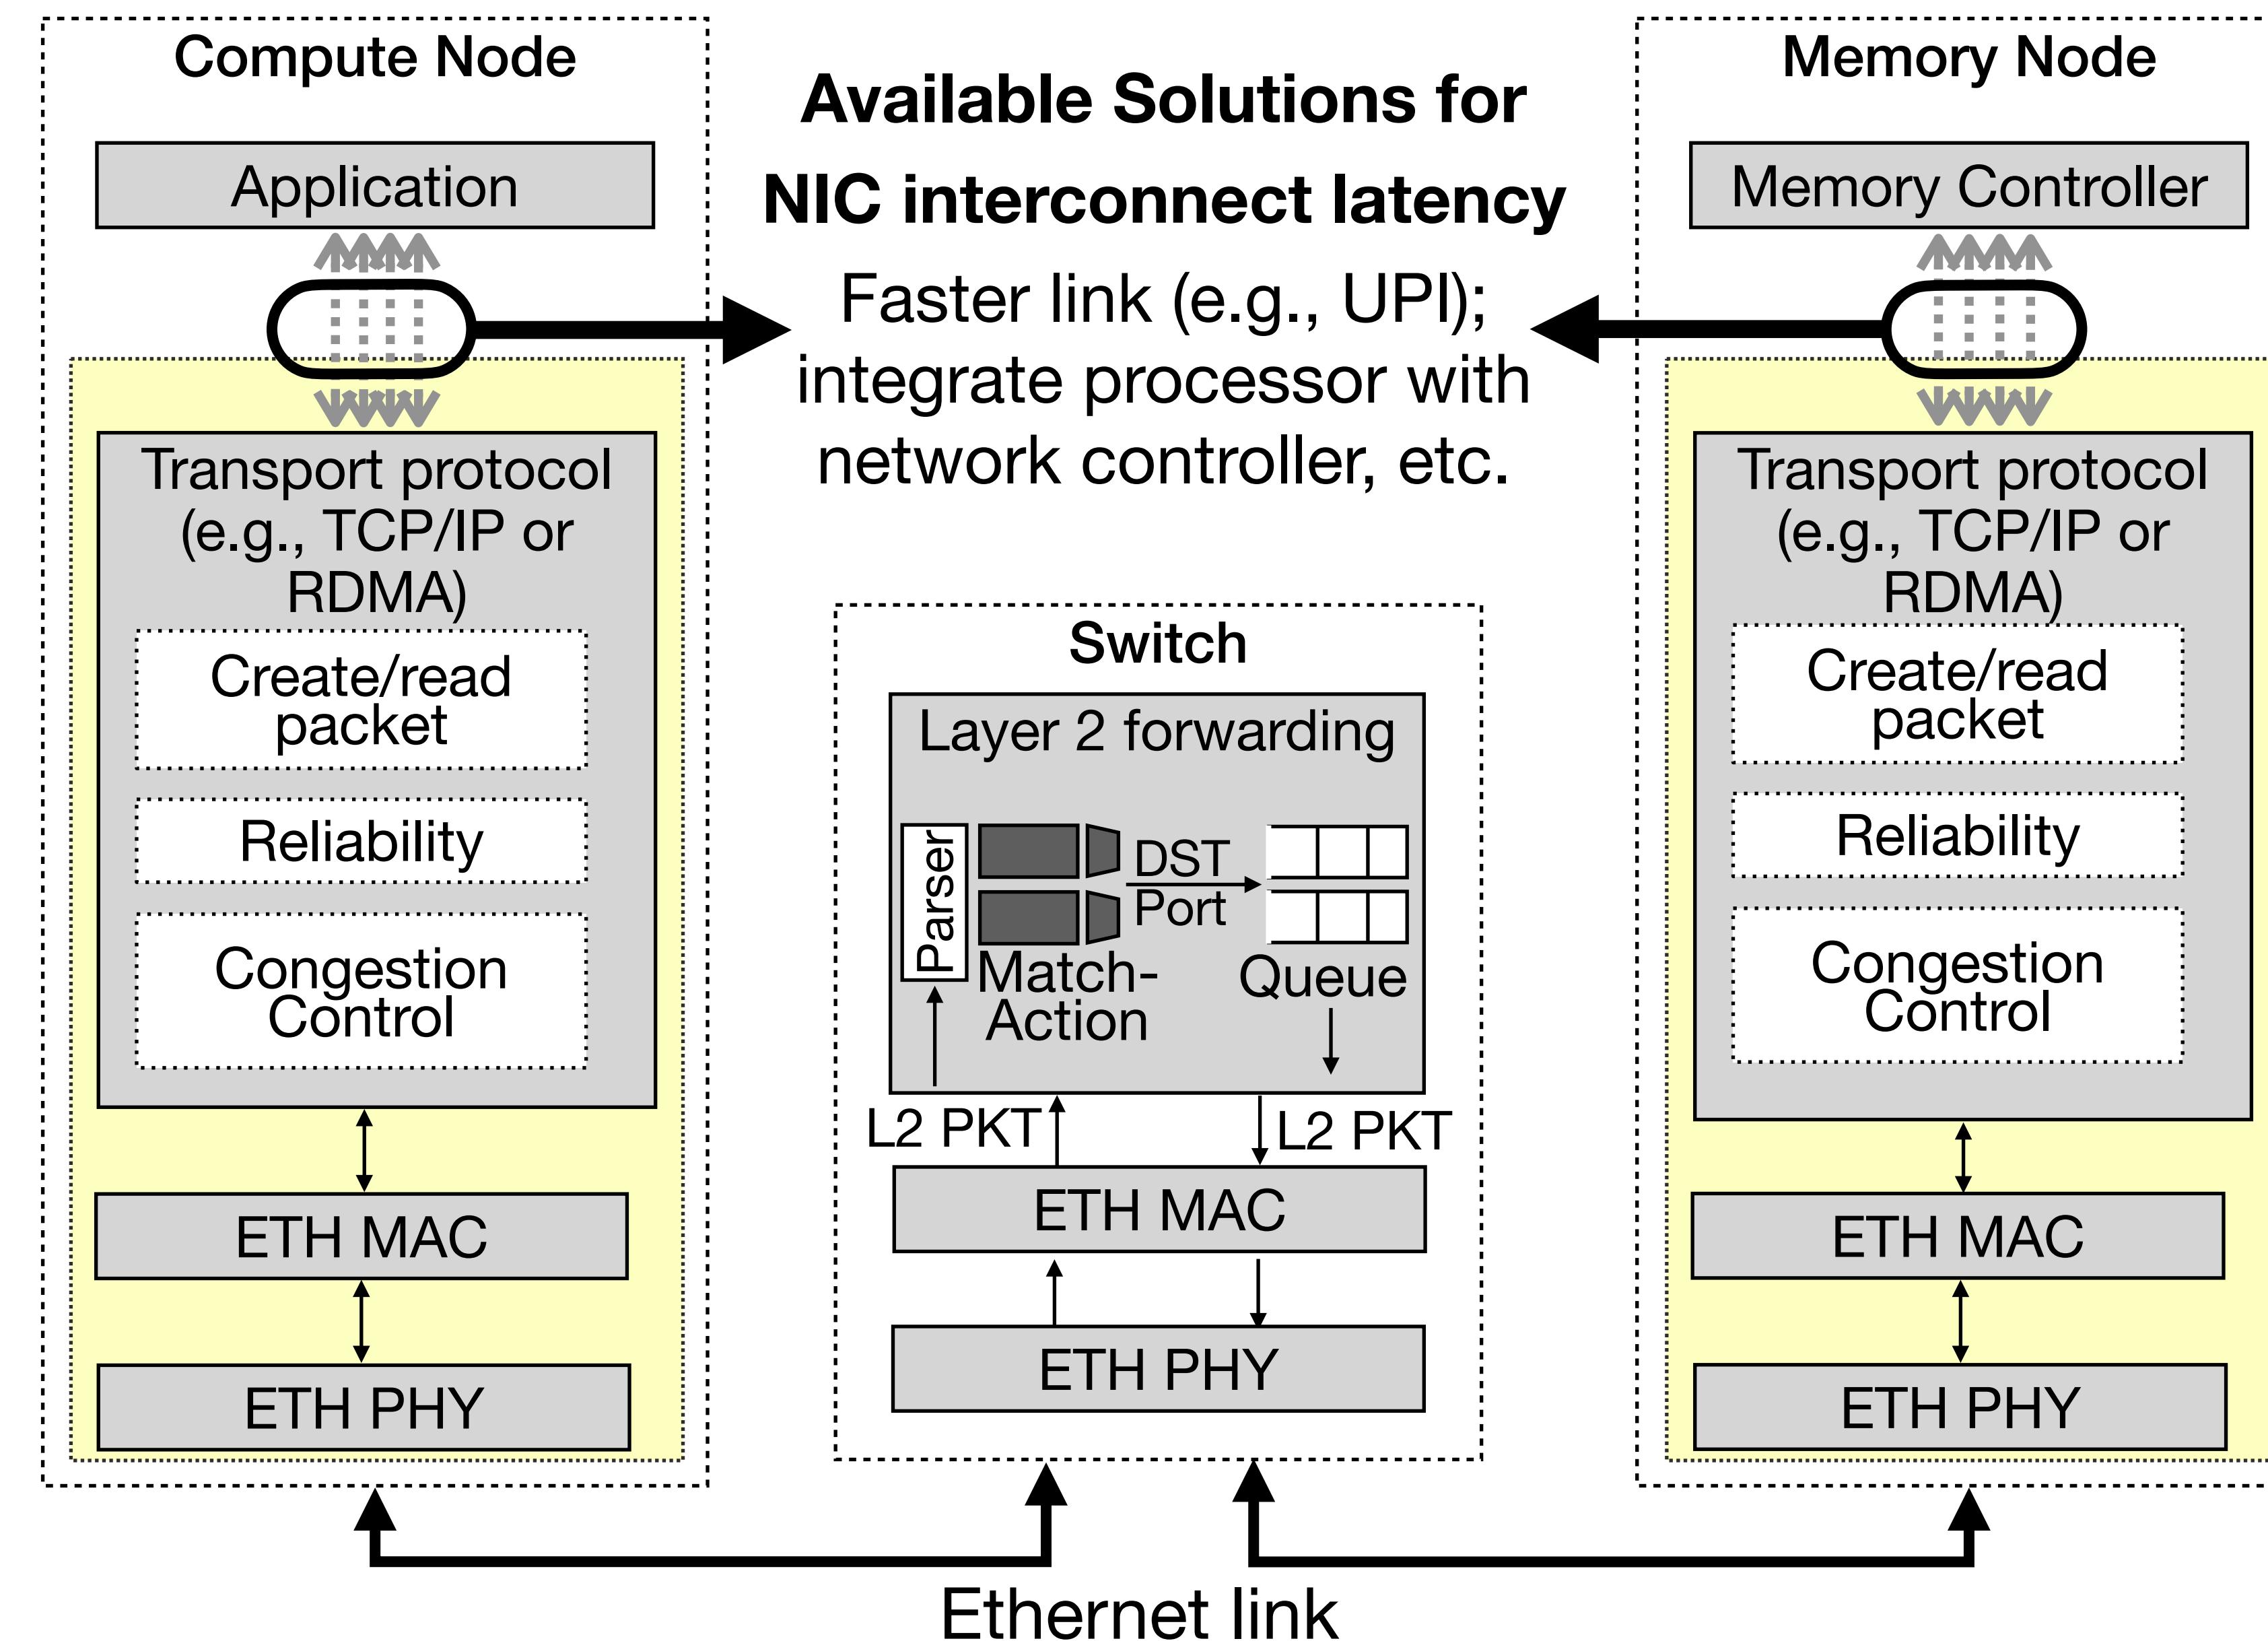
Ricardo Bianchini
Microsoft Azure
USA

But, separate fabrics for different traffic makes the network **costly** and **harder to manage**

Memory Disaggregation over Ethernet



Memory Disaggregation over Ethernet

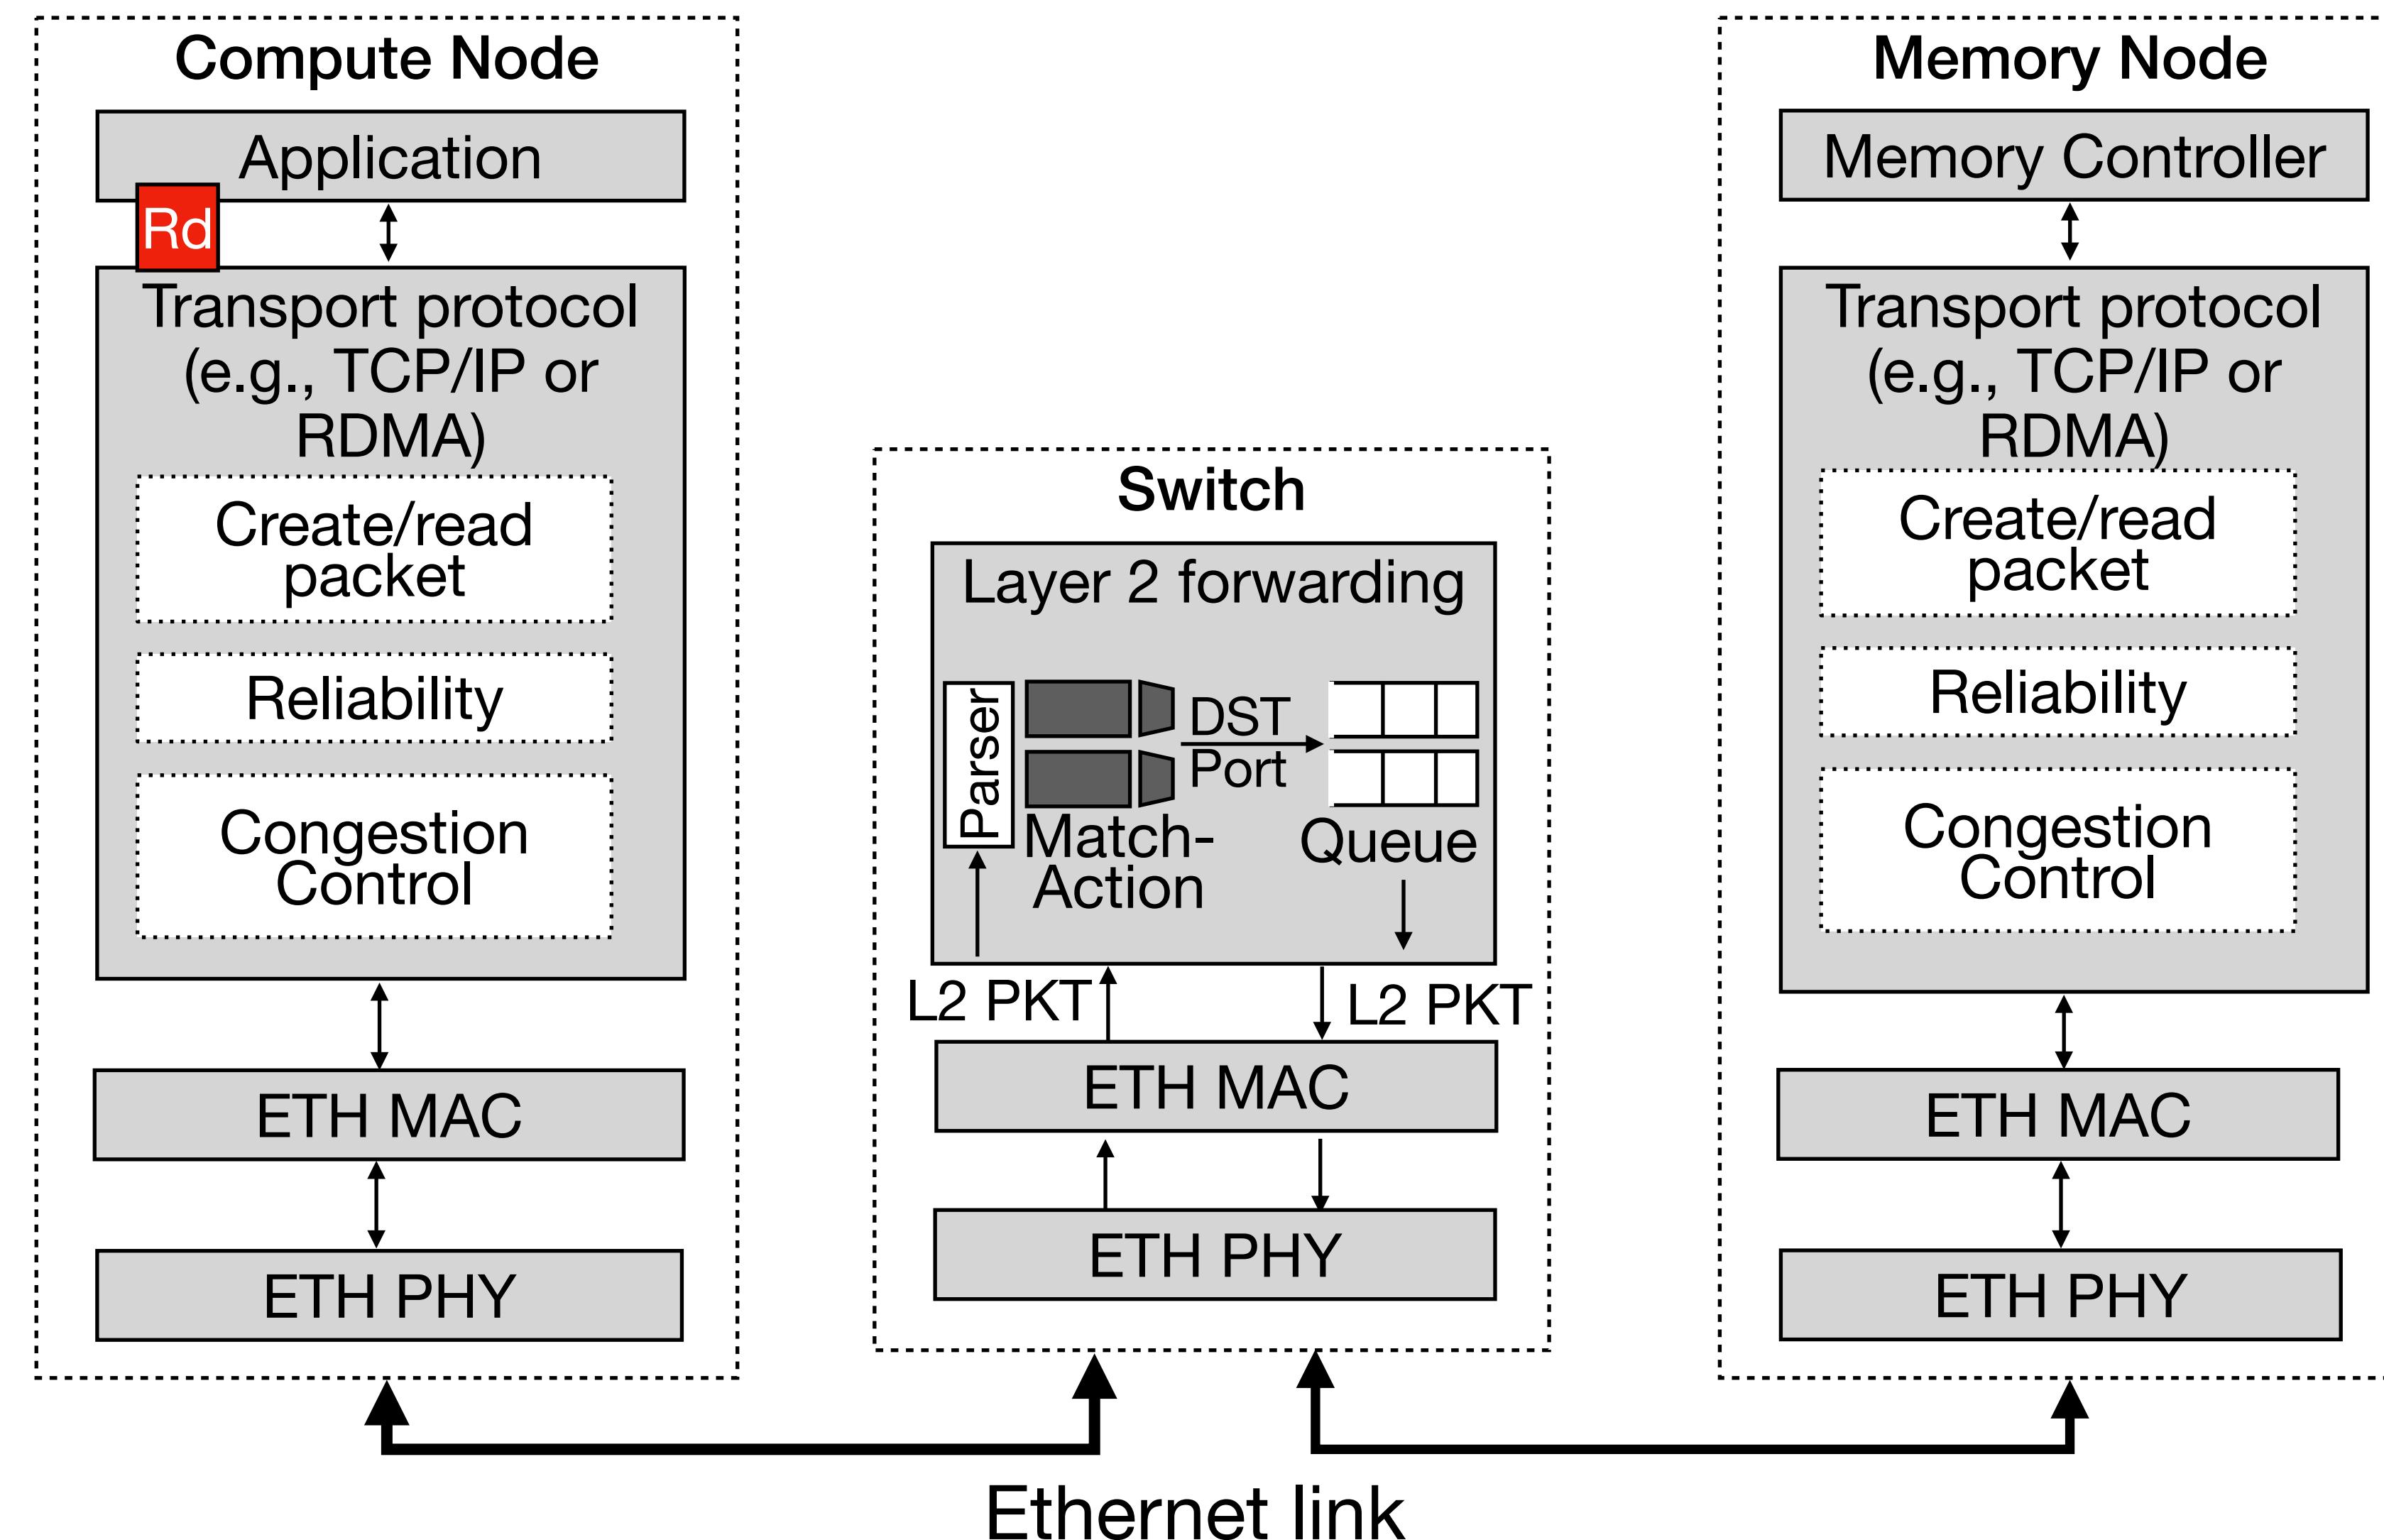


Research goal

Achieving near **intra-server memory access latency** over rack-scale **Ethernet**
(while maintaining high bandwidth utilization)

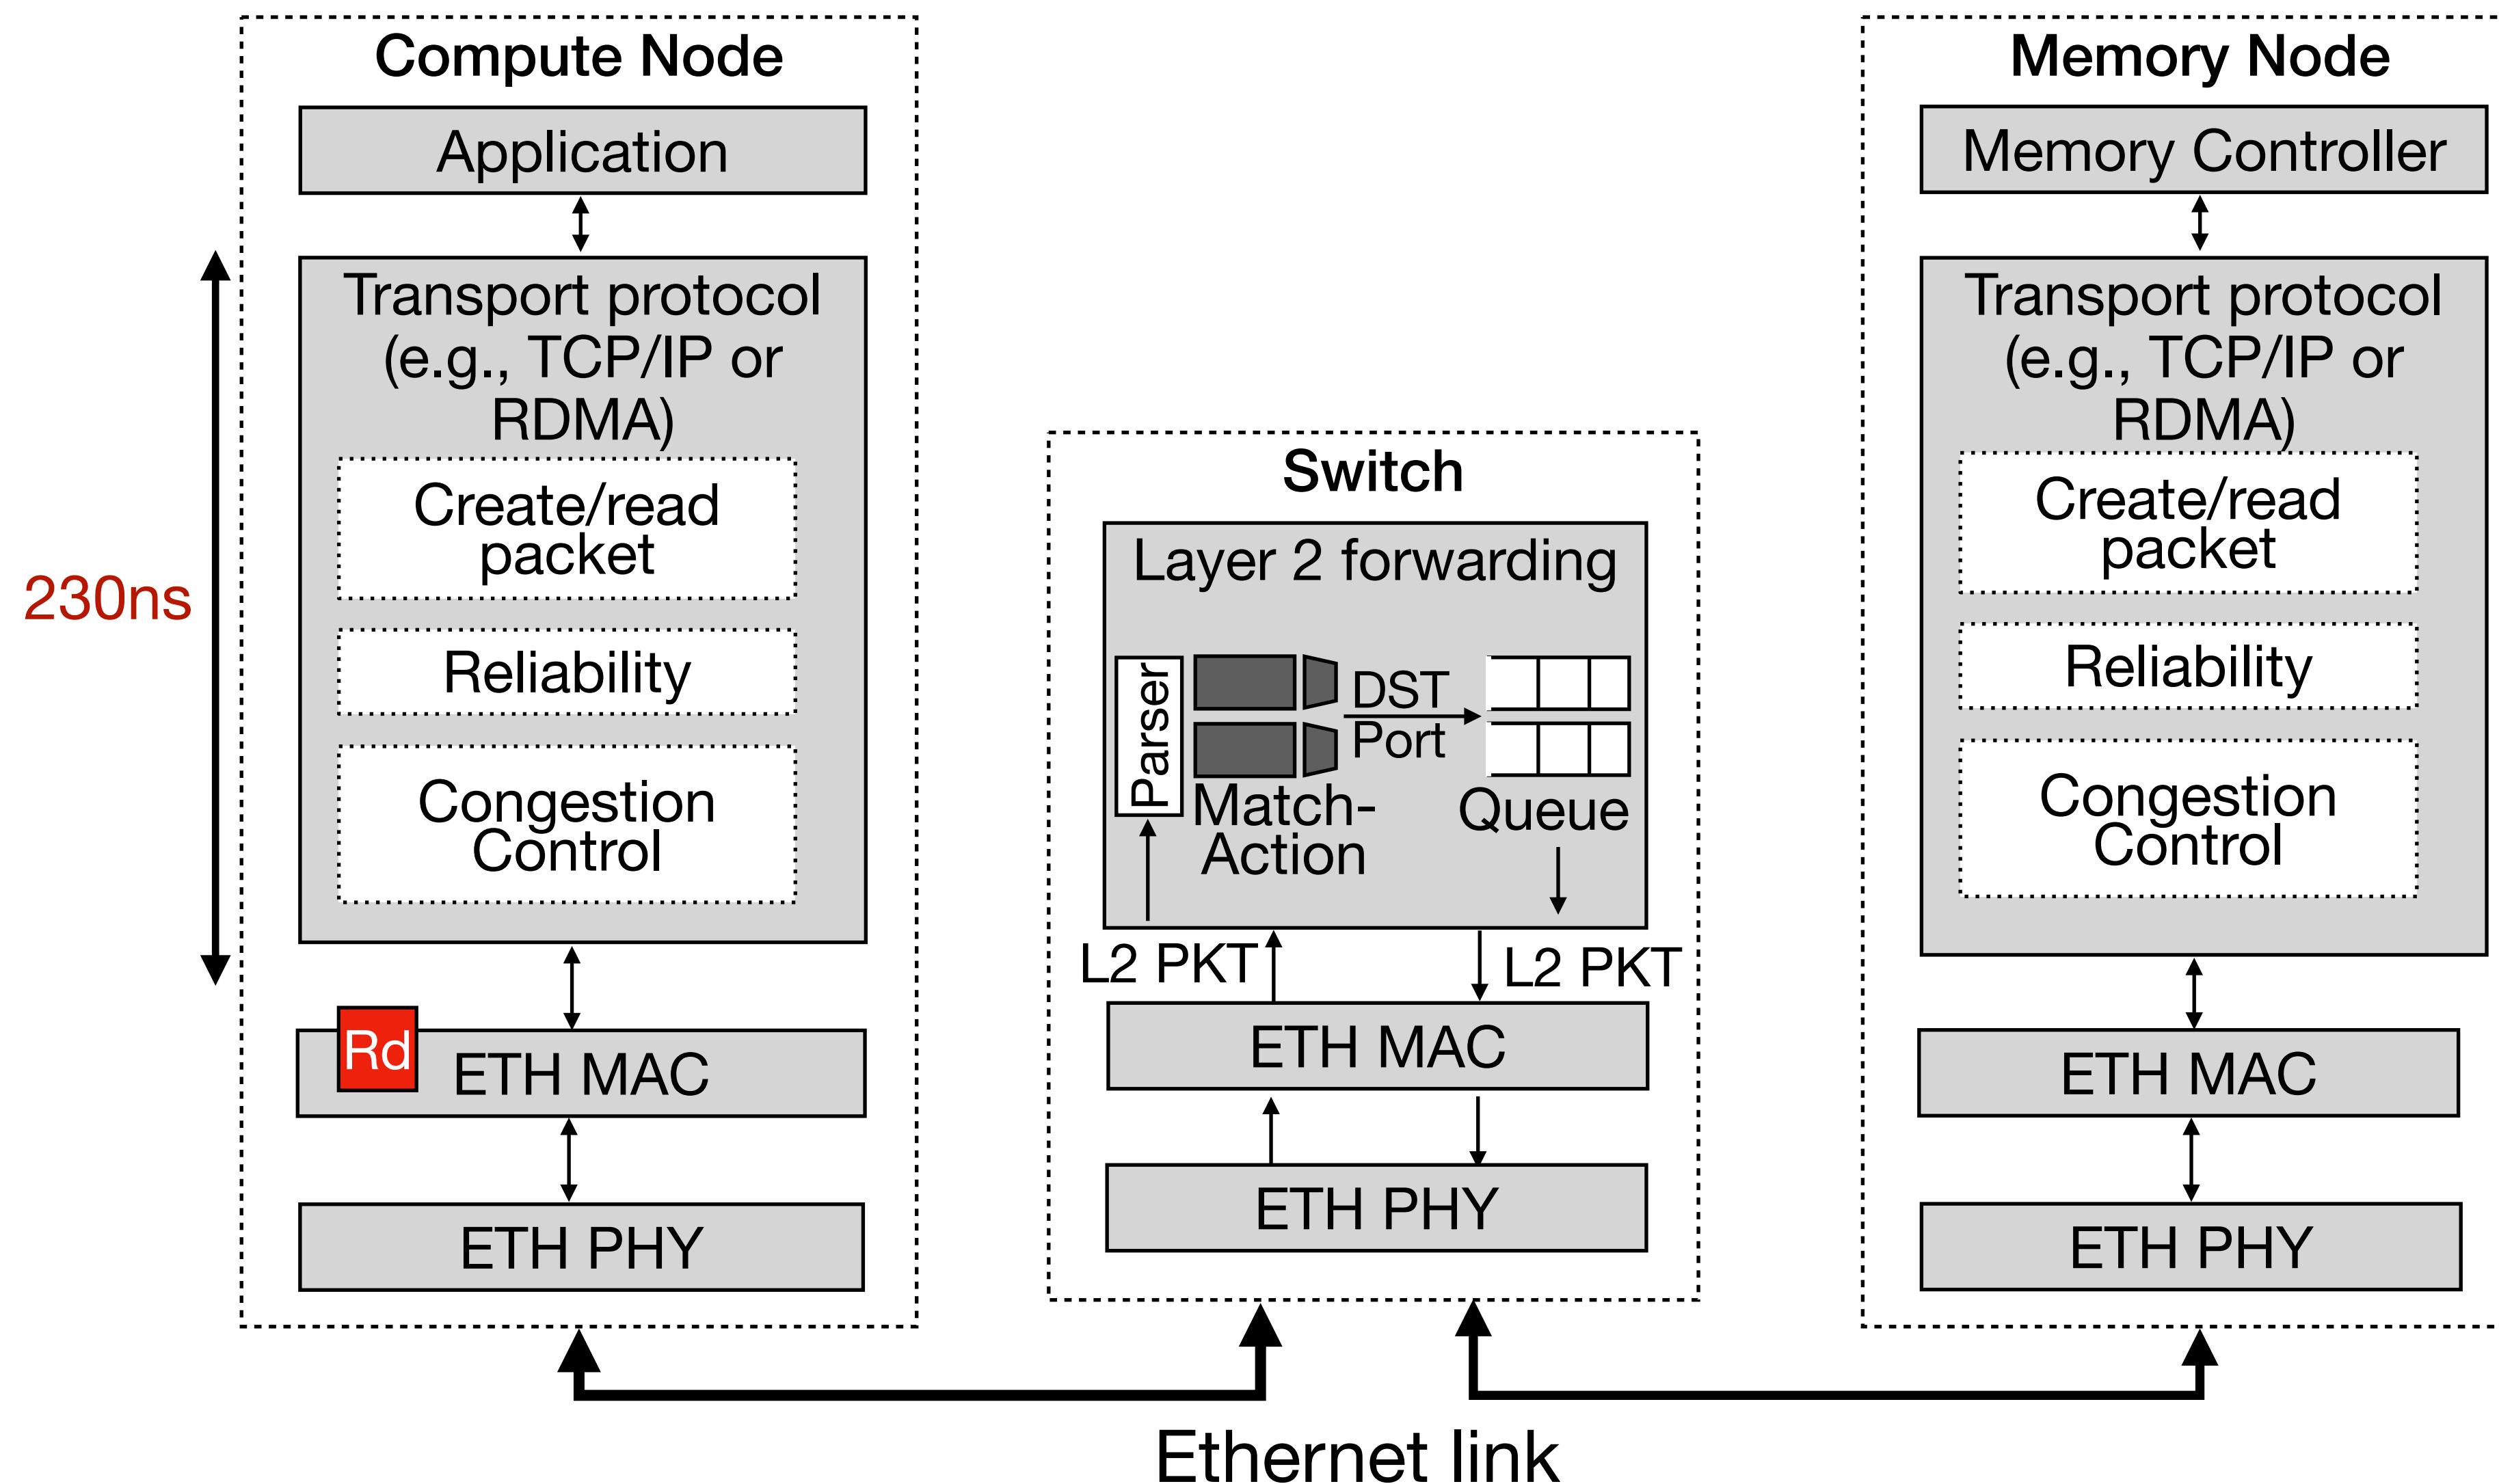
Latency Overheads of Existing Memory Disaggregation over Ethernet

An example of remote read request over RDMA



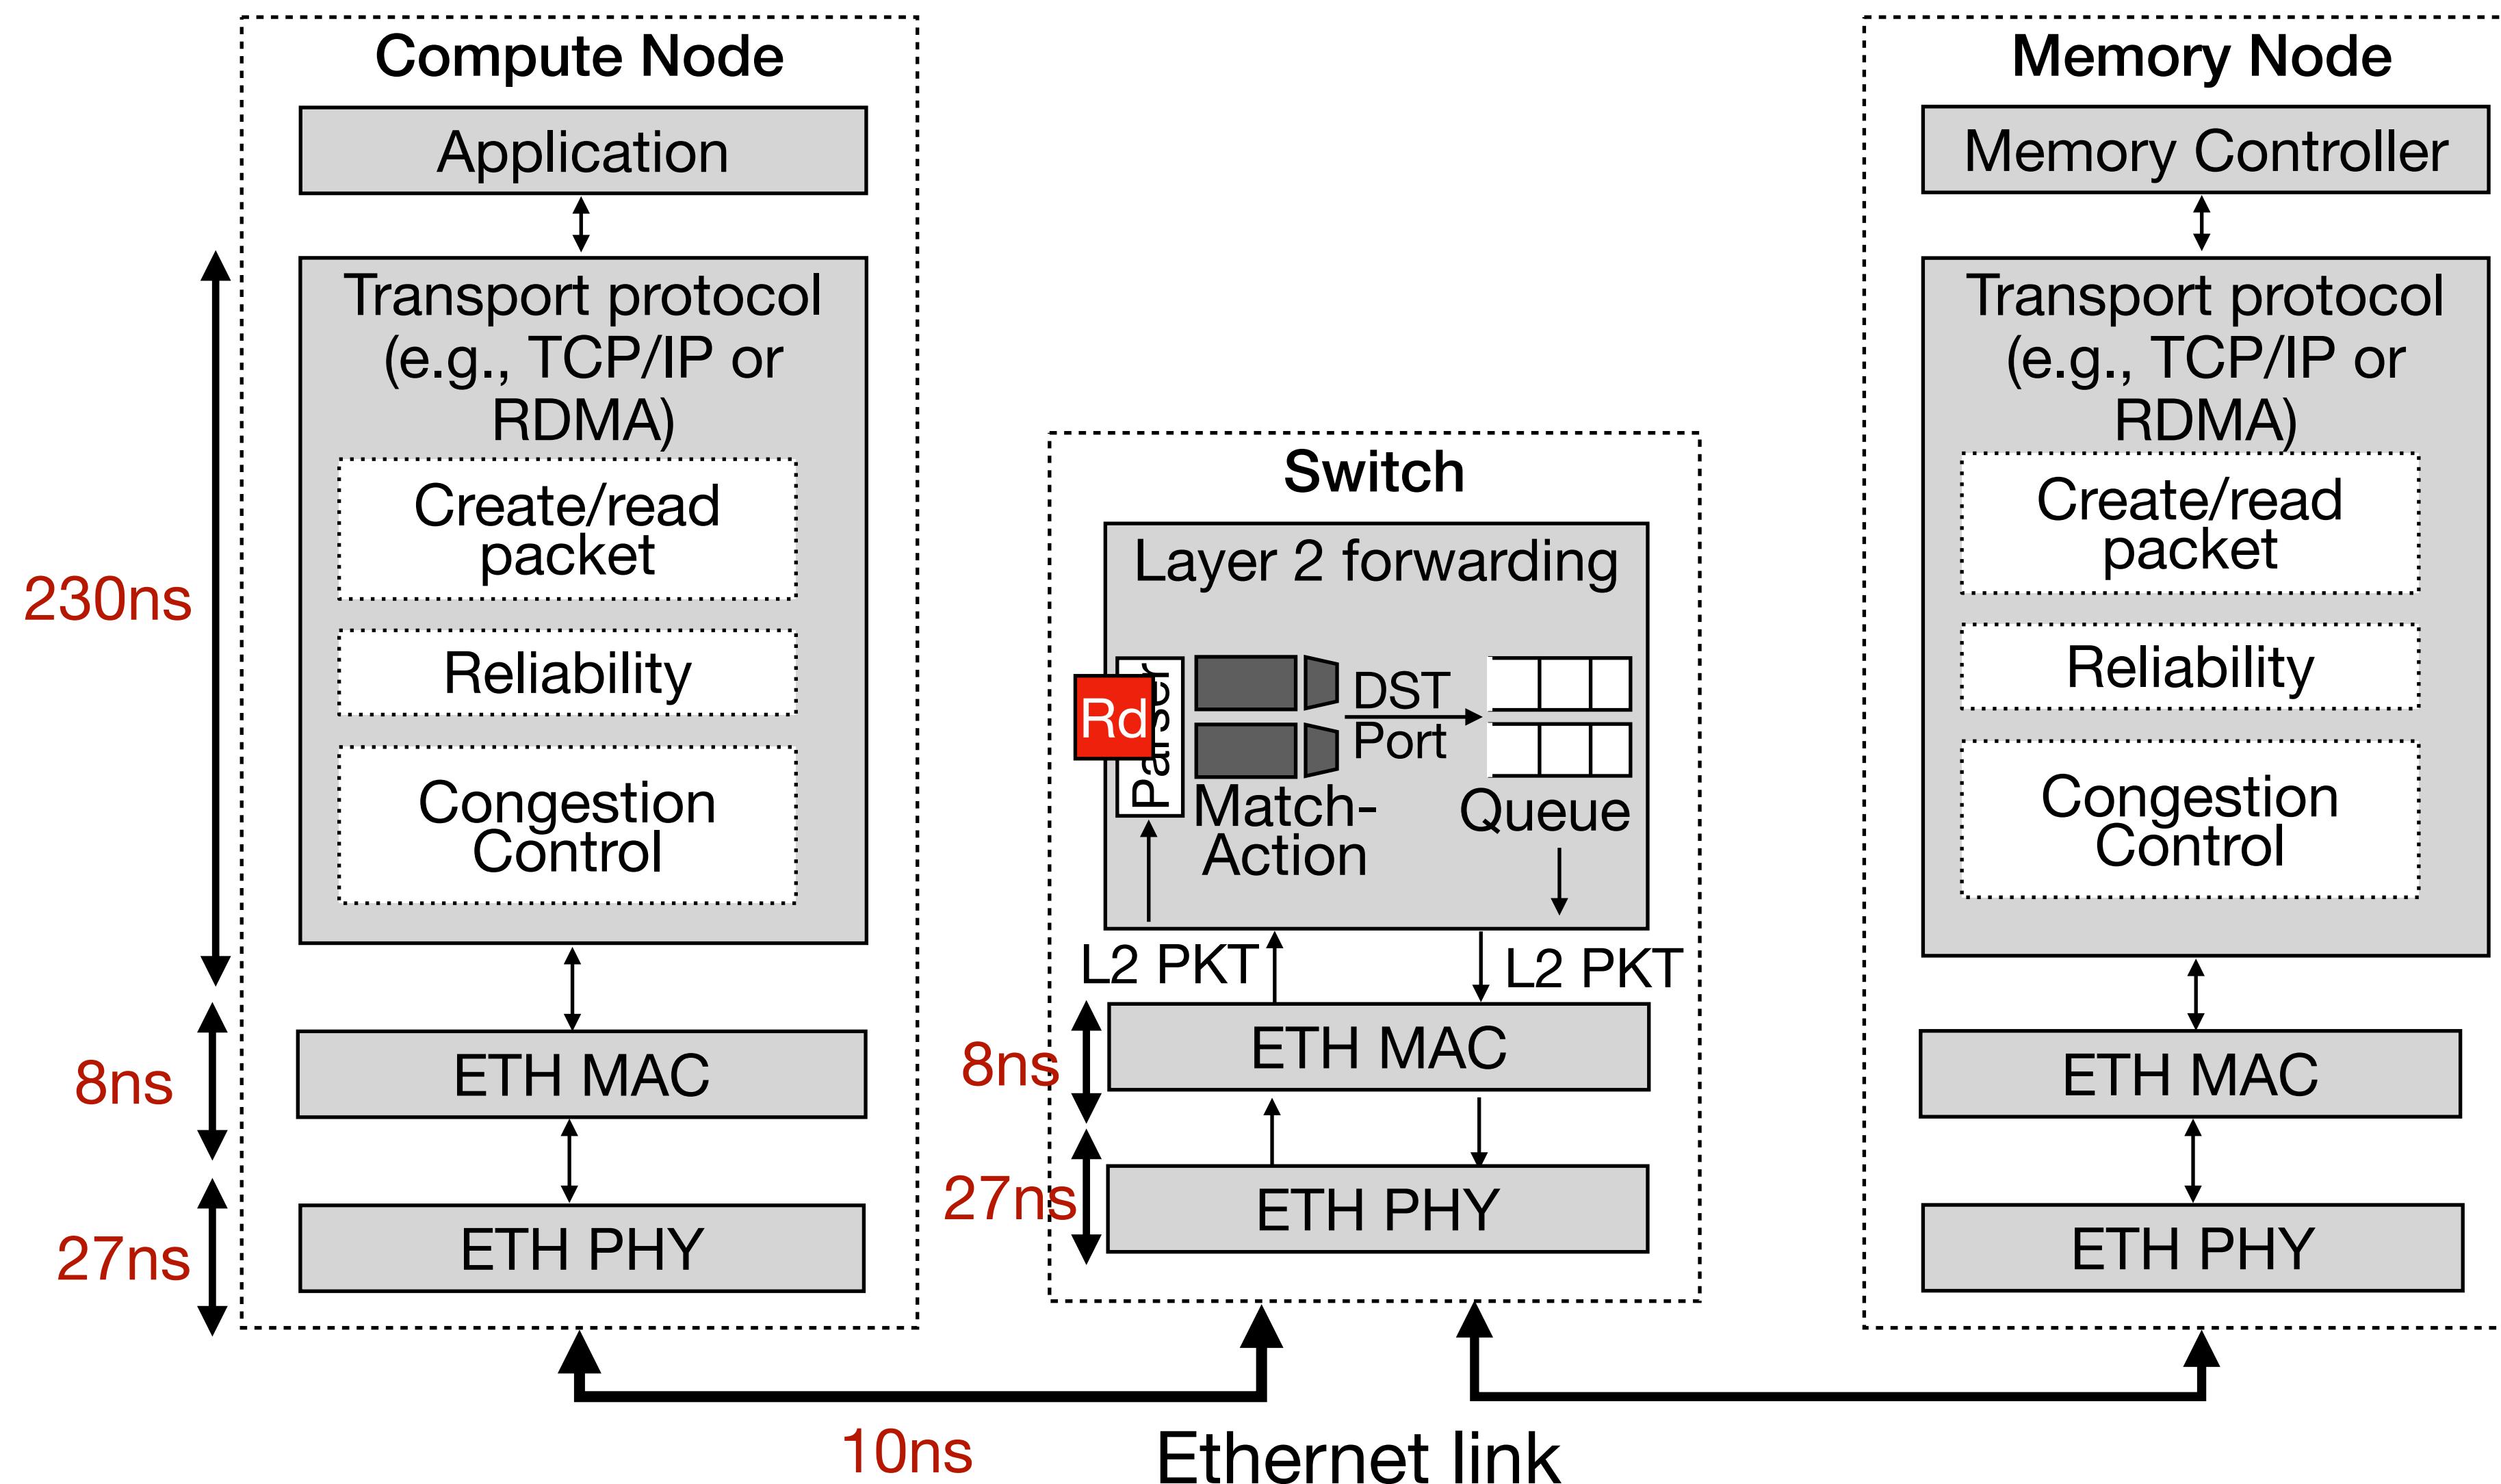
Latency Overheads of Existing Memory Disaggregation over Ethernet

An example of remote read request over RDMA



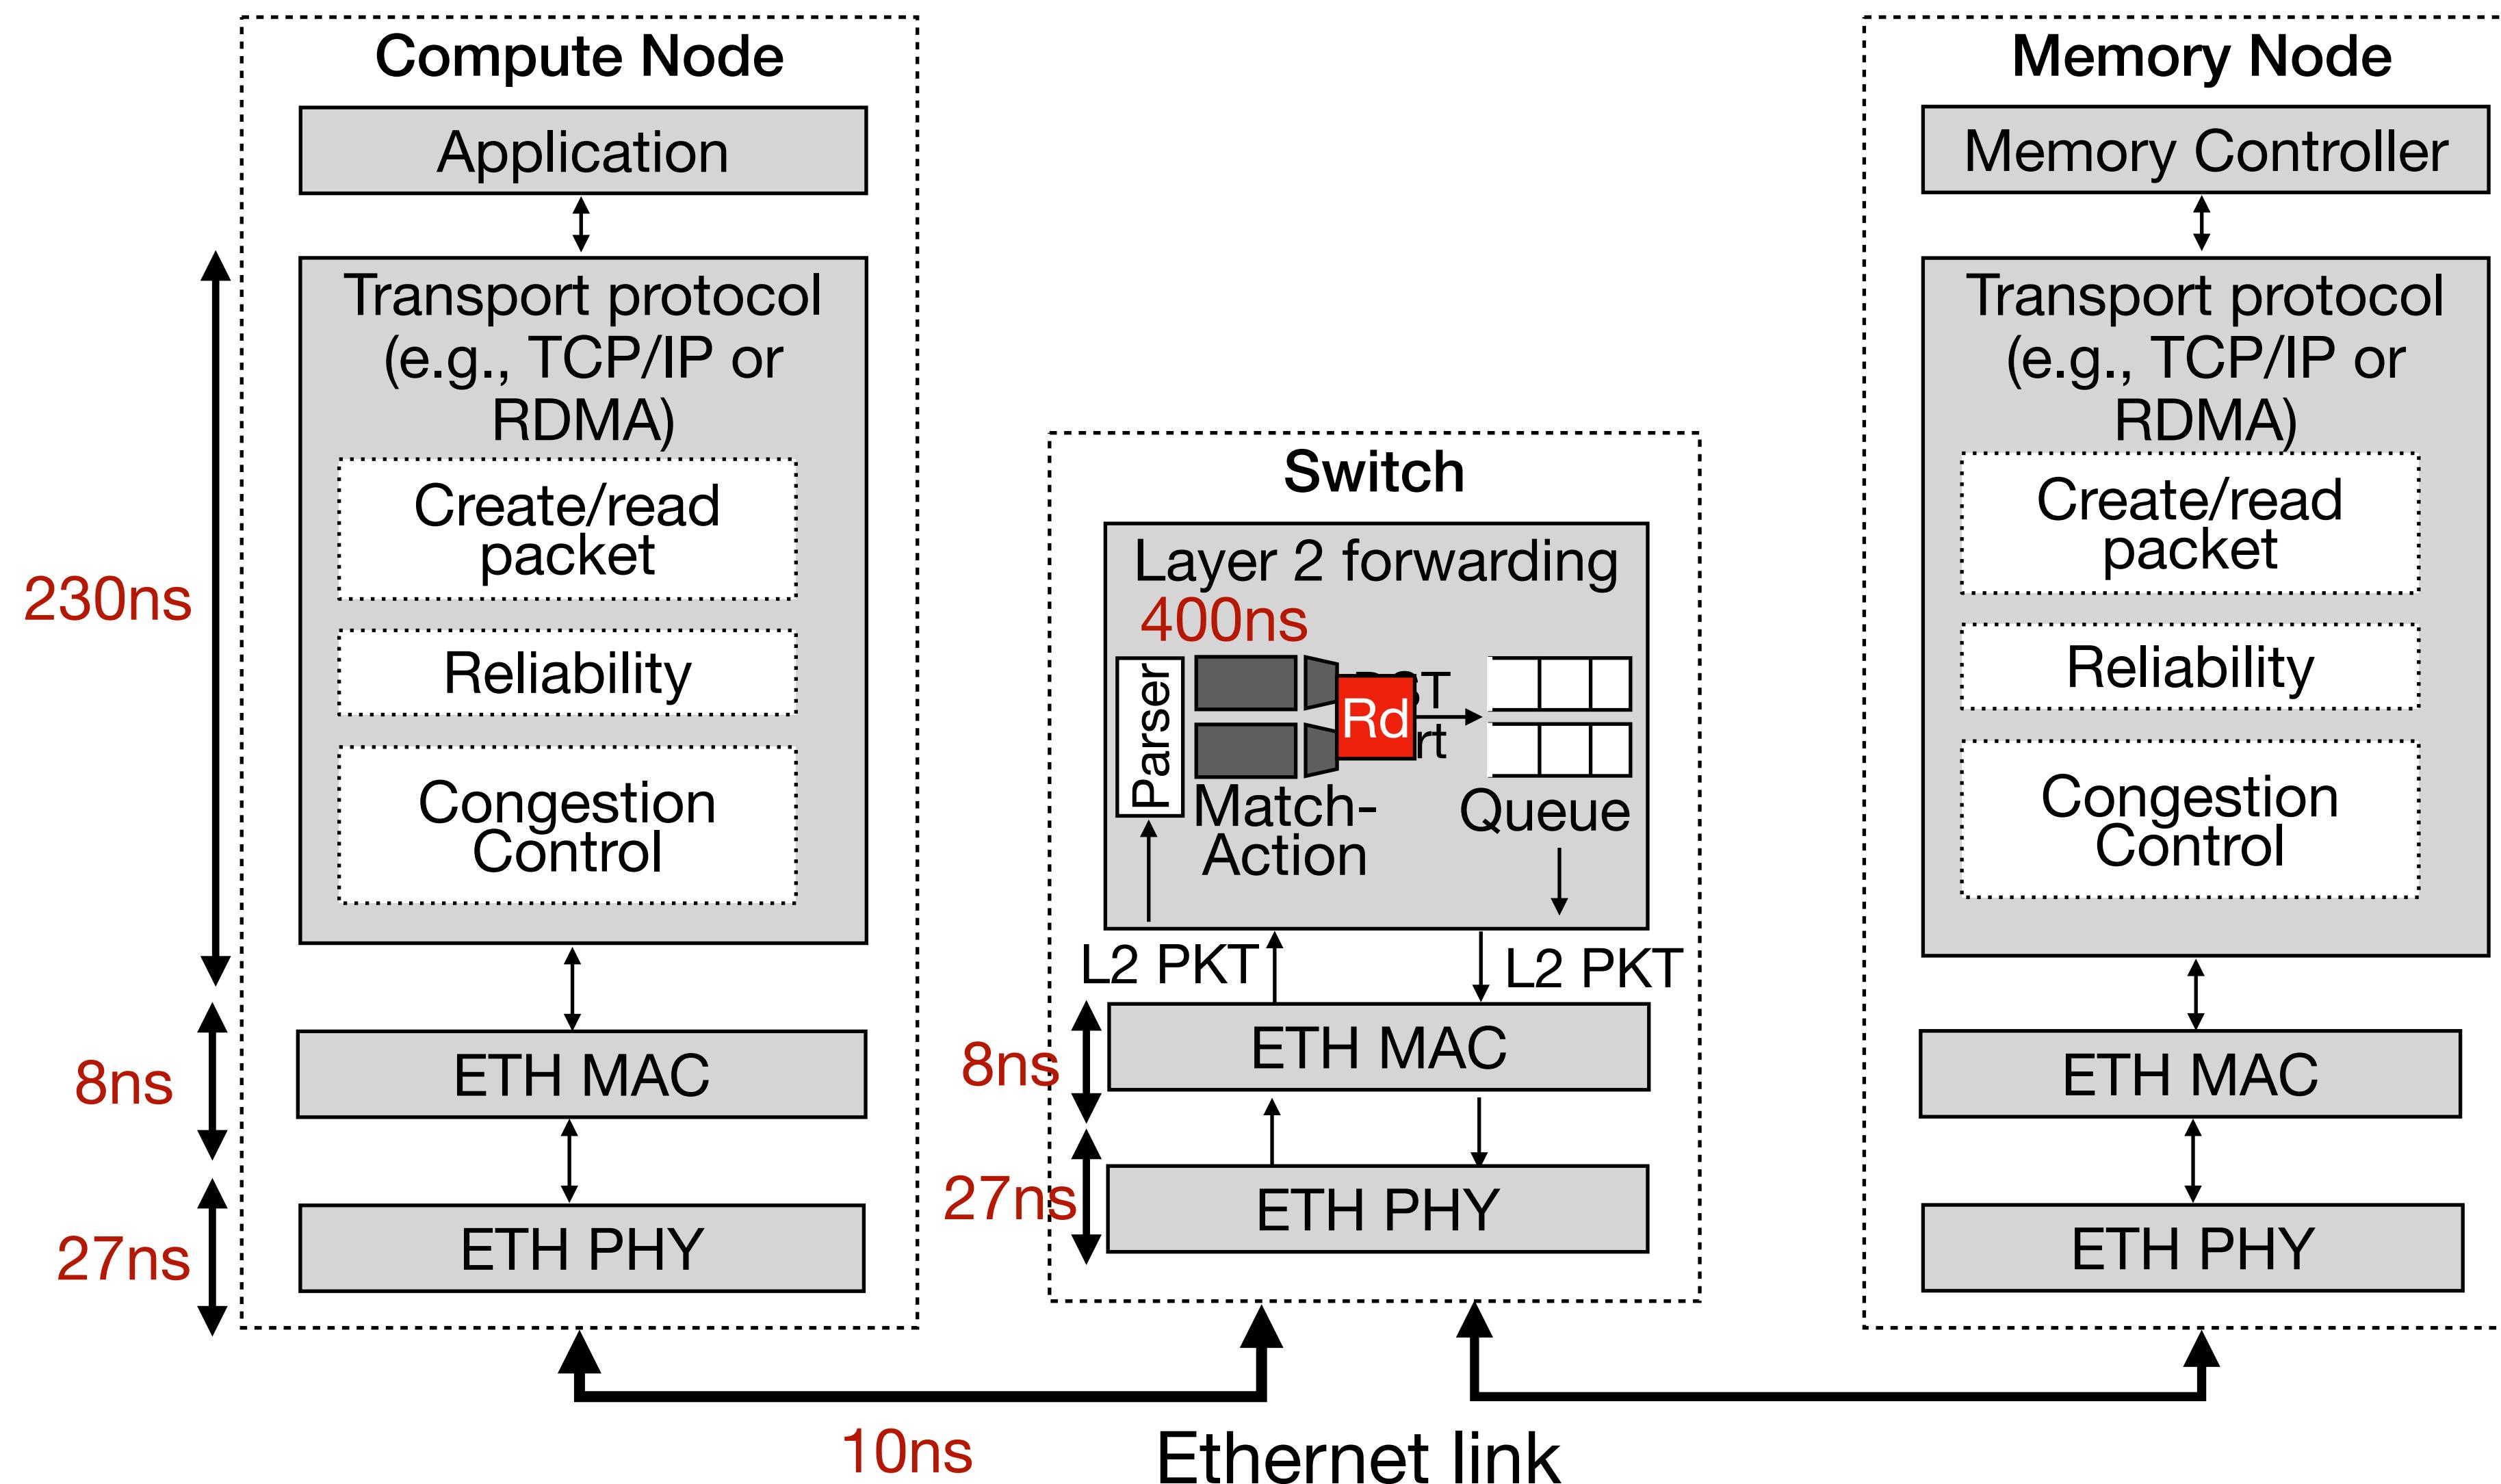
Latency Overheads of Existing Memory Disaggregation over Ethernet

An example of remote read request over RDMA



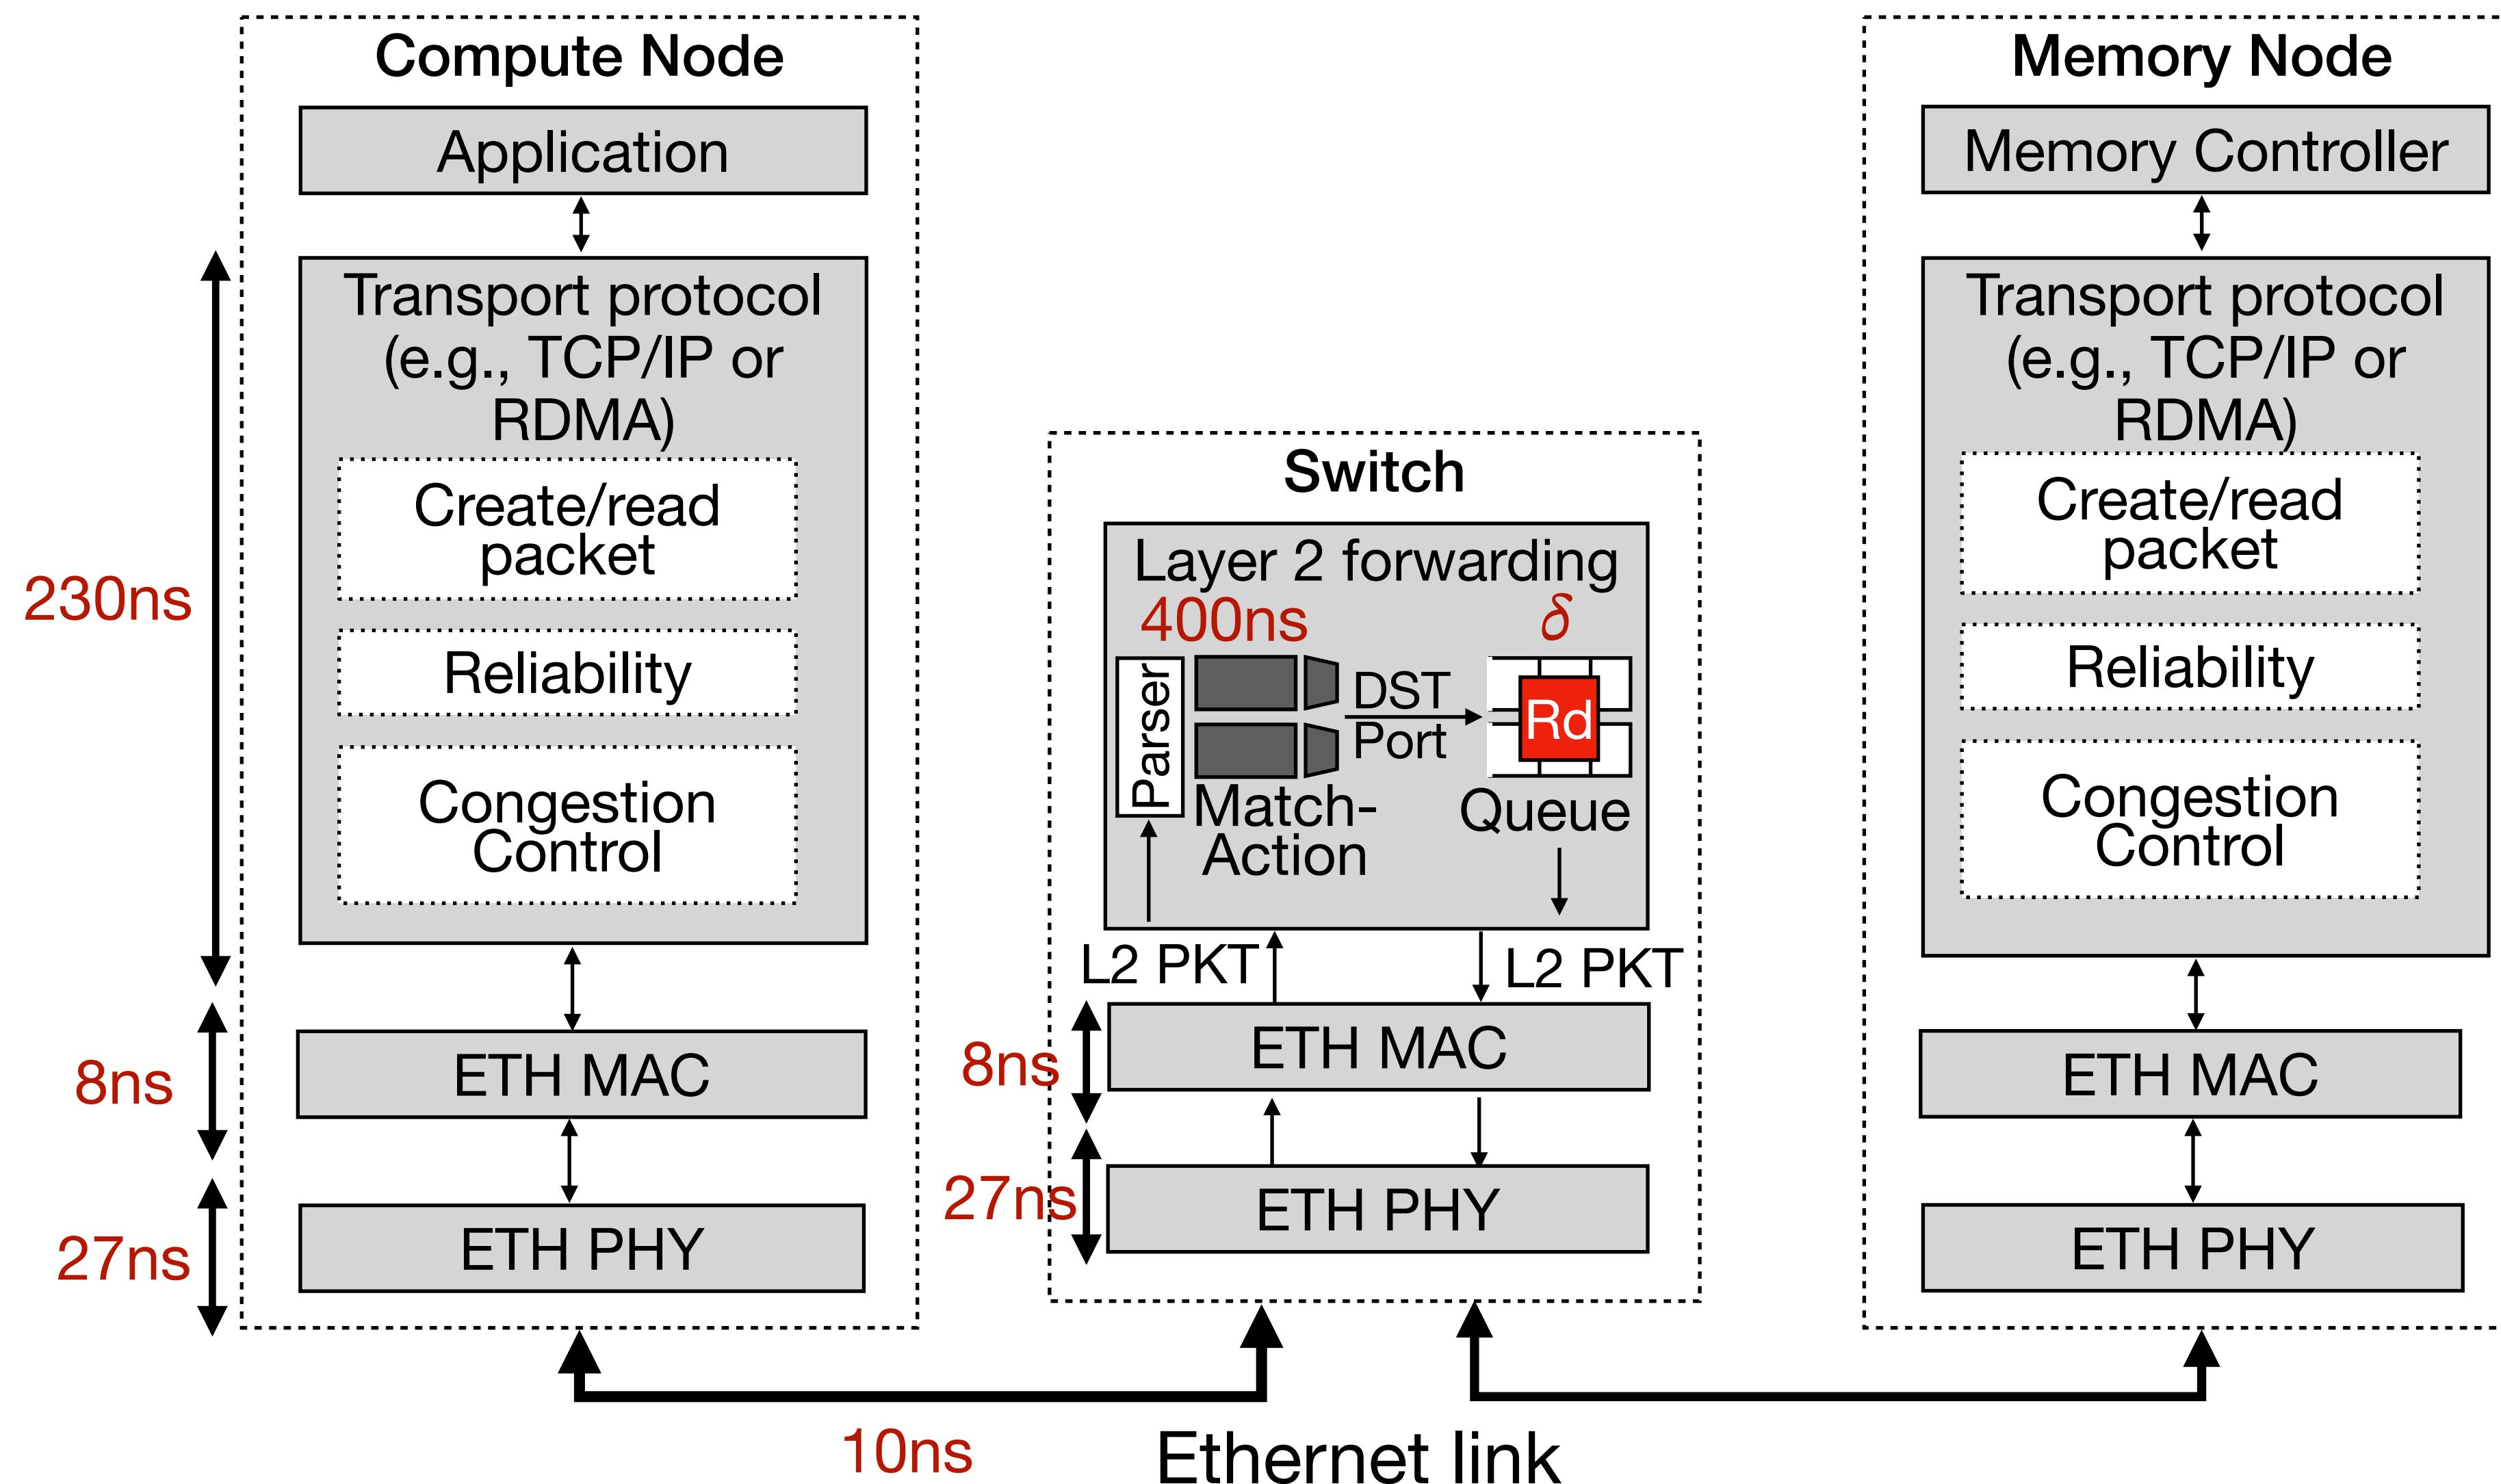
Latency Overheads of Existing Memory Disaggregation over Ethernet

An example of remote read request over RDMA



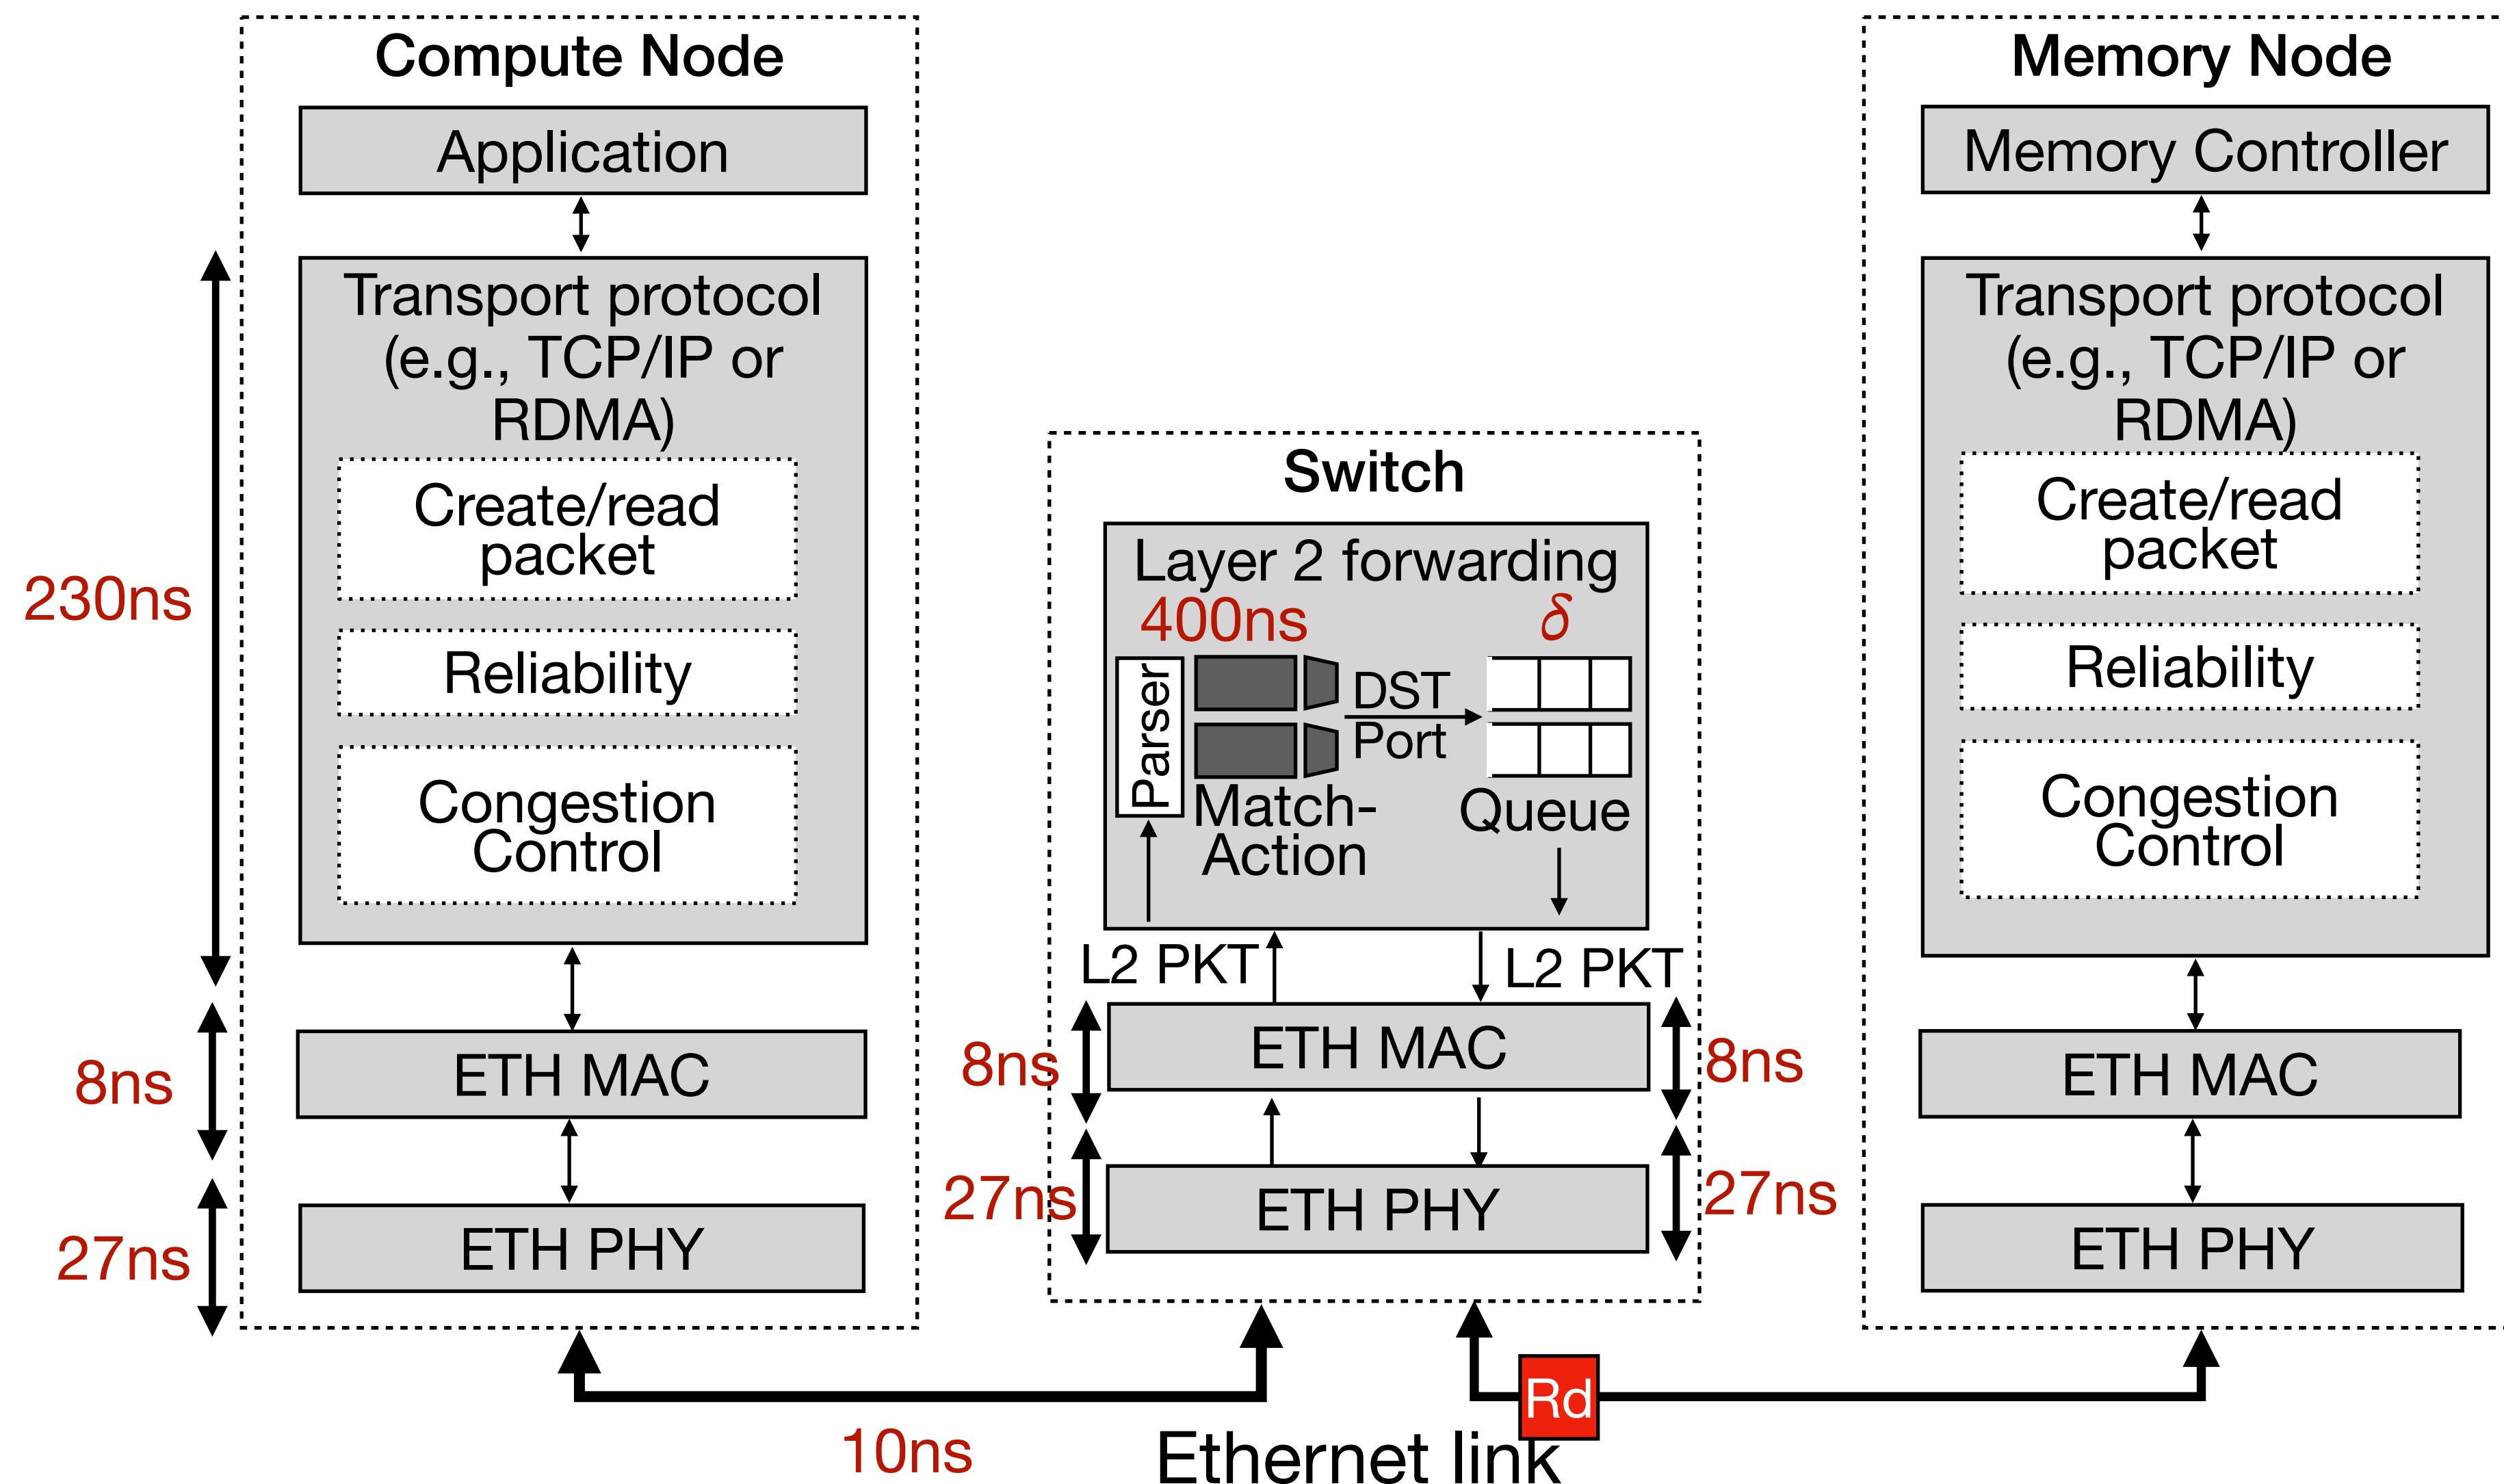
Latency Overheads of Existing Memory Disaggregation over Ethernet

An example of remote read request over RDMA



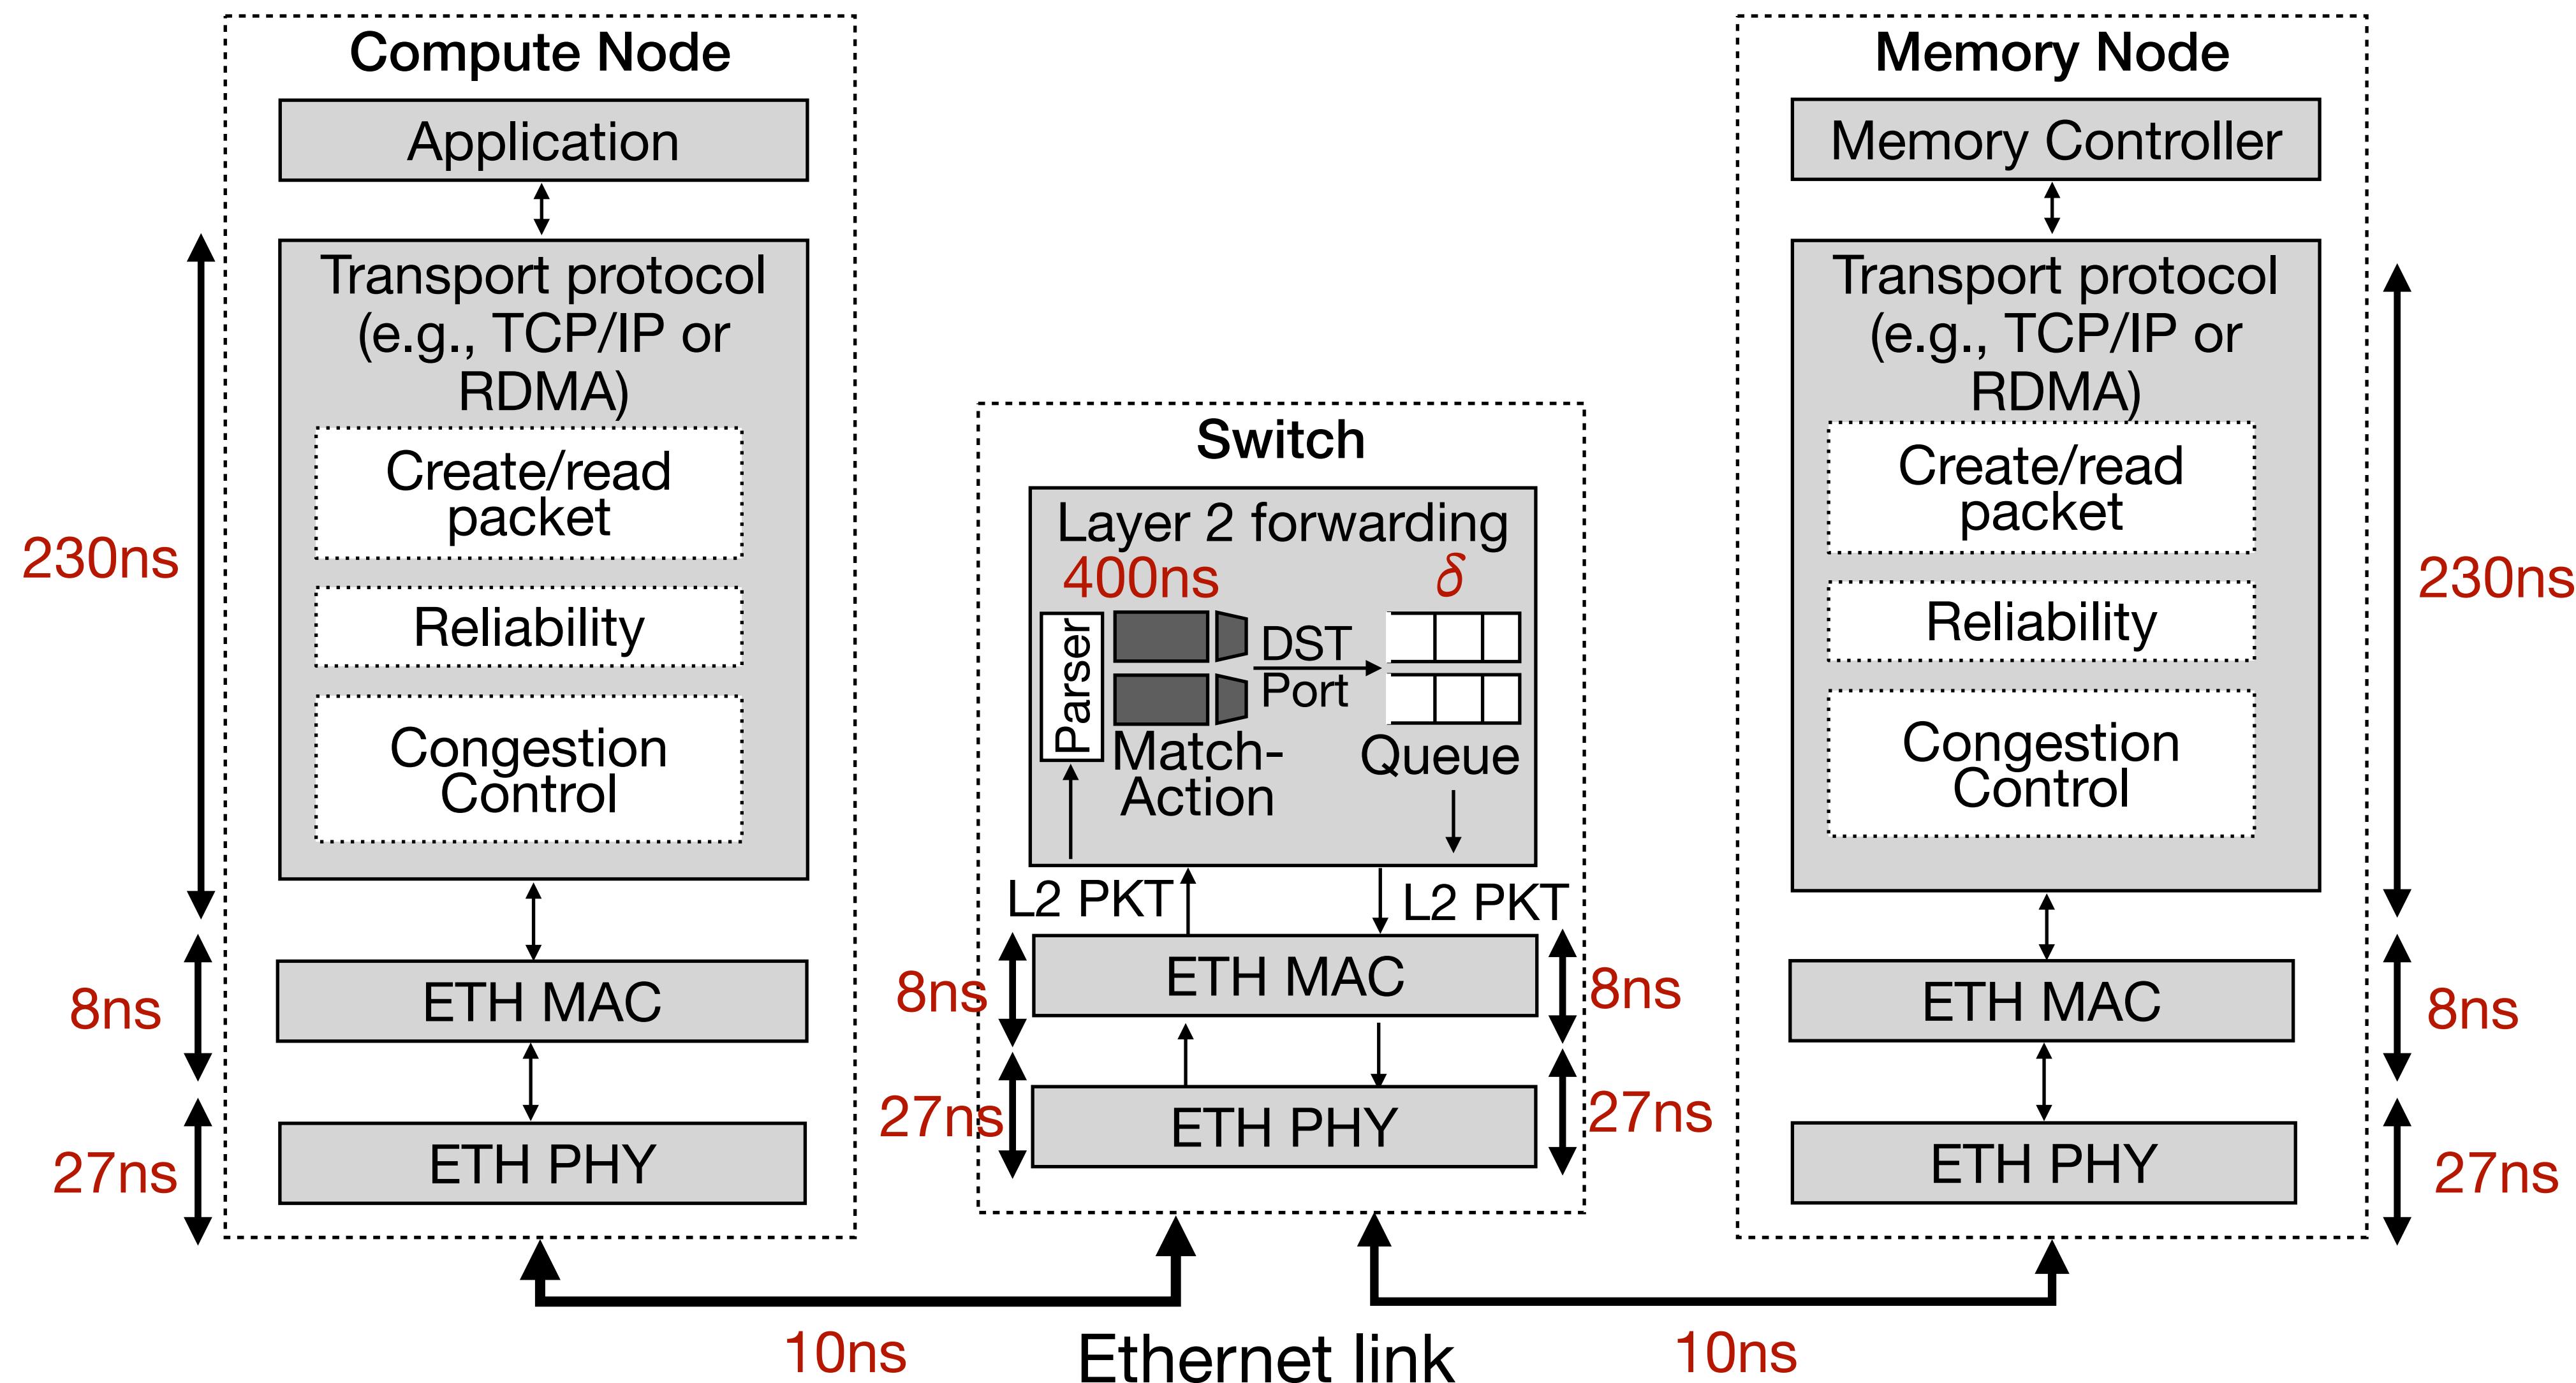
Latency Overheads of Existing Memory Disaggregation over Ethernet

An example of remote read request over RDMA



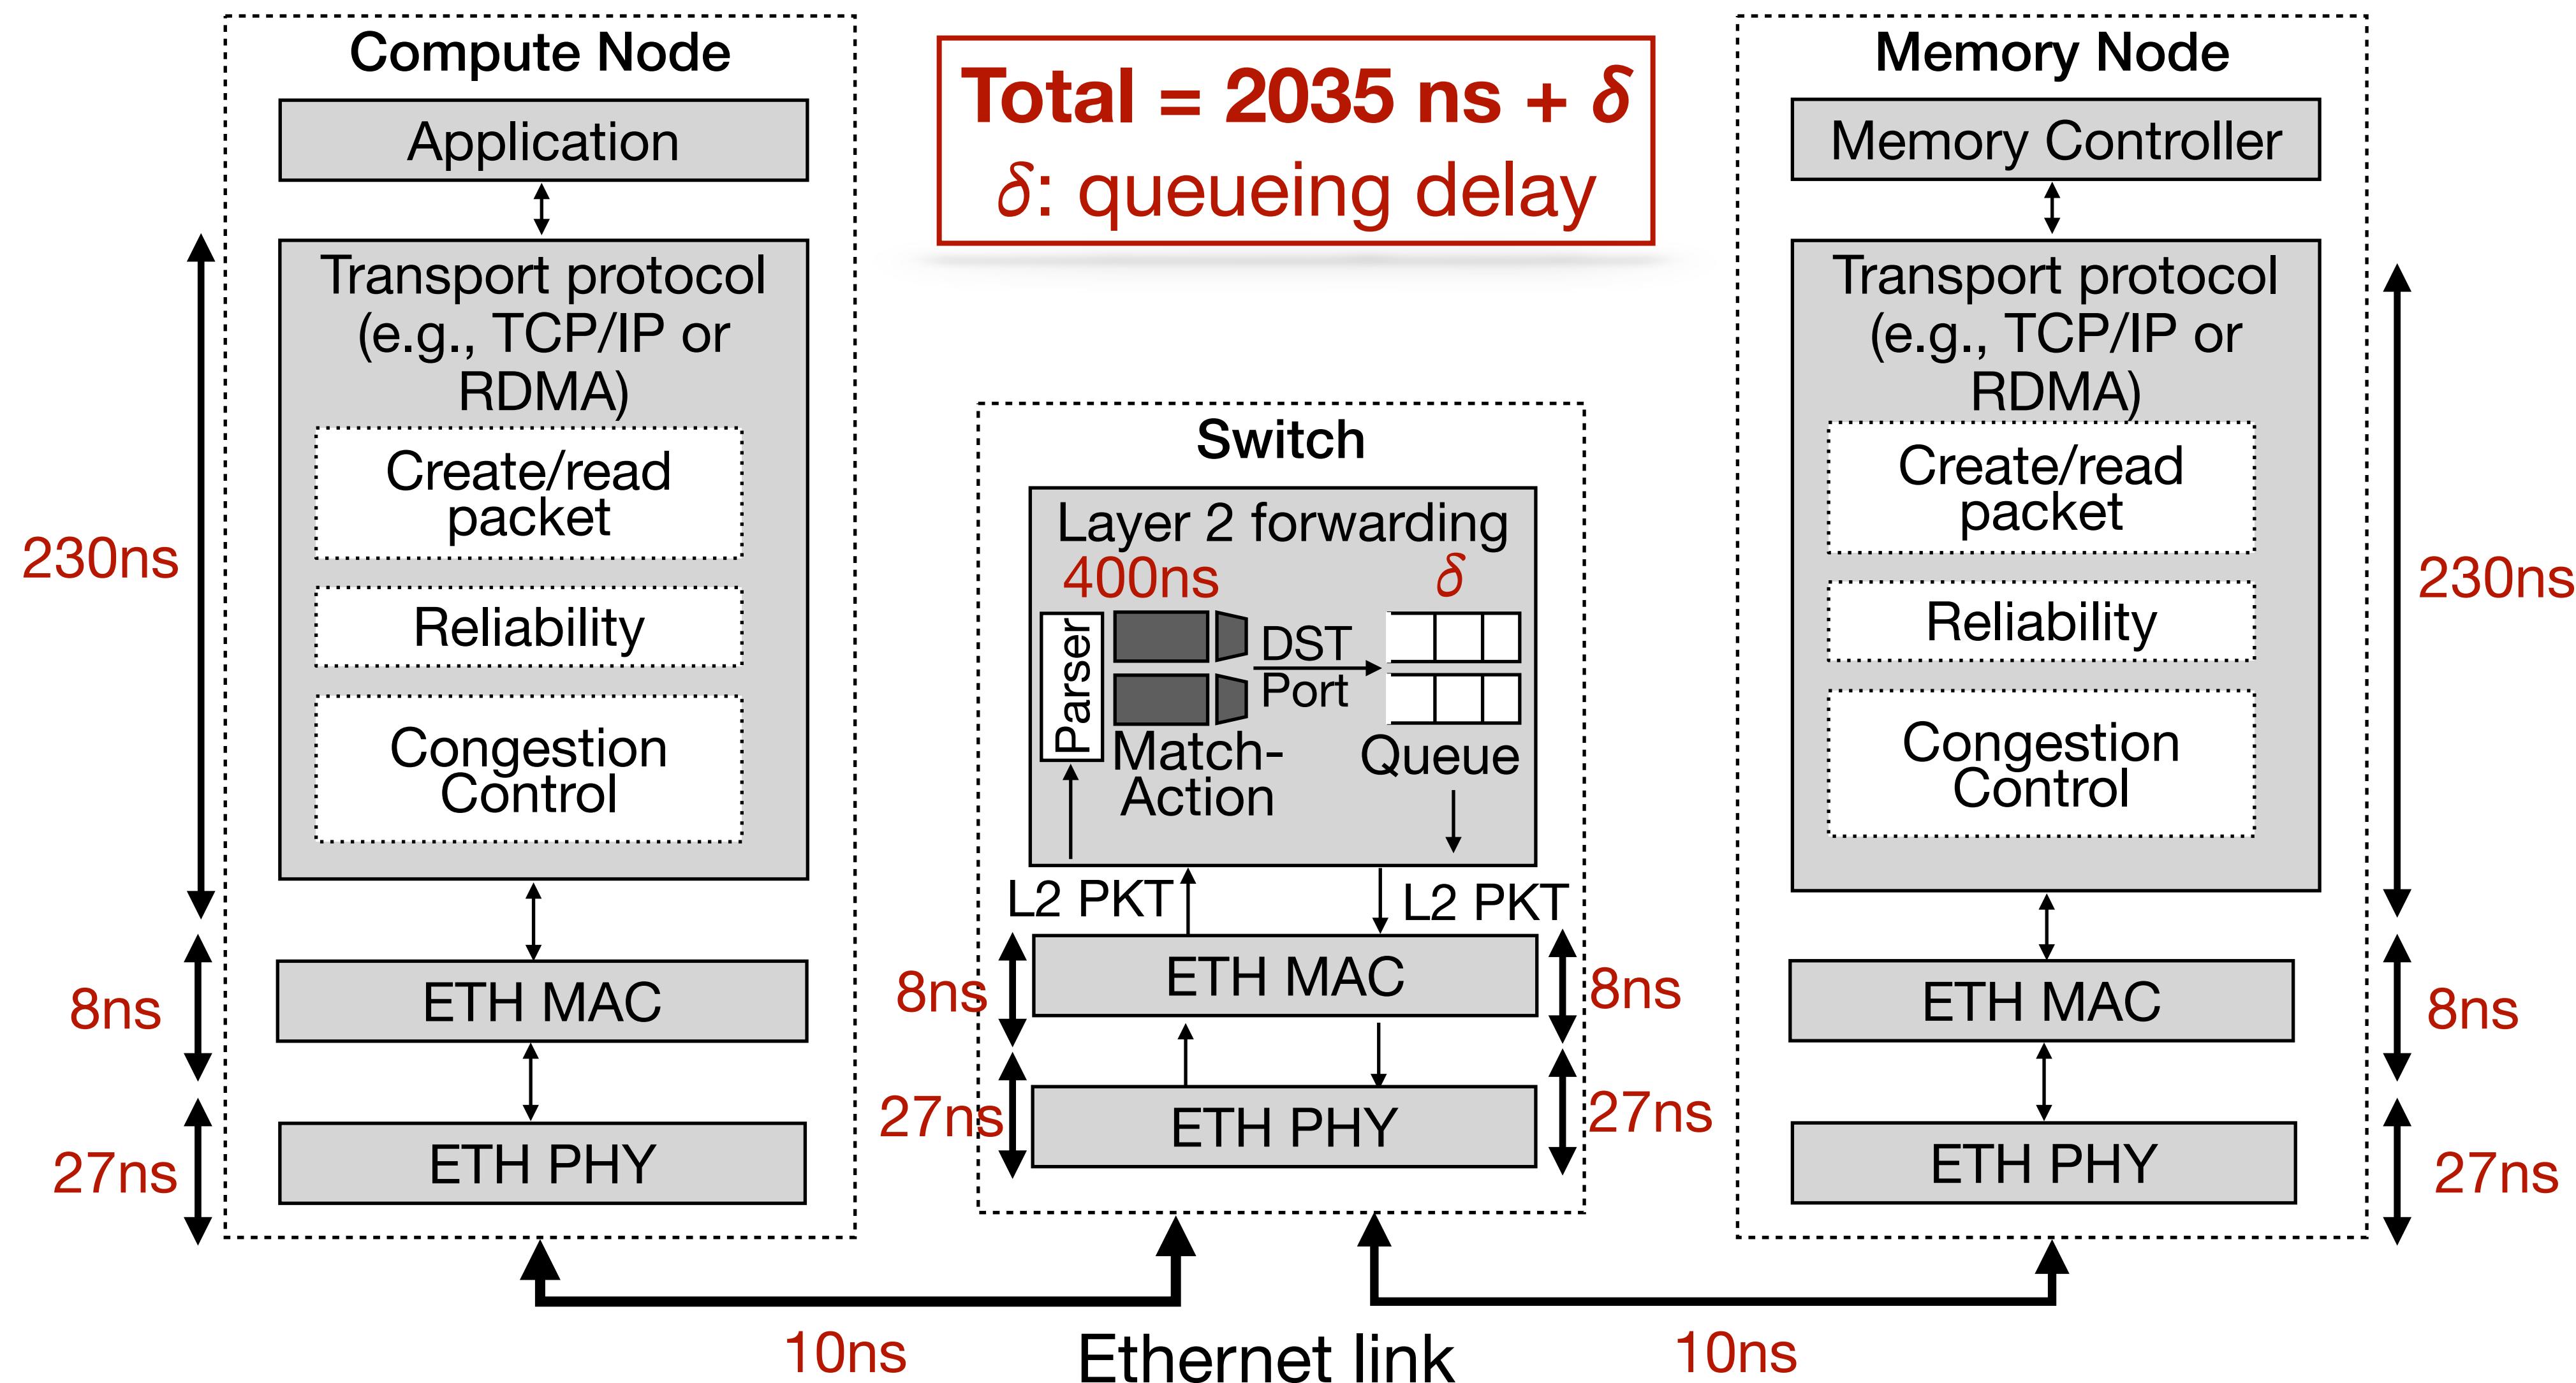
Latency Overheads of Existing Memory Disaggregation over Ethernet

An example of remote read request over RDMA

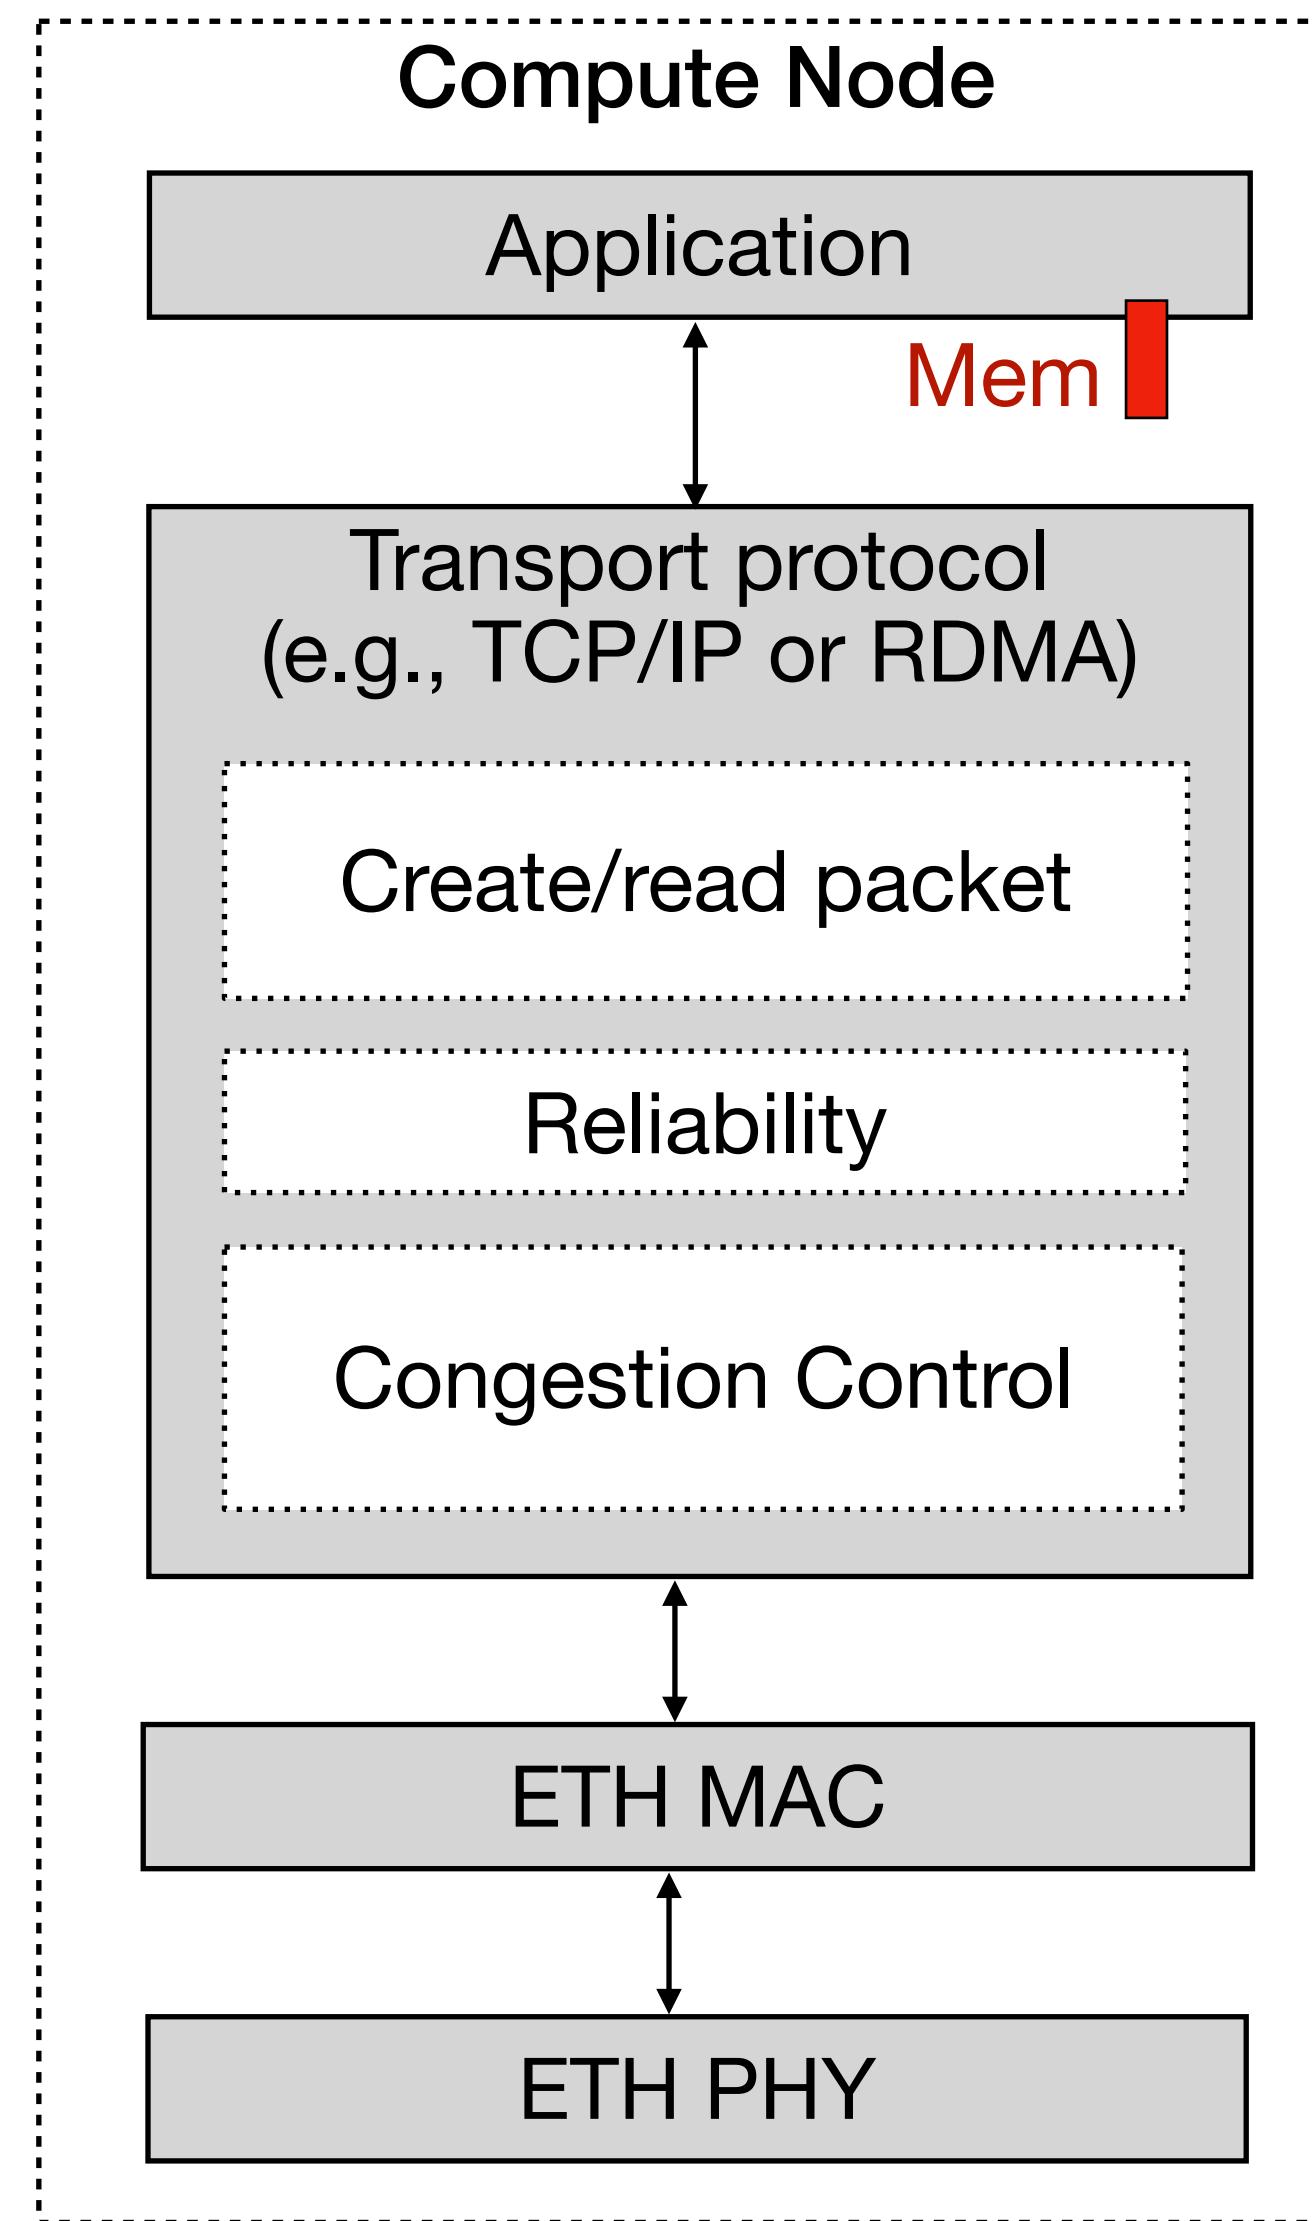


Latency Overheads of Existing Memory Disaggregation over Ethernet

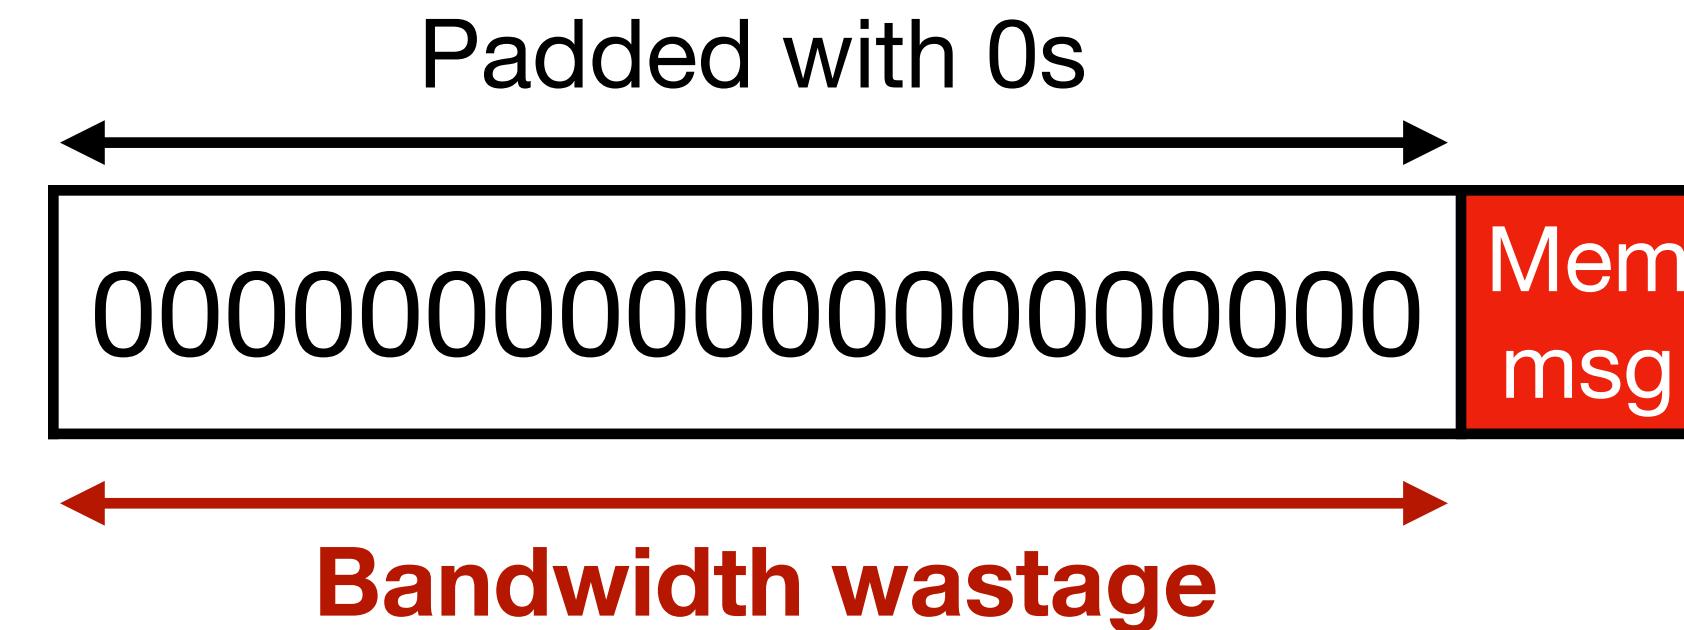
An example of remote read request over RDMA



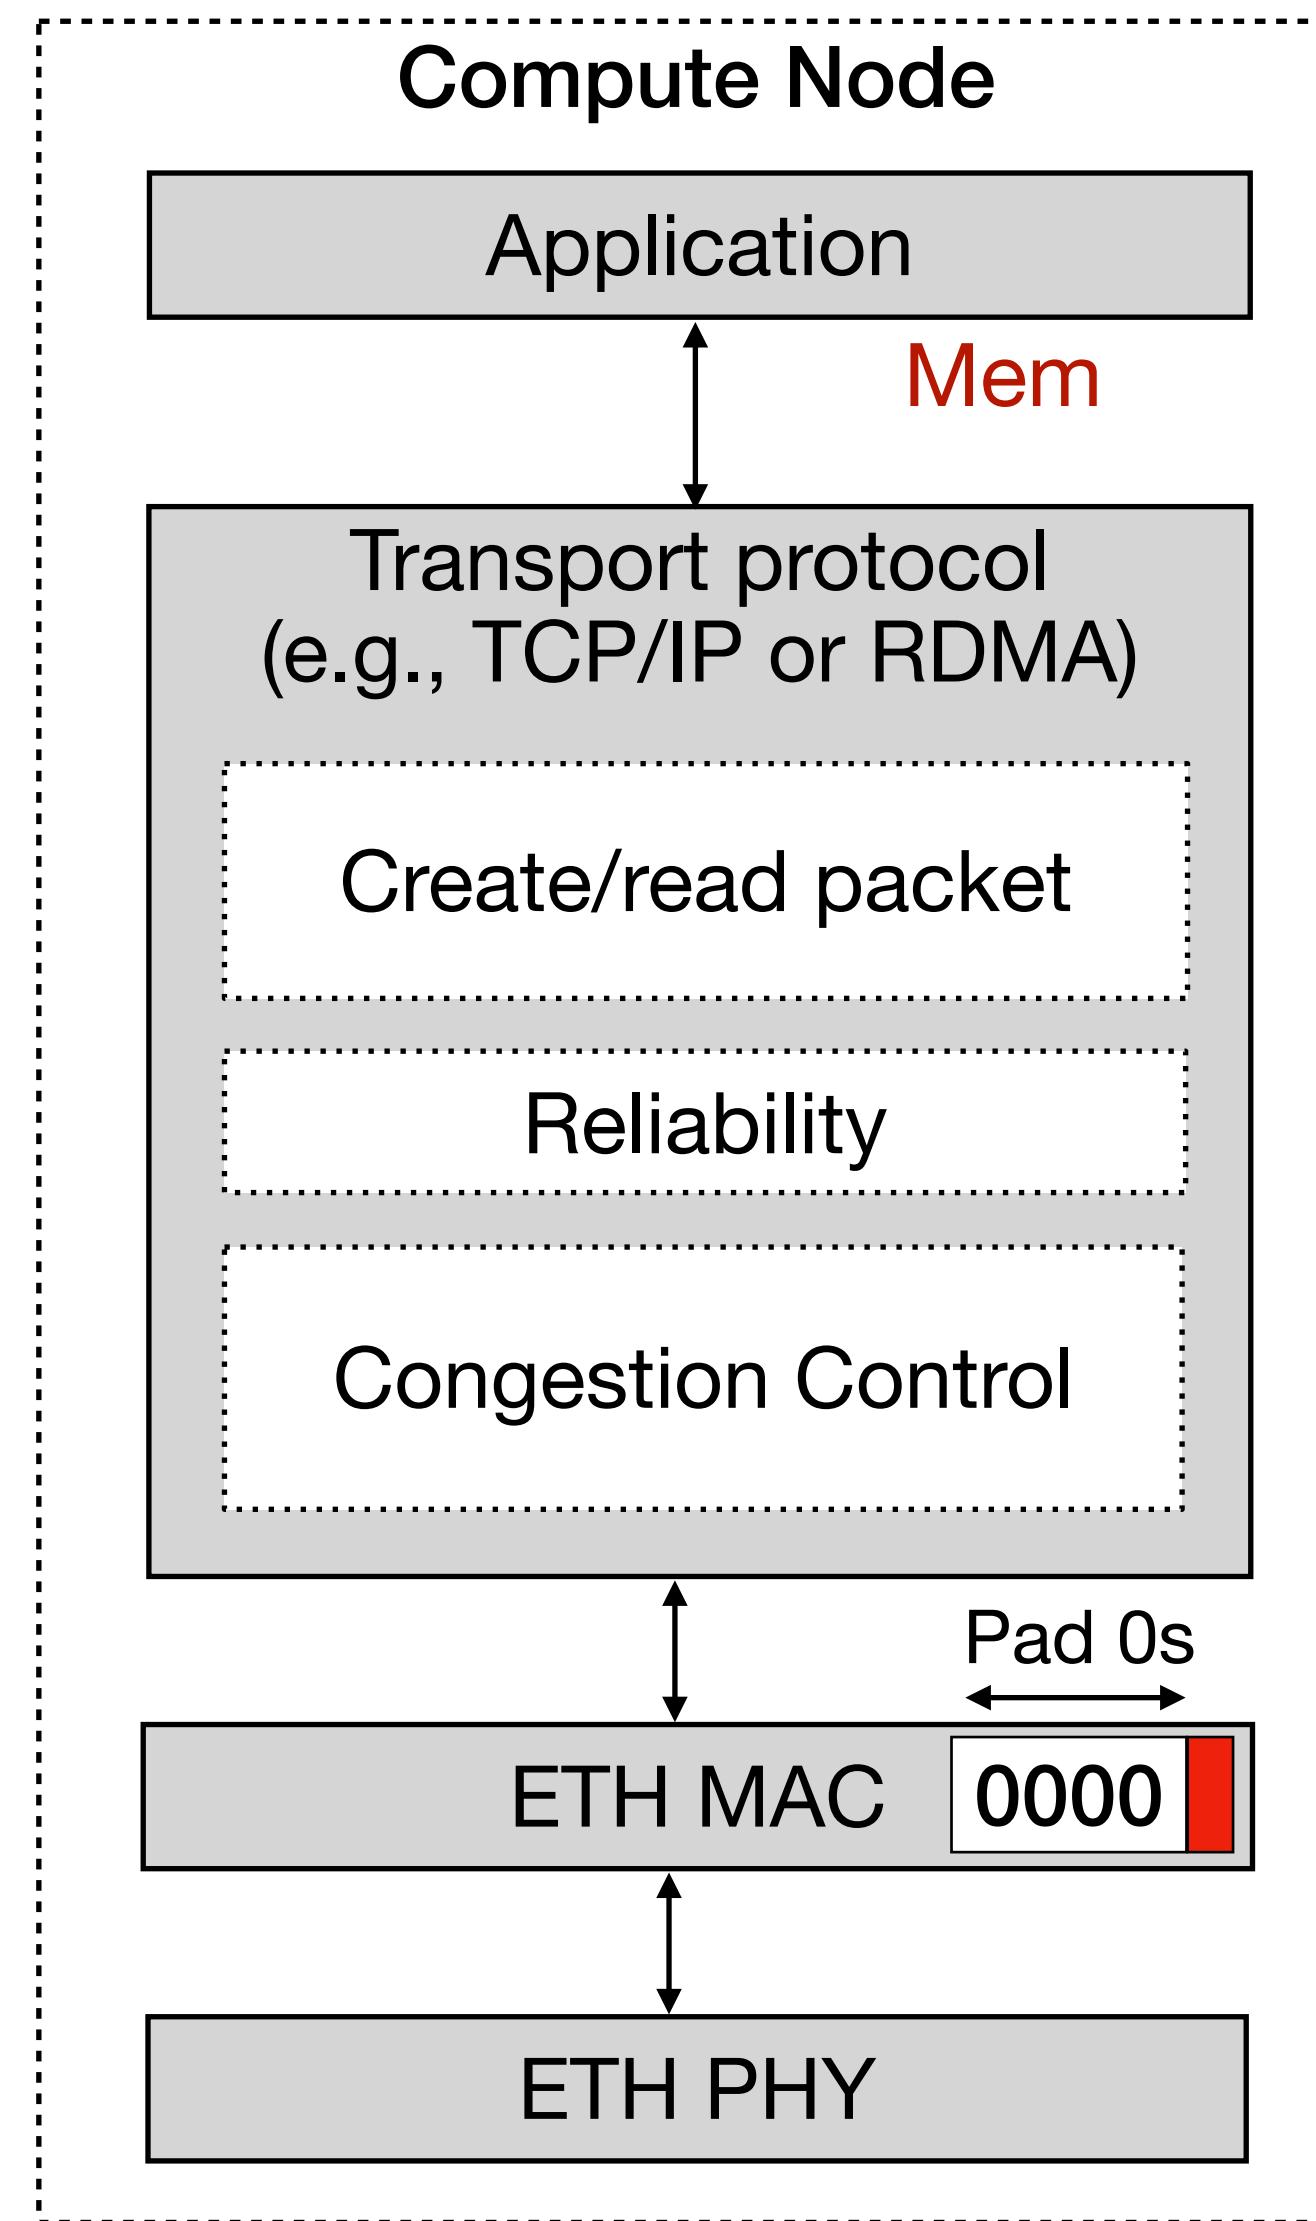
Other Overheads



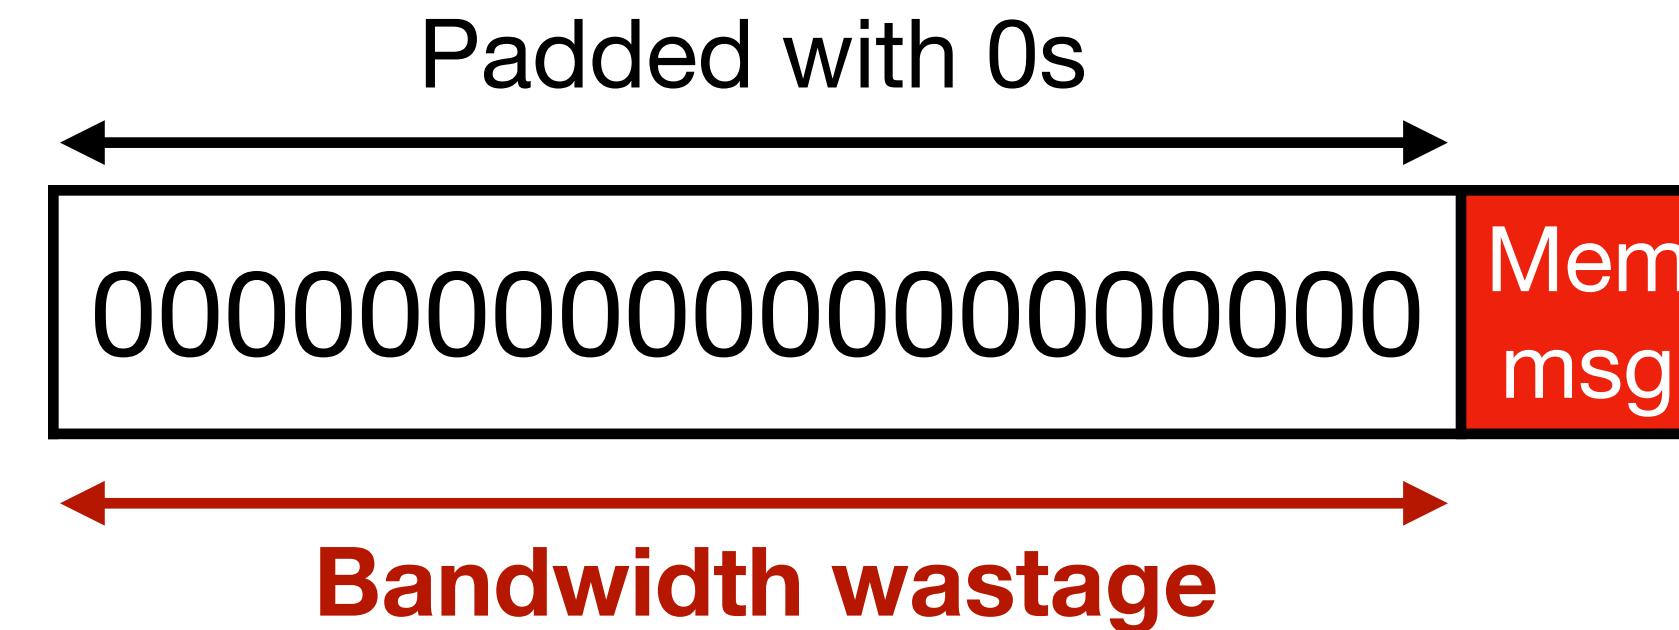
1. **Ethernet MAC enforces minimum 64B frame**
... but memory messages can be much smaller
(e.g., read requests are typically 8-16B)



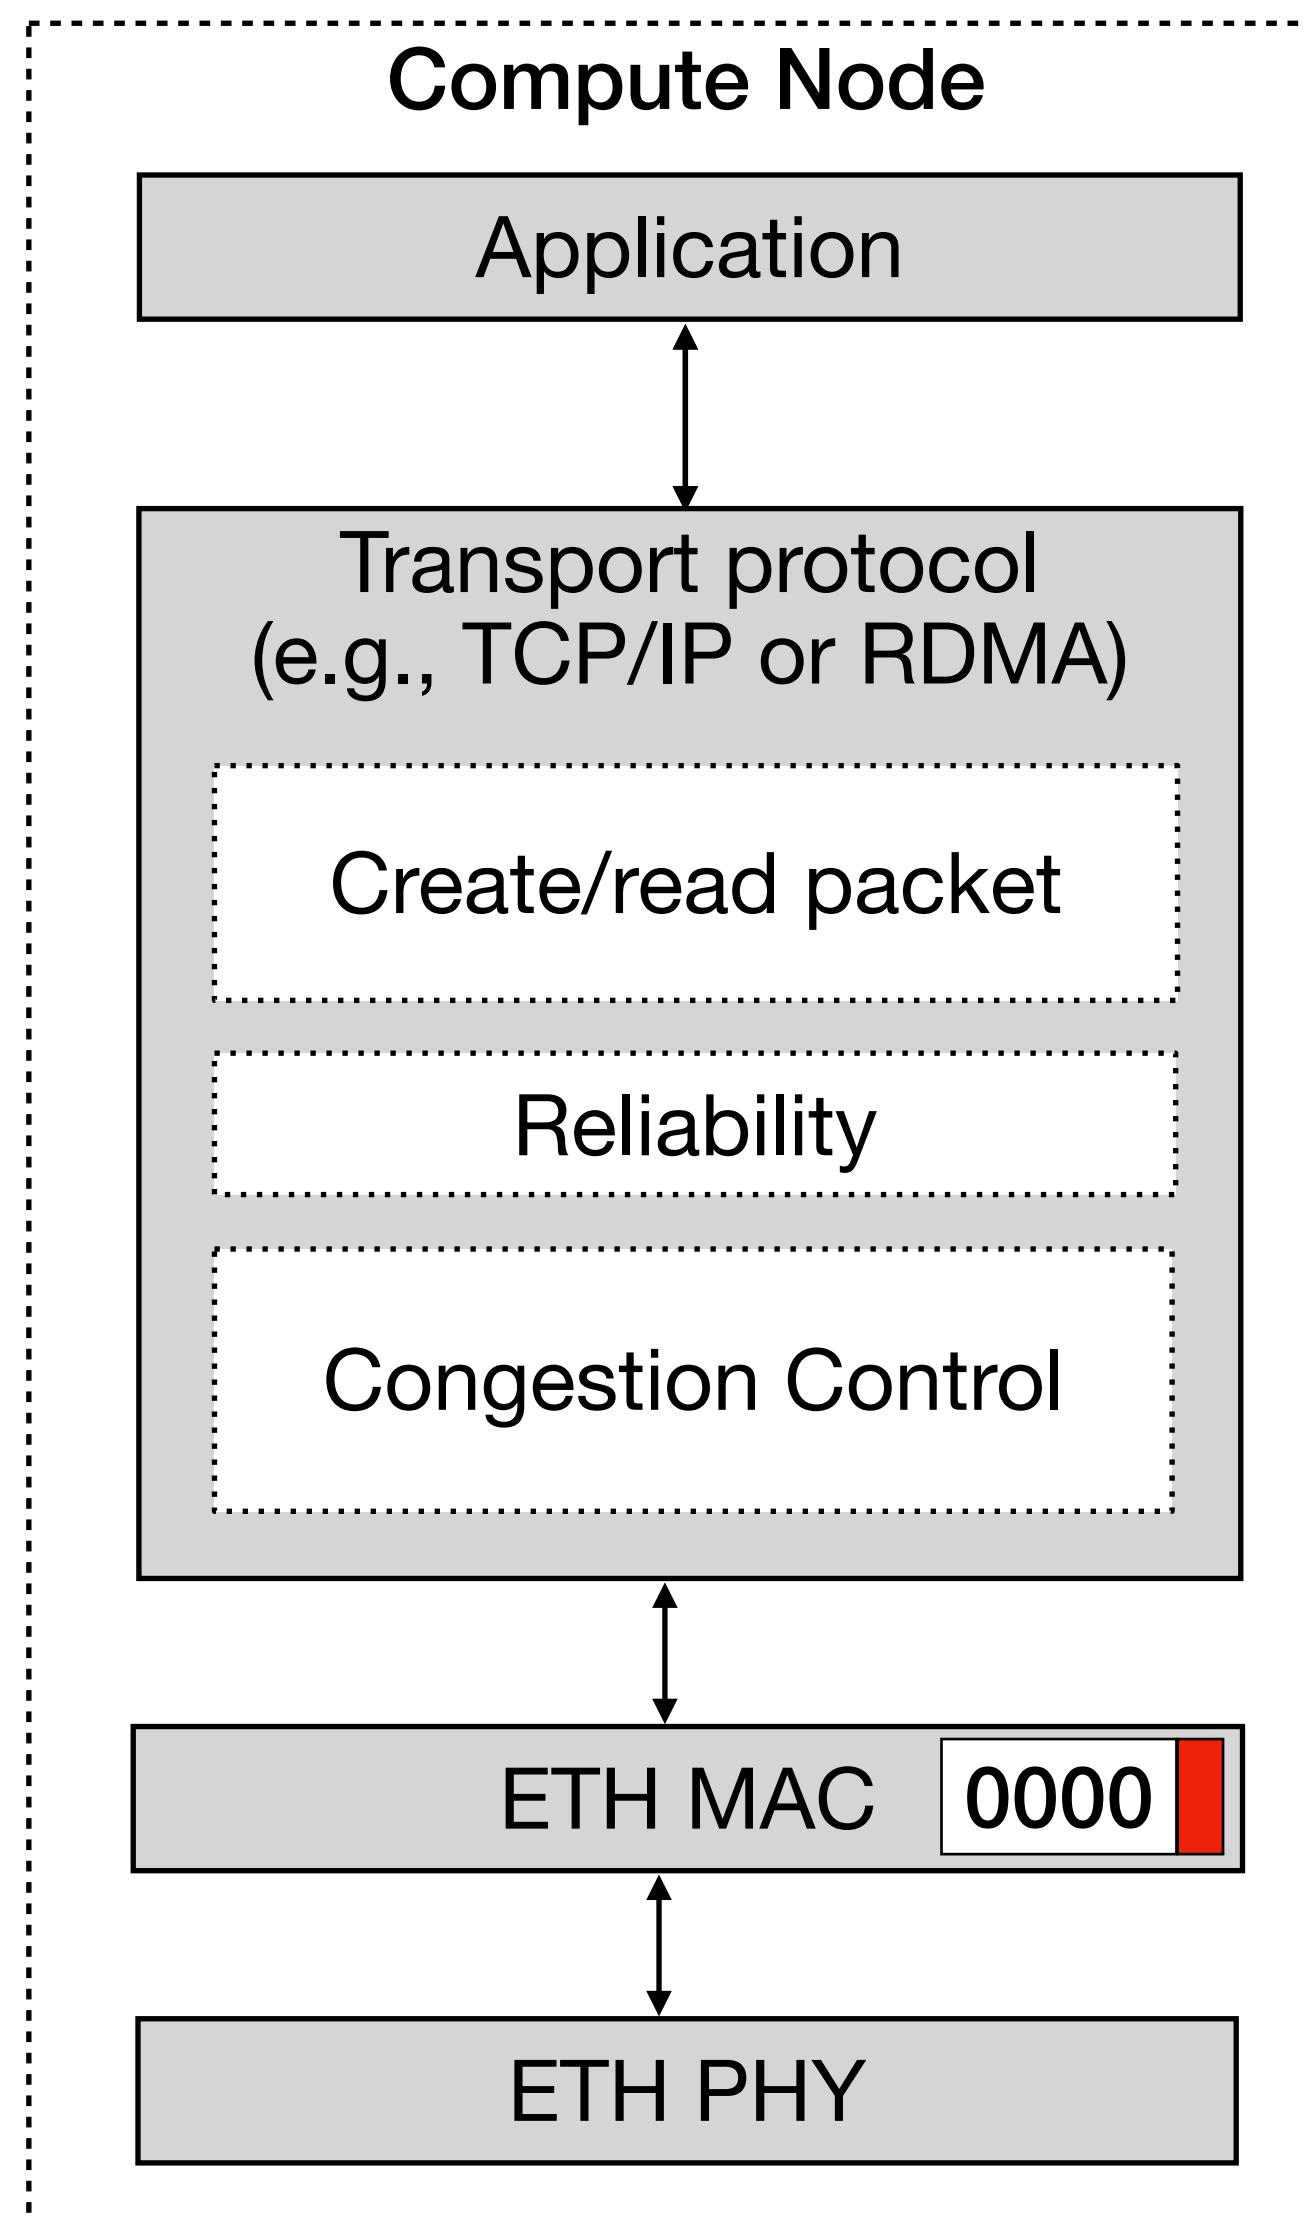
Other Overheads



1. **Ethernet MAC enforces minimum 64B frame**
... but memory messages can be much smaller
(e.g., read requests are typically 8-16B)



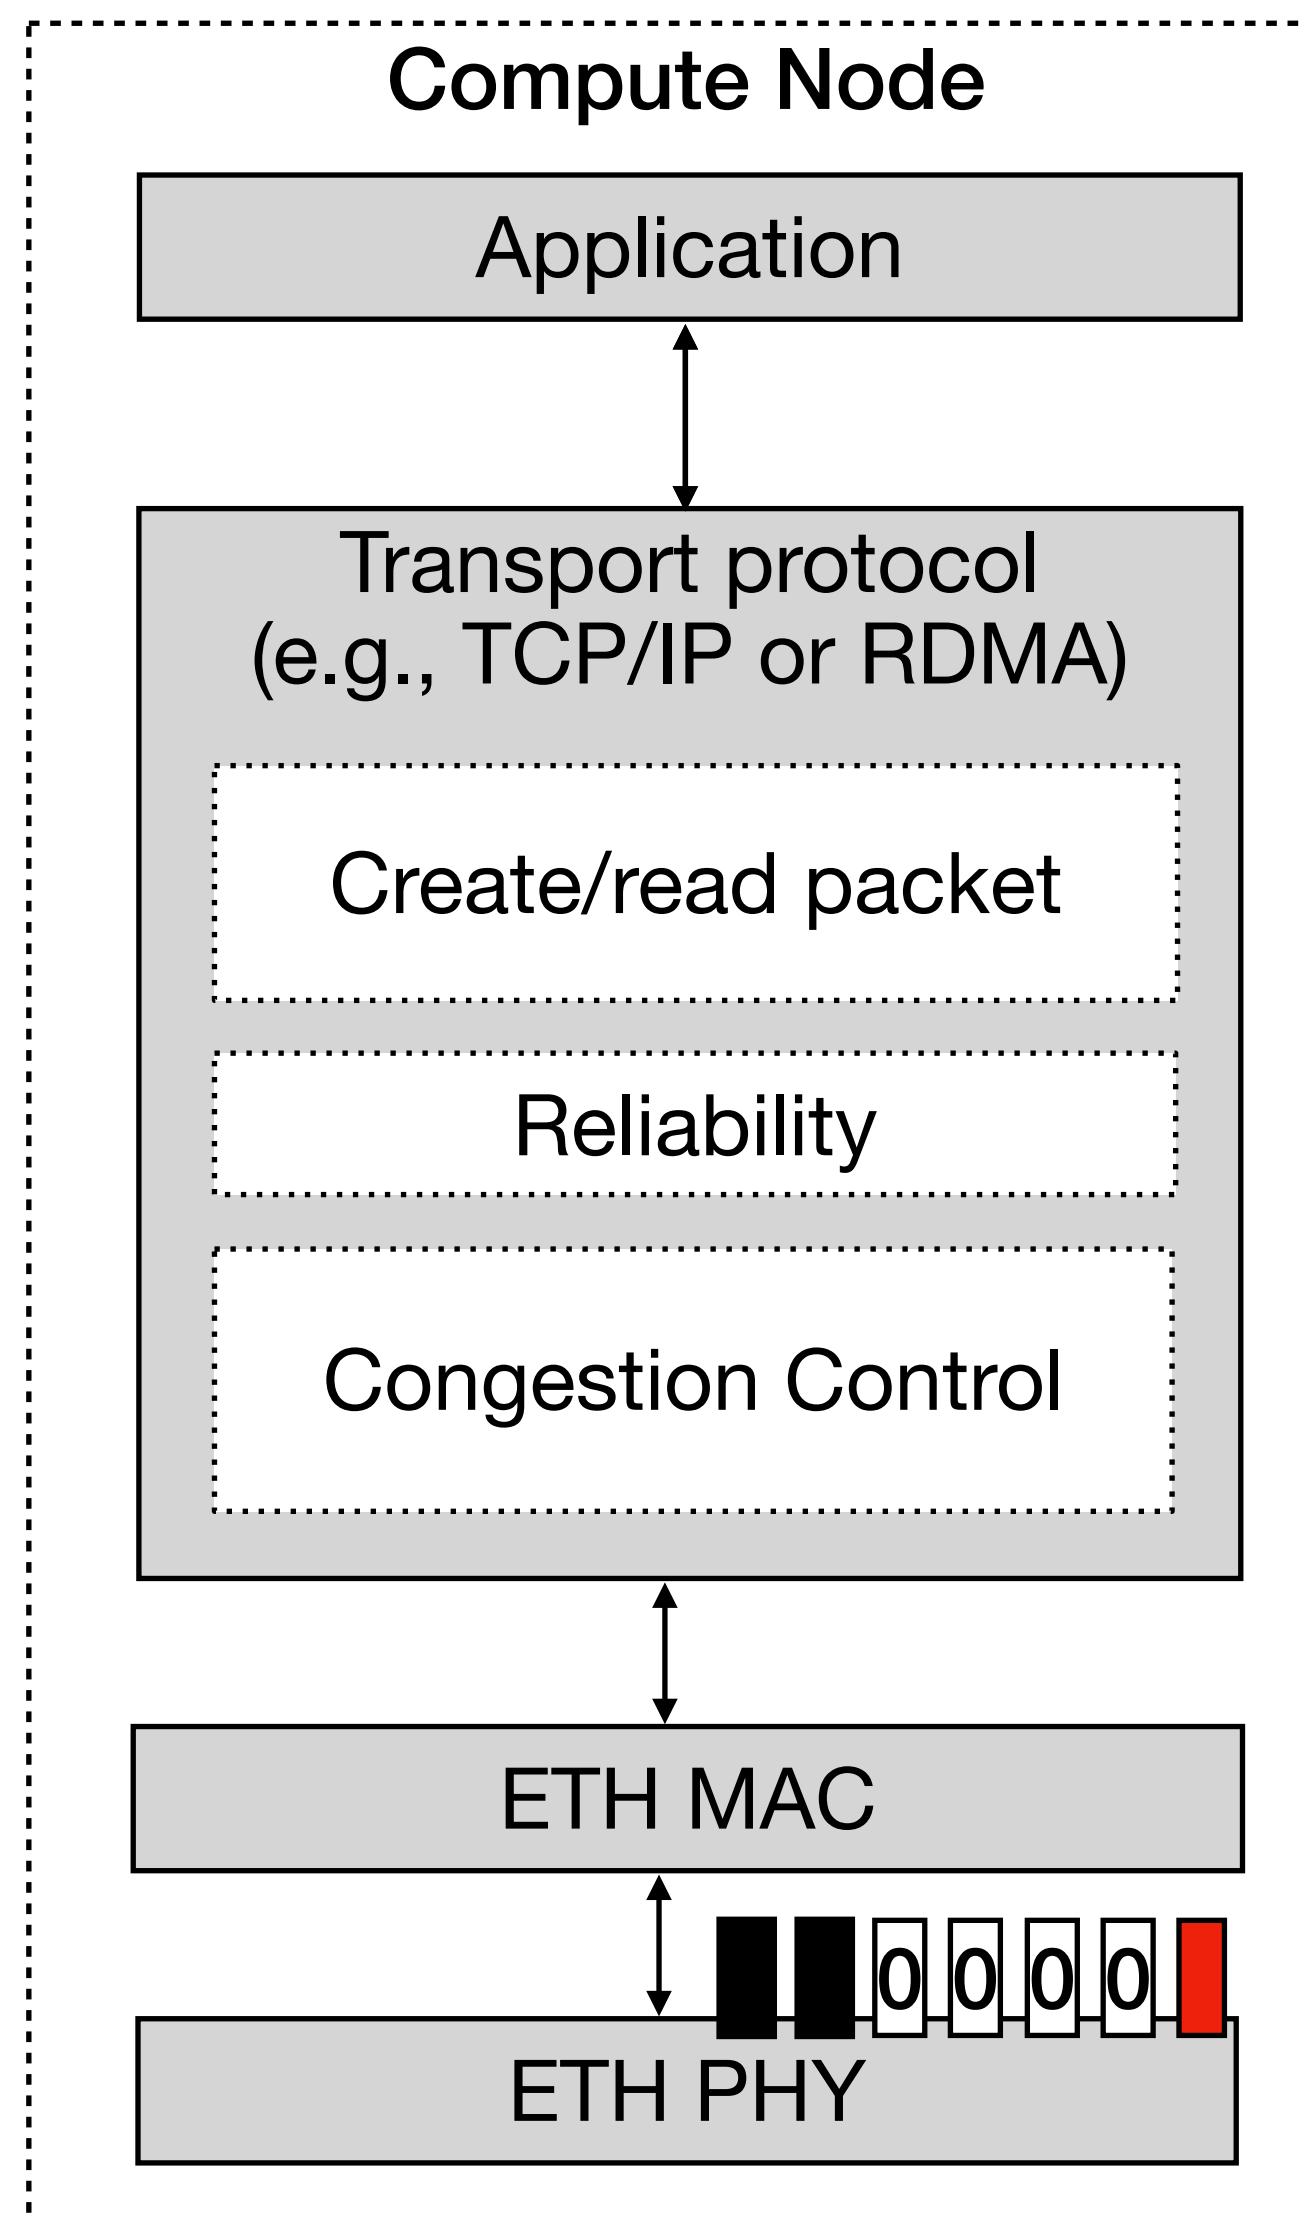
Other Overheads



1. **Ethernet MAC enforces minimum 64B frame**
... but memory messages can be much smaller
(e.g., read requests are typically 8-16B)
2. **Ethernet MAC enforces minimum of 12 bytes Inter-frame gap (IFG)**
... high overhead for small memory messages



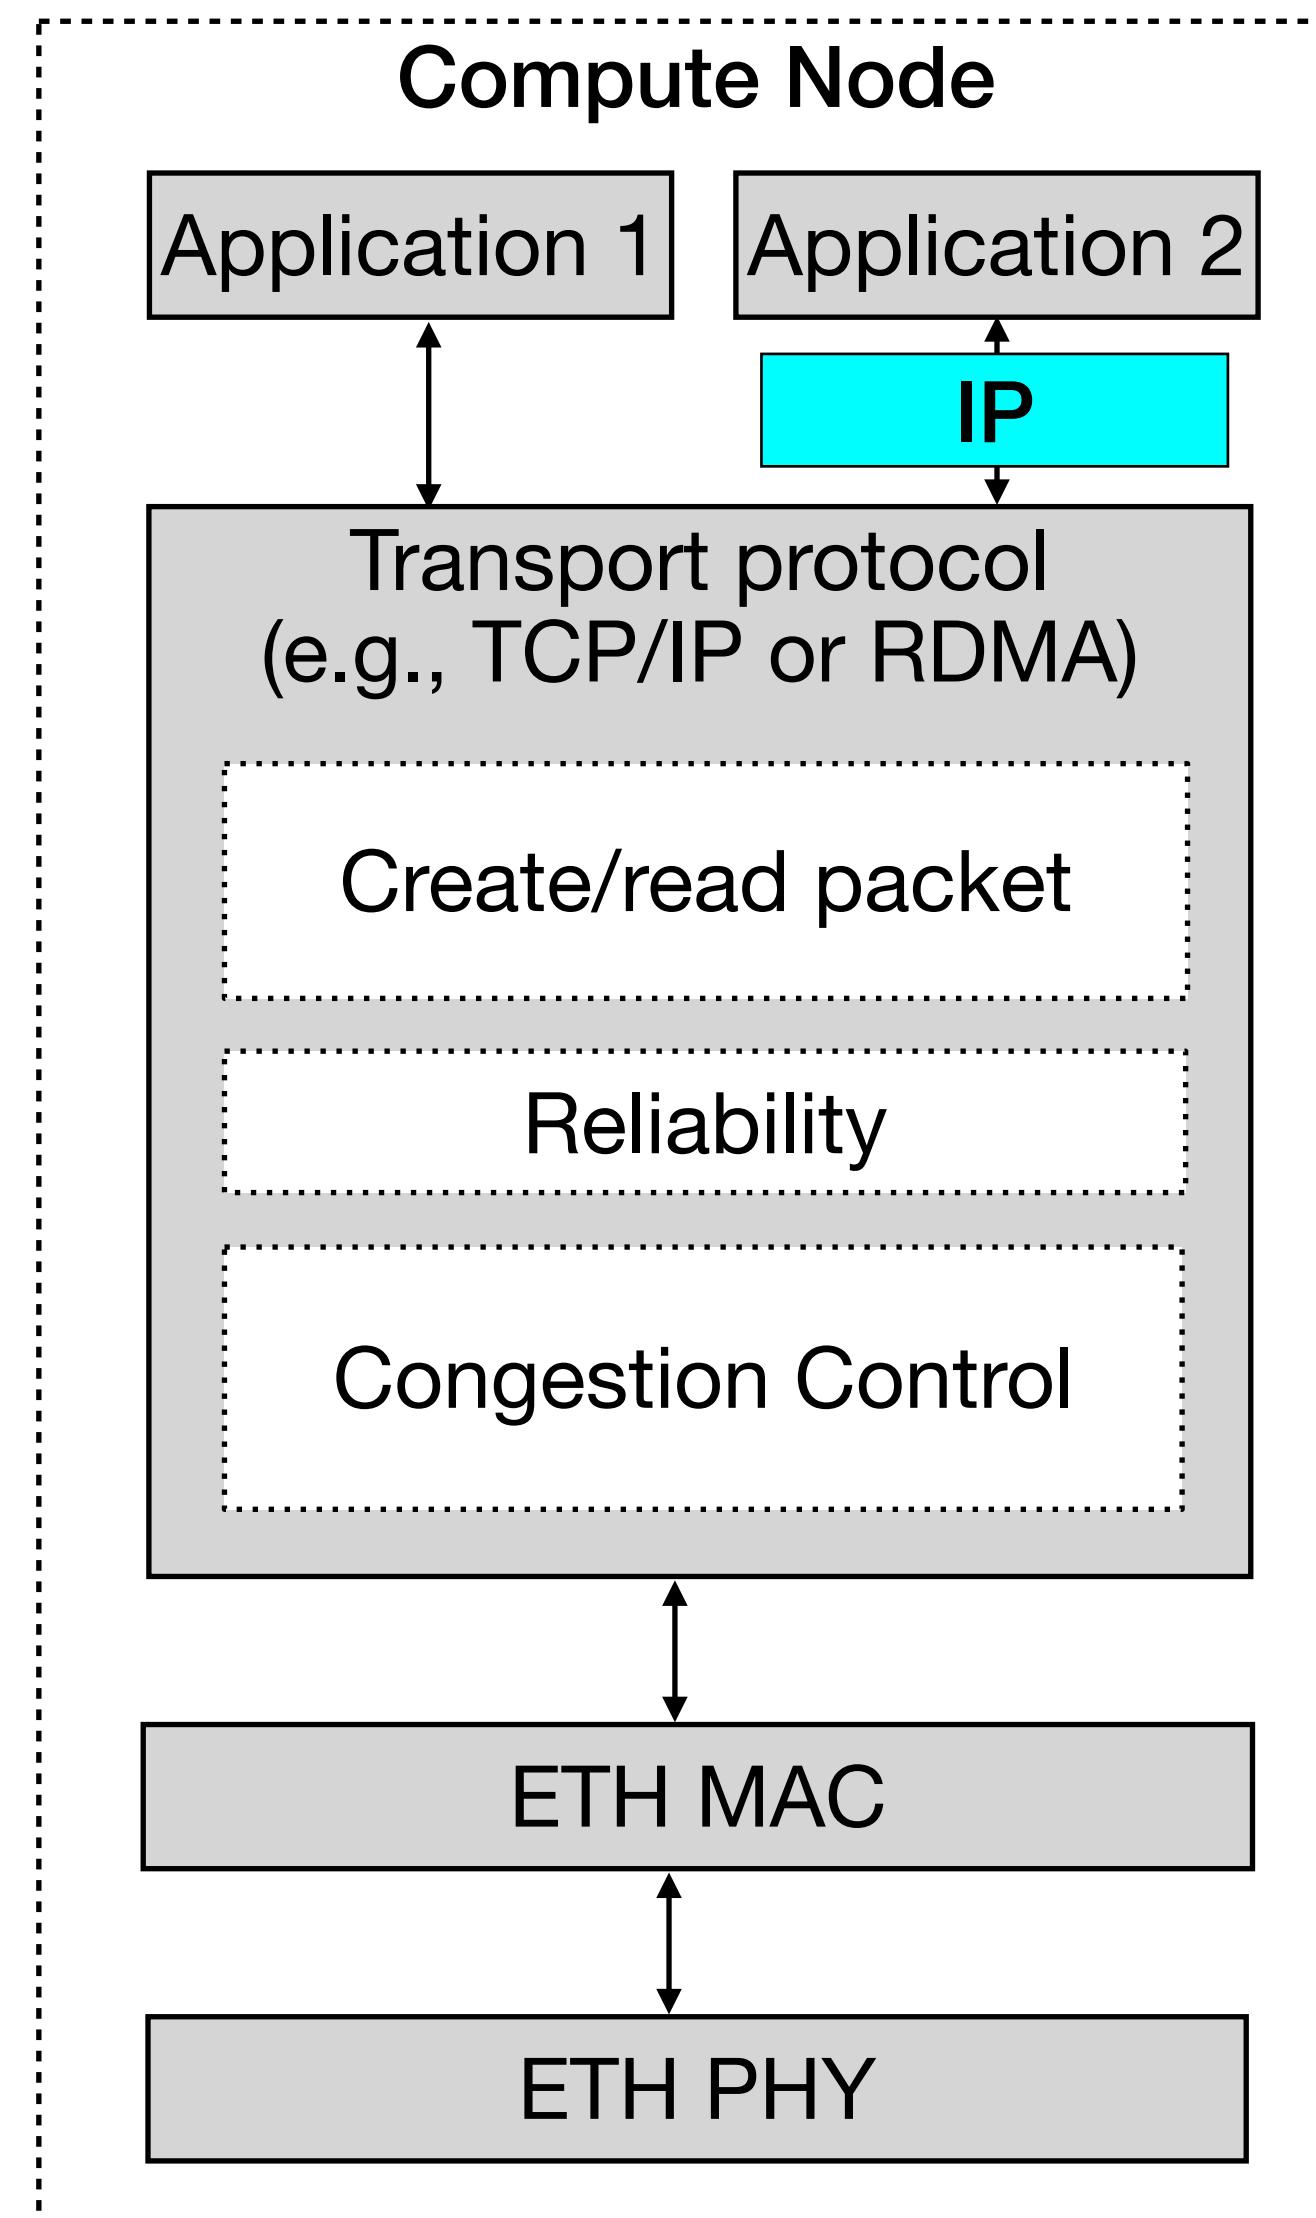
Other Overheads



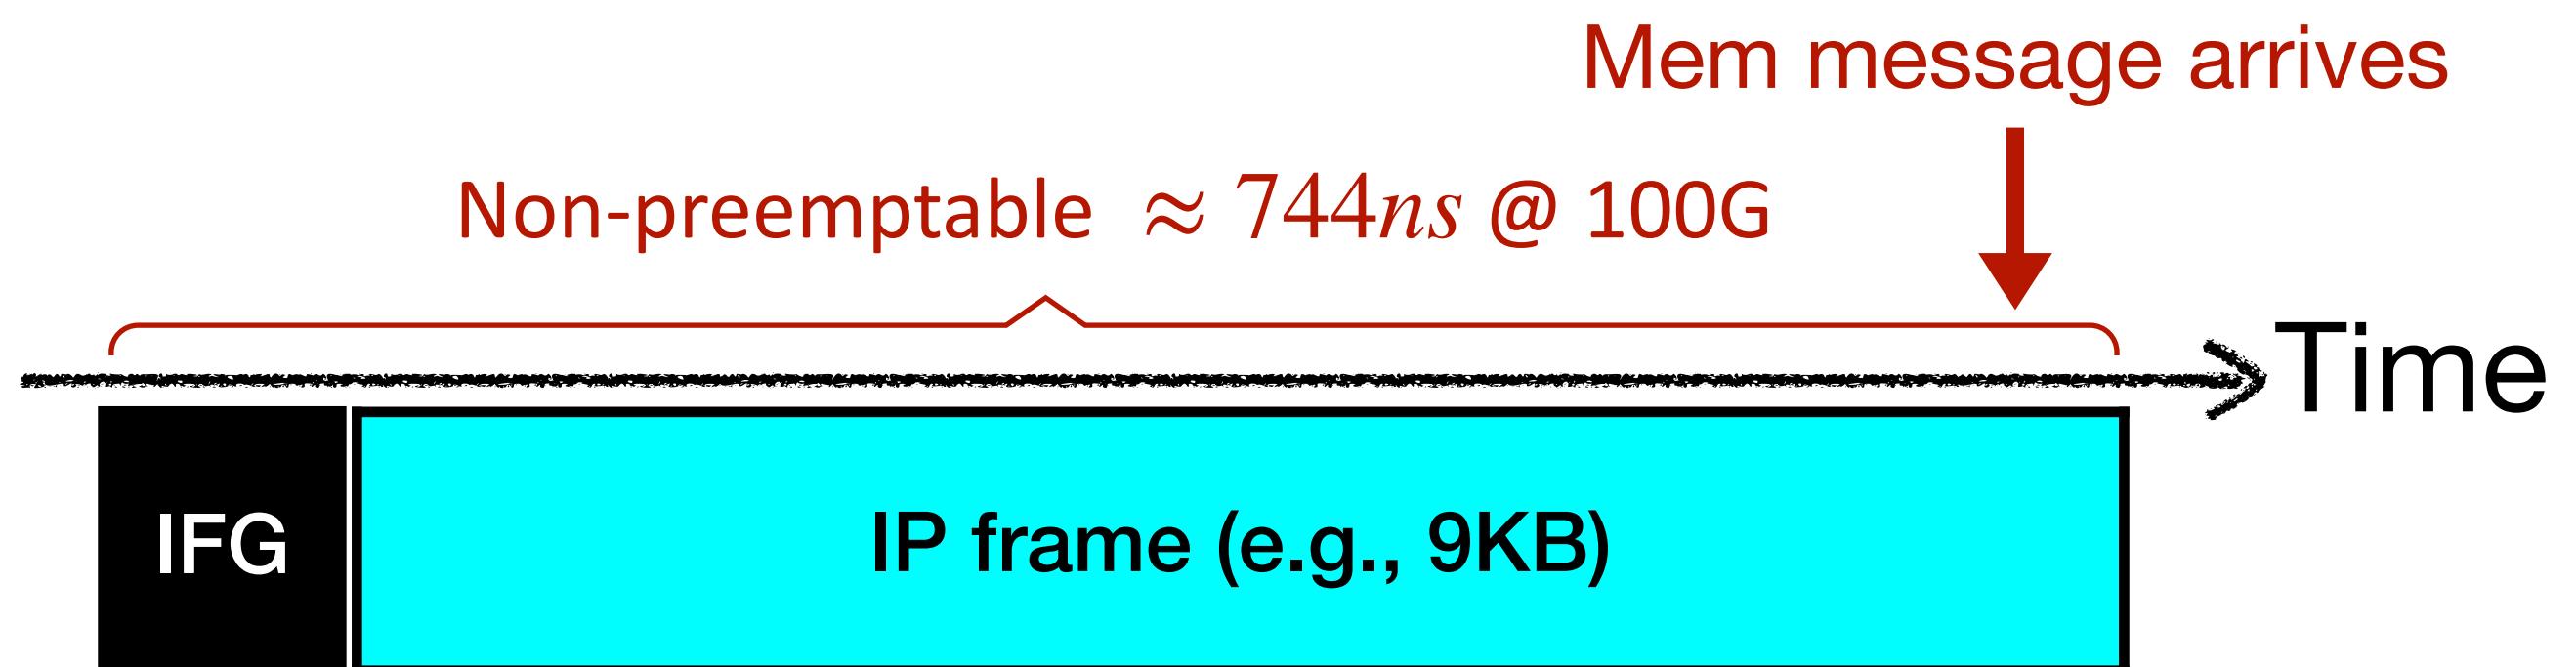
1. **Ethernet MAC enforces minimum 64B frame**
... but memory messages can be much smaller
(e.g., read requests are typically 8-16B)
2. **Ethernet MAC enforces minimum of 12 bytes Inter-frame gap (IFG)**
... high overhead for small memory messages



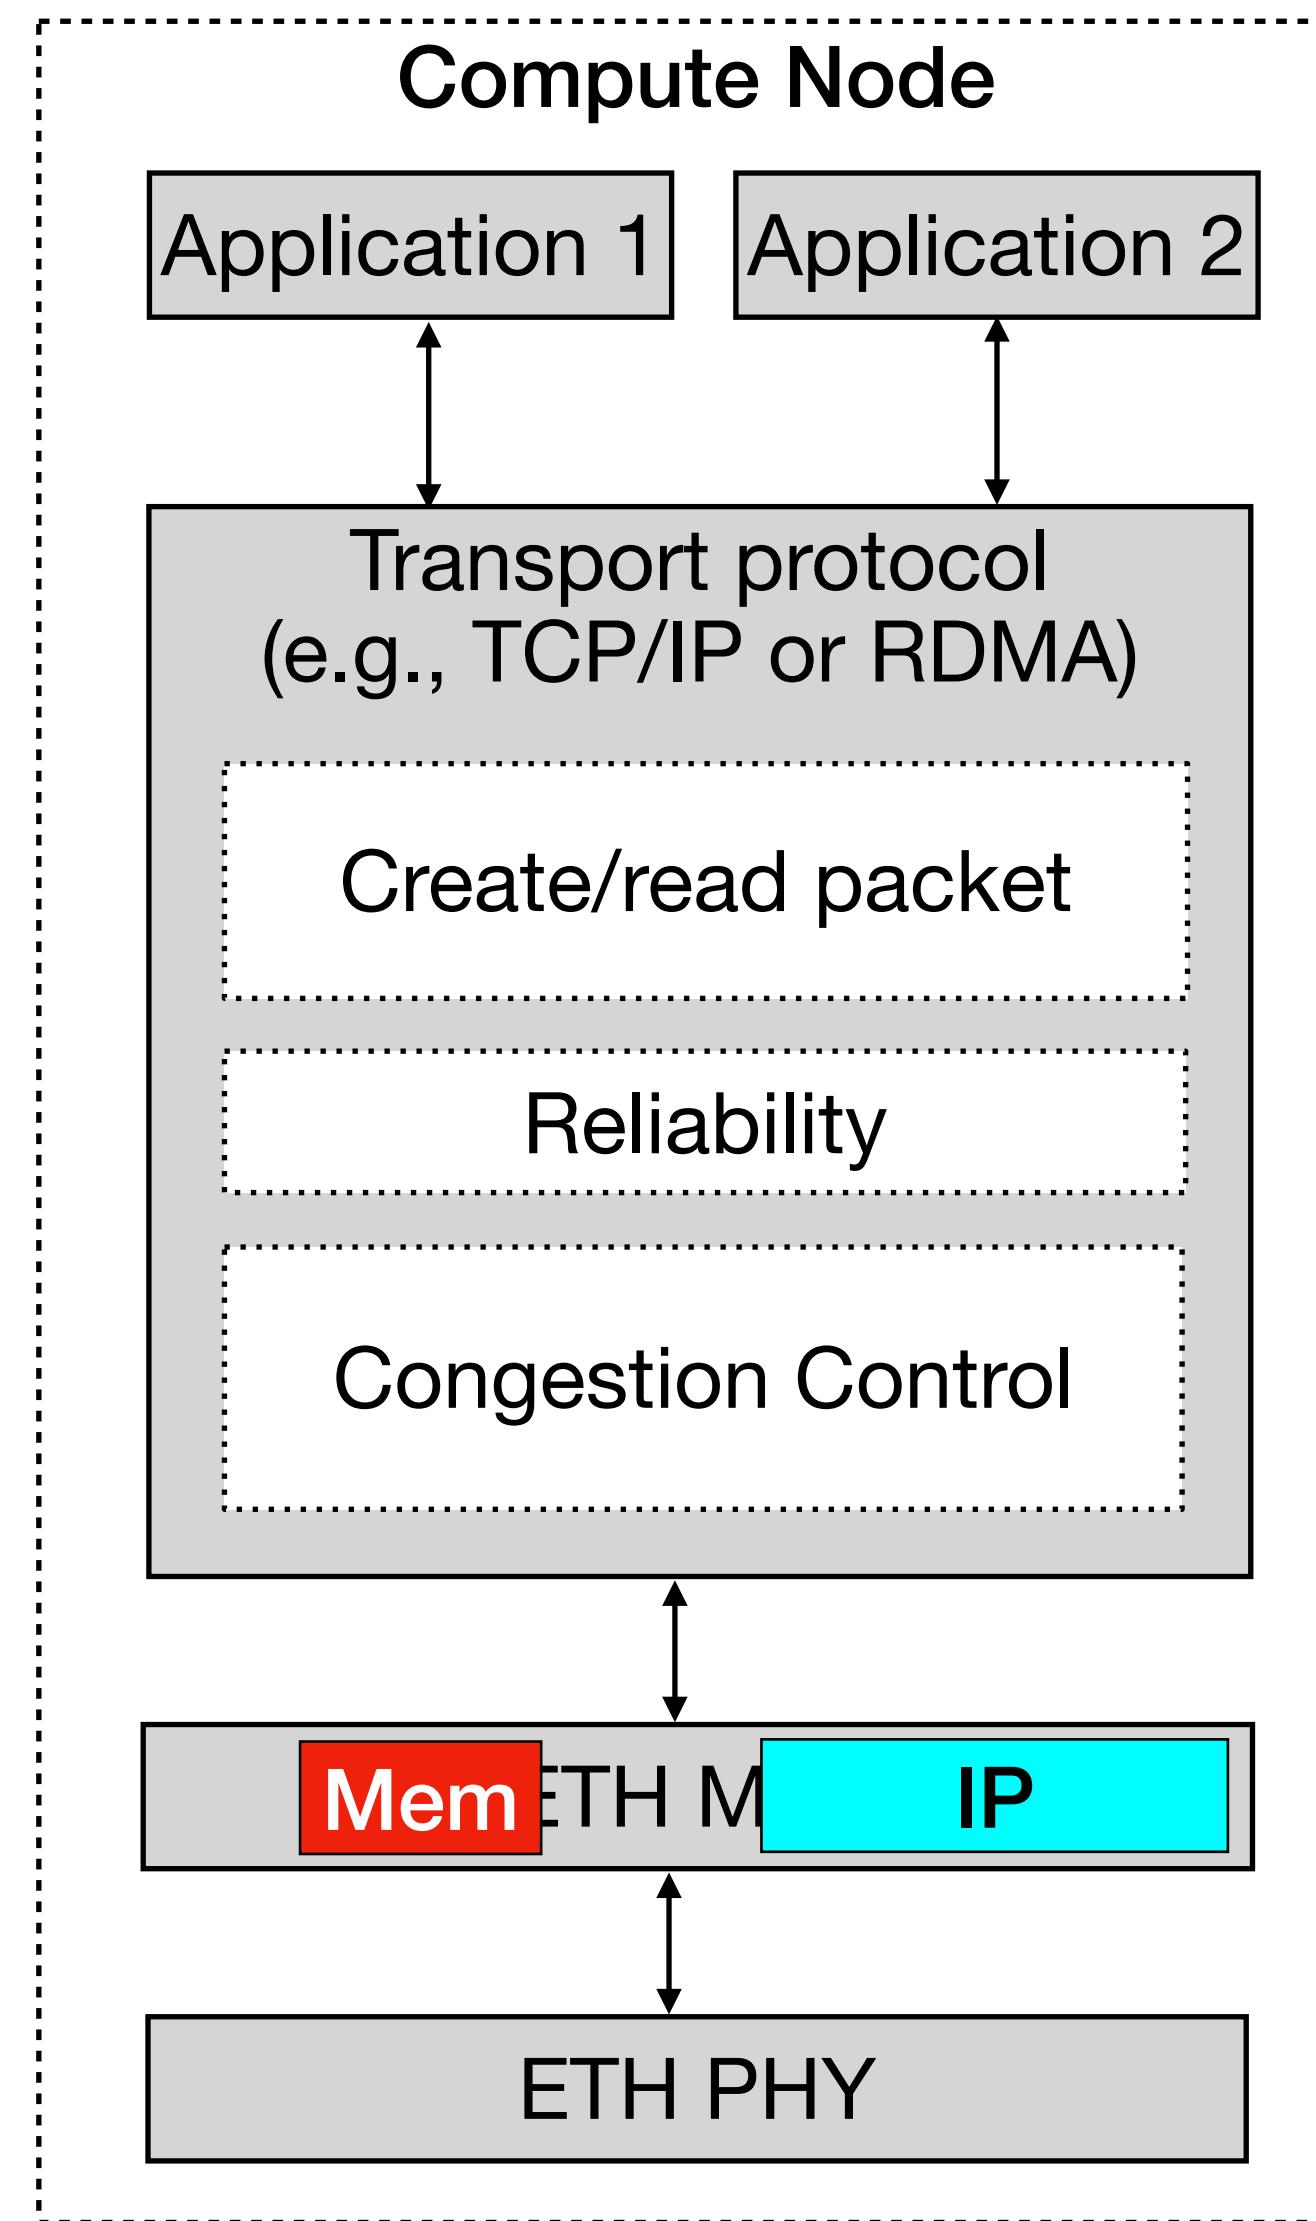
Other Overheads



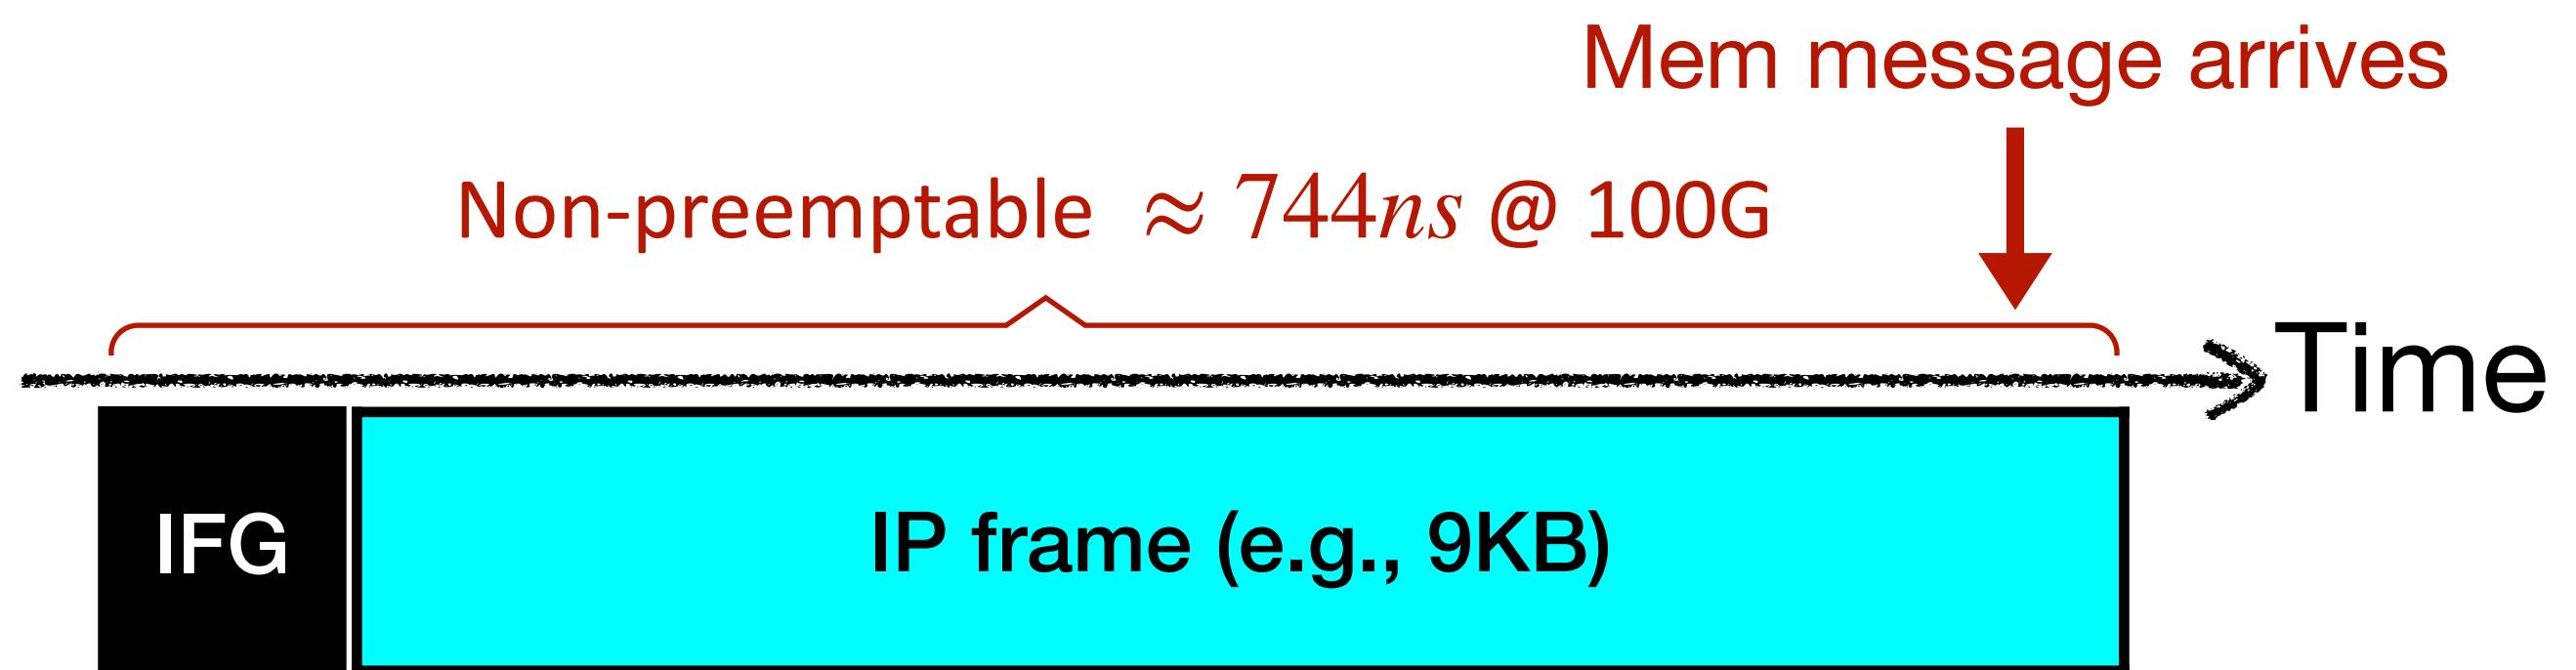
- Ethernet MAC enforces minimum 64B frame**
... but memory messages can be much smaller
(e.g., read requests are typically 8-16B)
- Ethernet MAC enforces minimum of 12 bytes Inter-frame gap (IFG)**
... high overhead for small memory messages
- Ethernet MAC does not allow intra-frame preemption**
... a large non-memory frame may block the transmission of a small memory message



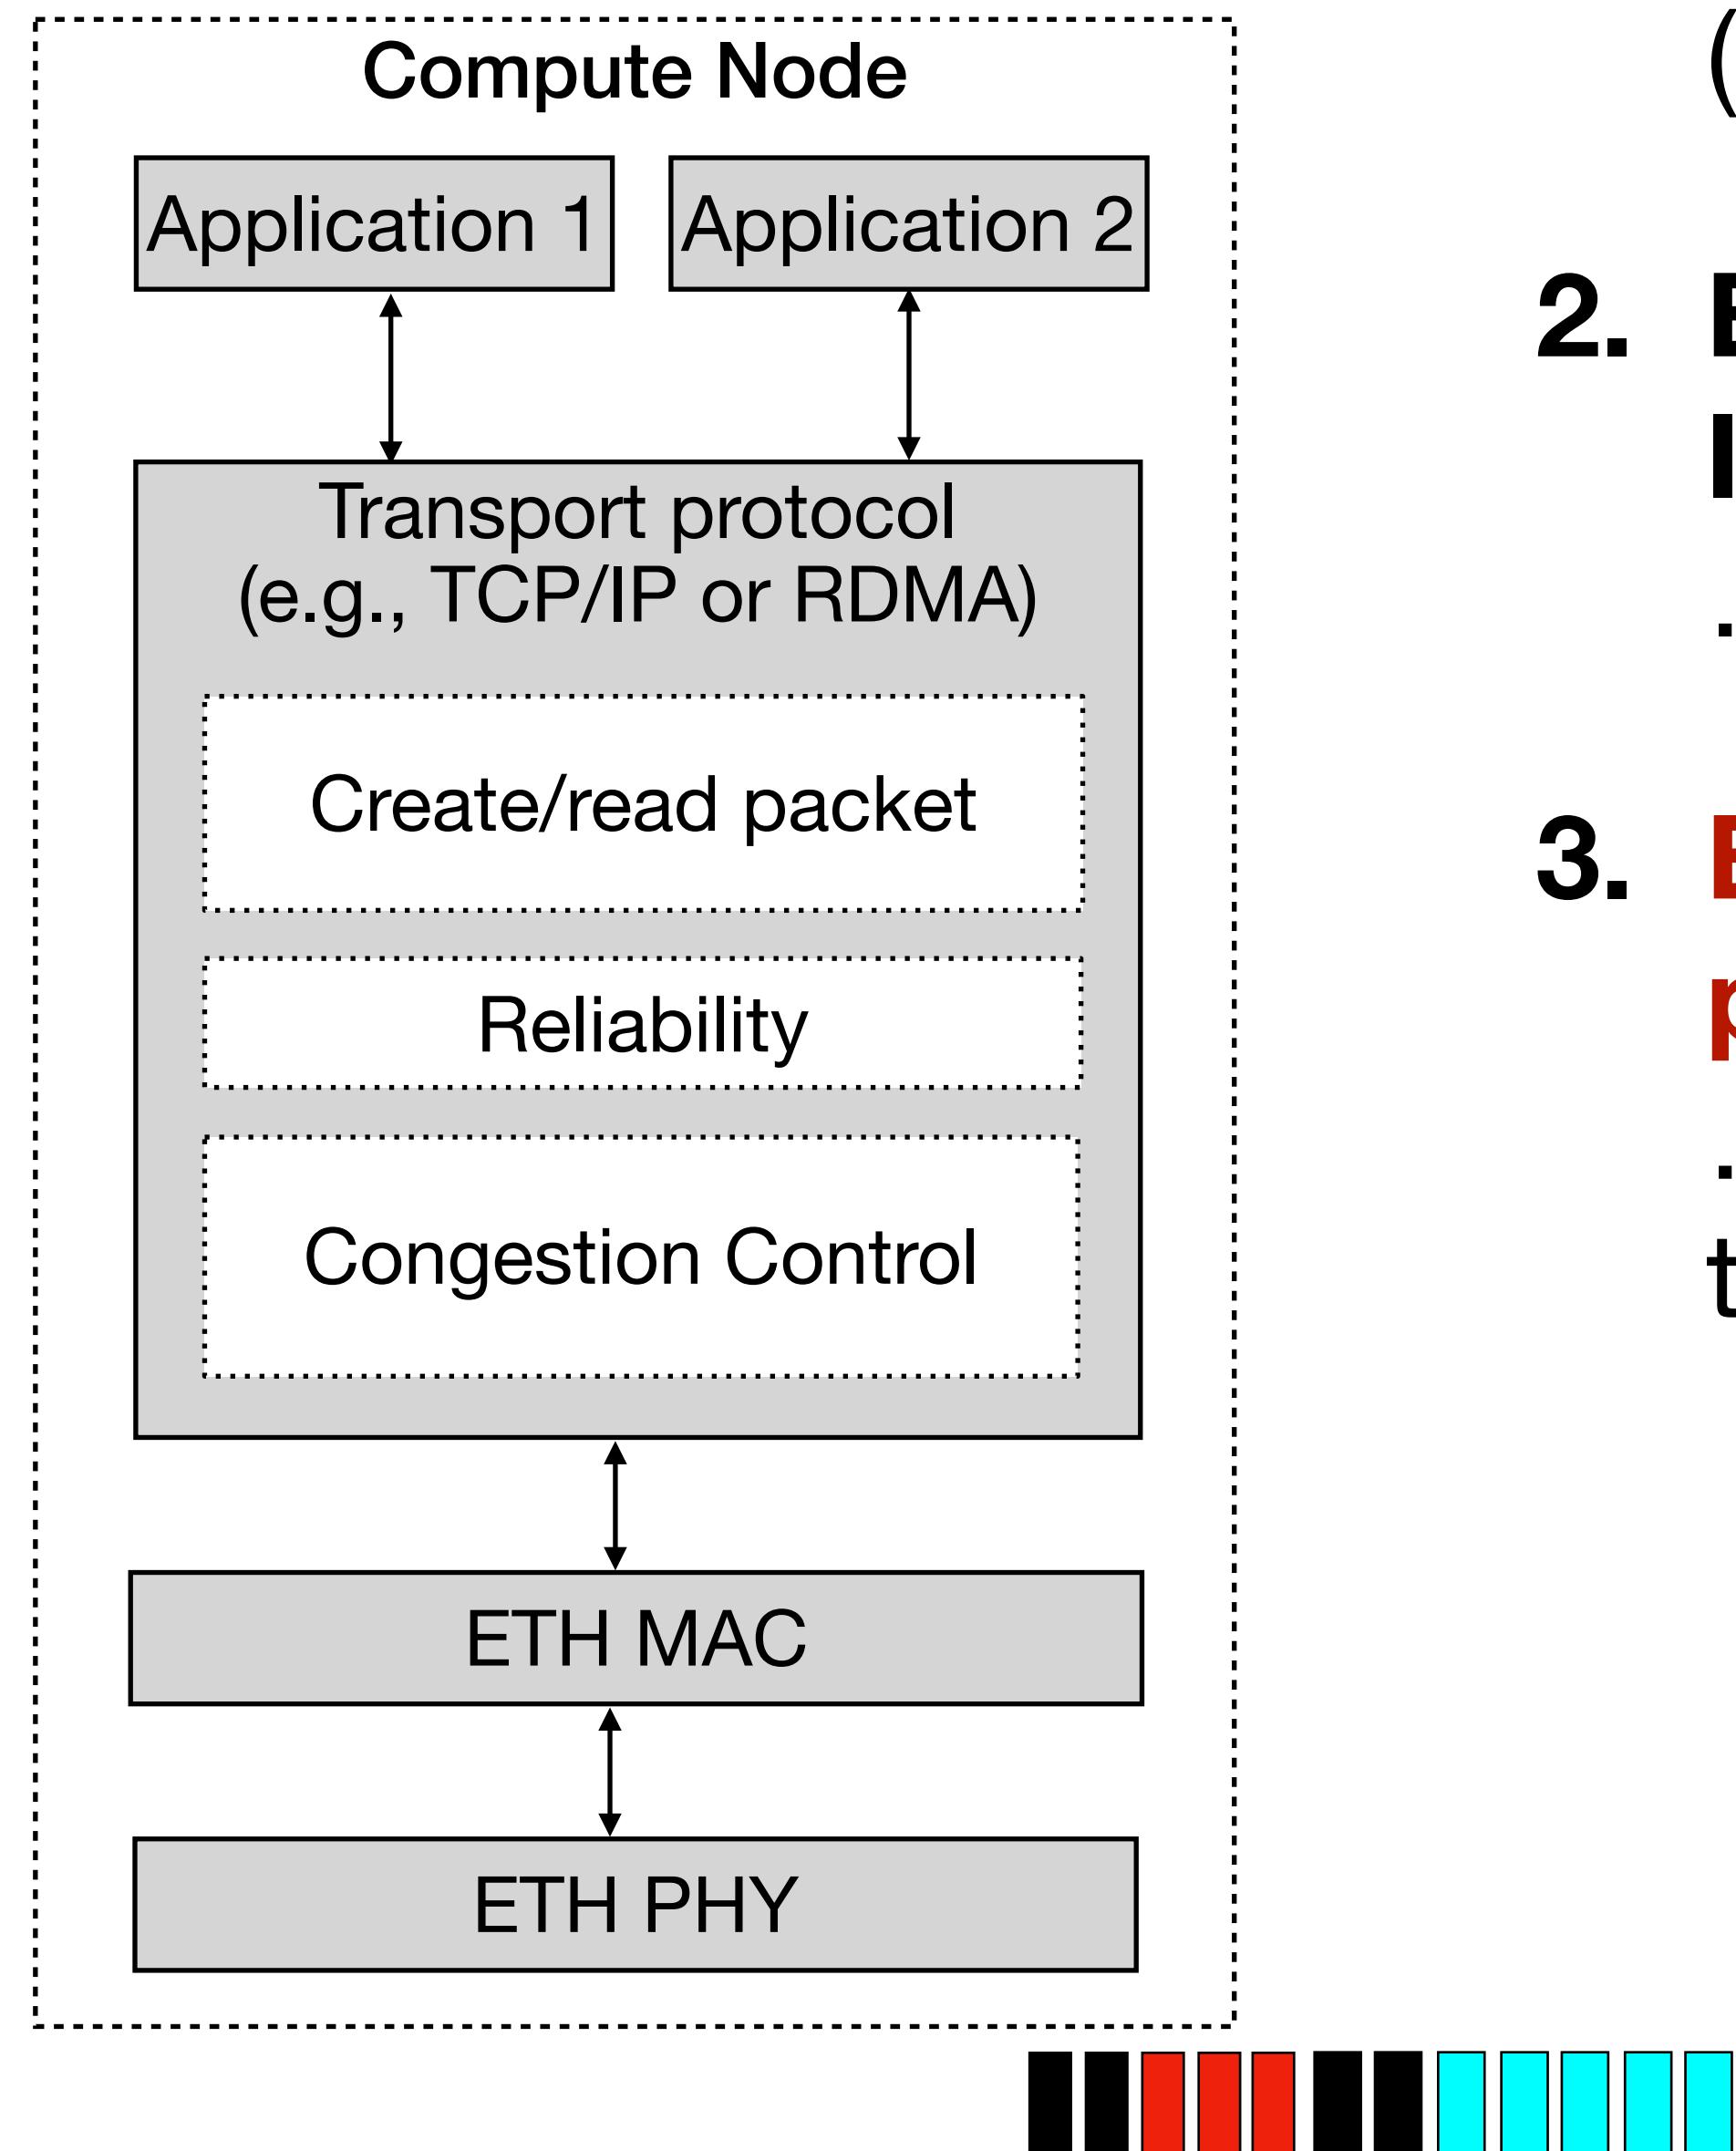
Other Overheads



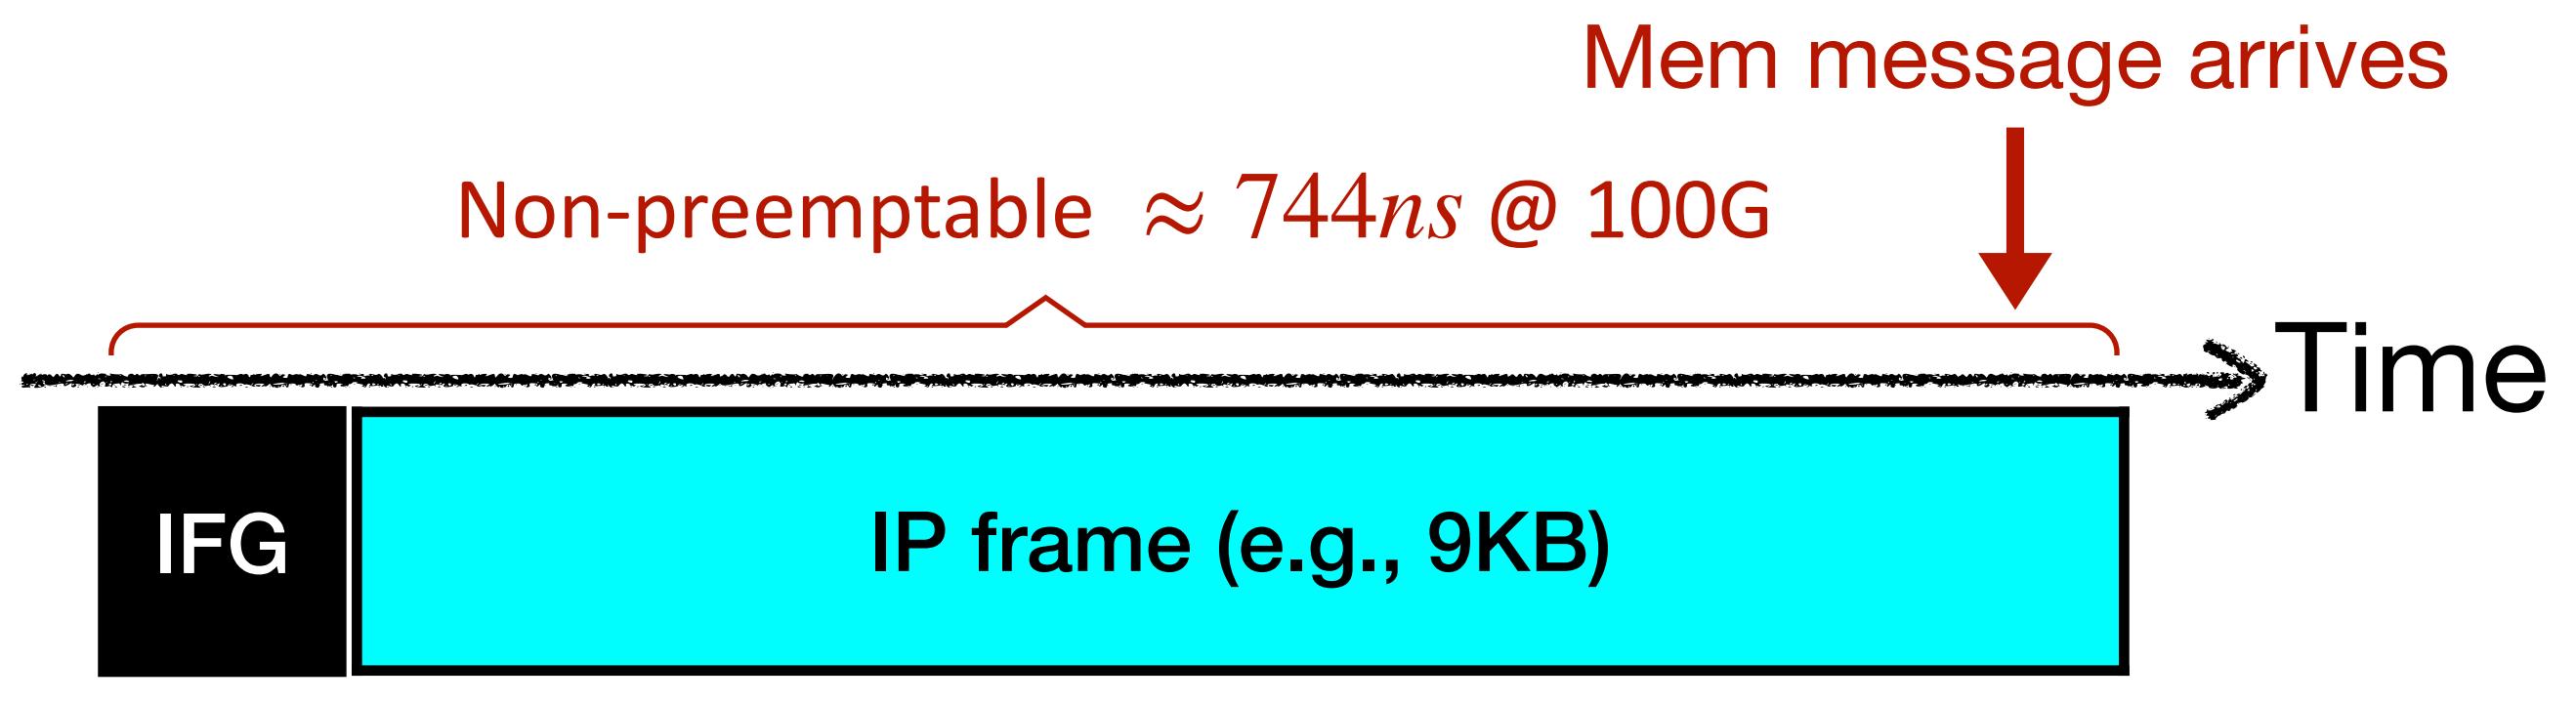
- Ethernet MAC enforces minimum 64B frame**
... but memory messages can be much smaller
(e.g., read requests are typically 8-16B)
- Ethernet MAC enforces minimum of 12 bytes Inter-frame gap (IFG)**
... high overhead for small memory messages
- Ethernet MAC does not allow intra-frame preemption**
... a large non-memory frame may block the transmission of a small memory message



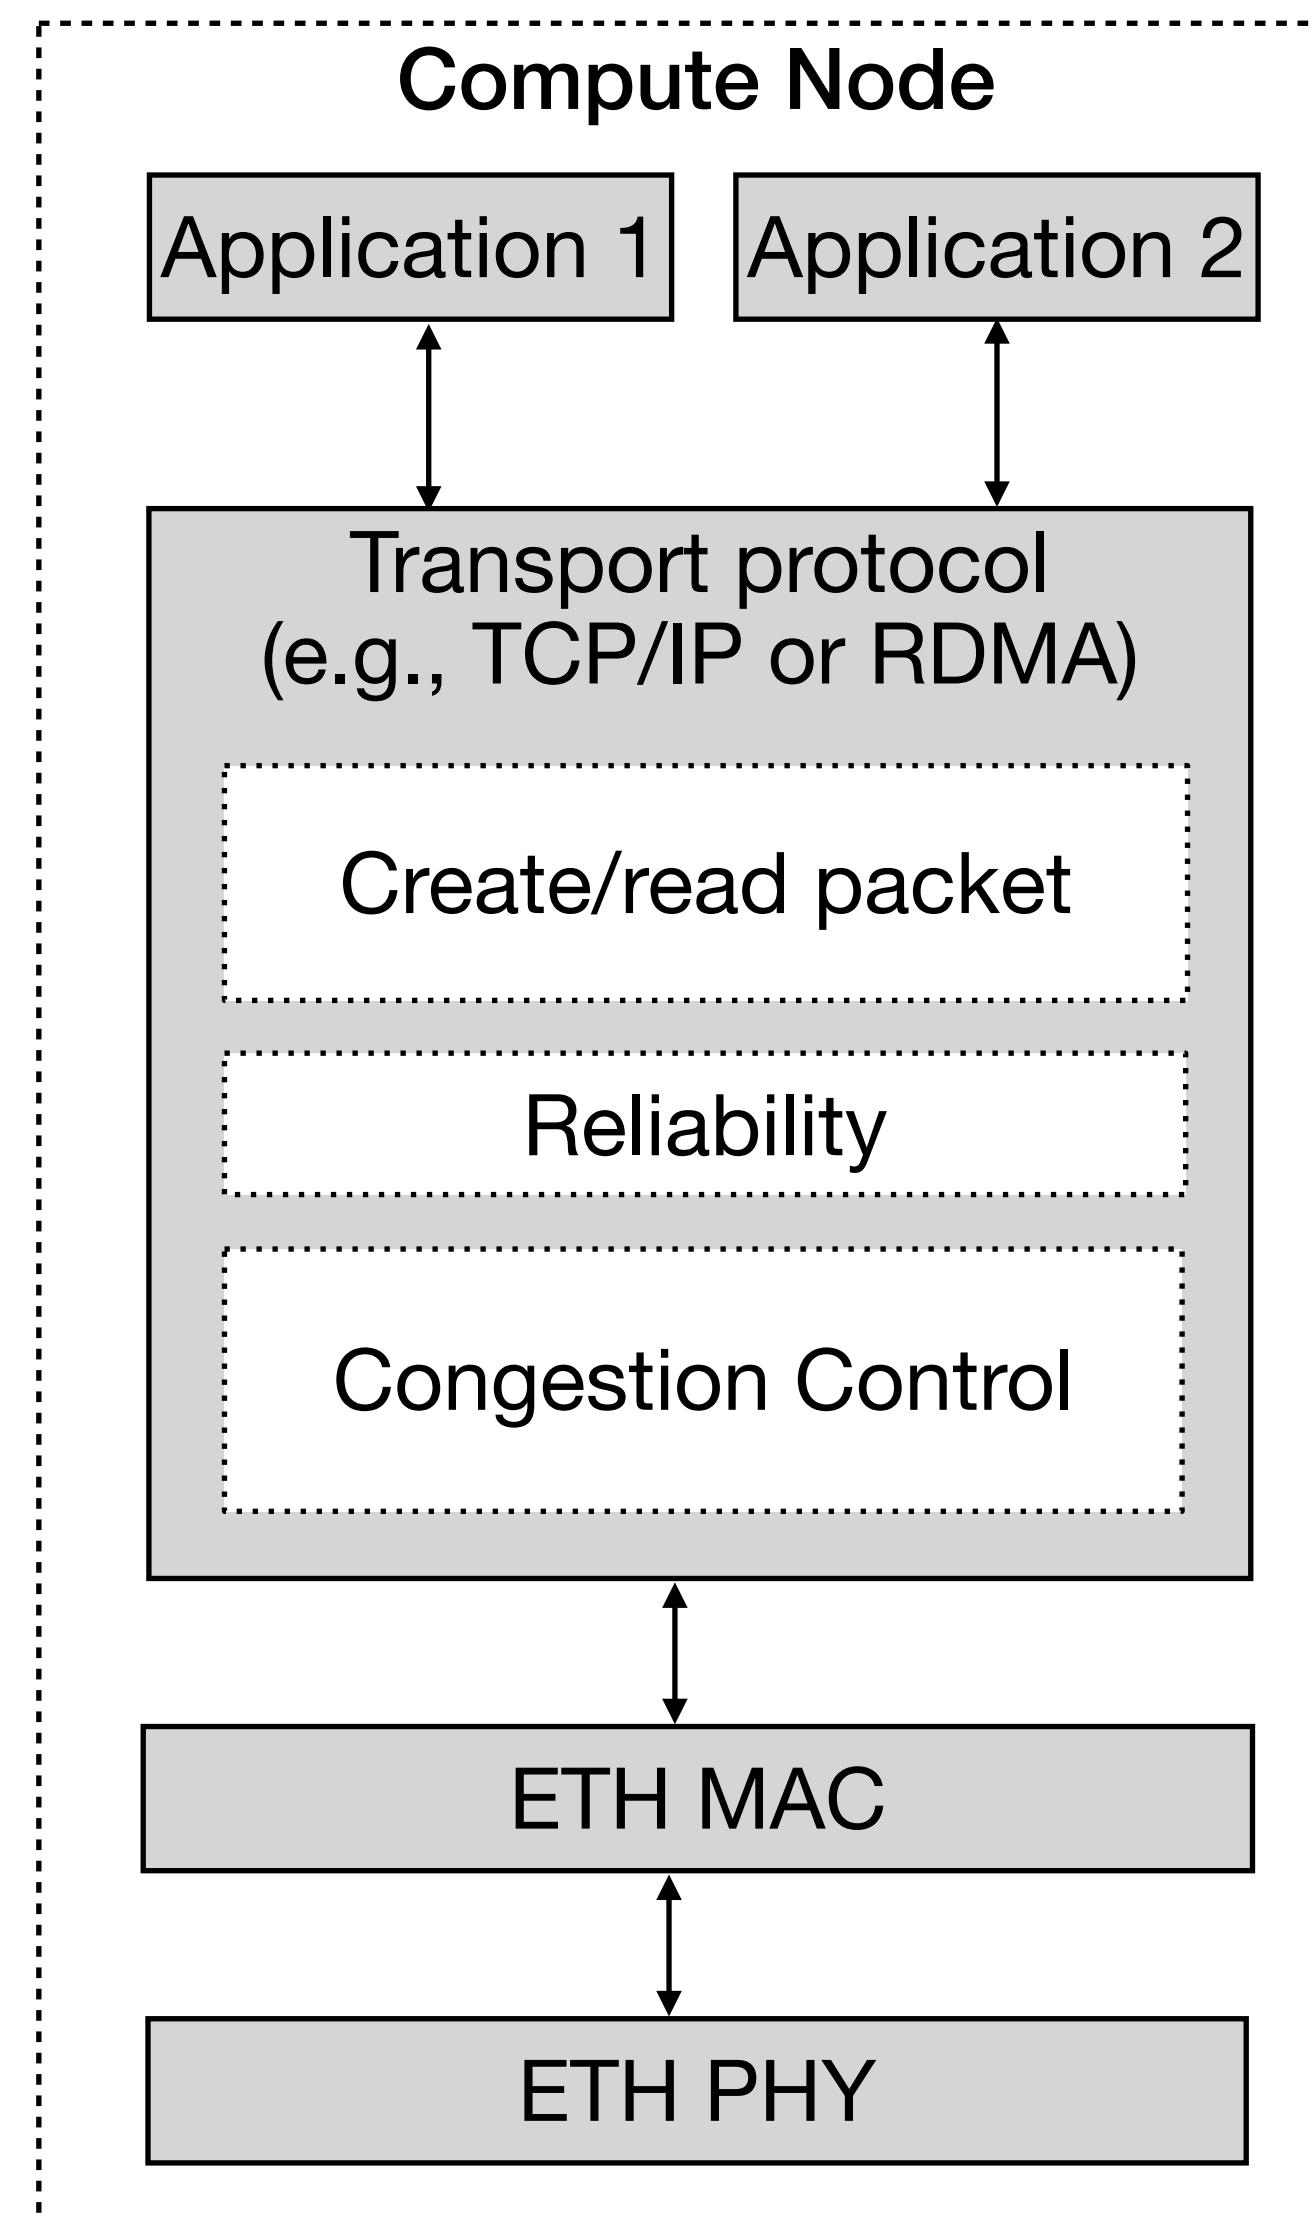
Other Overheads



- Ethernet MAC enforces minimum 64B frame**
... but memory messages can be much smaller
(e.g., read requests are typically 8-16B)
- Ethernet MAC enforces minimum of 12 bytes Inter-frame gap (IFG)**
... high overhead for small memory messages
- Ethernet MAC does not allow intra-frame preemption**
... a large non-memory frame may block the transmission of a small memory message



Other Overheads



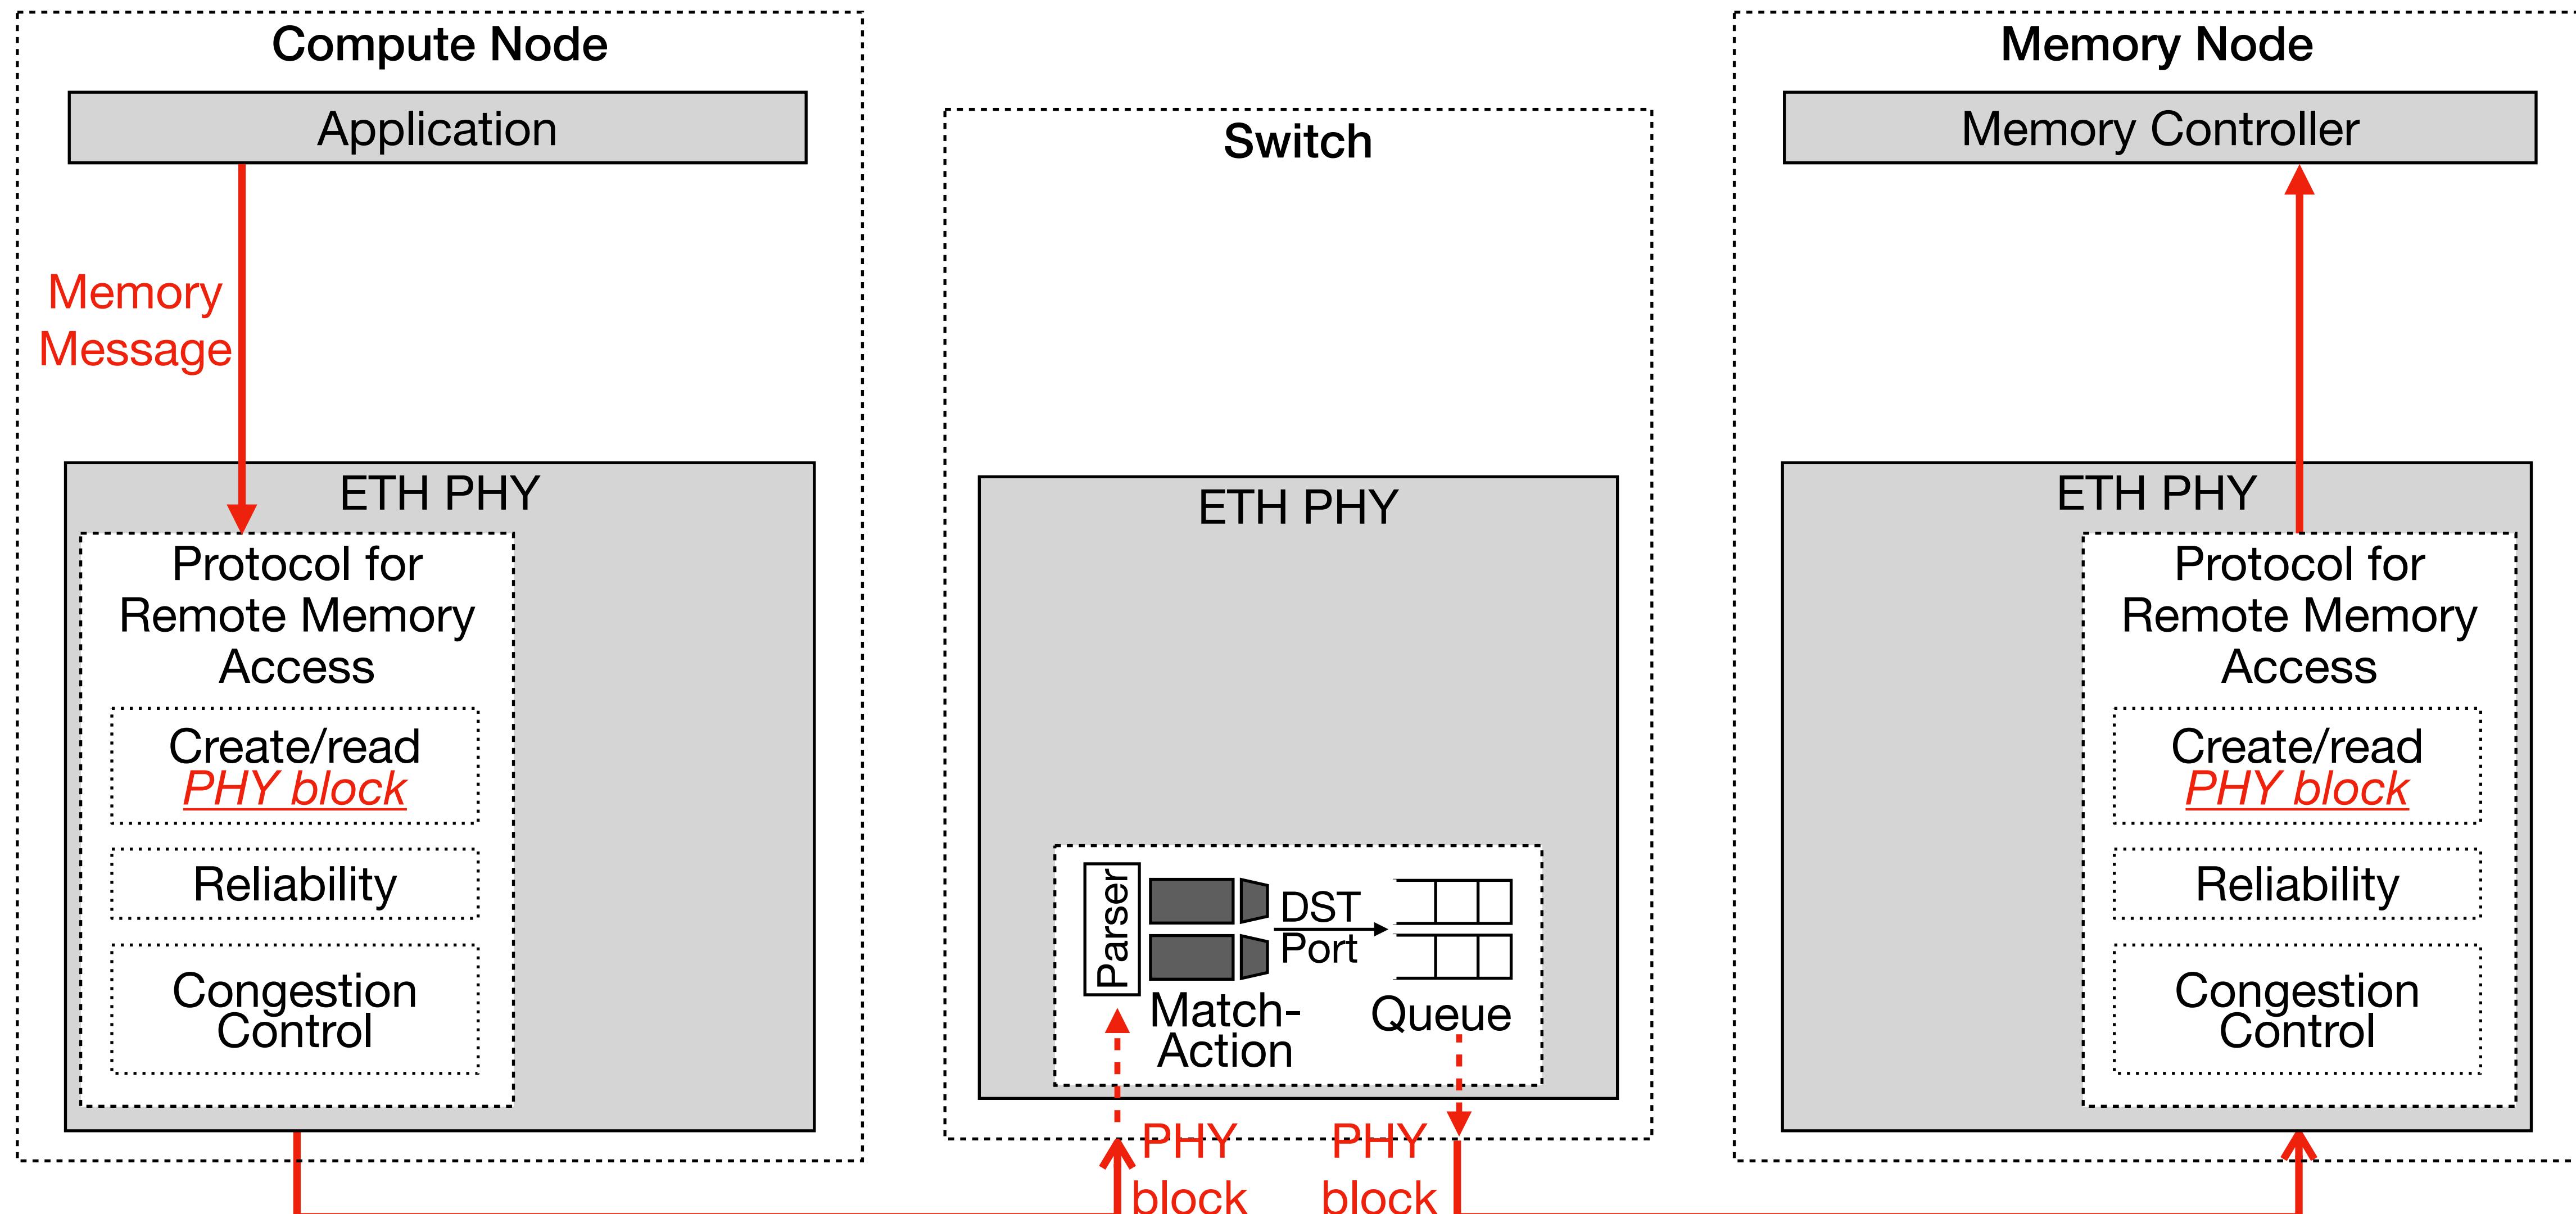
1. **Ethernet MAC enforces minimum 64B frame**
... but memory messages can be much smaller
(e.g., read requests are typically 8-16B)
2. **Ethernet MAC enforces minimum of 12 bytes Inter-frame gap (IFG)**
... high overhead for small memory messages
3. **Ethernet MAC does not allow intra-frame preemption**
... a large non-memory frame may block the transmission of a small memory message

Root cause: MAC layer processing

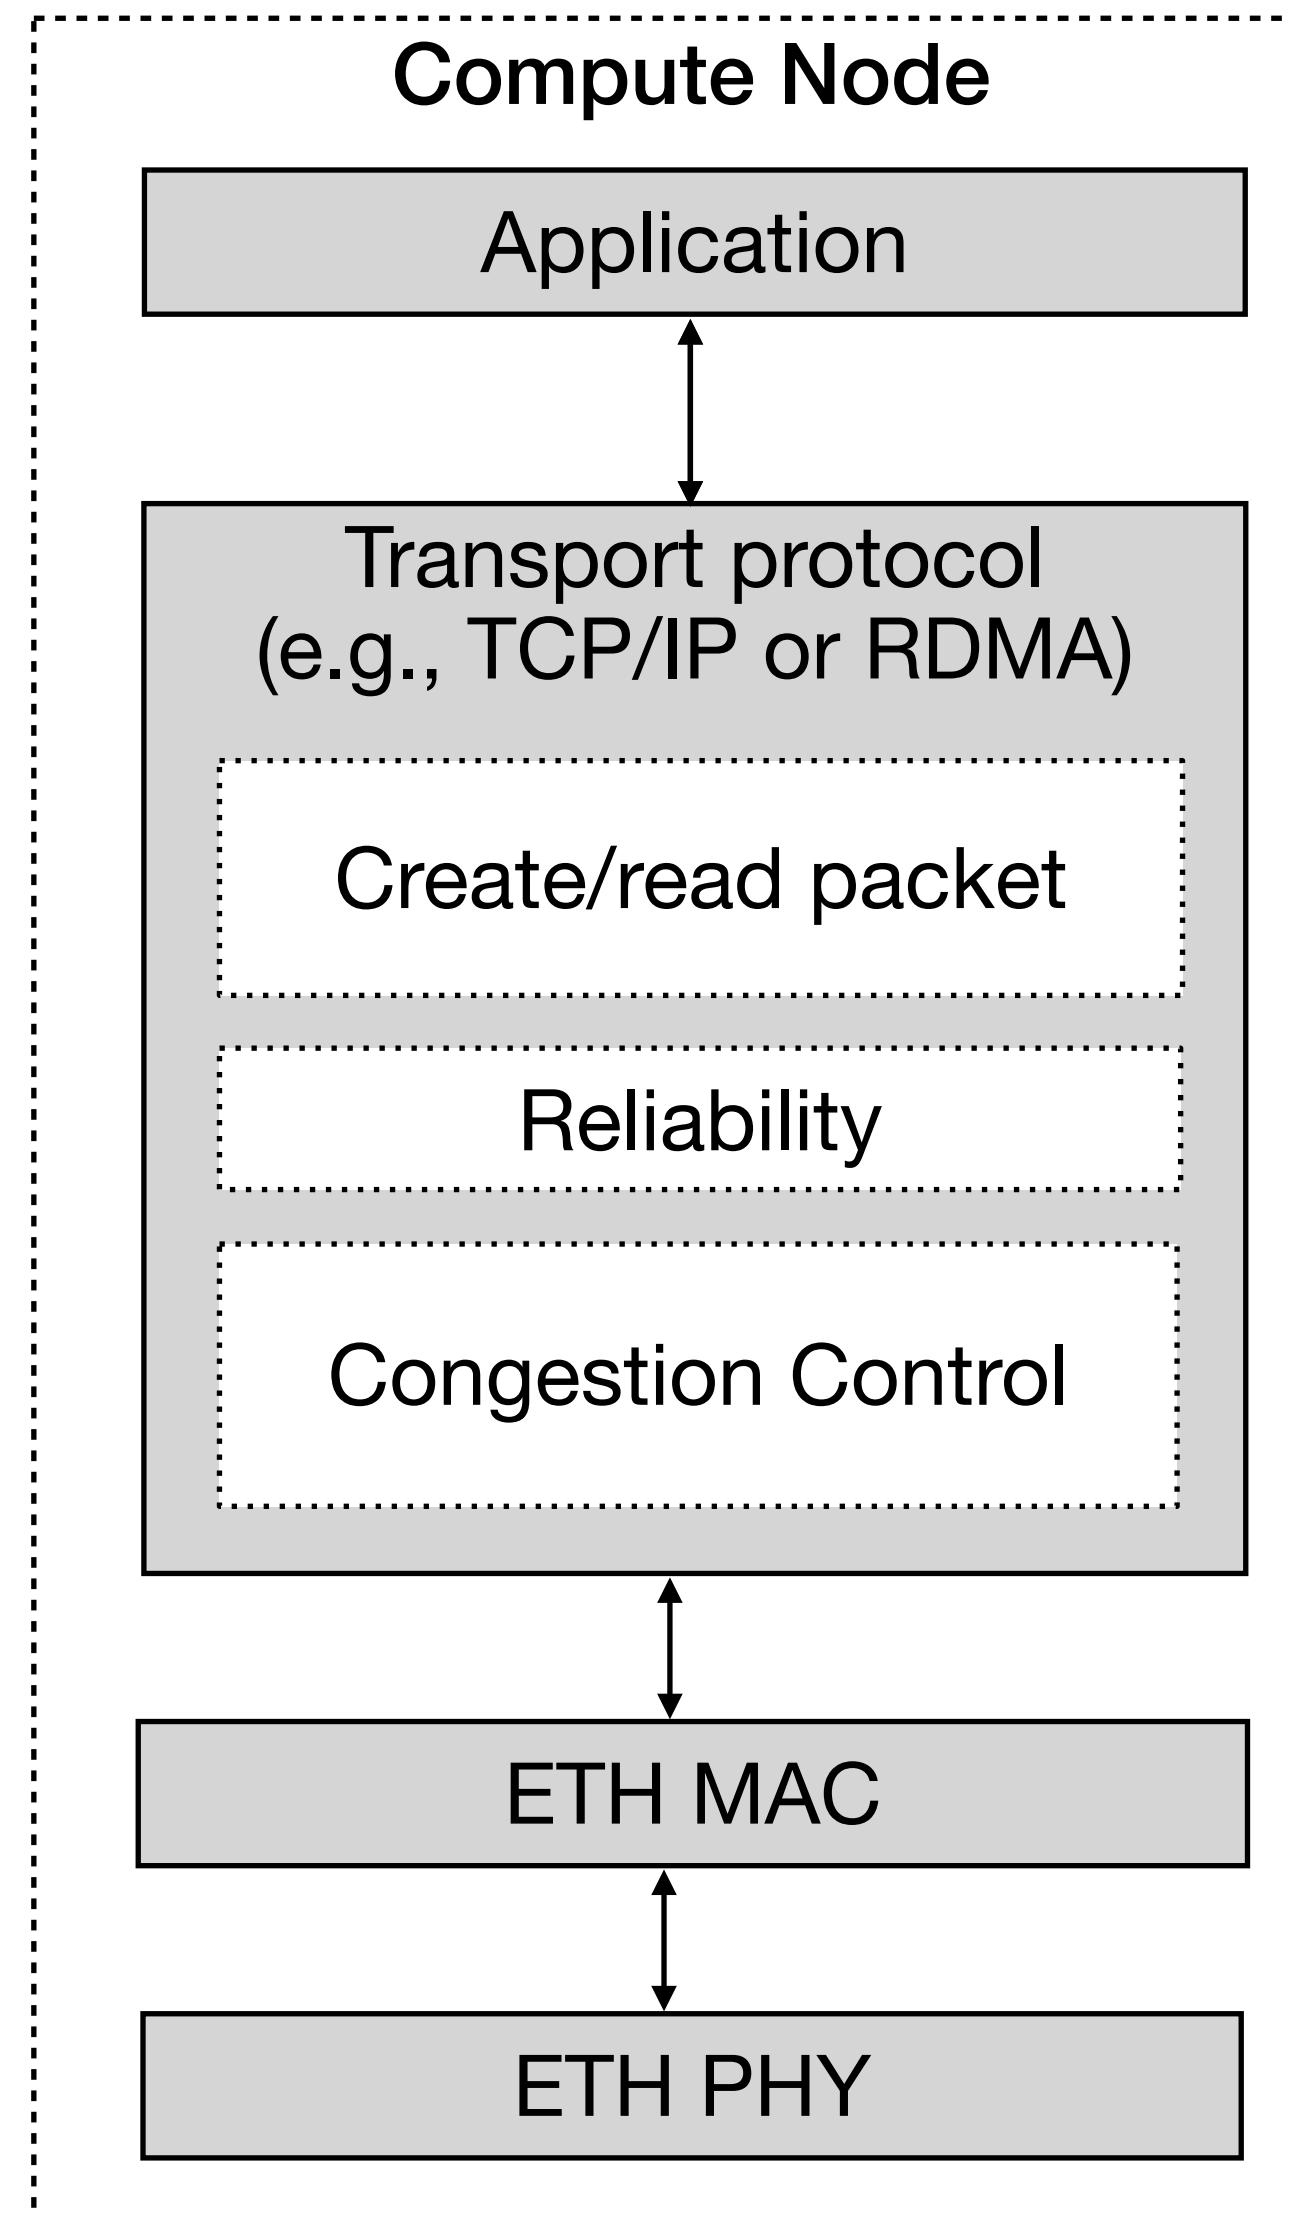
Design Choice # 1:

Implement the entire
protocol for remote memory access
within Ethernet's Physical Layer (PHY)

Architecture of Remote Memory Protocol in the PHY



Rationale for Remote Memory Protocol in PHY

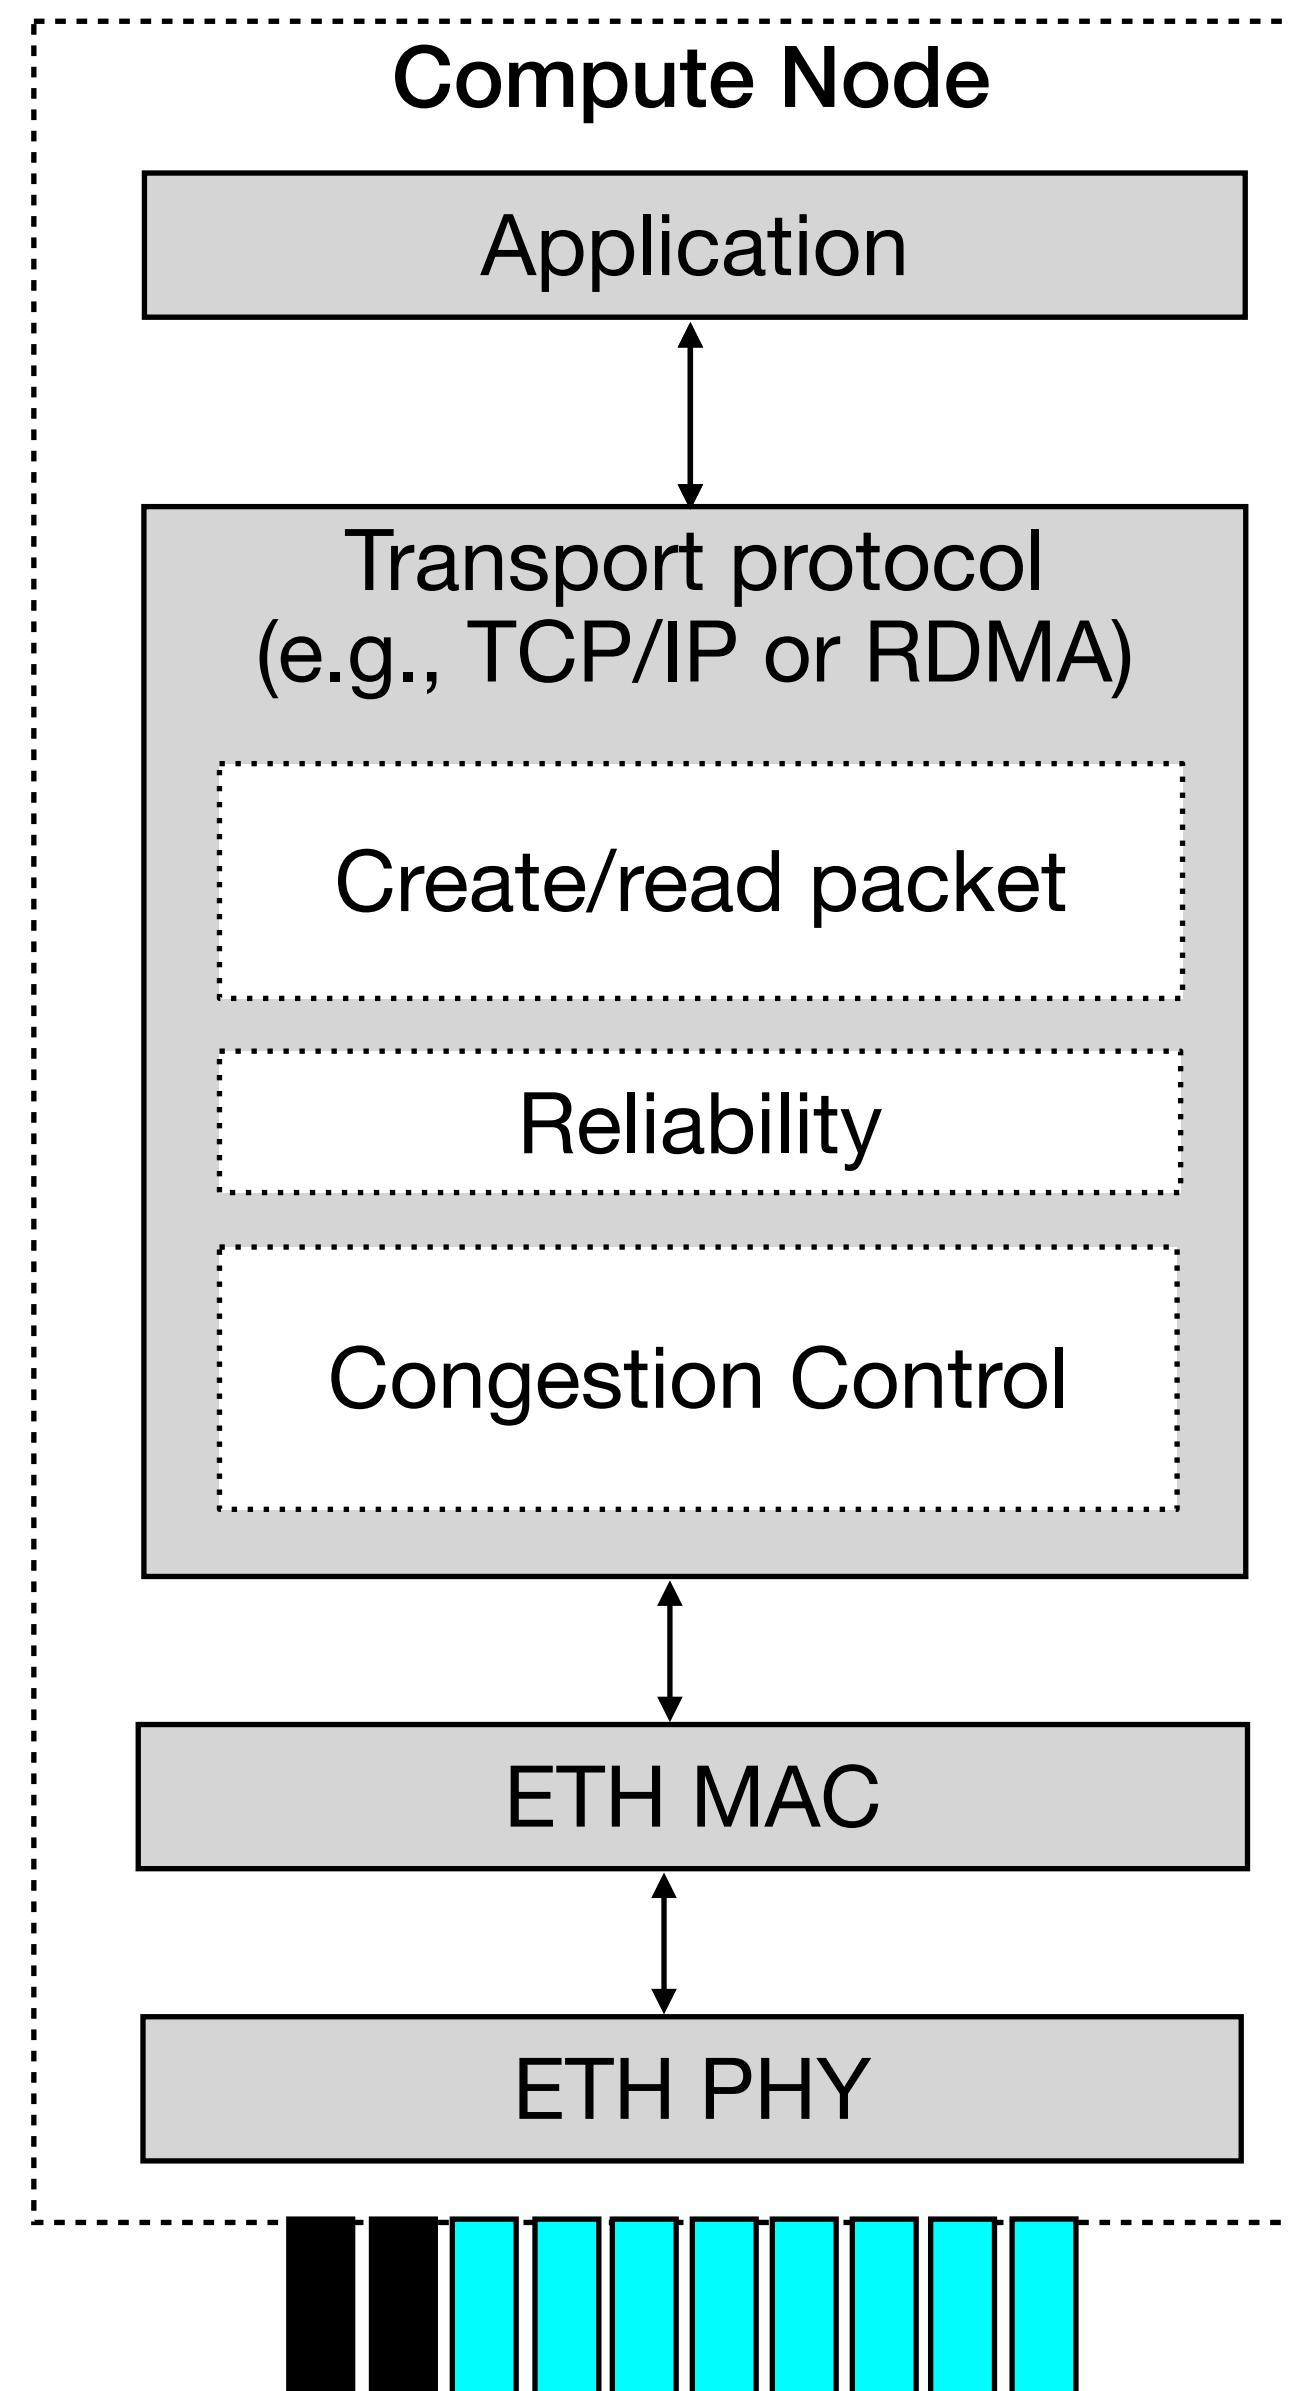


Ethernet PHY already reformats a MAC layer frame into a series of 66-bit PHY blocks

... thus, unlike the MAC layer that works at a frame granularity, PHY works at fine-grained block granularity

- 66 bit PHY block vs. 64 byte minimum MAC frame size
- Message interleaving can be done at block granularity in PHY rather than at frame granularity in MAC
- PHY also has access to IFG blocks

Rationale for Remote Memory Protocol in PHY

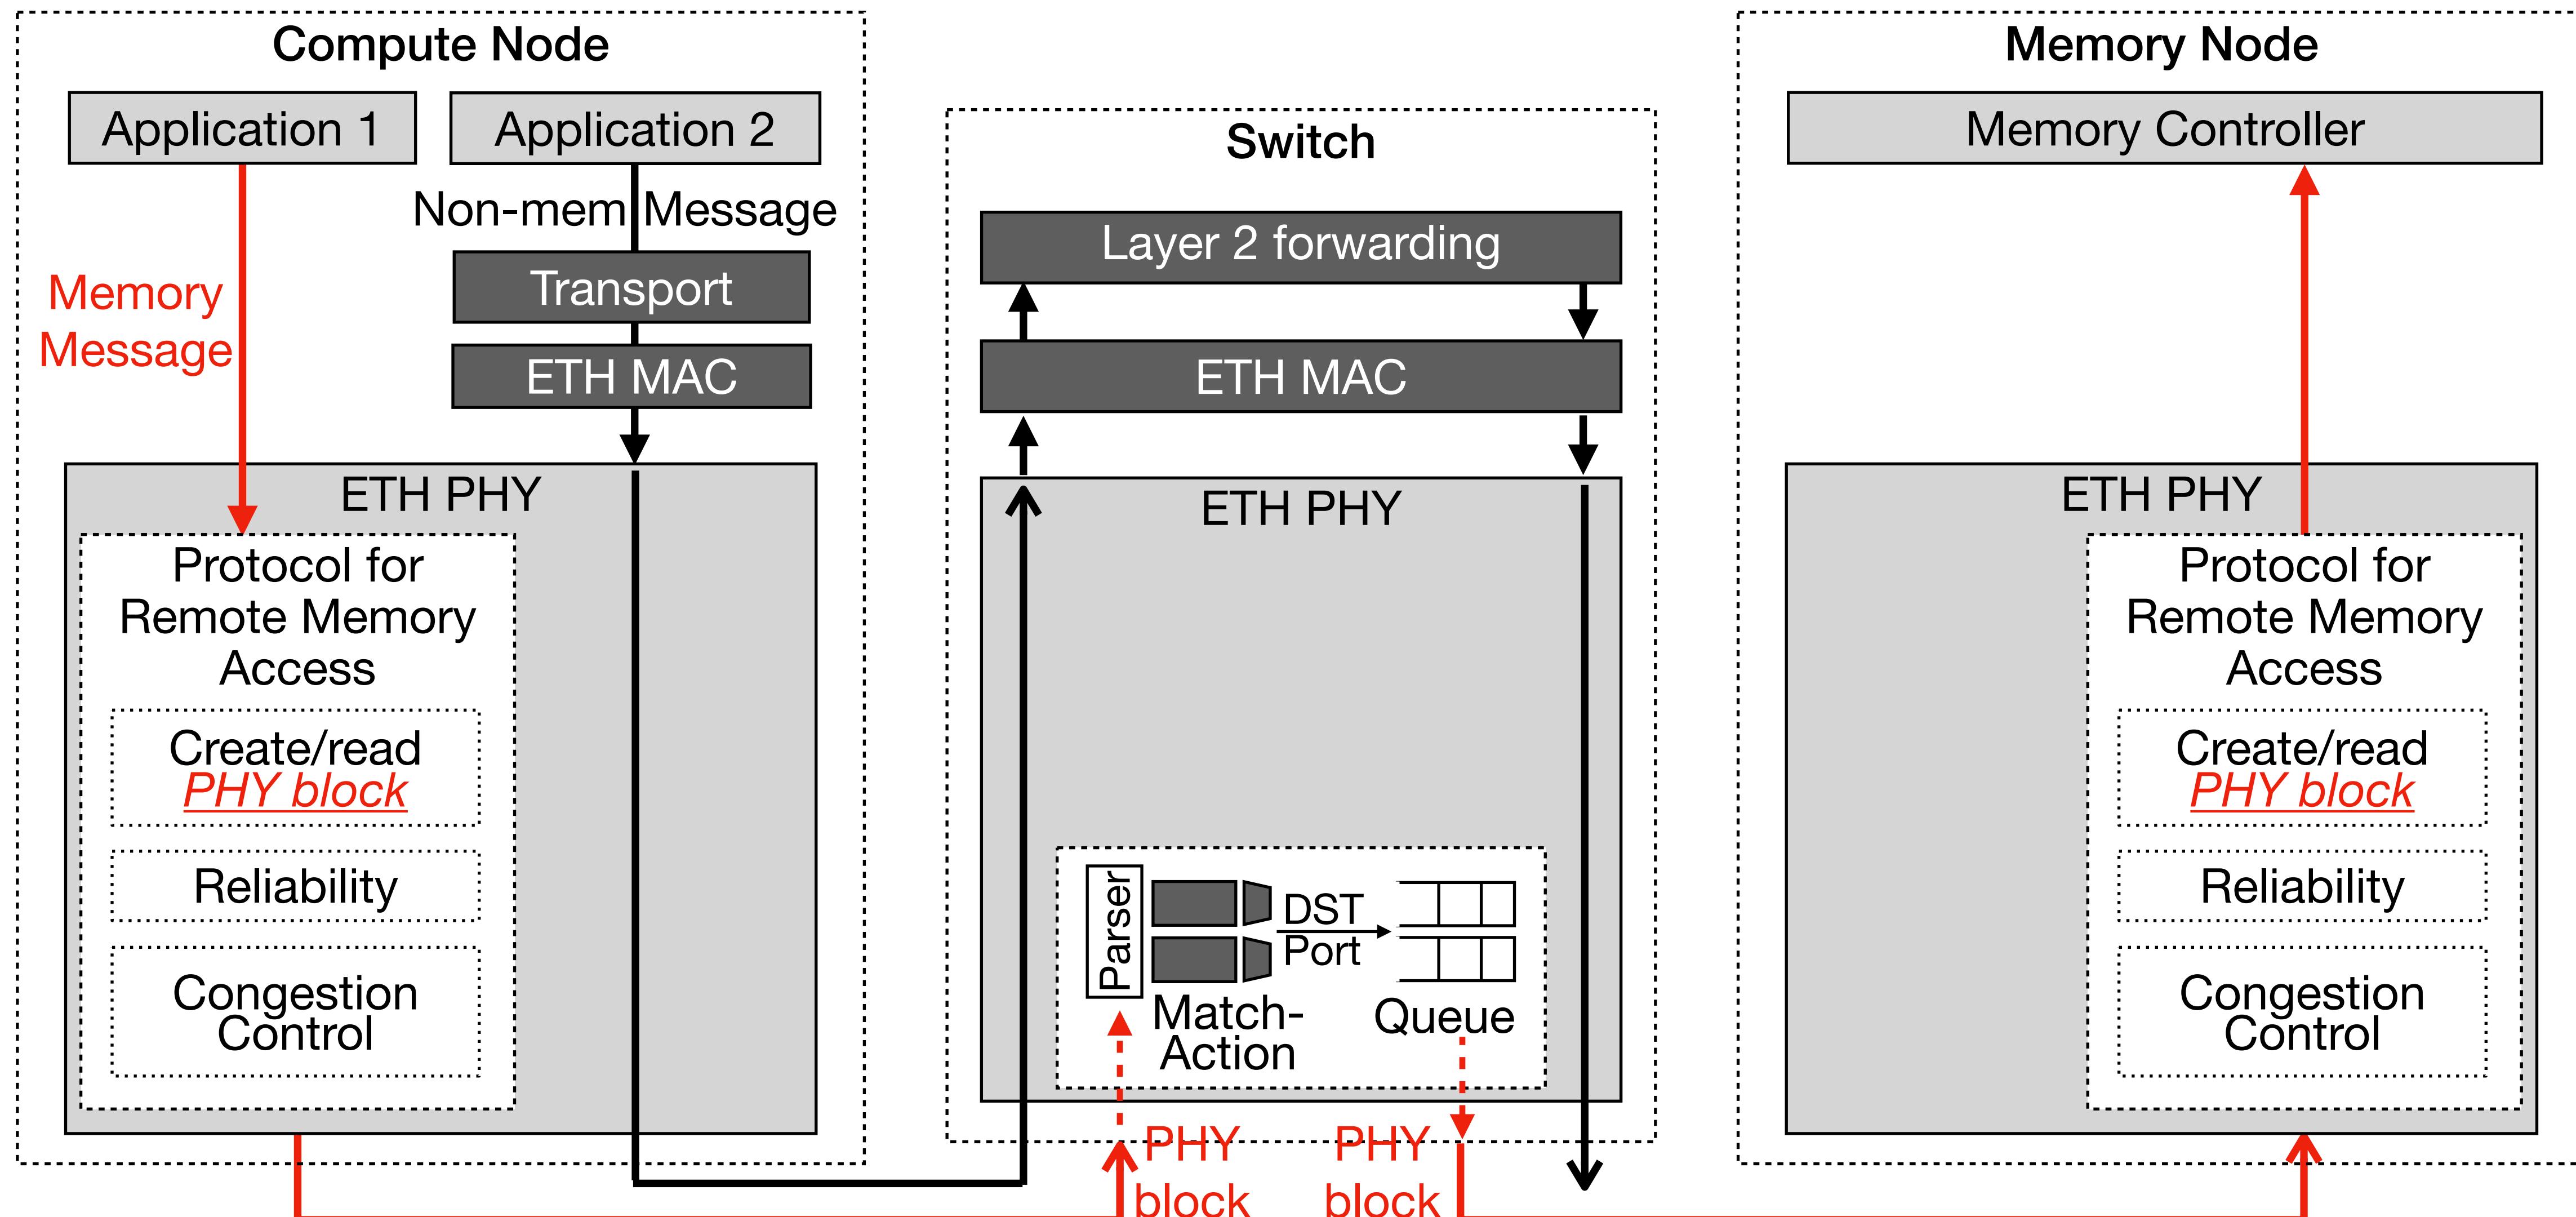


Ethernet PHY already reformats a MAC layer frame into a series of 66-bit PHY blocks

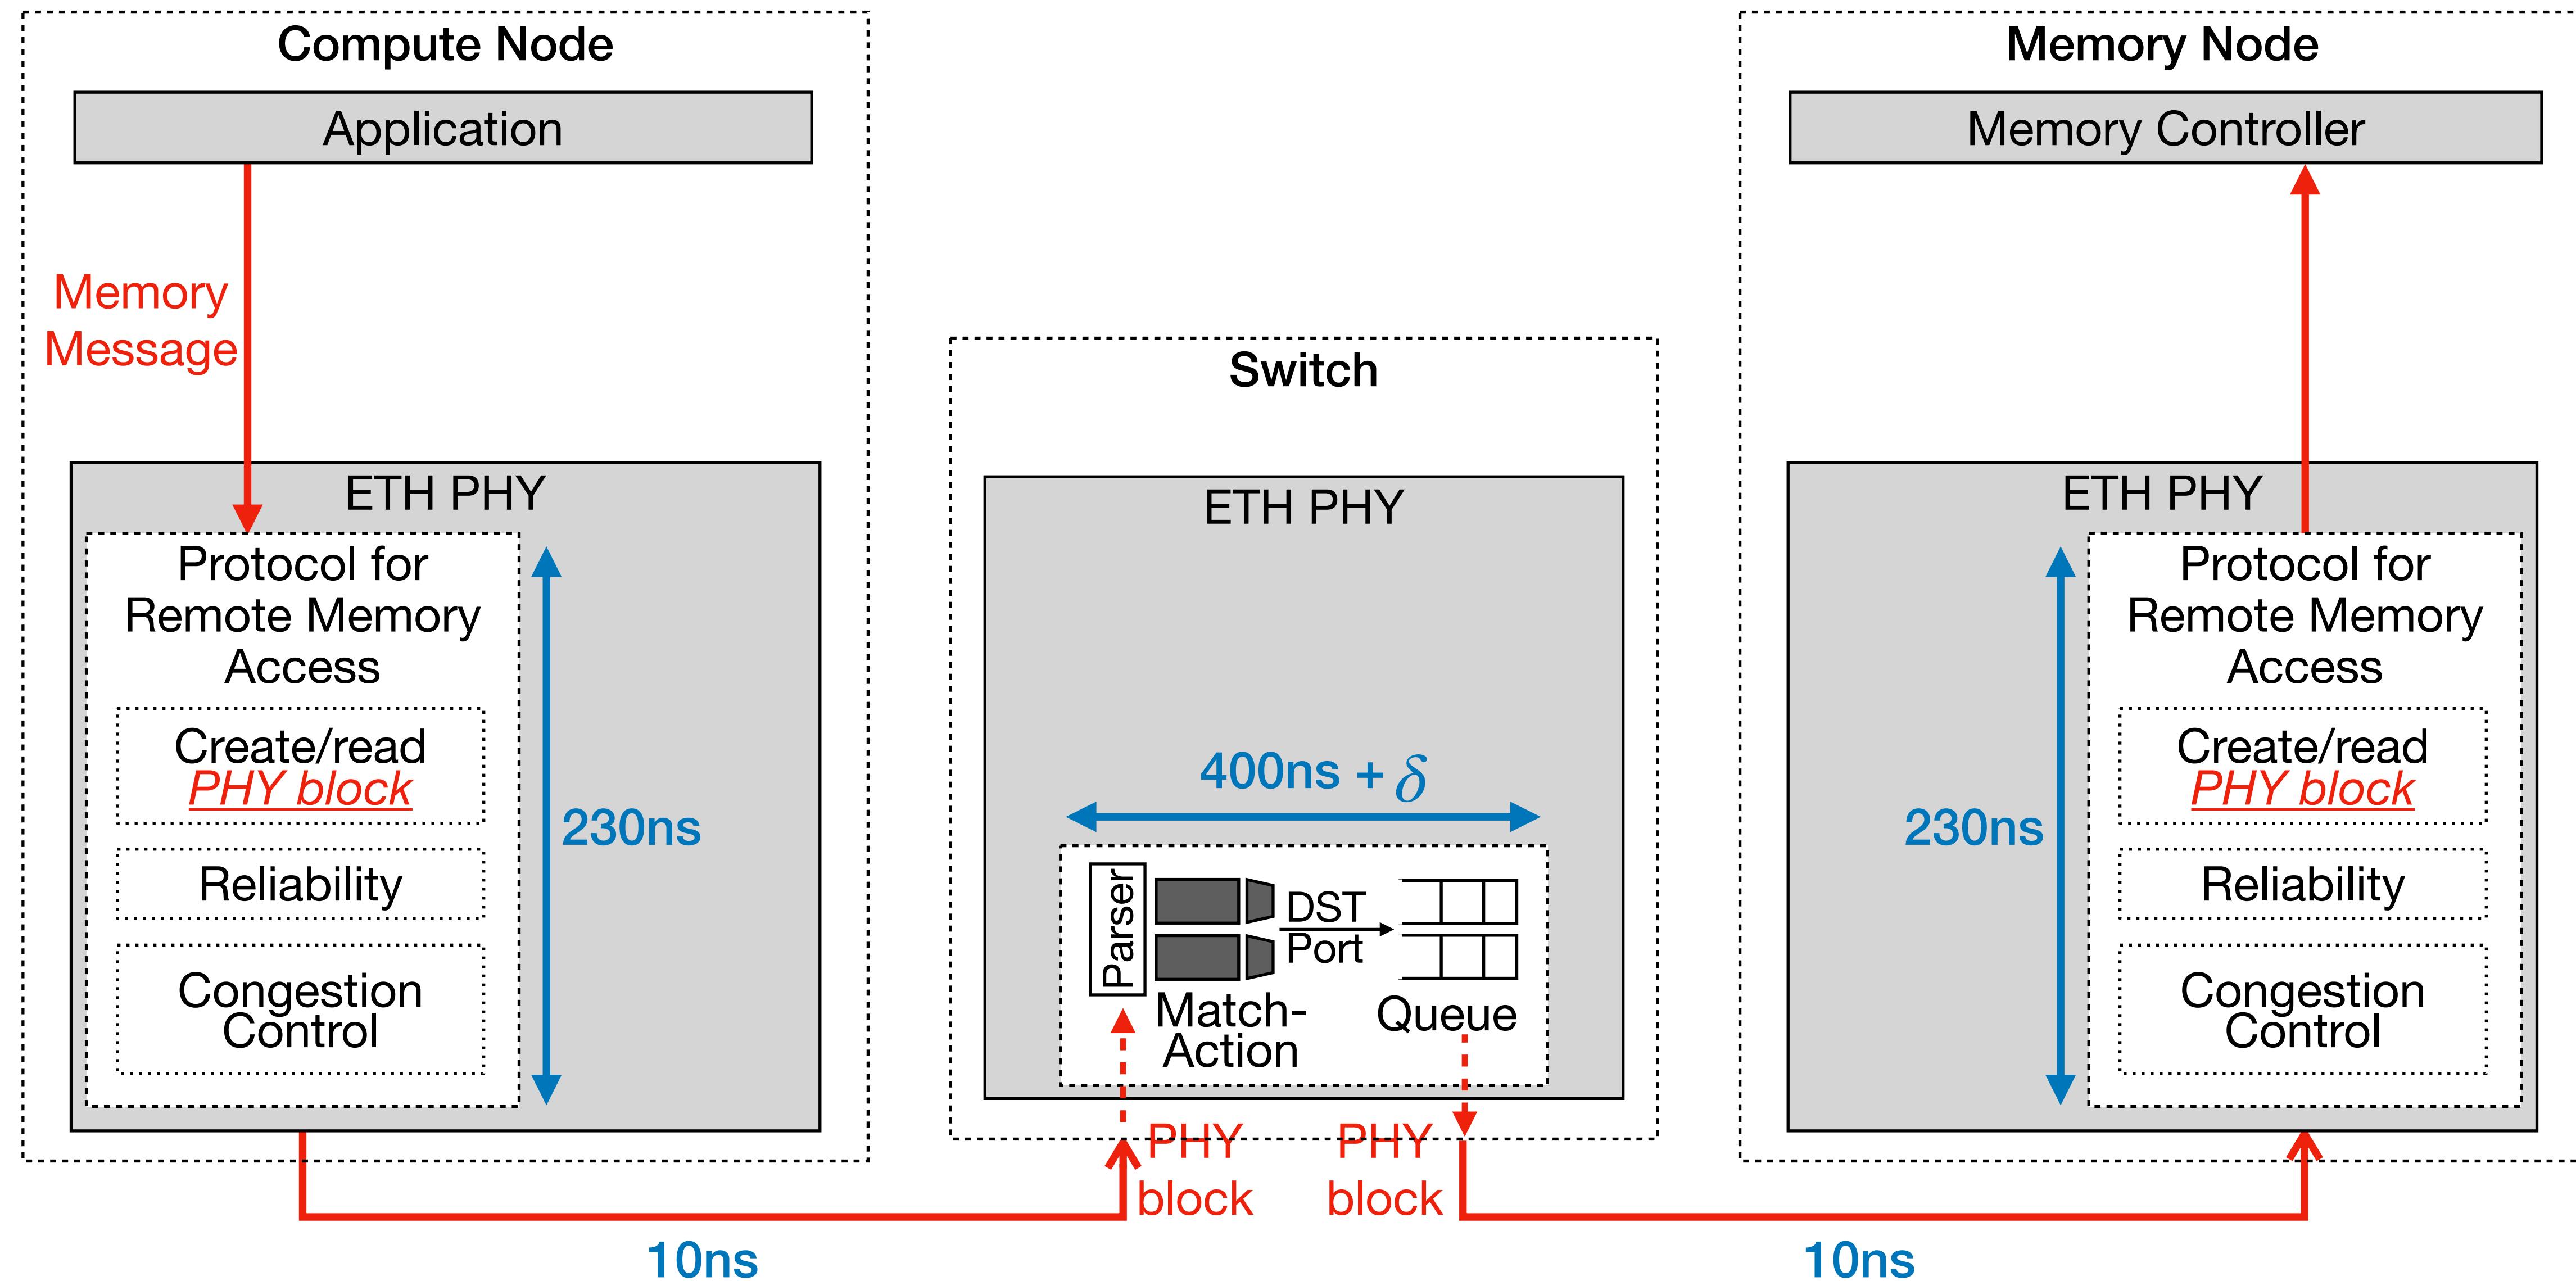
... thus, unlike the MAC layer that works at a frame granularity, PHY works at fine-grained block granularity

- 66 bit PHY block vs. 64 byte minimum MAC frame size
- Message interleaving can be done at block granularity in PHY rather than at frame granularity in MAC
- PHY also has access to IFG blocks

Architecture of Remote Memory Protocol in the PHY



Remote Memory Protocol in the PHY : What about latency?

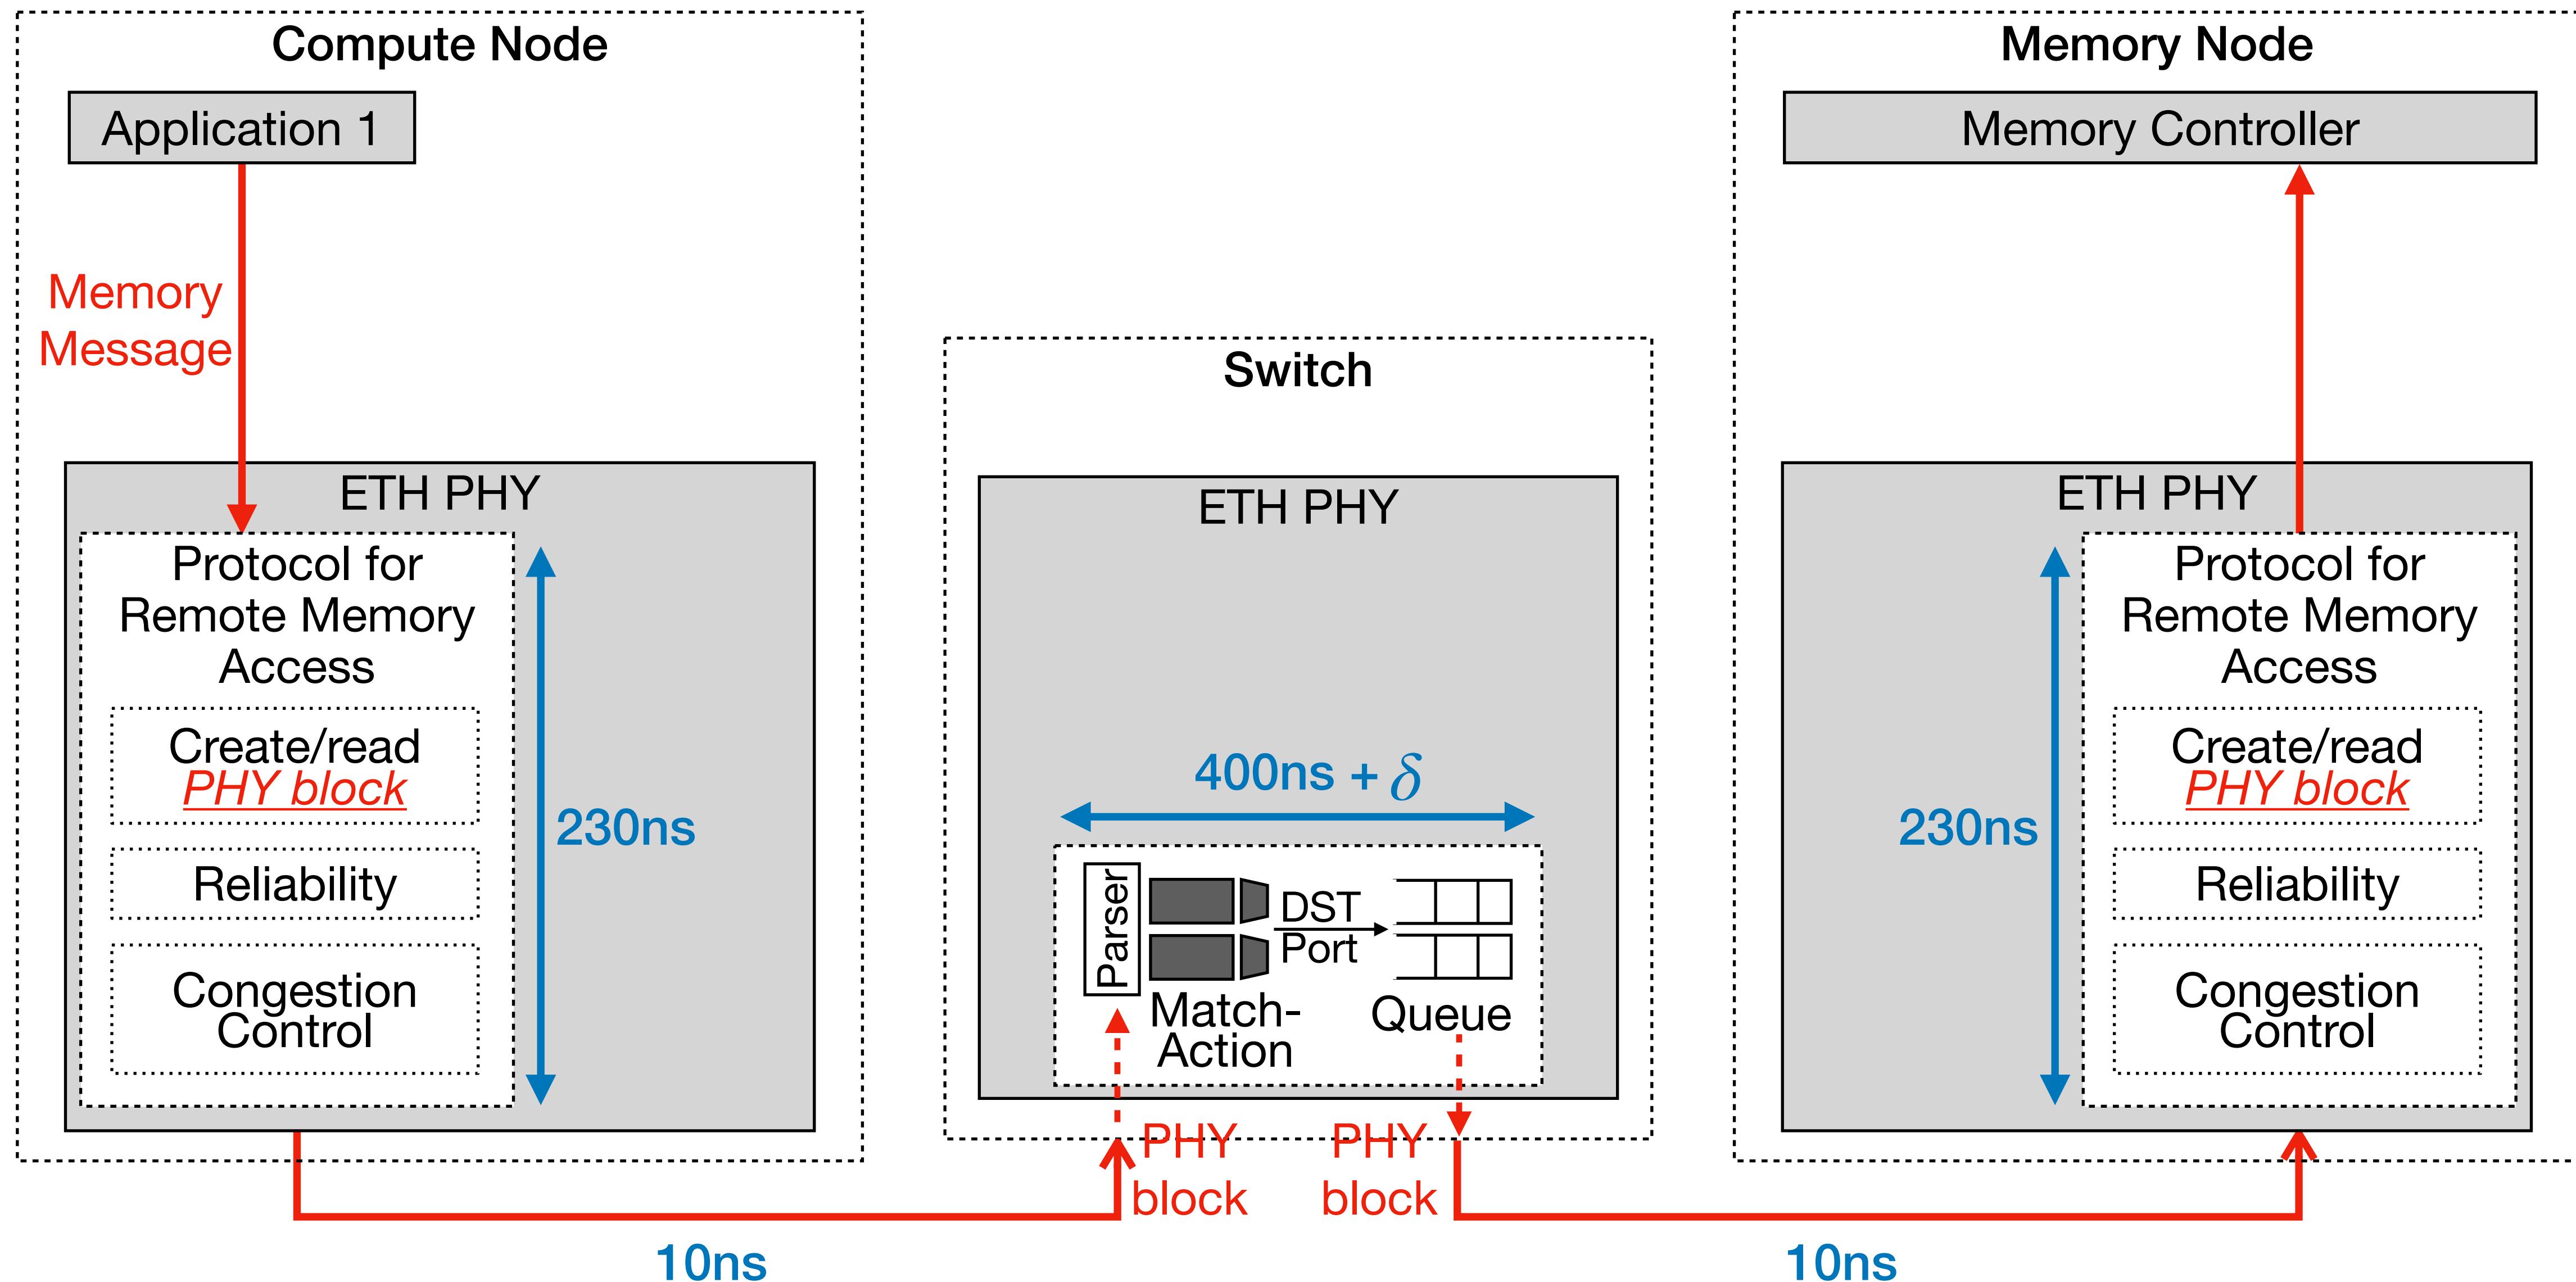


Design Choice # 2:

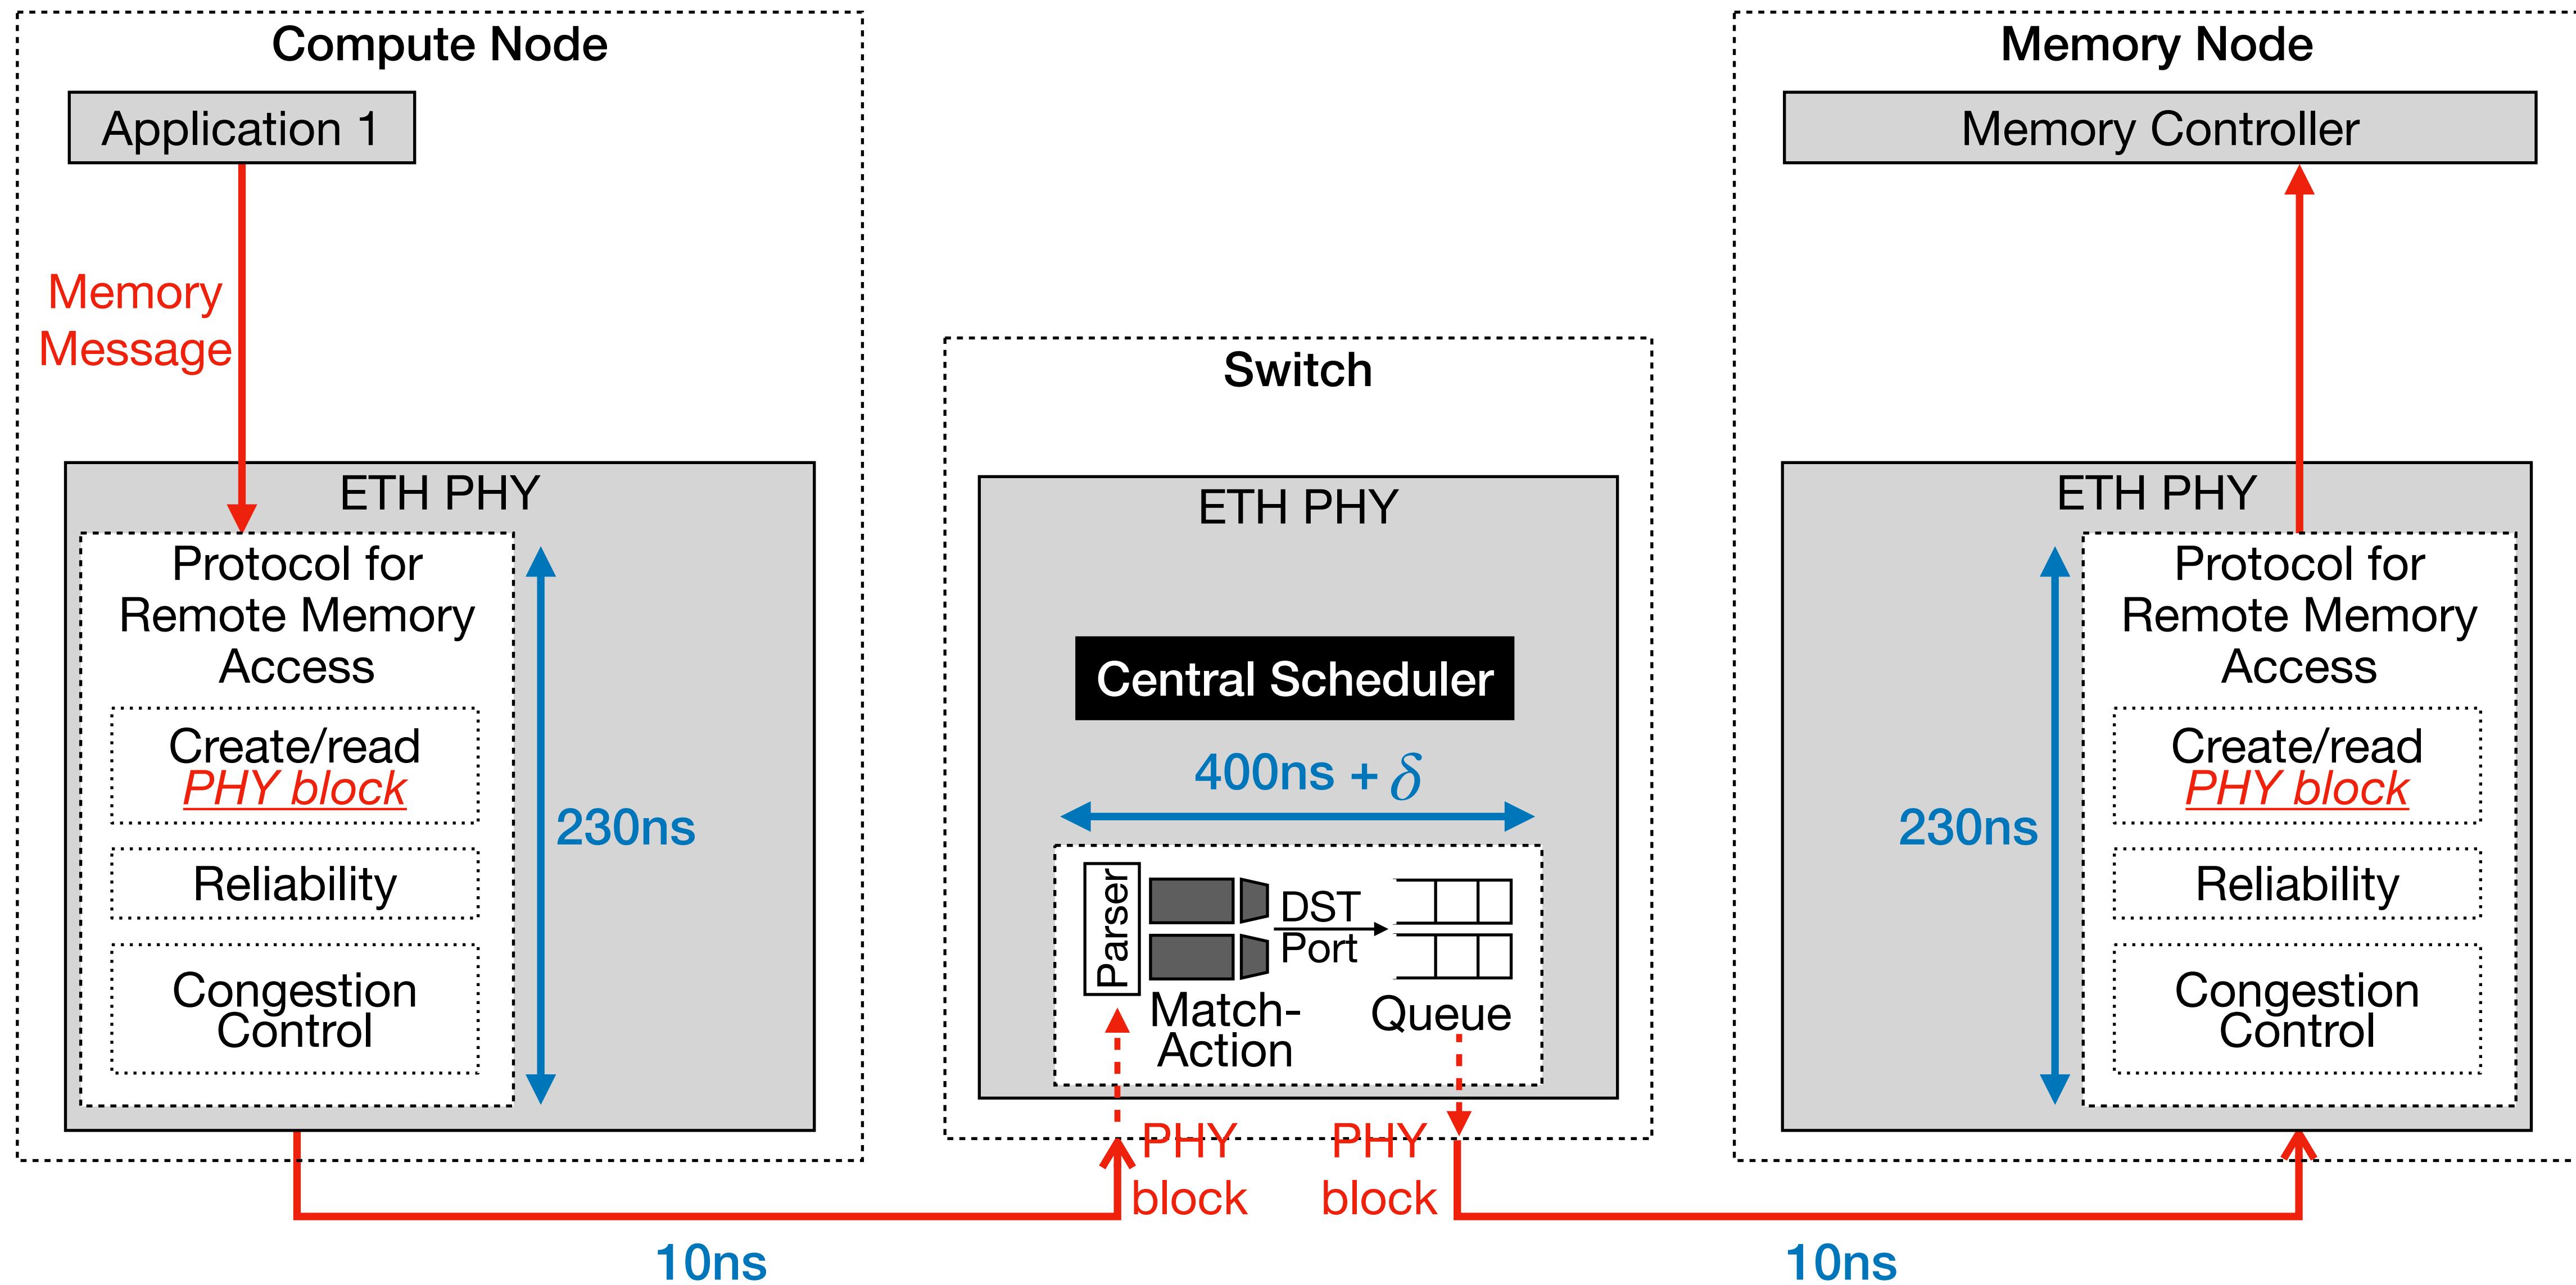
Packet Switching → Reconfigurable
(Circuit) Switching

Using a **centralized memory traffic scheduler**
implemented in the PHY of the Ethernet switch

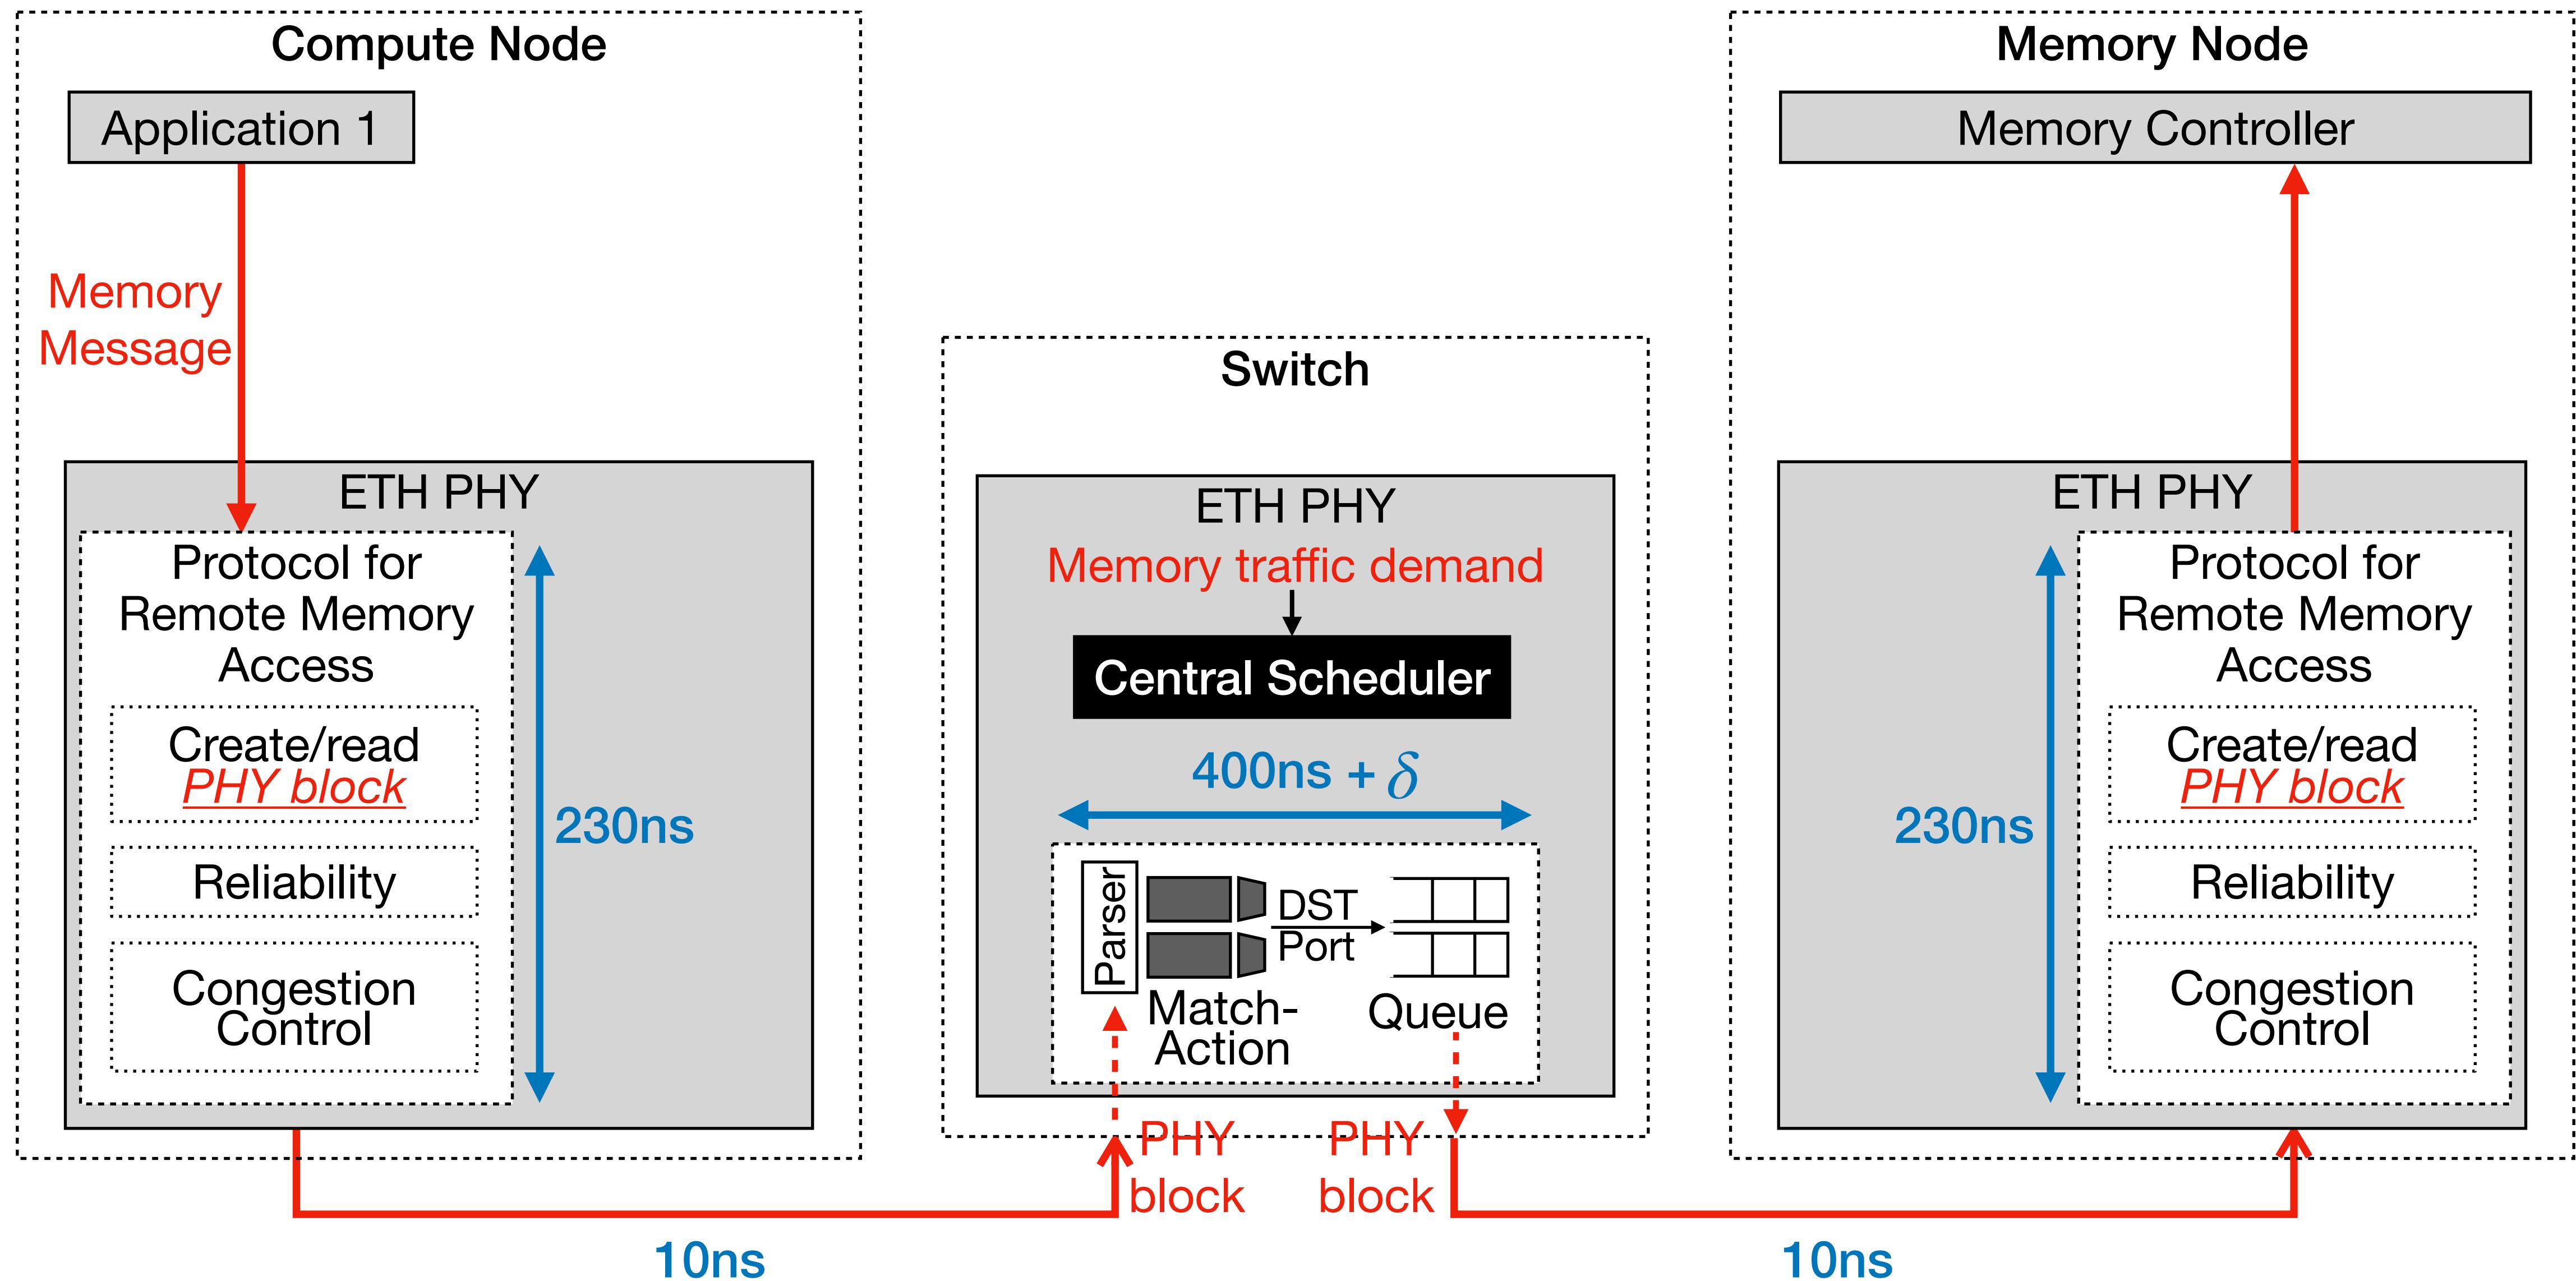
Central Scheduler in the Switch PHY



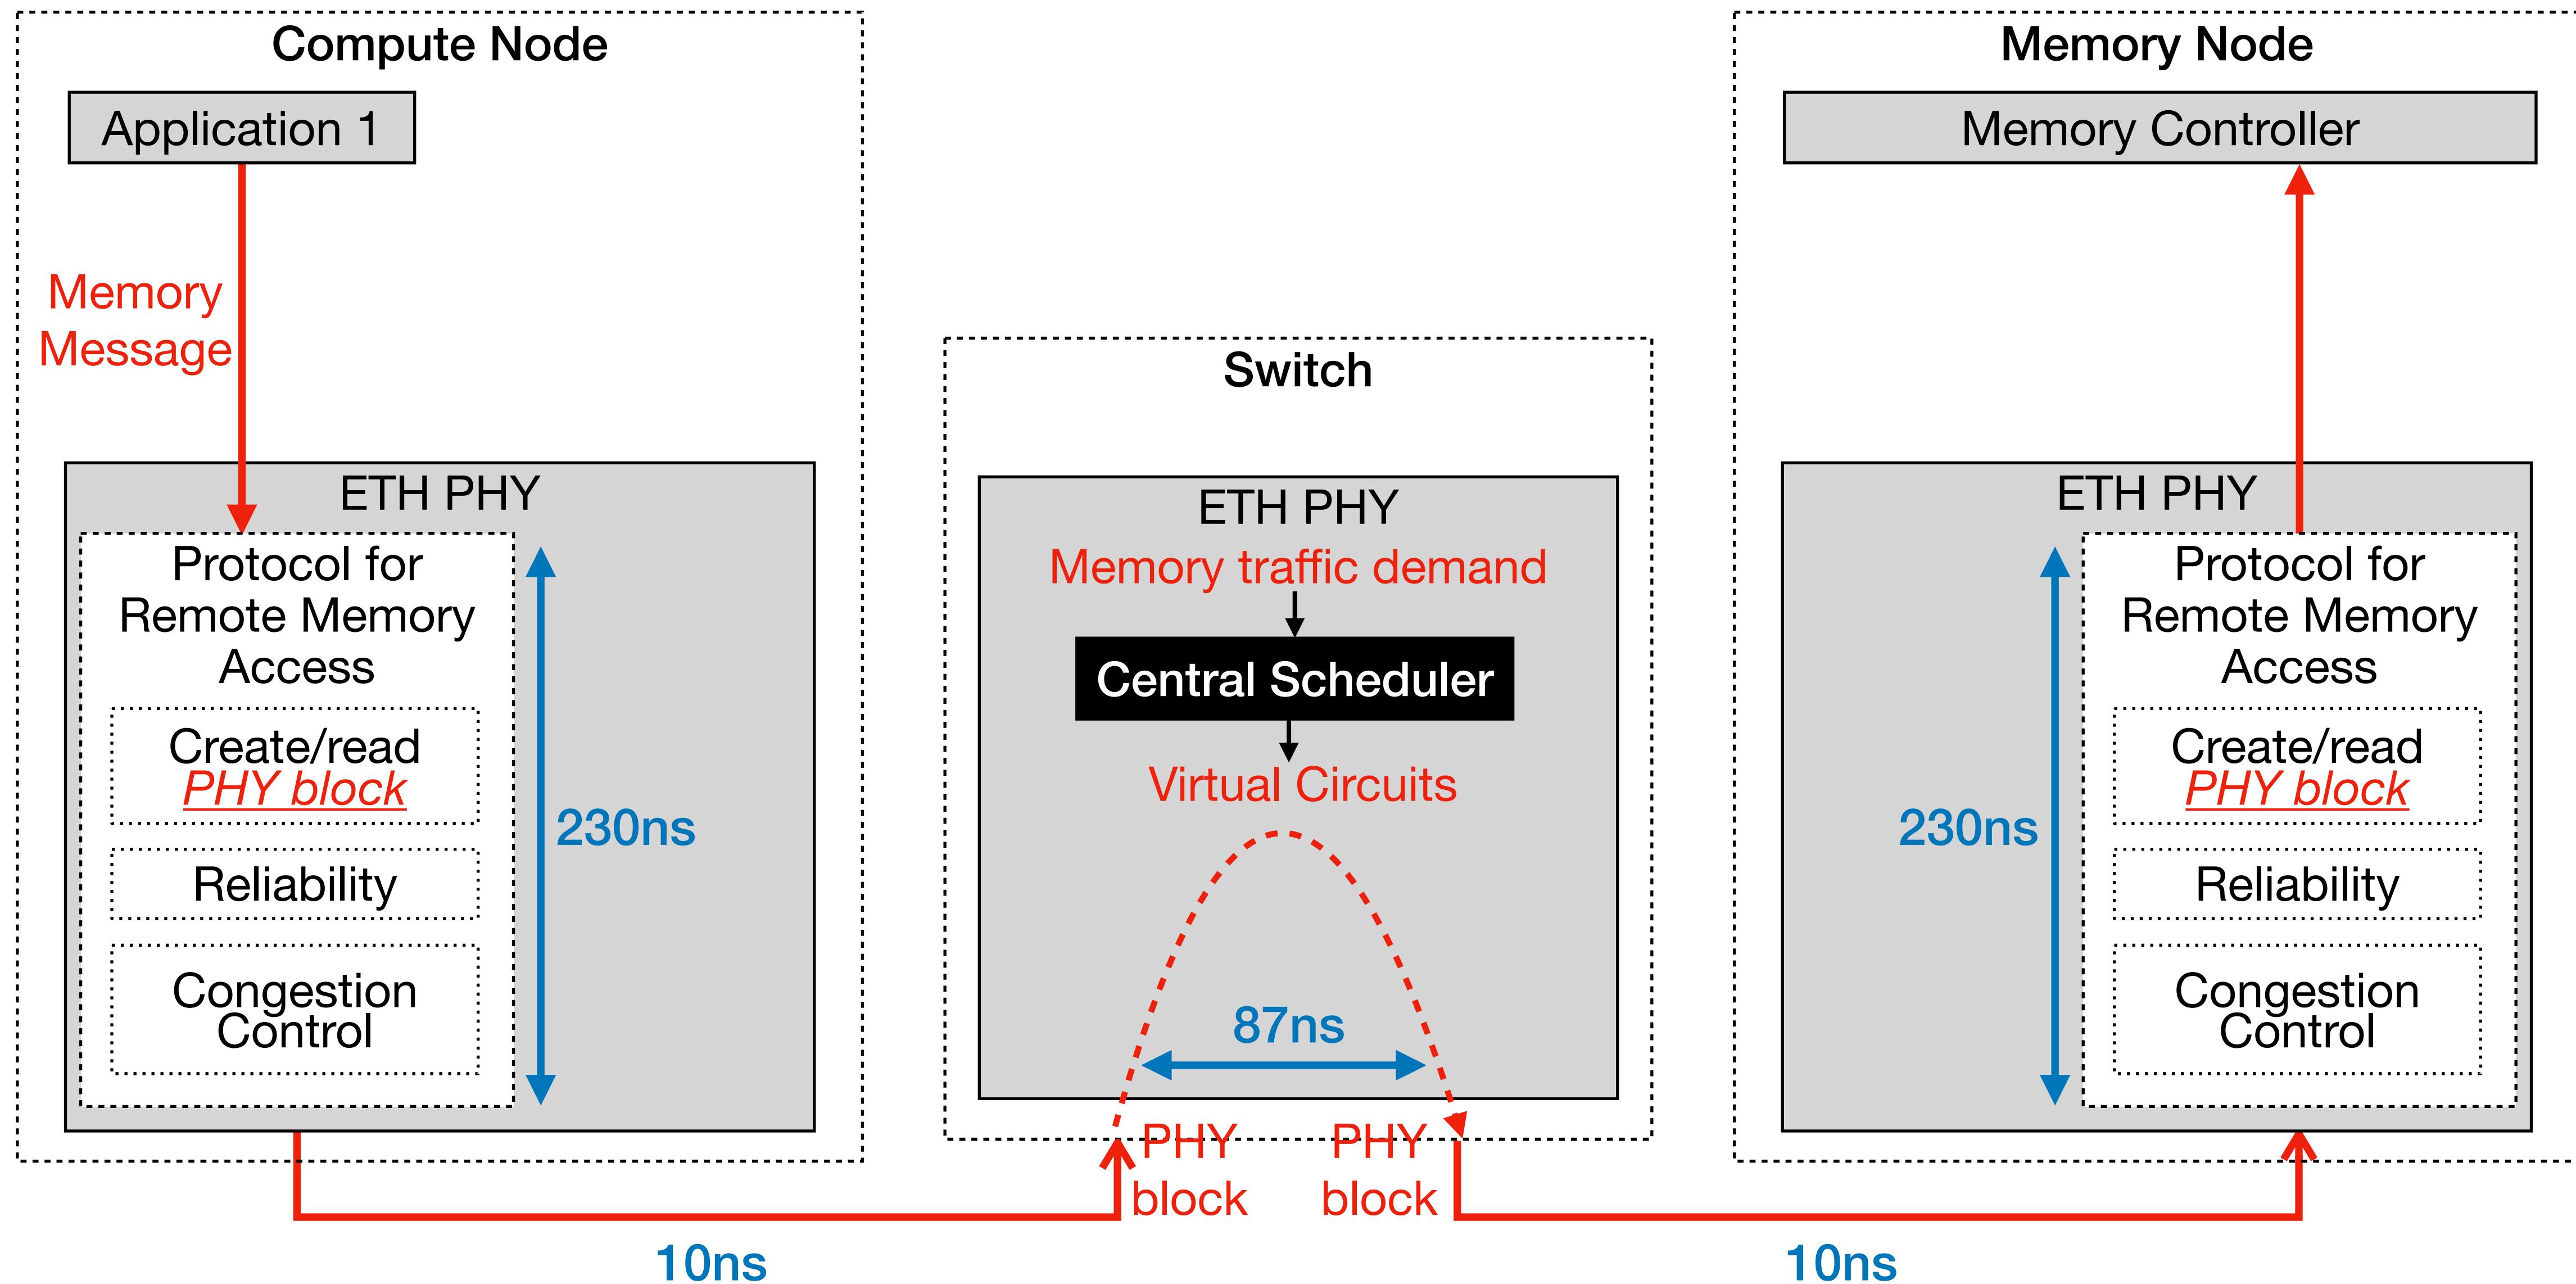
Central Scheduler in the Switch PHY



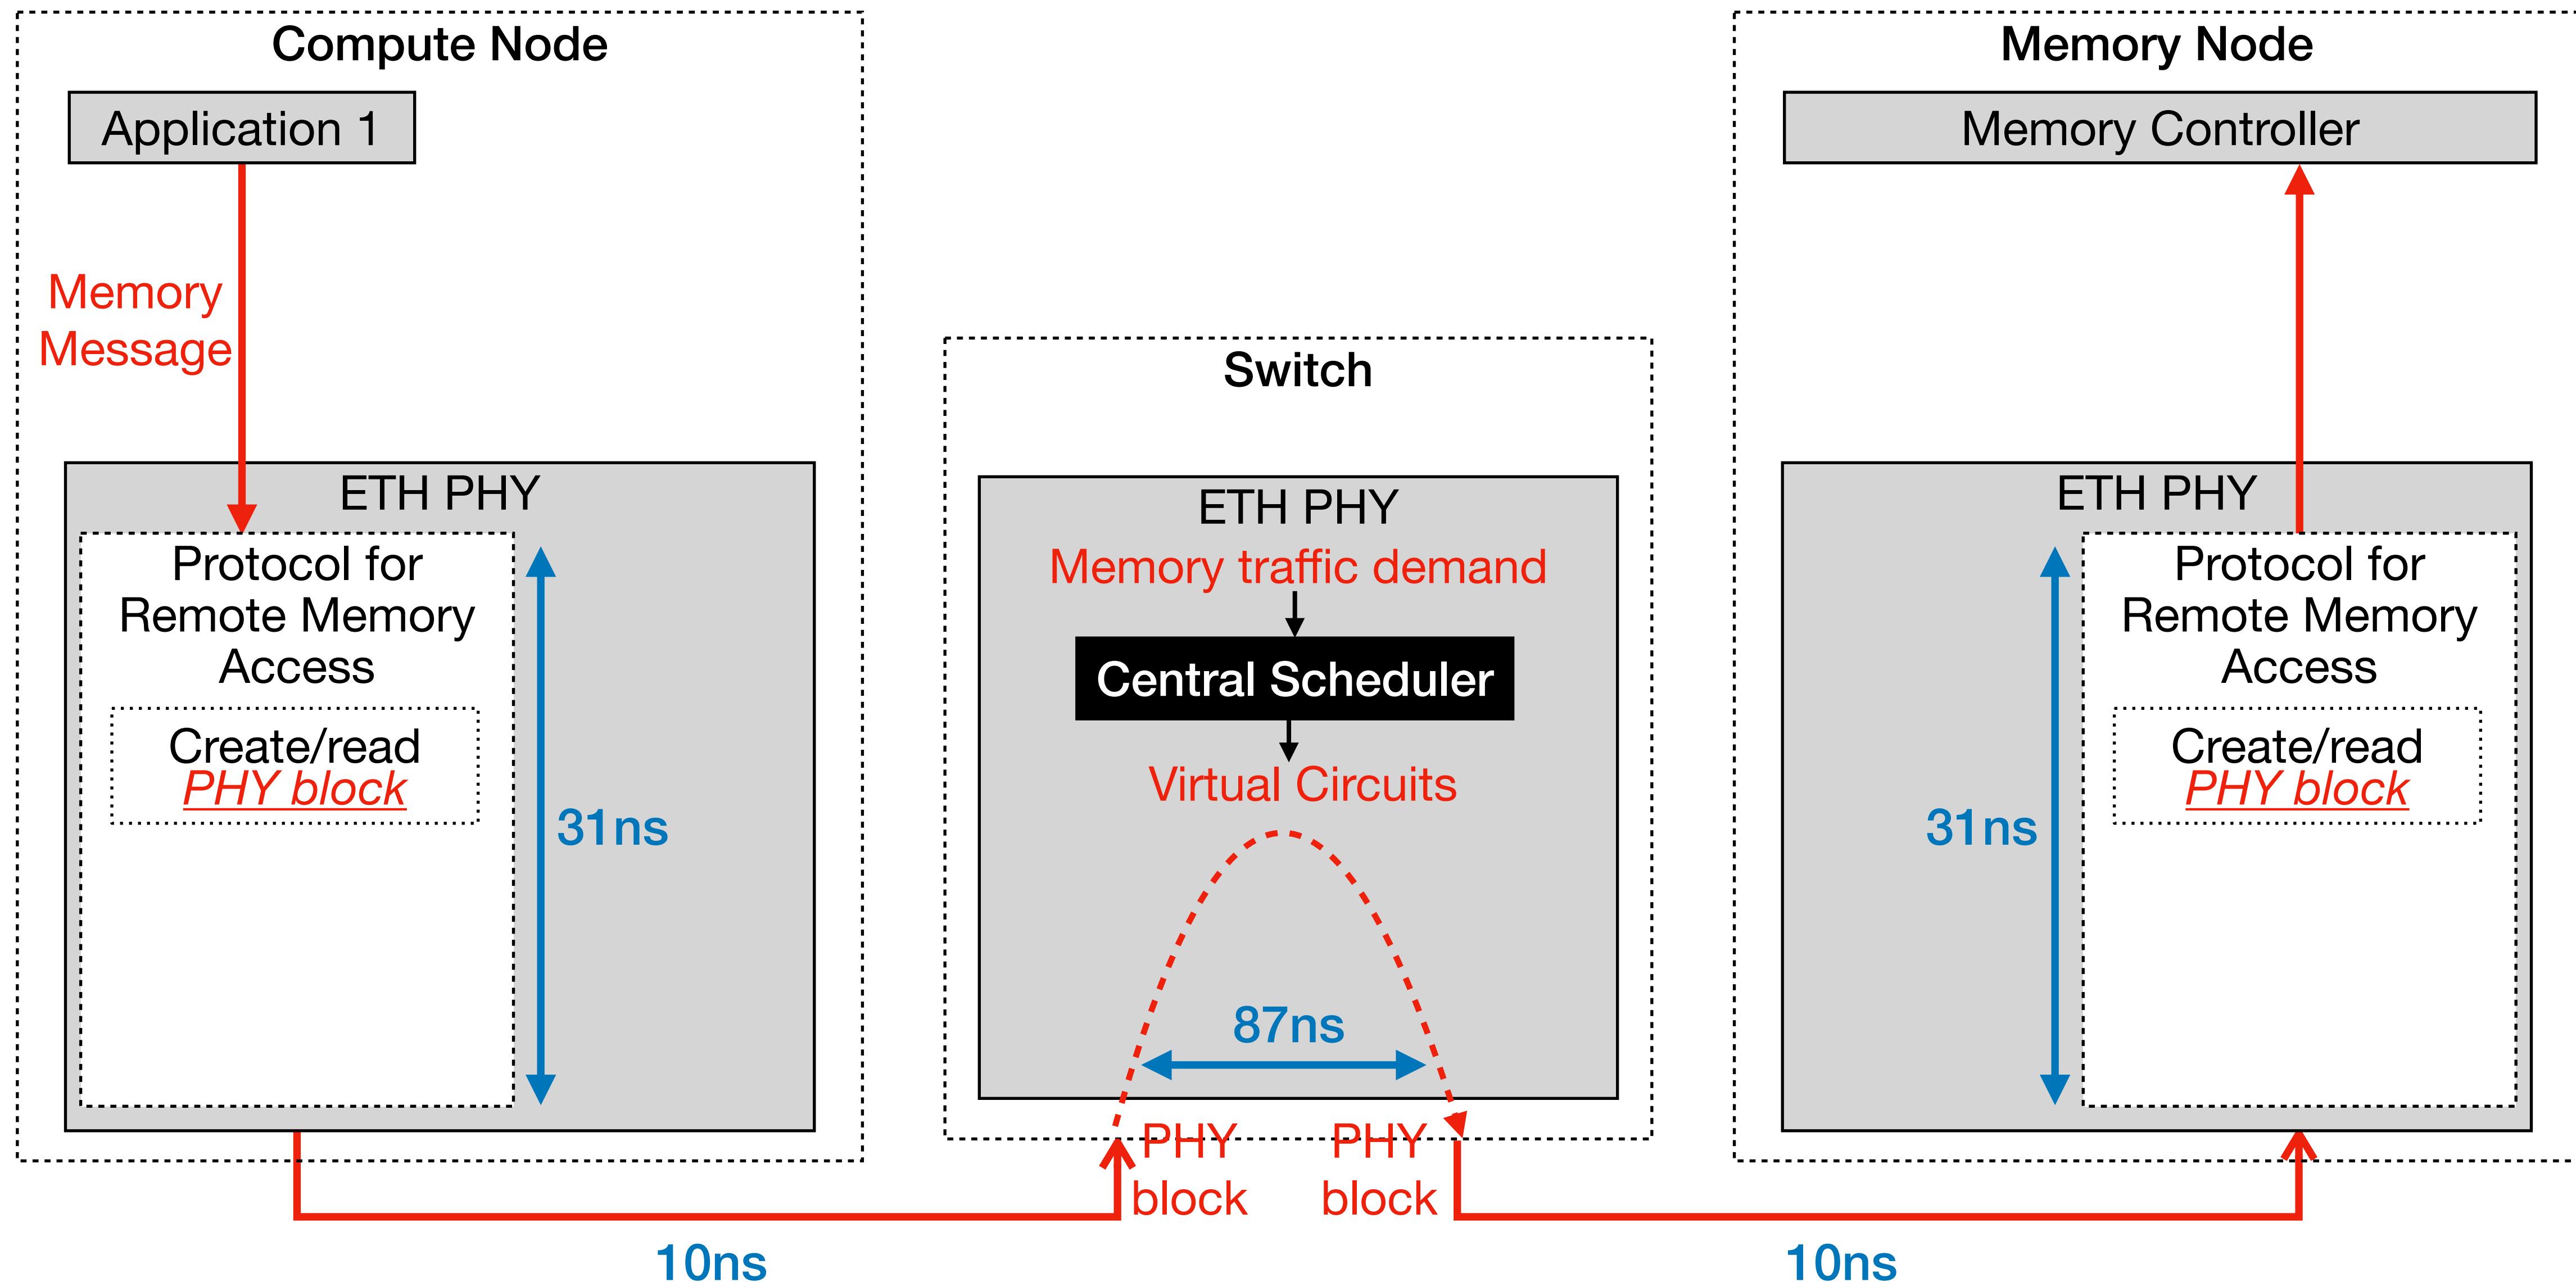
Central Scheduler in the Switch PHY



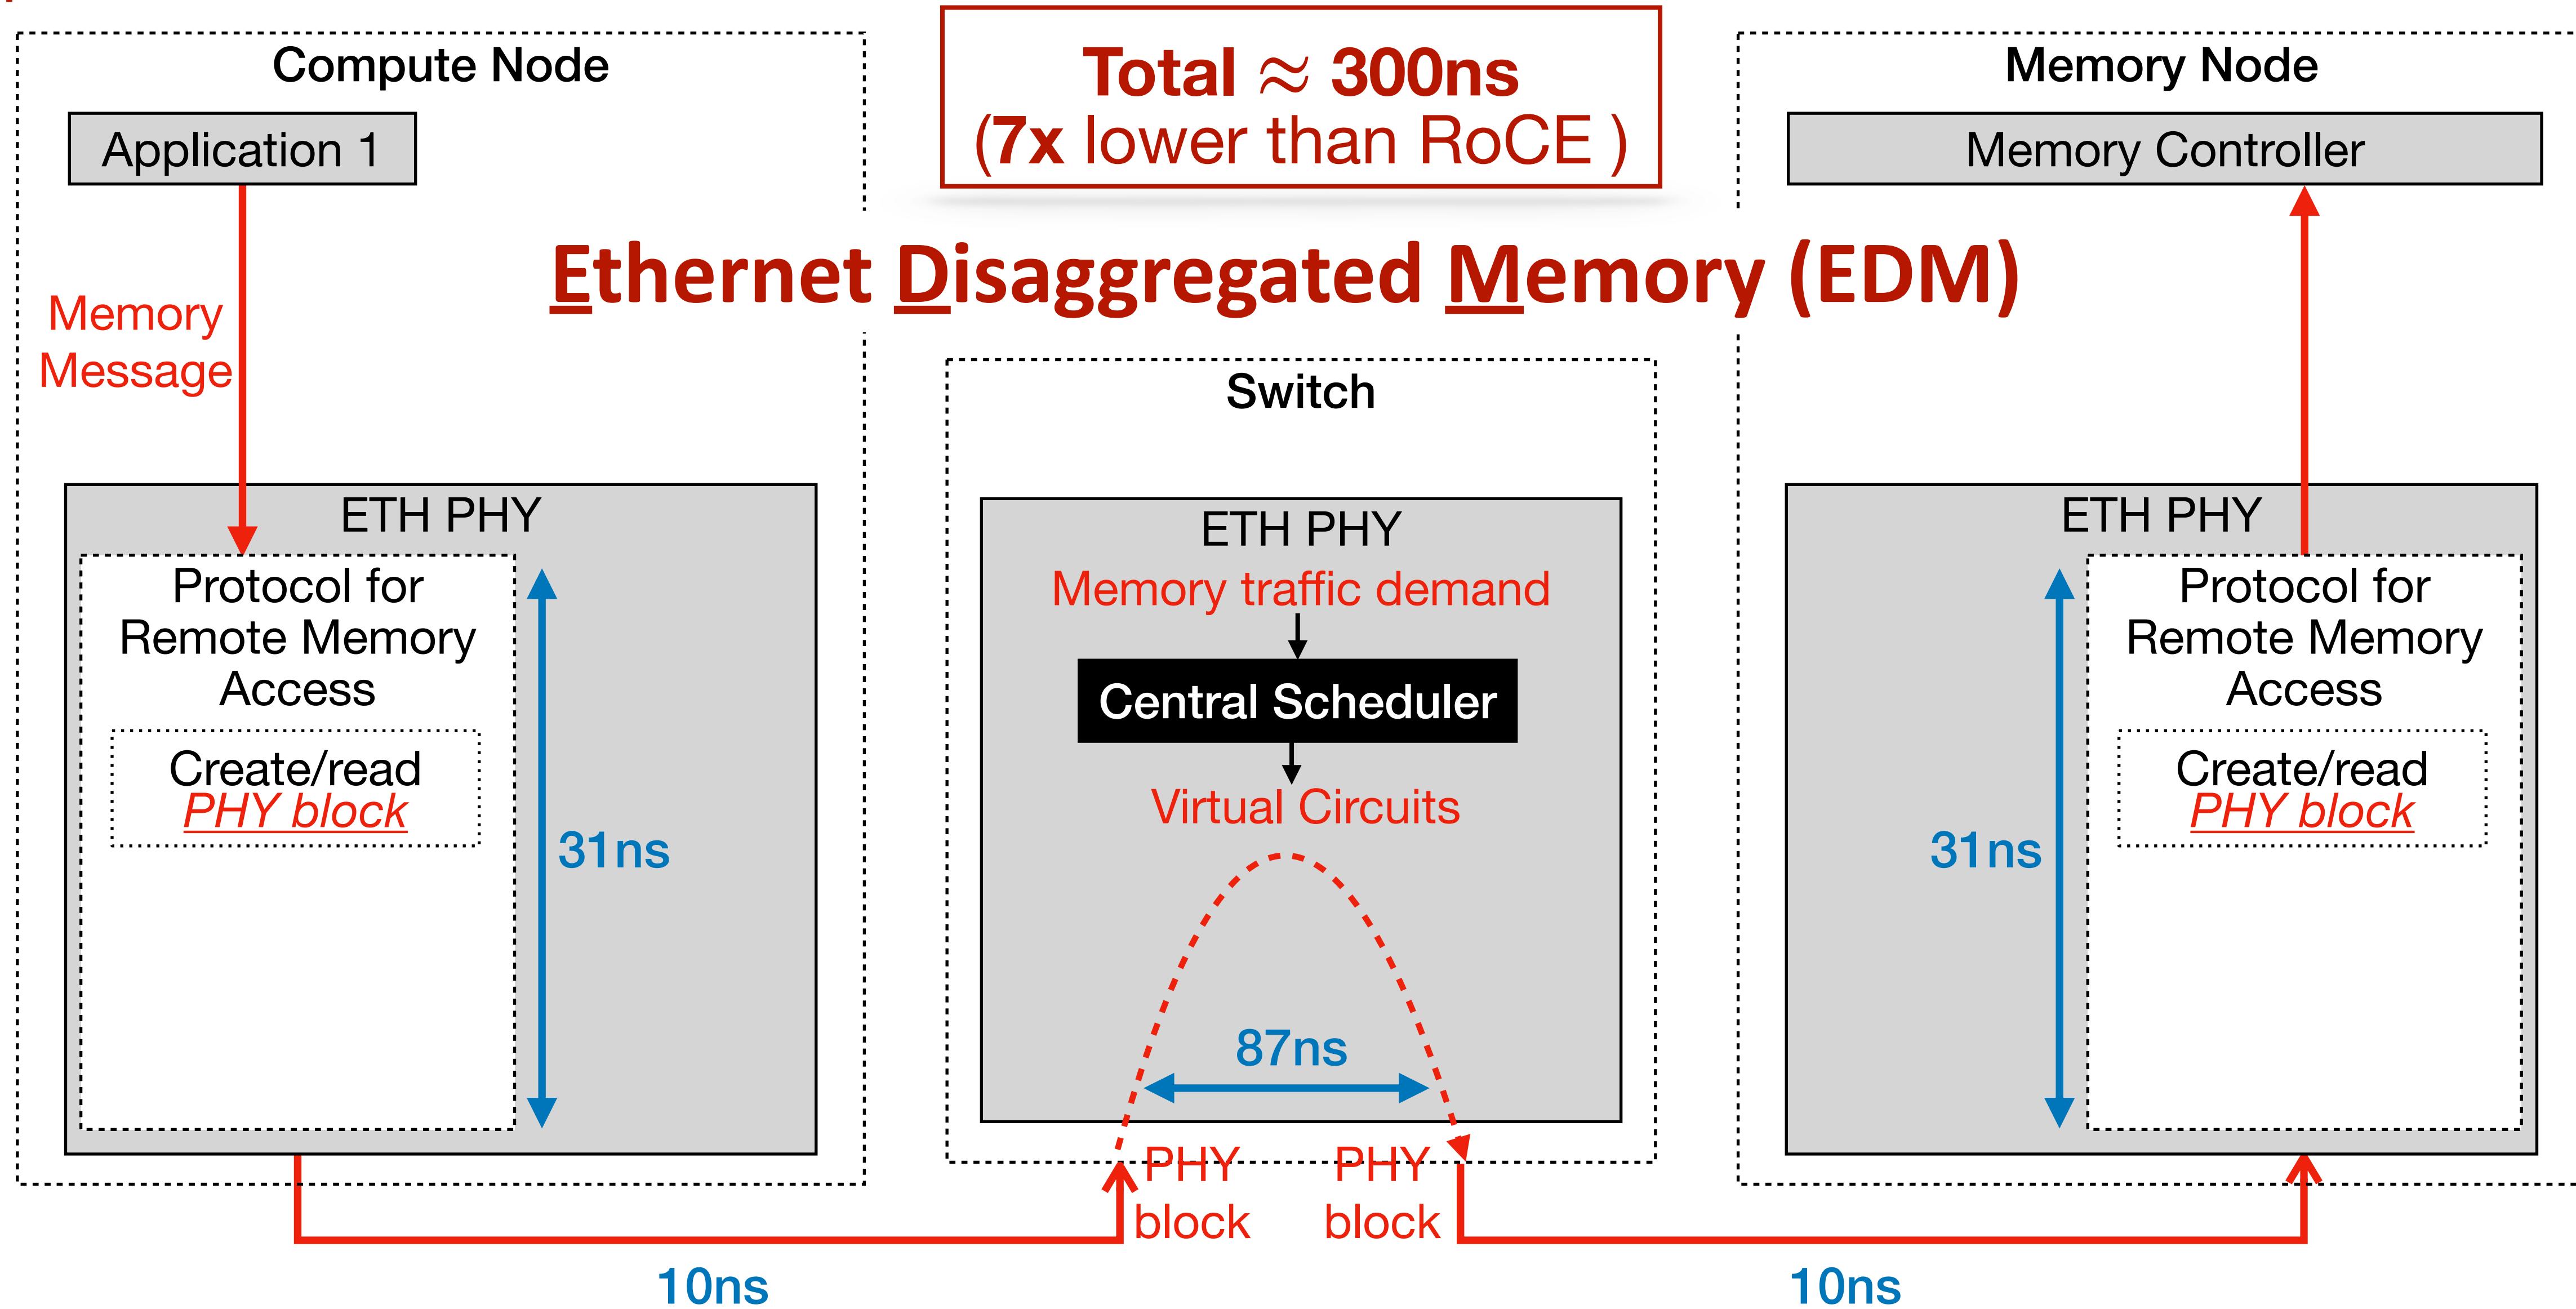
Central Scheduler in the Switch PHY



Central Scheduler in the Switch PHY



Central Scheduler in the Switch PHY



Challenges w/ Central Scheduler

- **Challenge 1:** Accurate traffic demand estimation
- **Challenge 2:** Send demands to the switch with low bandwidth, latency overhead
- **Challenge 3:** Line rate, low latency scheduler pipeline

Challenges w/ Central Scheduler

- **Challenge 1:** Accurate traffic demand estimation
 - Solution: Leverage the interface to memory controller
- **Challenge 2:** Send demands to the switch with low bandwidth, latency overhead
- **Challenge 3:** Line rate, low latency scheduler pipeline

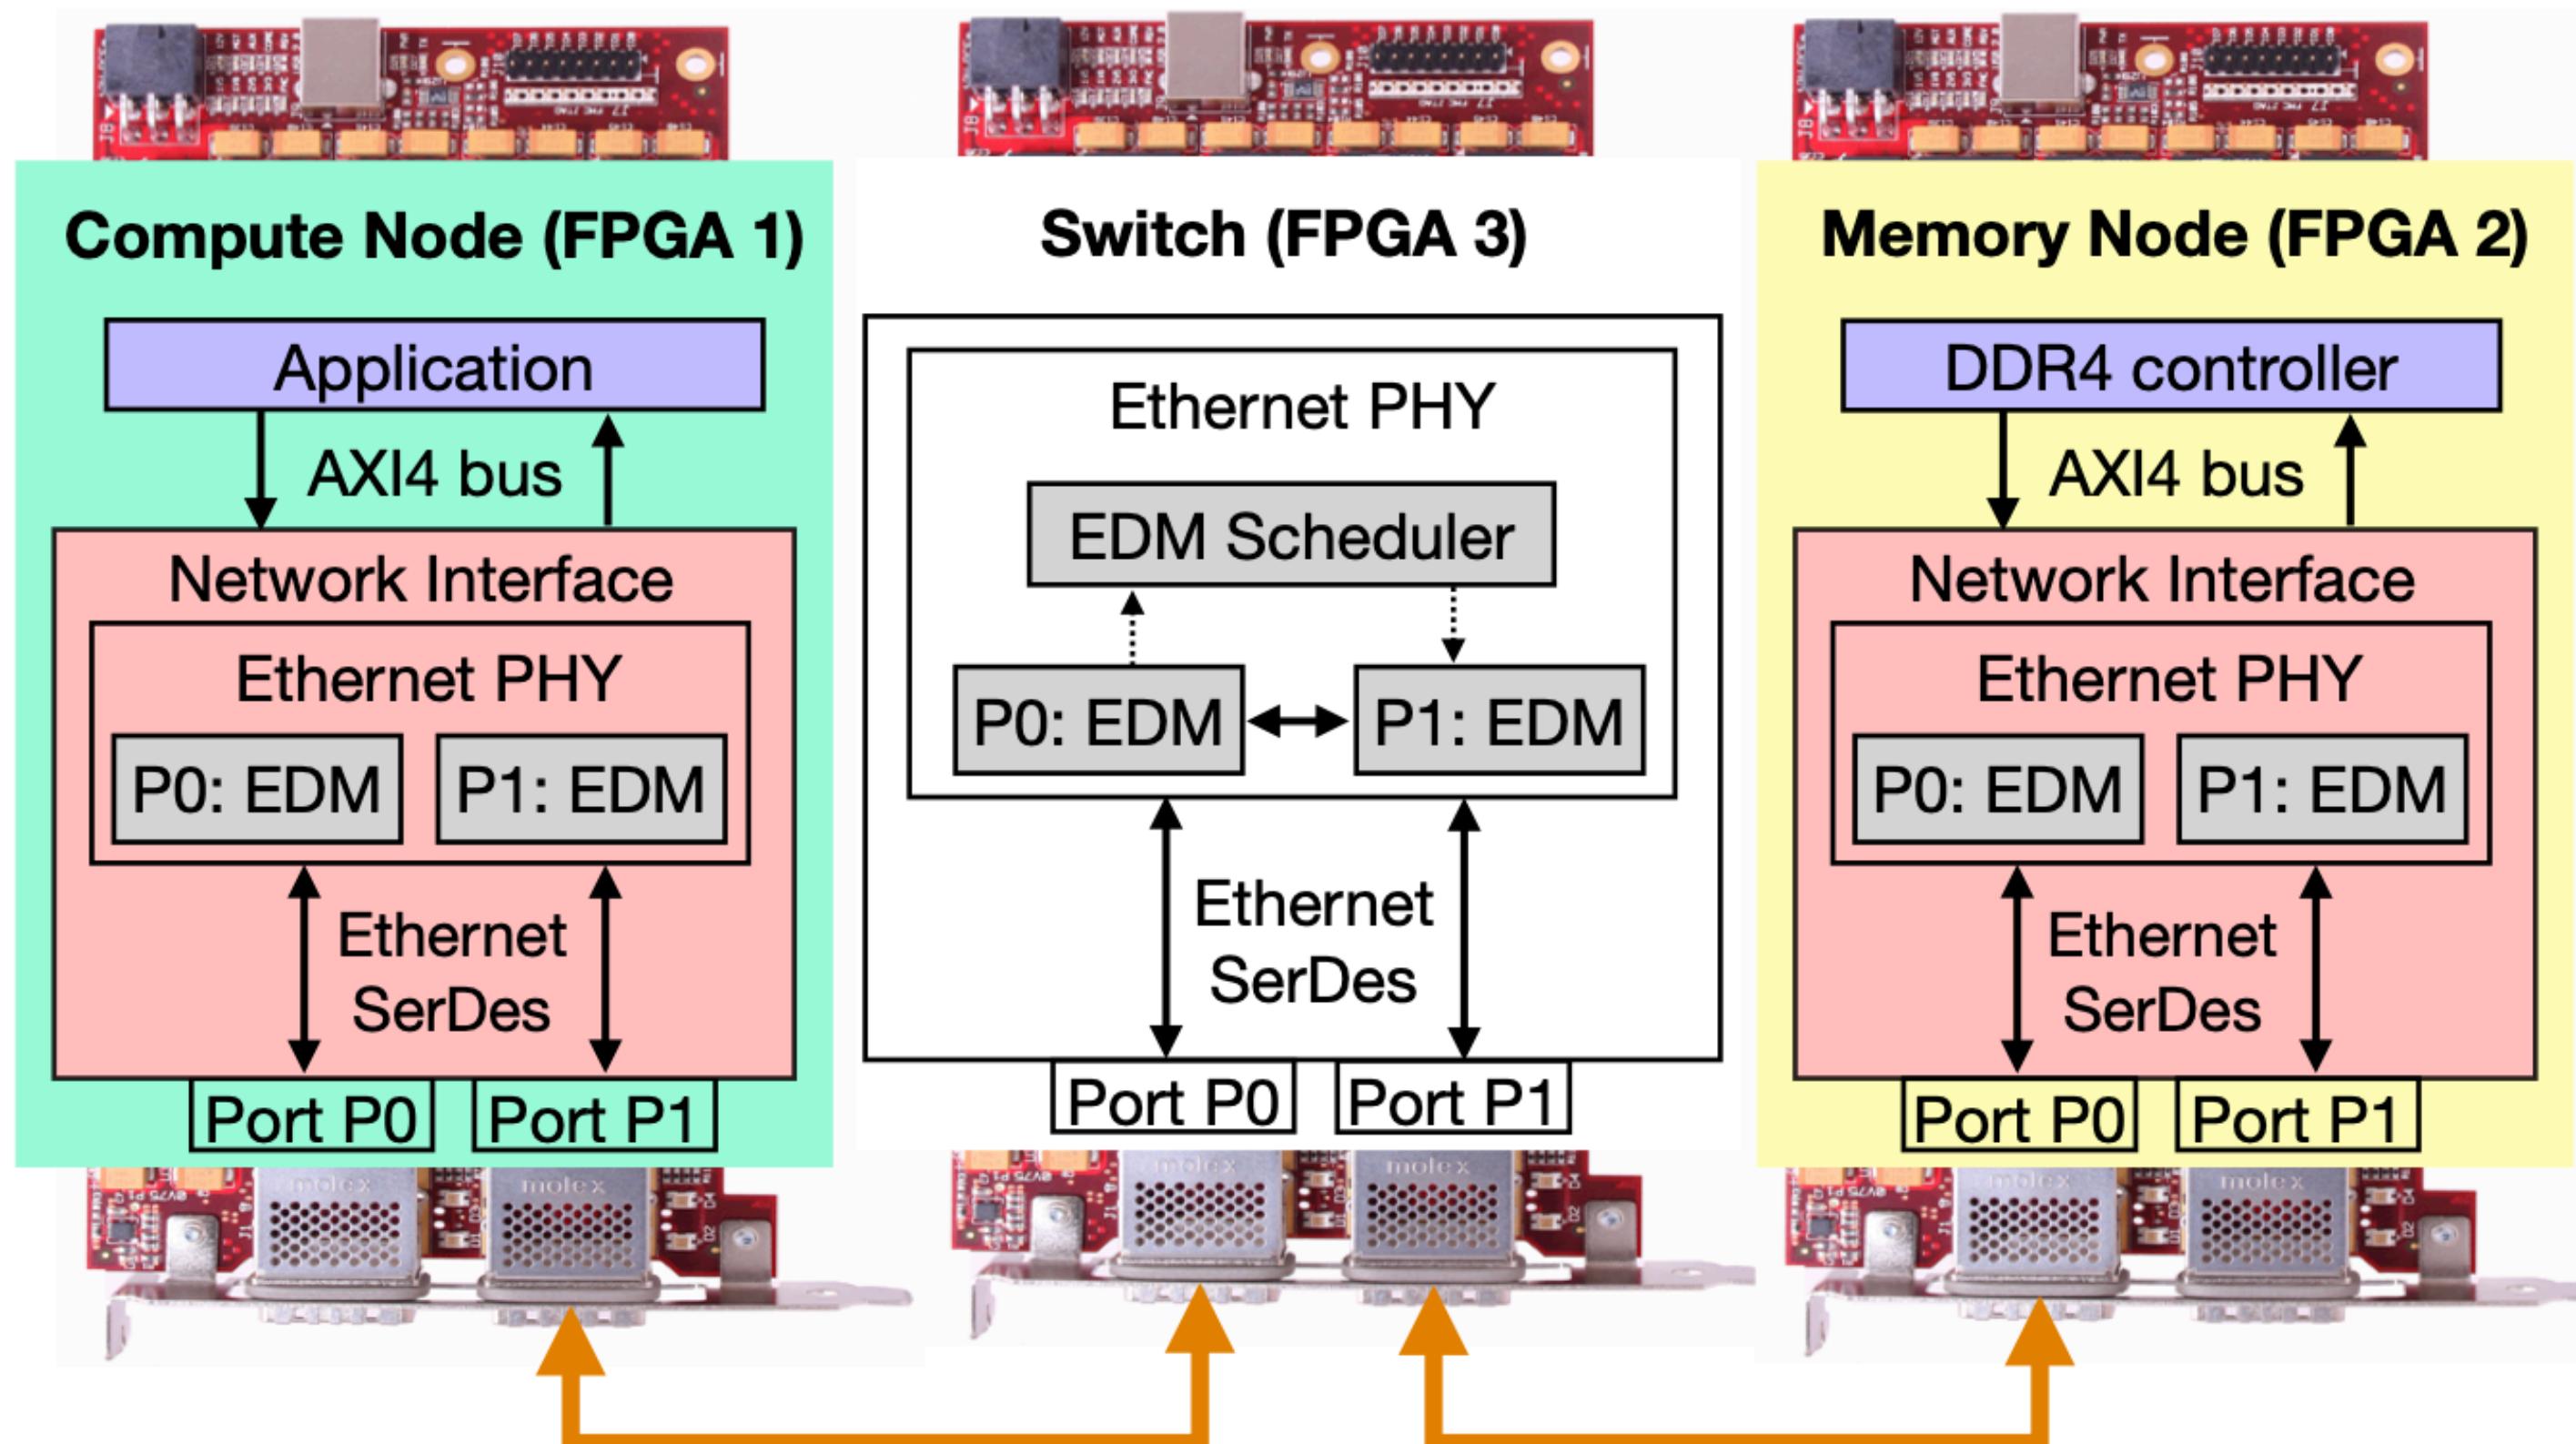
Challenges w/ Central Scheduler

- **Challenge 1:** Accurate traffic demand estimation
 - Solution: Leverage the interface to memory controller
- **Challenge 2:** Send demands to the switch with low bandwidth, latency overhead
 - Solution: Leverage request-reply nature of memory access
- **Challenge 3:** Line rate, low latency scheduler pipeline

Challenges w/ Central Scheduler

- **Challenge 1:** Accurate traffic demand estimation
 - Solution: Leverage the interface to memory controller
- **Challenge 2:** Send demands to the switch with low bandwidth, latency overhead
 - Solution: Leverage request-reply nature of memory access
- **Challenge 3:** Line rate, low latency scheduler pipeline
 - Solution: Leverage hardware parallelism in switch's PHY

Implementation & Evaluation

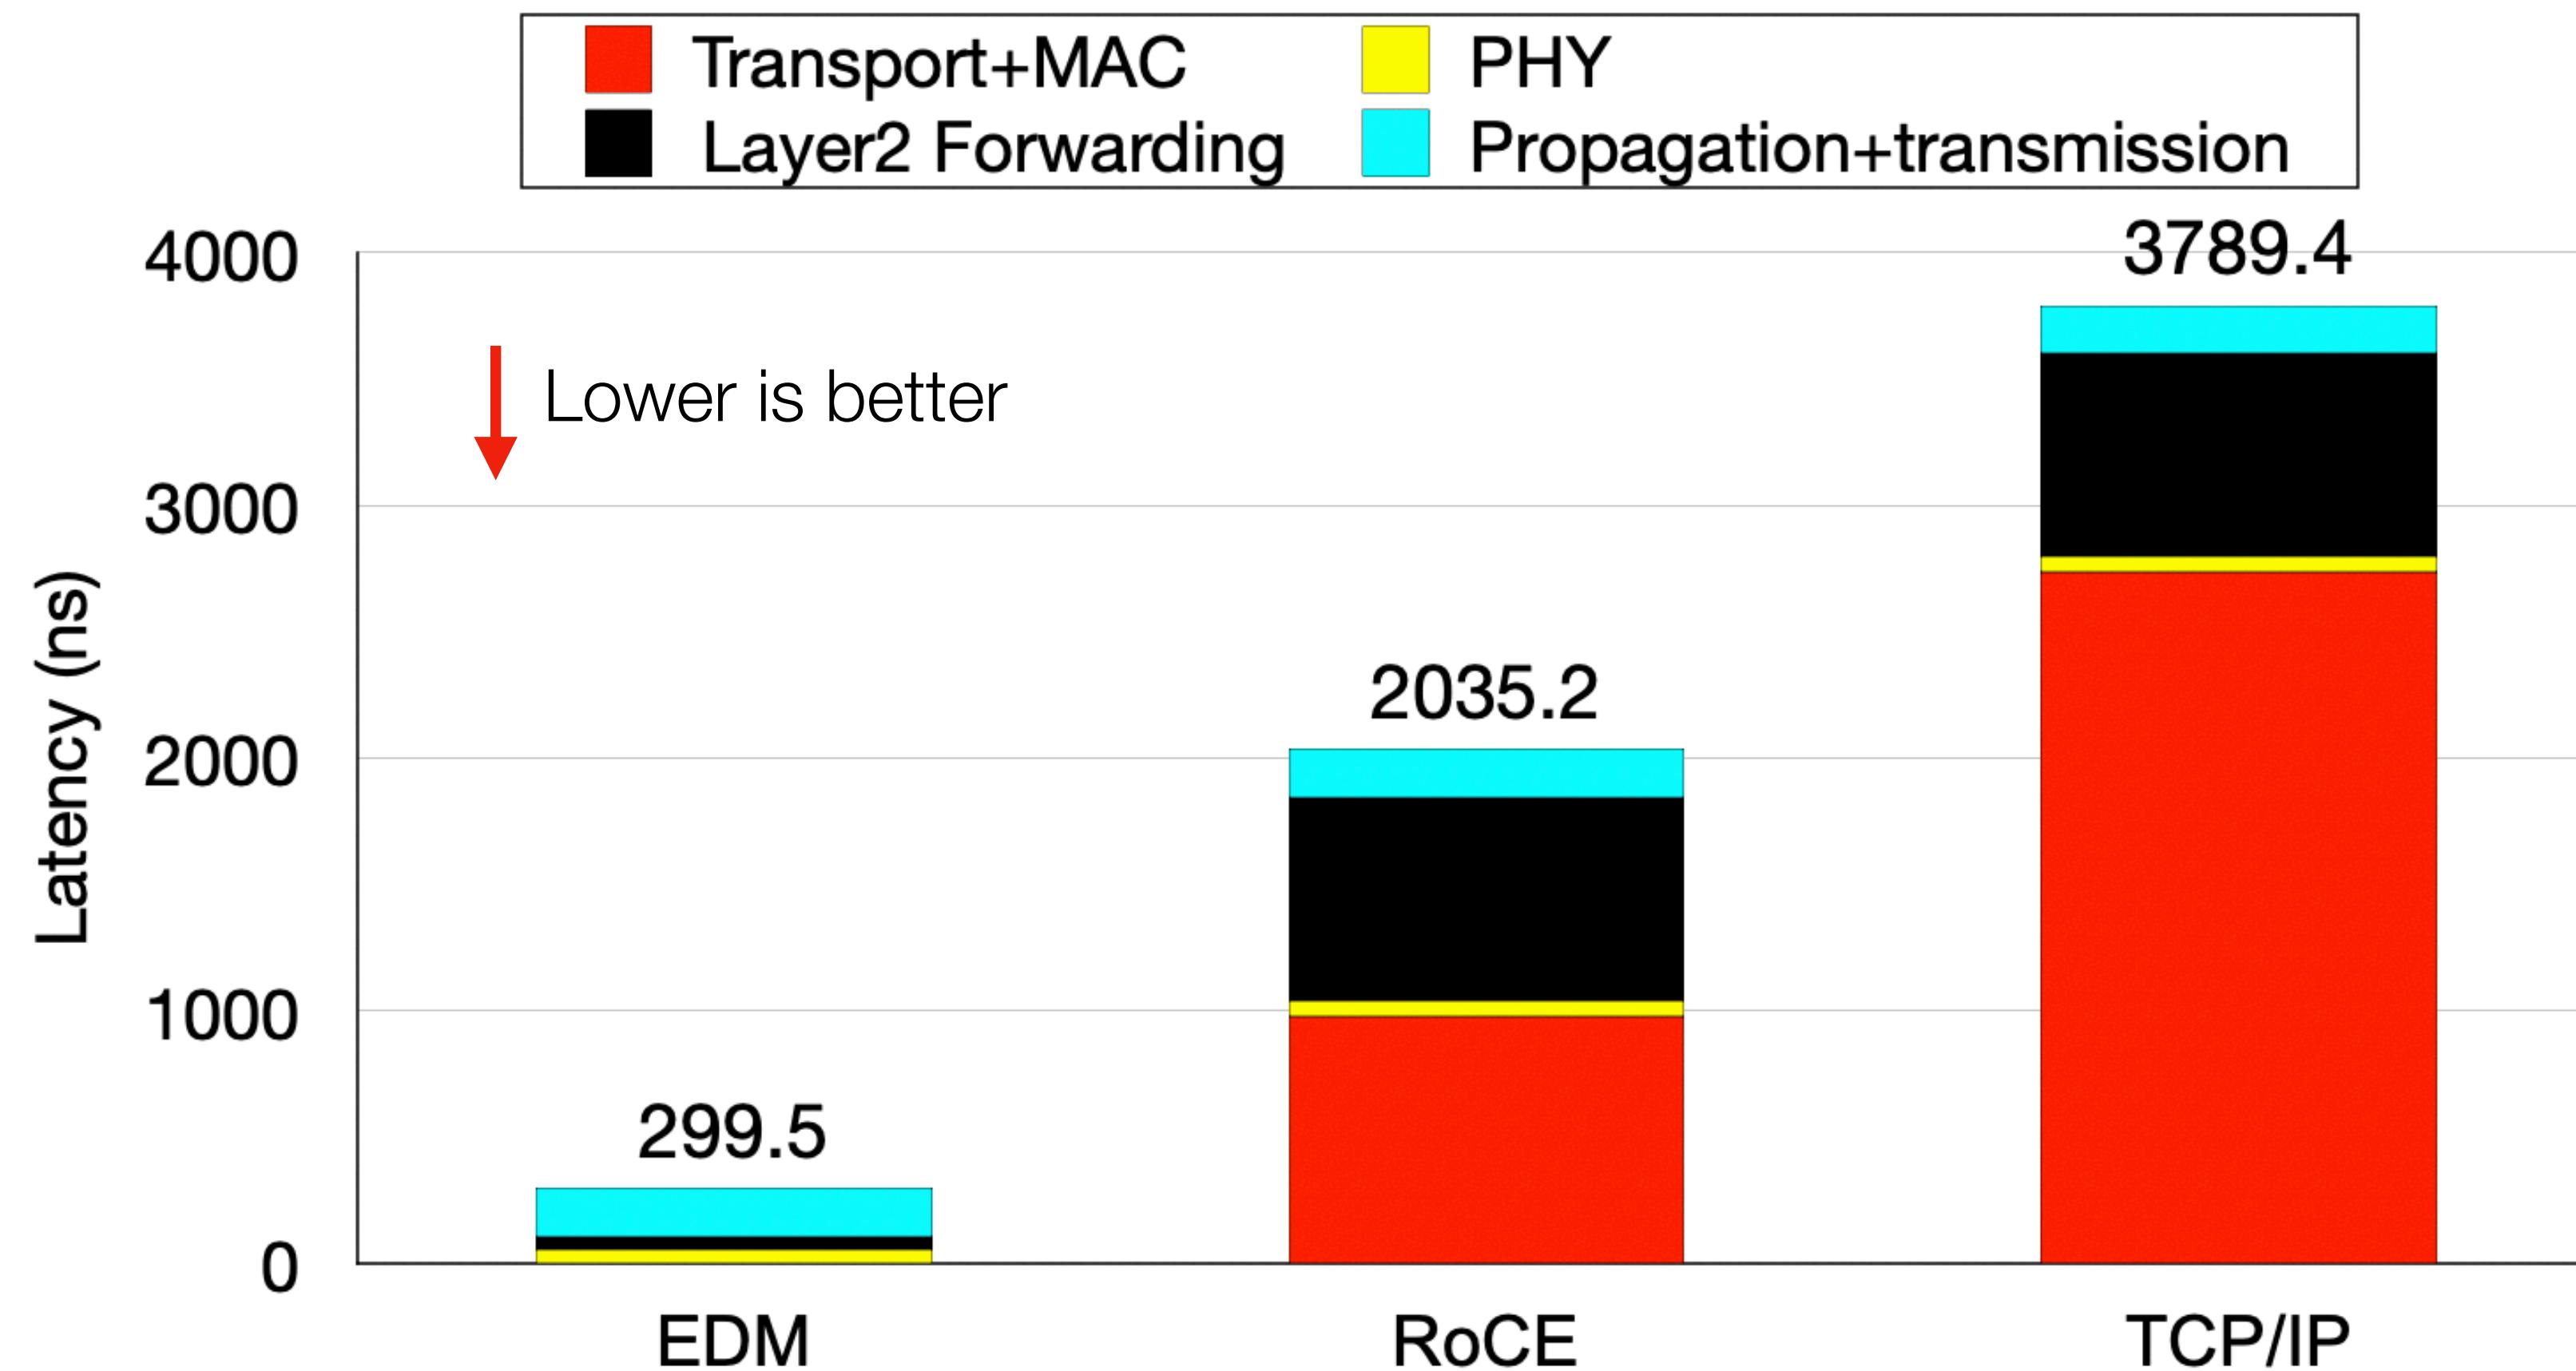


Hardware Testbed

- Three Xilinx Alveo U200 FPGAs
- Open-source 25GbE (Corundum)
- Synopsys ASIC RTL compiler

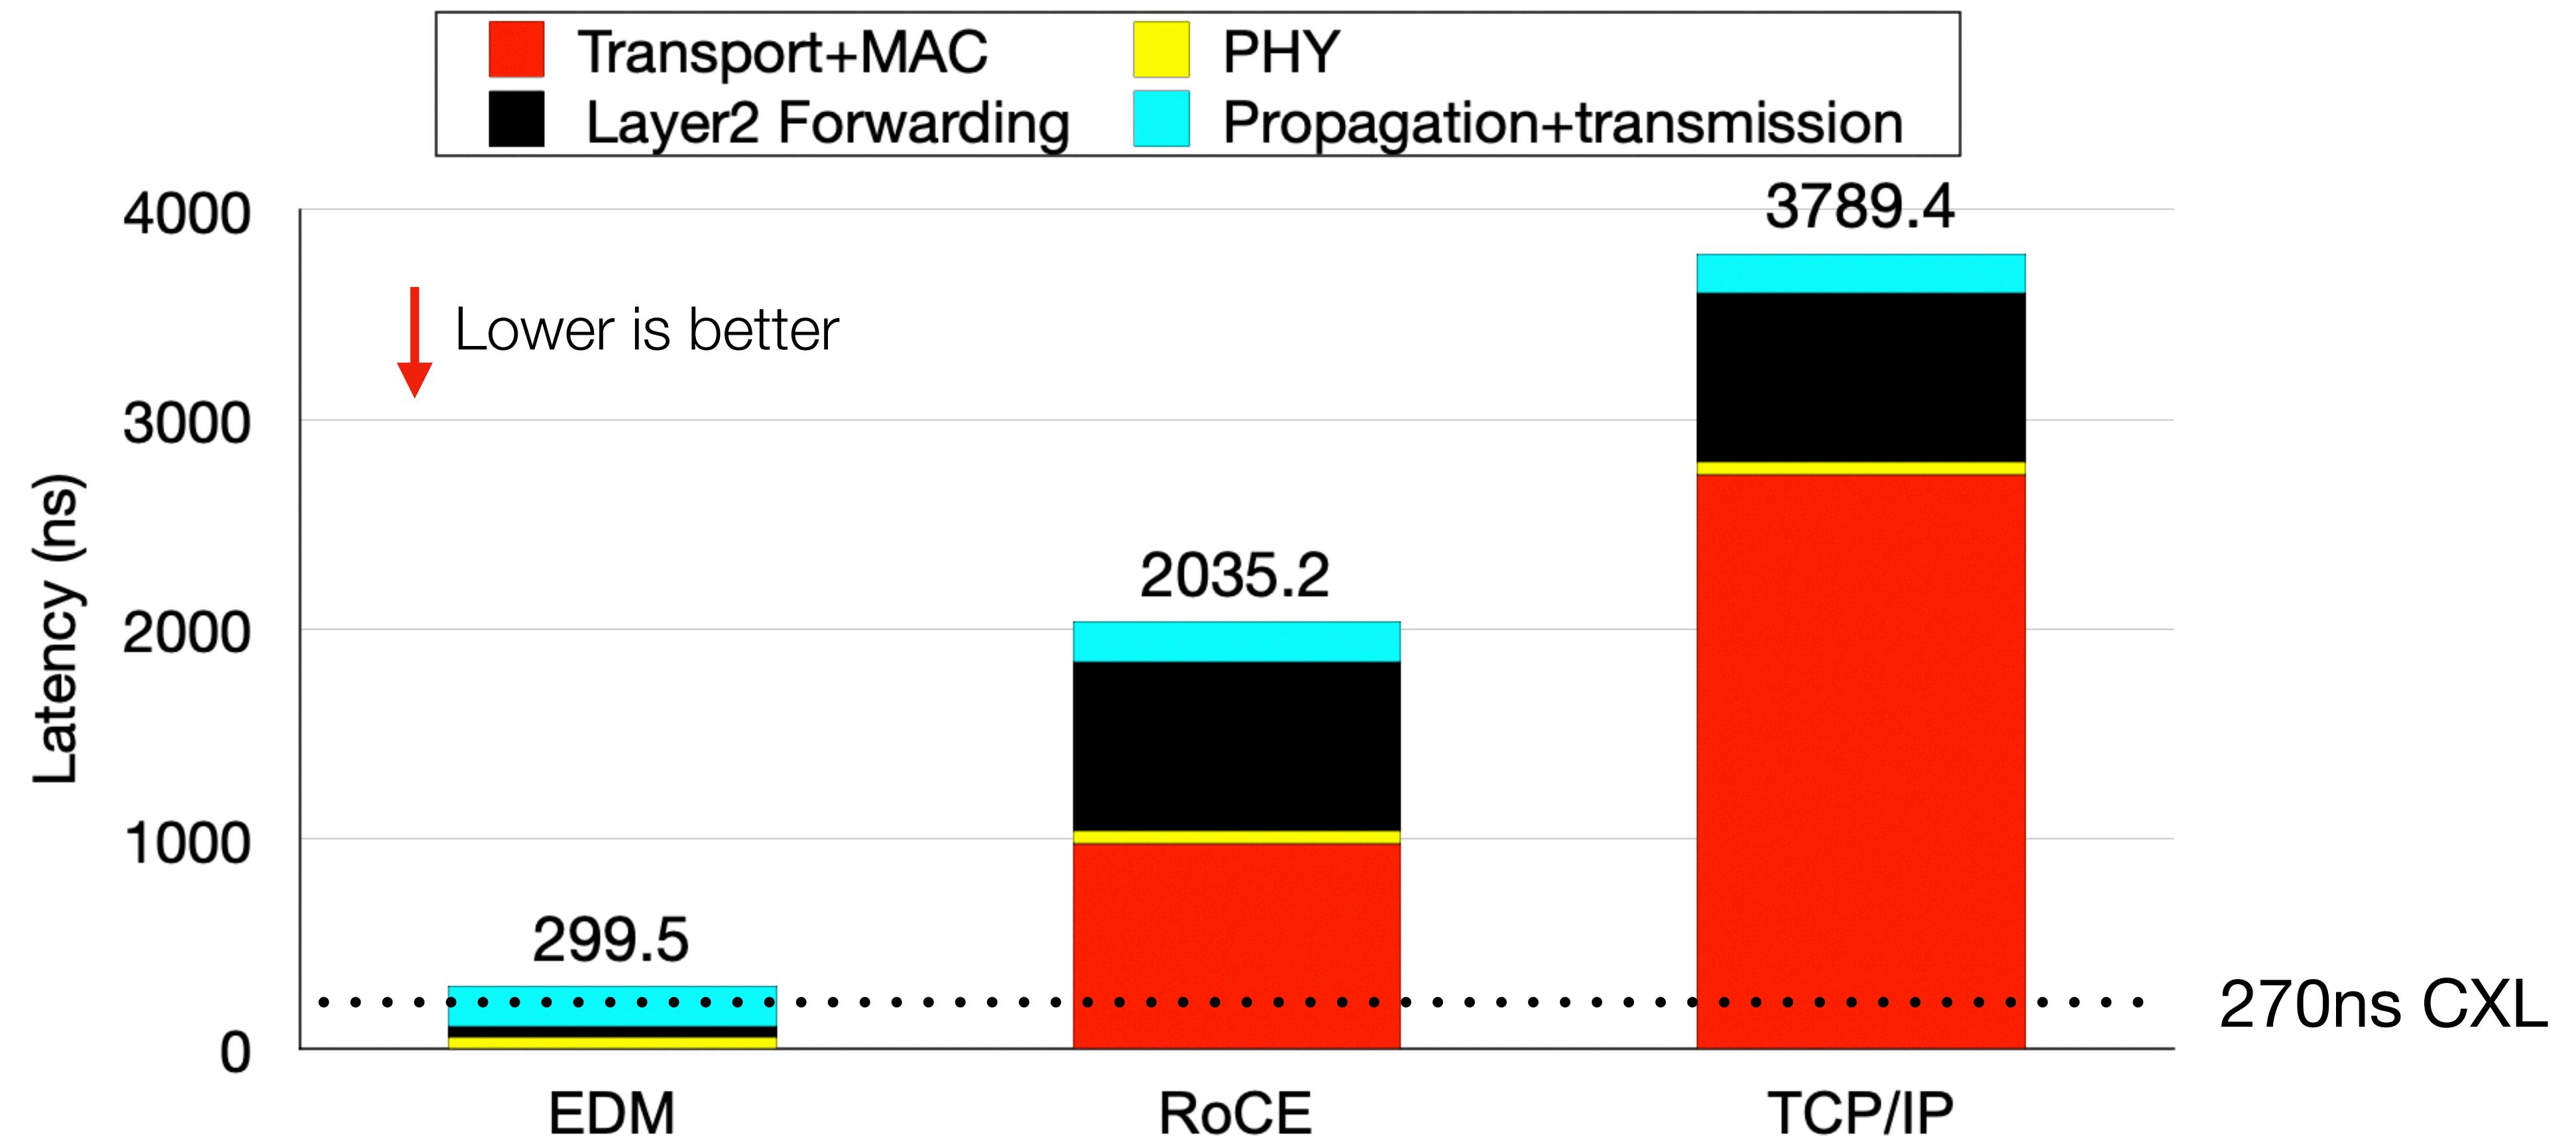
Evaluation Result

- End-to-end unloaded latency

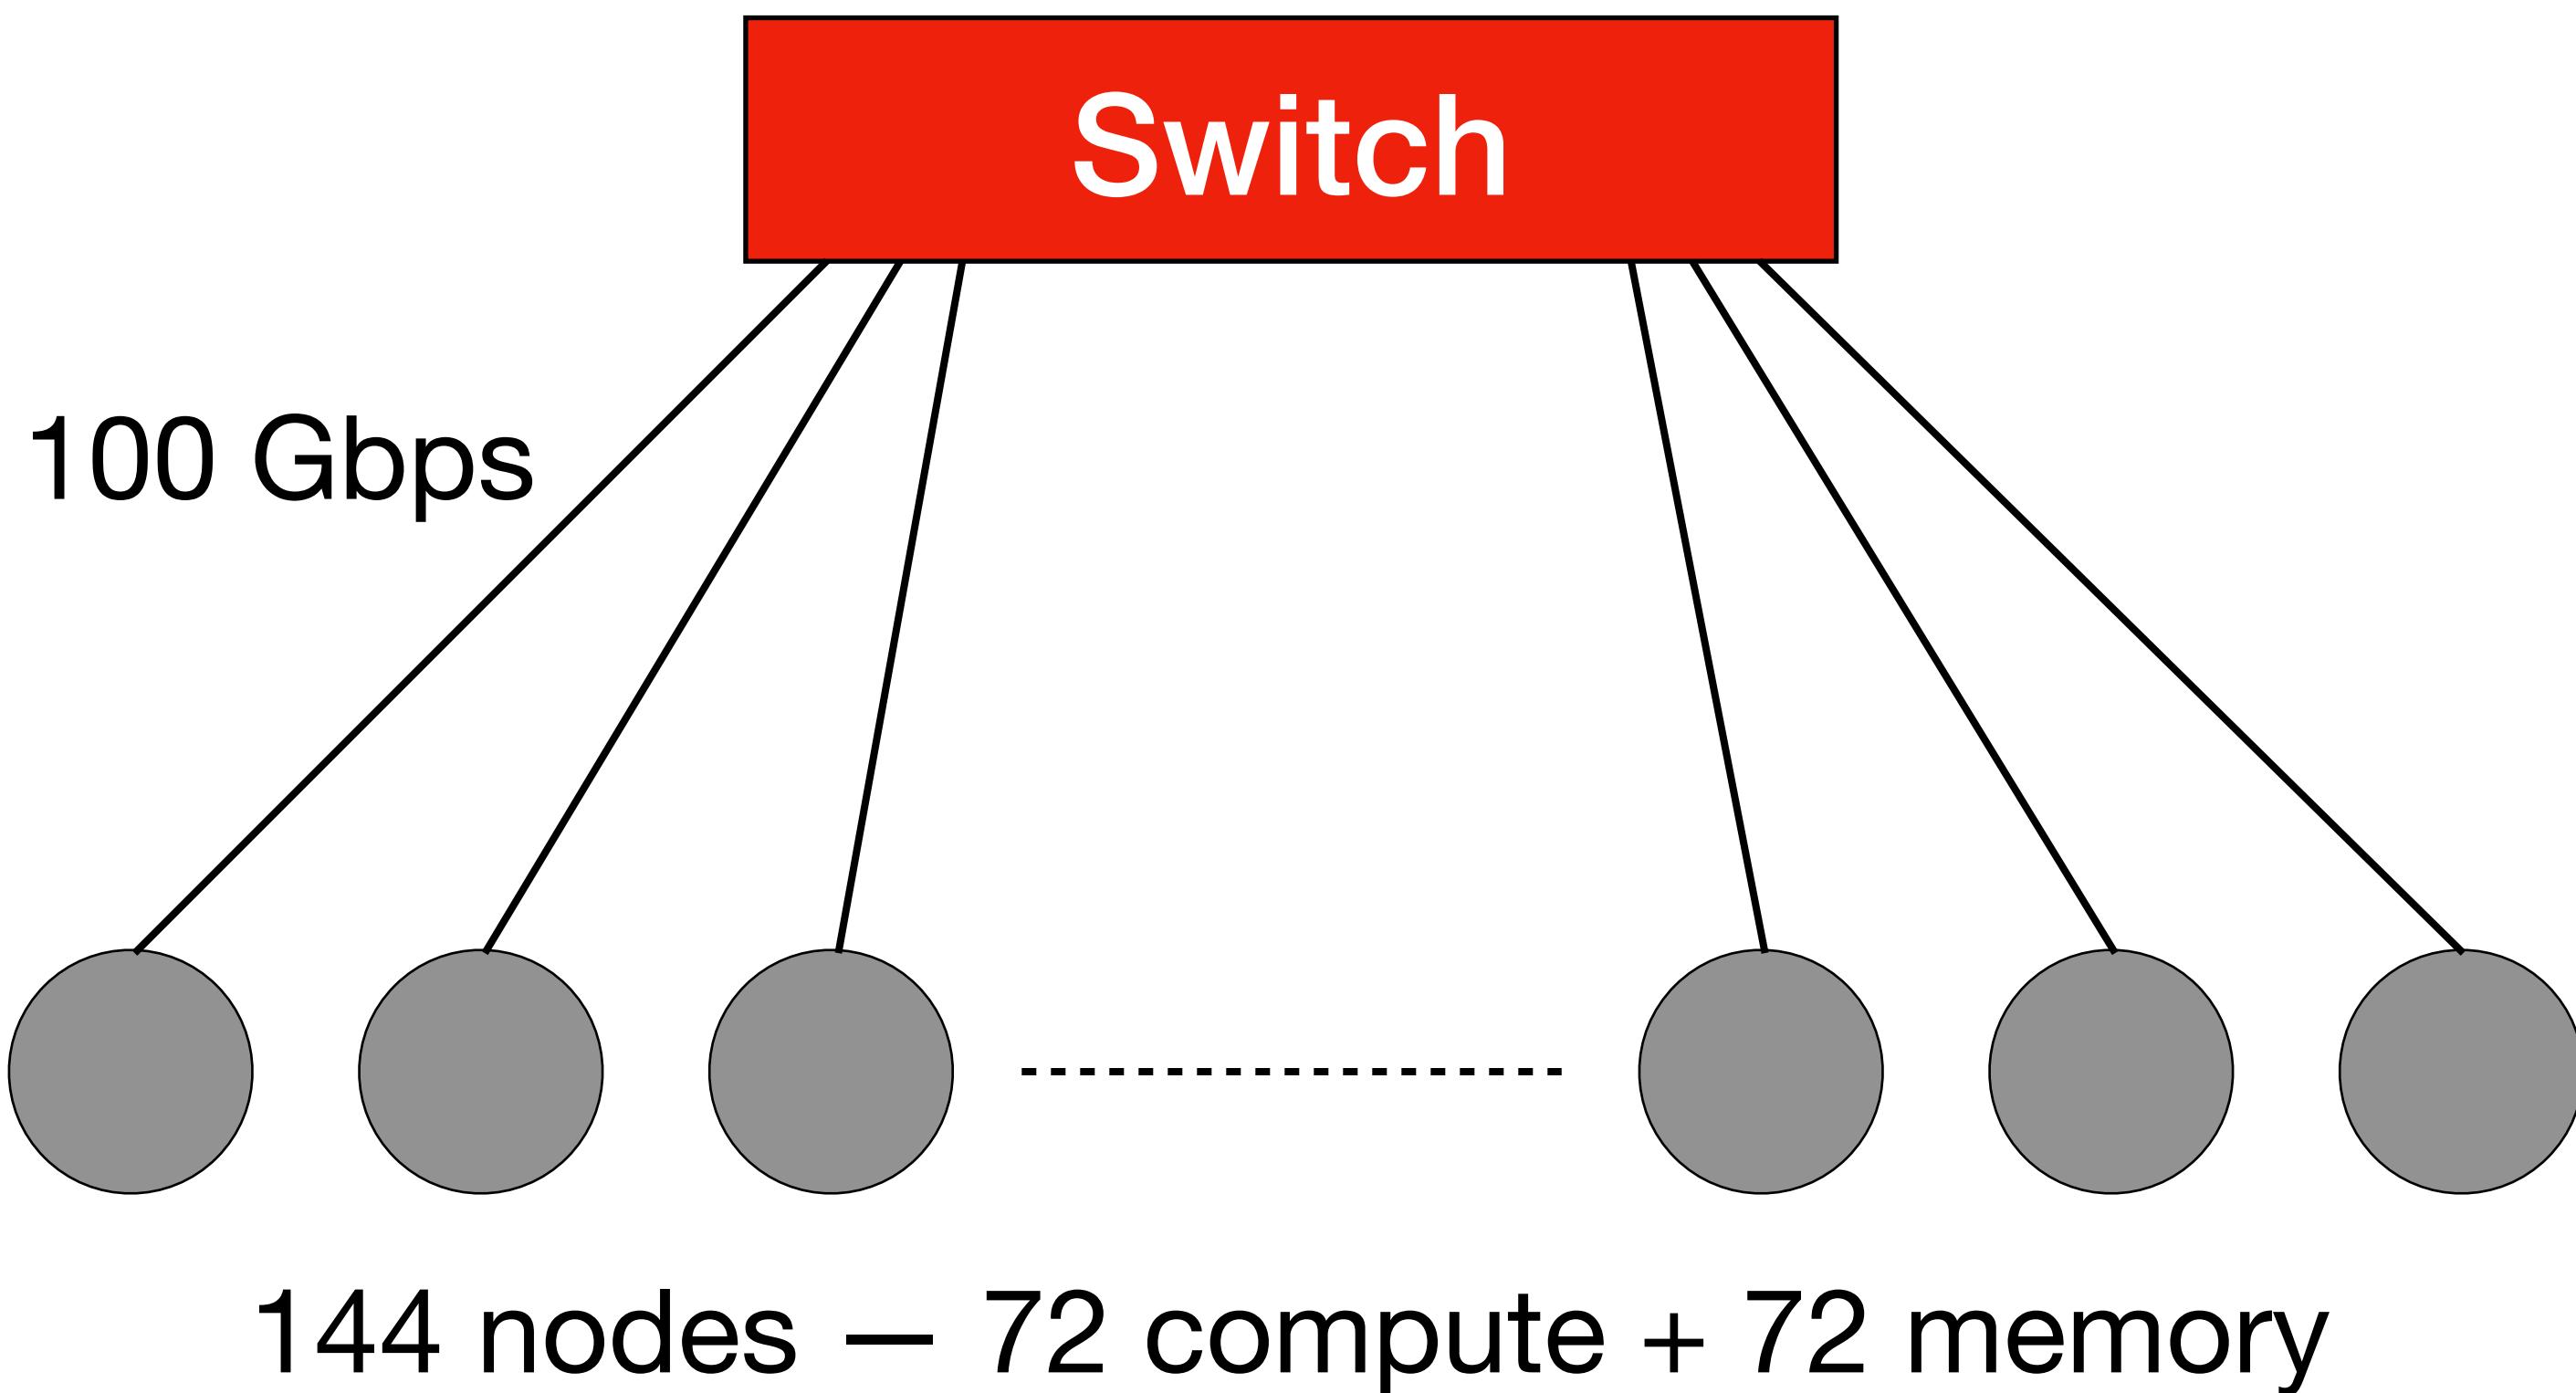


Evaluation Result

- End-to-end unloaded latency



Implementation & Evaluation



Hardware Testbed

- Three Xilinx Alveo U200 FPGAs
- Open-source 25GbE (Corundum)
- Synopsys ASIC RTL compiler

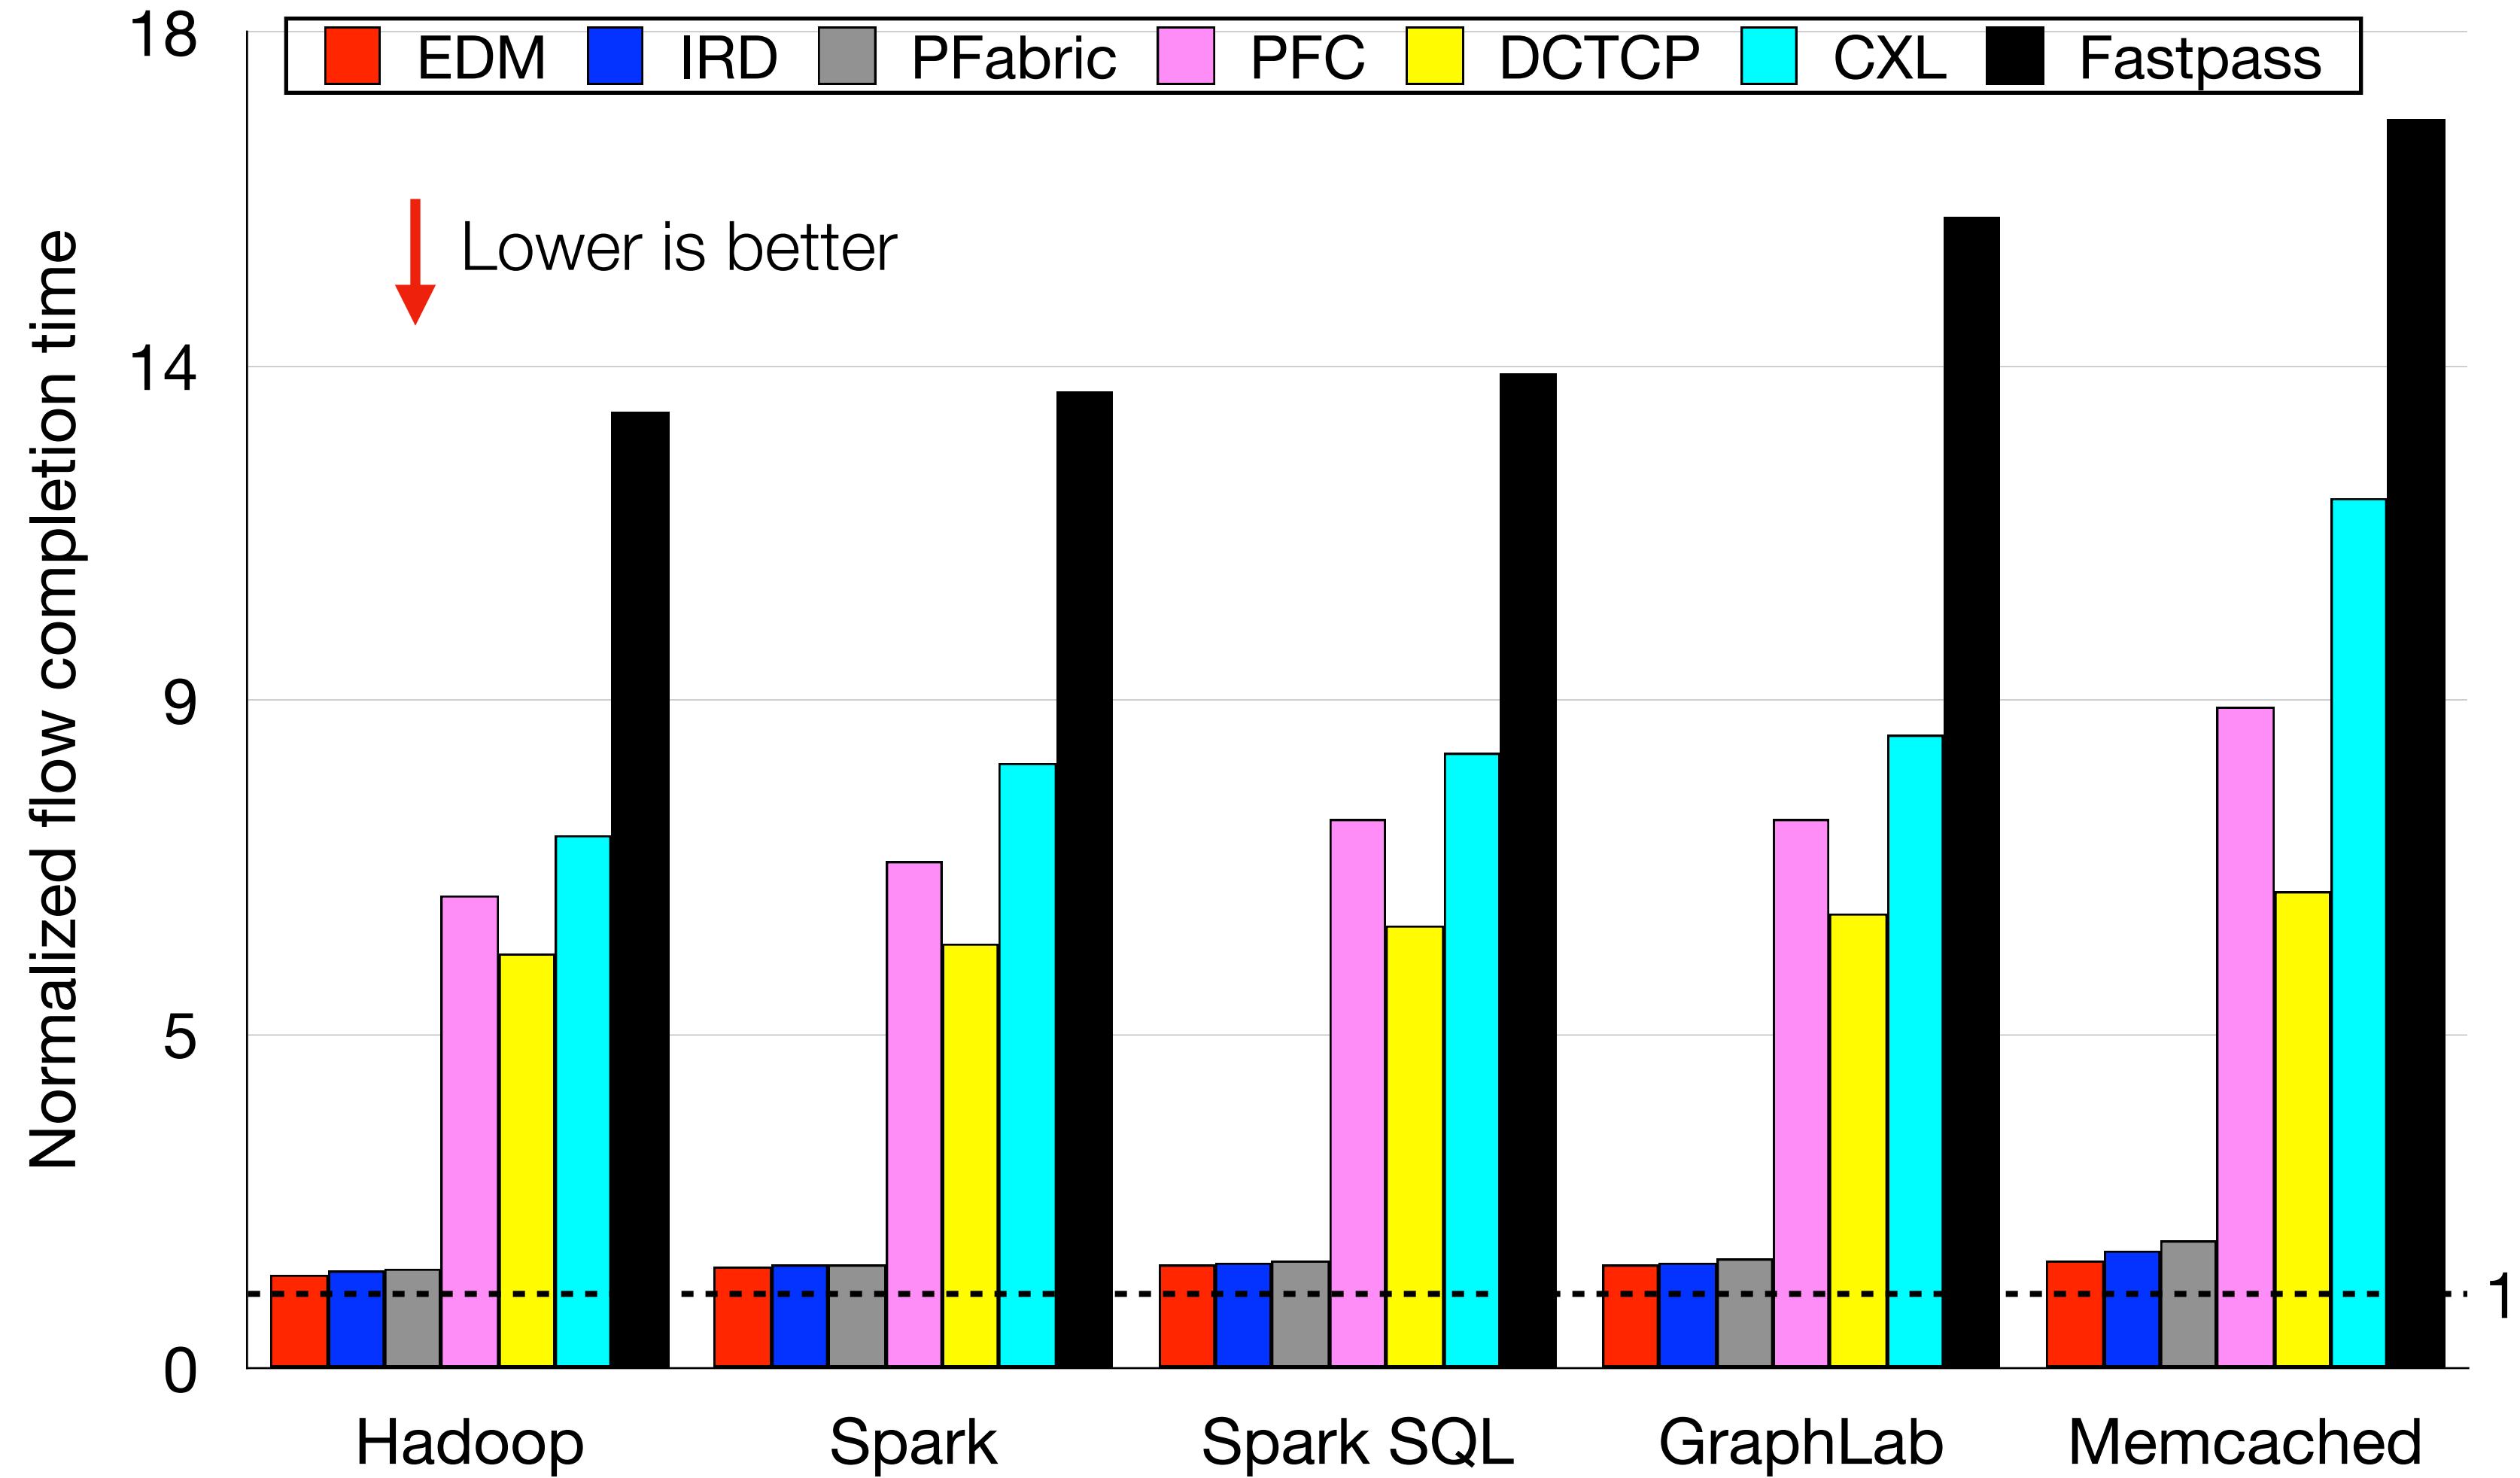
Network Simulator

- A single rack with 144 nodes
- Fed with real-world traces
- Compare against 6 classes of scheduling / congestion control

Evaluation Result

- Disaggregated workloads in a loaded network

Experiment name	Dataset
Hadoop, Spark	Generator @ sortbenchmark.org
Spark SQL	Big Data Benchmark@Berkeley
GraphLab	Movie rating data @ Netflix
Memcached	KV-store@ YCSB



Summary

- EDM is a low latency Ethernet fabric for memory disaggregation.
- EDM uses two ideas for low latency w/ high bandwidth utilization:
 - EDM implements the **protocol for remote memory access** entirely in the **Ethernet PHY**.
 - EDM implements a **fast, centralized memory traffic scheduler** in the switch's PHY.
- EDM incurs a latency of **~300ns (7x lower than RoCE)** in an unloaded network, and **< 1.3x** its unloaded latency under heavy network loads.

Thank you!

Code: <https://github.com/wegul/EDM>